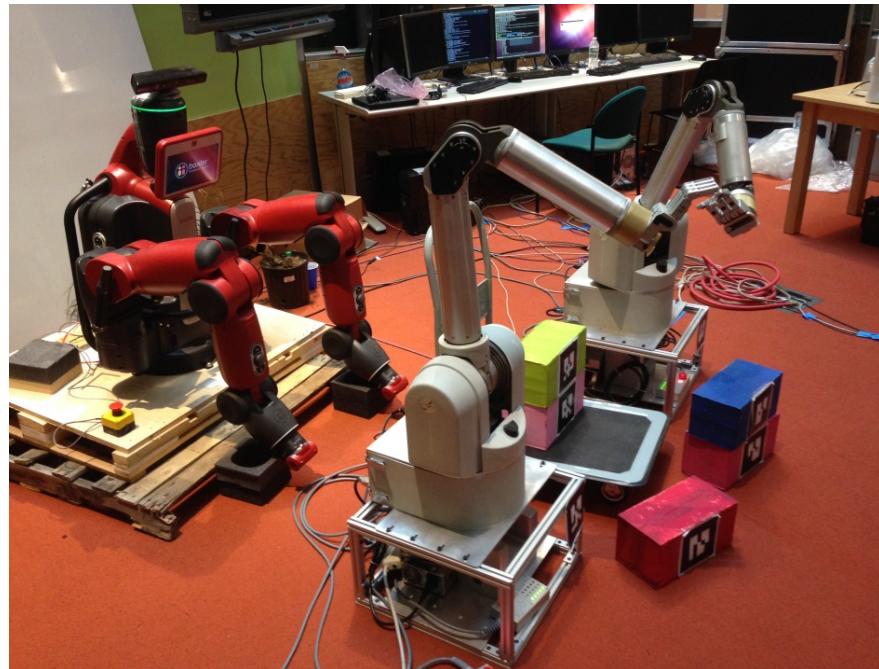


1st Summer School on Cognitive Robotics

i.e., How to create “thinking” robots



Introduction: Architectures for Autonomy

Brian Williams and Tiago Vaquero

June 12th, 2017

Outline

- **Summer School in a Nutshell**
- Architectures for Autonomy
- Programming Cognitive Robots

Summer School is about Creating Goal-directed Systems



1. Decision-making components for generating, refining, executing and monitoring plans on line ([15 Tutorials](#)).
2. Architectures and principles for creating goal-directed systems from decision-making components
(3T, Remote Agent, Claraty, MDS, Idea, [Enterprise/ROS](#), ROSplan).
3. Dynamic languages, elevated to the goal-level through partial specification and operation on state
(PRS, RAPS, TCC, Golog, [RMPL](#)).

Day Structure

8:30am - 10:00am Tutorial

10:00am - 10:30am Coffee break

10:30am - 12:00pm Tutorial

12:00pm - 1:00pm Lunch

1:00pm - 2:30pm Tutorial

2:30pm - 3:00pm Coffee break

3:00pm - 6:00pm Labs

6:00pm - 7:00pm Open lab hours

Today

9:30-10:20

10:20-10:30

Tutorials and Labs

Day 1: Robust Execution Introduction 1: Architectures for Autonomy Tutorial 2: Self-Monitoring, Self-Diagnosing Systems Tutorial 3: Temporal Networks for Dynamic Scheduling	Lab: Enterprise/ROS Familiarization and Robust Execution
Day 2: Motion Planning Tutorial 4: Sampling-based Motion Planning Tutorial 5: Single and Multi-Robot Path Planning with Quality Guarantees Tutorial 6: Trajectory Optimization for Underactuated Robots	Lab: Incorporating Motion Planning for Autonomous Vehicles
Day 3: Activity Planning Tutorial 7: Classical Planning@Robotics: Methods and Challenges Tutorial 8: Planning in Hybrid Domains Tutorial 9: Planning of Concurrent Timelines	Lab: Incorporating Activity Planning
Day 4: Perception and Manipulation Tutorial 10: Multi-vehicle Routing with Time Windows Tutorial 11: Generative Models for Perception Tutorial 12: Fundamentals of Robotic Manipulation and Grasping	Lab: Manipulation and Multi-vehicle Routing
Day 5: Planning with Uncertainty and Risk Tutorial 13: Probabilistic Planning Tutorial 14: Localization and Mapping Tutorial 15: Risk-bounded Planning and Scheduling	Lab: Challenge, as it all comes together

Speakers



Andrew Coles
King's College London



John W. Fisher
MIT



Joerg Hoffmann
Saarland University



Luke Hunsberger
Vassar College



Lydia Kavraki
Rice University



Phil Kilby
CSIRO Data61 / ANU



Sven Koenig
USC



Scott Kuindersma
Harvard University



Mark Moll
Rice University



Alberto Rodriguez
MIT



Nicholas Roy
MIT



David Wang
NuVu



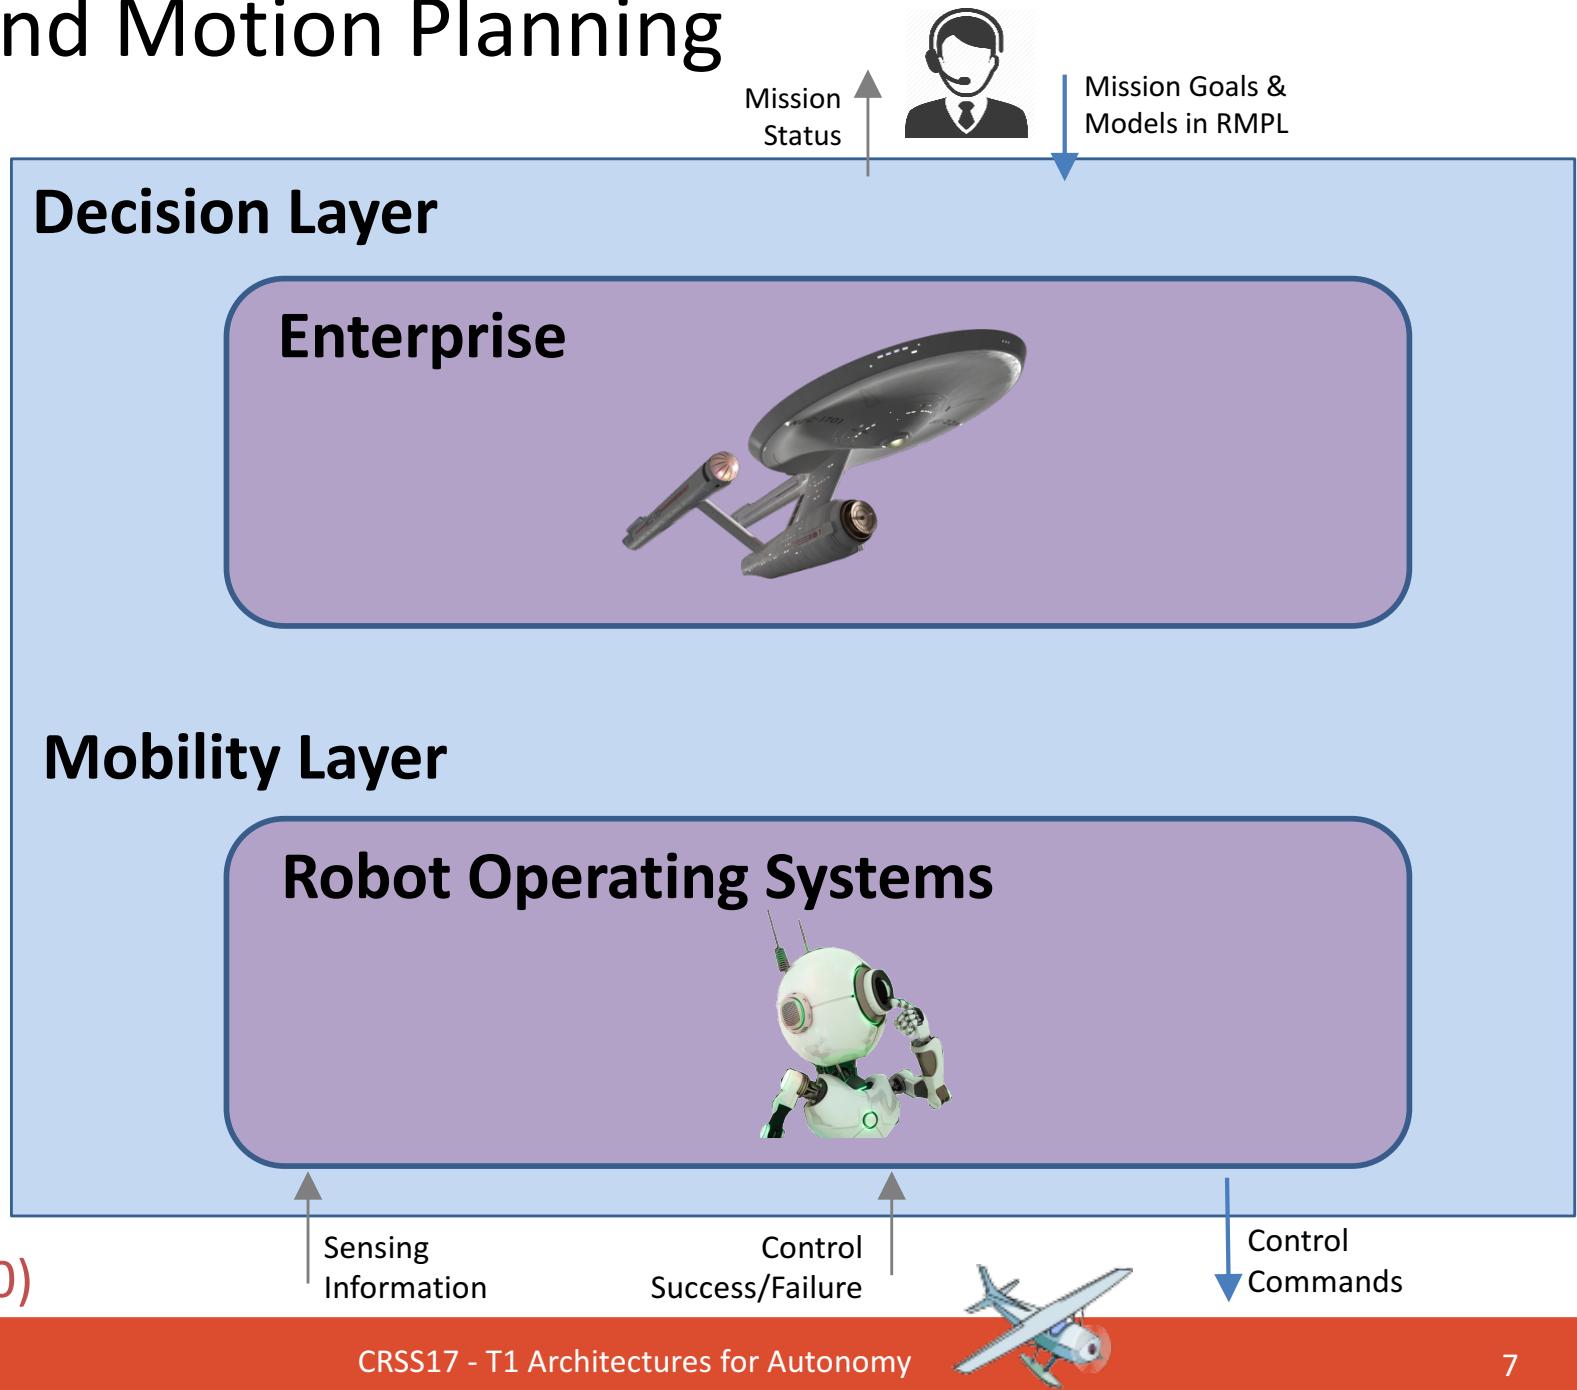
Brian Williams
MIT



Shlomo Zilberstein
UMass Amherst

Labs use a Goal-directed Executive for Task and Motion Planning

ANERS



Labs embedded in simulations and hardware

ANERS

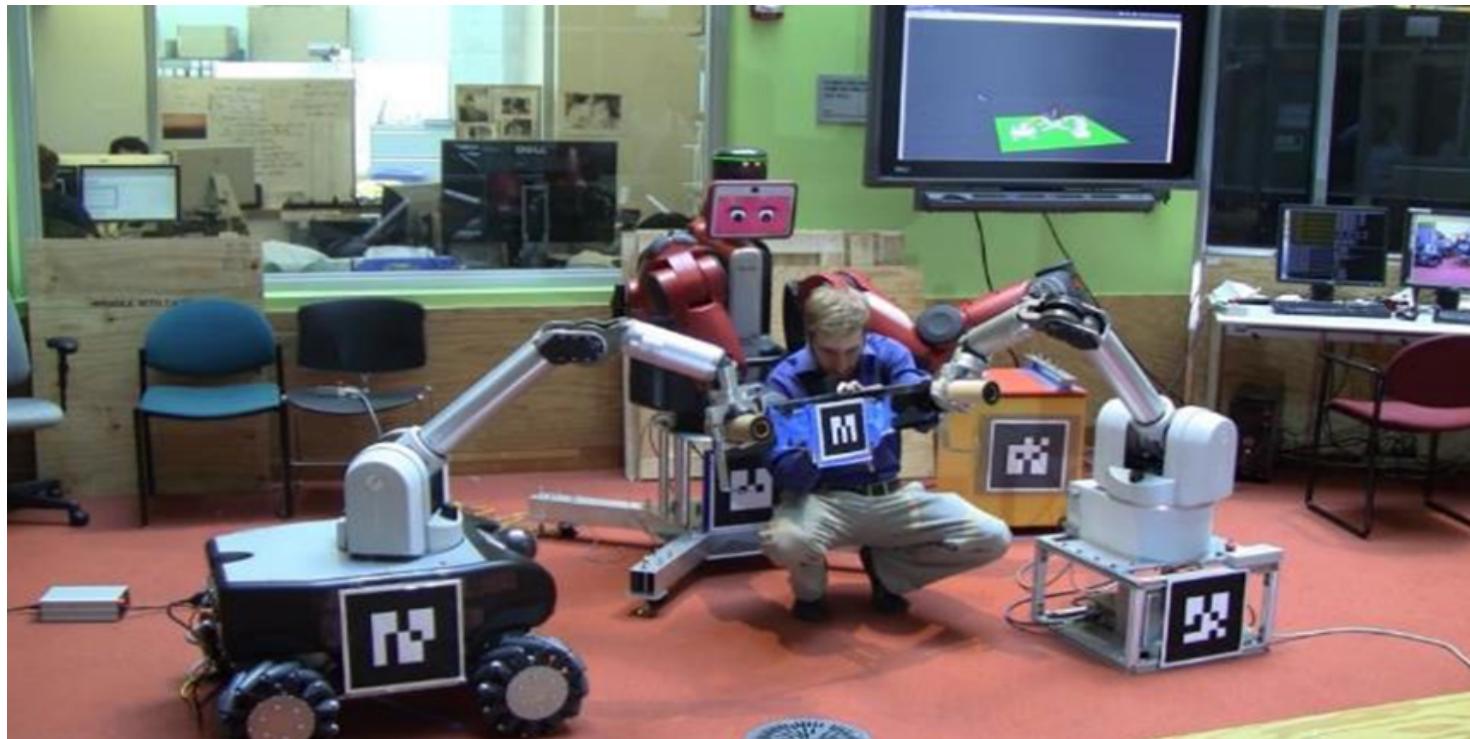


Working out-of-the-box with your '15 holiday gift

Labs Driven by a Grand Challenge



Labs



**Model-based Embedded and Robotic Systems
Group (MERS)**

Room 226, 2nd Floor, The Stata Center, Building 32

Be warned – bleeding edge!

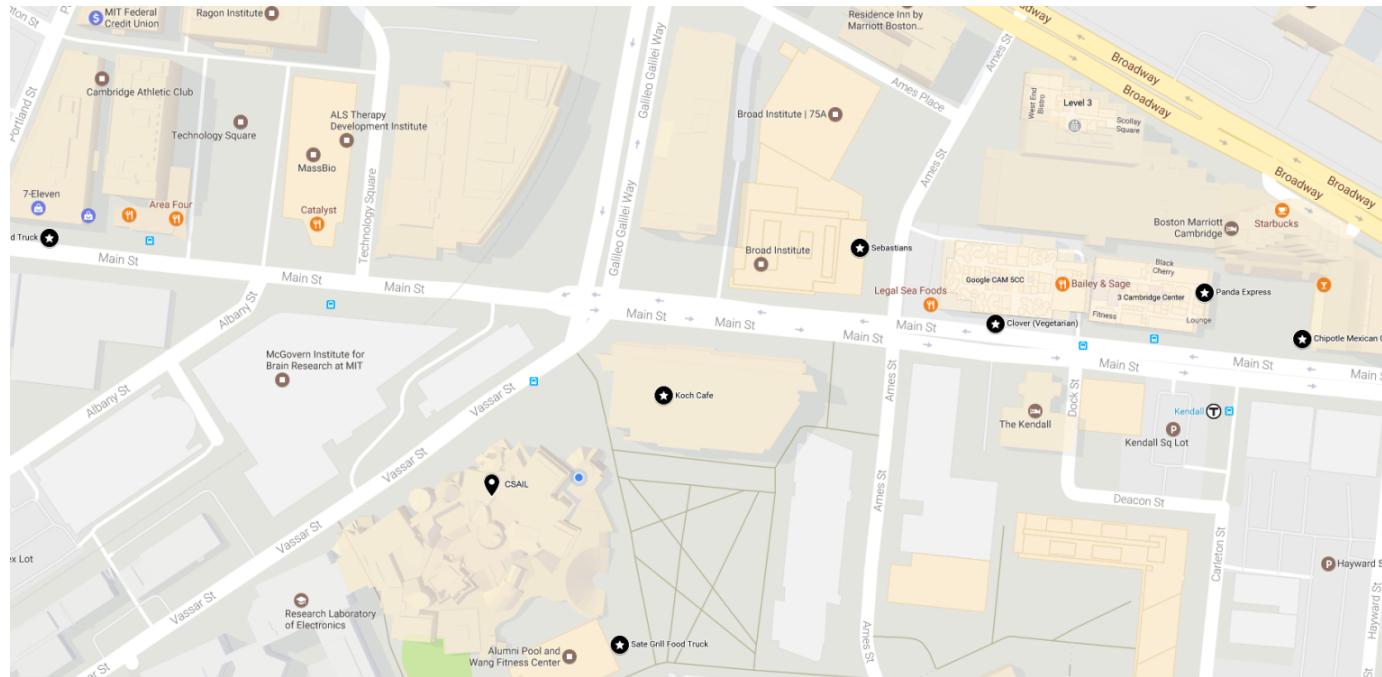


Lunch



You supply! Many options around campus . . Explore!

Student Center include Cambridge Grill, Shawarma Shack (Middle Eastern Food), Cafe Spice, Laverde's, and Subway, food trucks, sub shops, take out, brew pubs and gourmet restaurants



Clover
(Vegetarian), Flour
Bakery, Koch Cafe,
Sebastians, Sata
Grill Food Truck,
Chinese Food
Truck, Chipotle
Mexican Grill, and
Panda Express.

Social Dinner – Summer Shack



Tuesday, June 13 @ 7pm

Transportation not provided. Taking subway recommended

Need Help or Have Questions?

Talk to one of the **organizers** or the **MERS group team**.
Look for folks with **blue lanyard**

Need directions?

Use the ***Summer School Booklet*** as a reference

Most Important

Learn and
Have Fun!

Tiago, Brian and the MERS Team

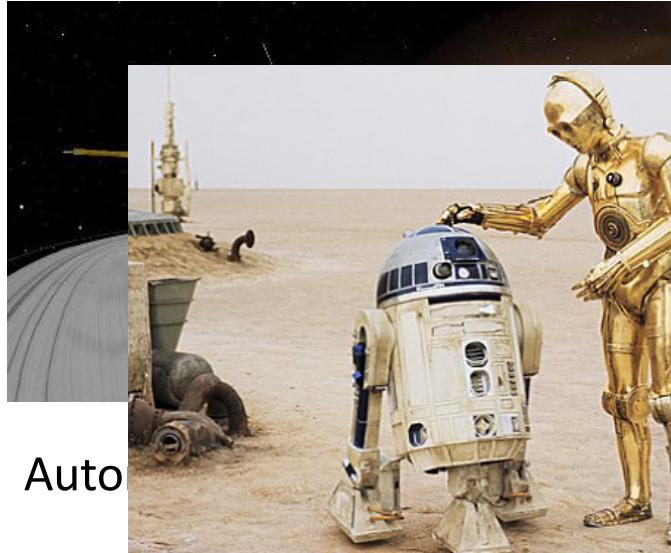
Outline

- Summer School in a Nutshell
- **Architectures for Autonomy**
- Programming Cognitive Robots

Some Papers to Read

- R. Alami, R. Chatila, S. Fleury, M. Ghallab, F. Ingrand, An Architecture for Autonomy, The International Journal of Robotics Research, Vol 17, Issue 4, 1998.
- R. Bonasso, D. Kortenkamp, D. Miller, M. Slack. Experiences with an architecture for intelligent, reactive agents, JETAI, 1996.
- G. Brown, D. Bernard, R. Rasmussen, Attitude and articulation control for Cassini spacecraft: A fault tolerance overview. In 14th AIAA/IEEE Digital Avionics Systems Conference, 1995.
- M. Cashmore, M. Fox, D. Long, D. Magazzeni, B. Ridder, A. Carrera, N. Palomeras, N. Hurtos, M. Carreras. ROSPlan: Planning in the Robot Operating System, In Twenty-Fifth International Conference on Automated Planning and Scheduling, 2015.
- A. Ceballos, S. Bensalem, A. Cesta, L. De Silva, S. Fratini, F. Ingrand, J. Ocon, A. Orlandini. A goal-oriented autonomous controller for space exploration. Proceedings of ESA Symposium on Advanced Space Technologies in Robotics and Automation (ASTRA), vol 11.
- K. Currie , A. Tate , S. Bridge. O-Plan: the Open Planning Architecture, Artificial Intelligence, Volume 52, Issue 1, pp 49-86, November 1991.
- R. Firby. Adaptive execution in complex dynamic worlds. Ph.D. Dissertation, Yale University. 1978.
- R. Firby, Task networks for controlling continuous processes, In Proceedings of the Second International Conference on AI Planning Systems Chicago IL, June 1994.
- E. Gat. Integrating planning and reacting in a heterogeneous asynchronous architecture for controlling real-world mobile robots. In Proc. of AAAI, 1992.
- E. Gat. On Three-Layer Architectures, in Artificial Intelligence and Mobile Robots. MIT Press, 1997.
- K. Haigh, M. Veloso. Interleaving Planning and Robot Execution for Asynchronous User Requests. Autonomous Agents, pp 79-95, 1998.
- B. Hayes-Roth. An Architecture for adaptive Intelligent Systems. Artificial Intelligence 72, 1995.
- J. Laird, A. Newell, P. Rosenbloom. SOAR: An architecture for general intelligence. Artificial Intelligence, 33(1), 1987.
- C. McGann, F. Py, K. Rajan, H. Thomas, R. Henthorn, R. McEwen. T-rex: A model-based architecture for AUV control, in 3rd Workshop on Planning and Plan Execution for Real-World Systems, vol 2007.
- D. Musliner, E. Durfee, K. Shin. CIRCA: A cooperative intelligent, real-time control architecture. IEEE Transactions and Systems, Man, and Cybernetics 23(6), 1993.
- N. Nilsson. Shakey The Robot, Technical Note 323. AI Center, SRI International, 333 Ravenswood Ave., Menlo Park, CA 94025, Apr 1984.
- N. Nilsson, Teleo-Reactive Programs for Agent Control, Journal of Artificial Intelligence, Volume 1, pages 139-158, 1994.
- D. Tran, S. Chien, R. Sherwood, R. Castano, B. Cichy, A. Davies, G. Rabideau. The Autonomous Sciencecraft Experiment Onboard the EO-1 Spacecraft. In Proc. of AAAI 2004.

Examples of Cognitive Robotic Systems



Autono-



mation gathering scouts



ected

tive
tive

rmation



A
and



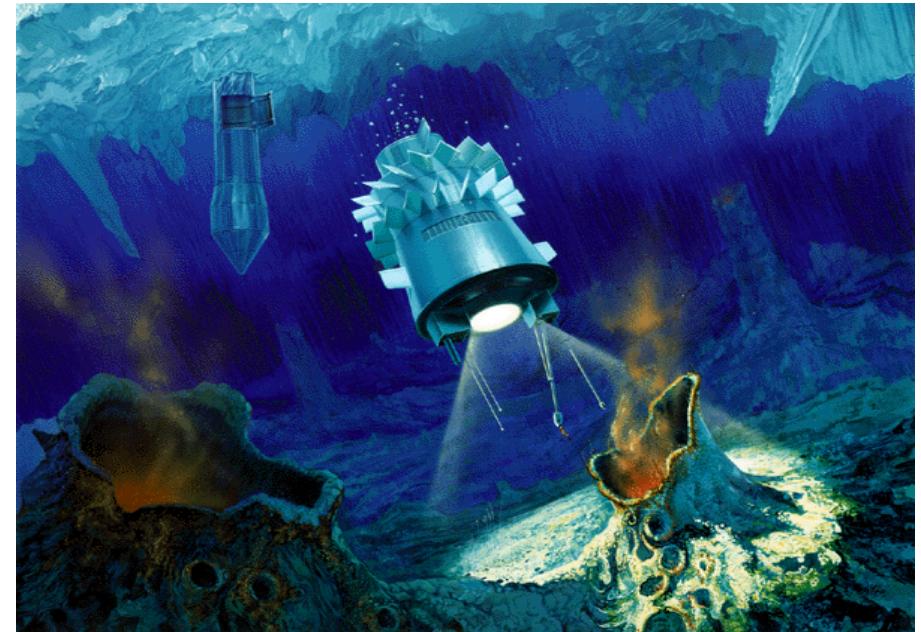
for Man
and Manipulation tasks



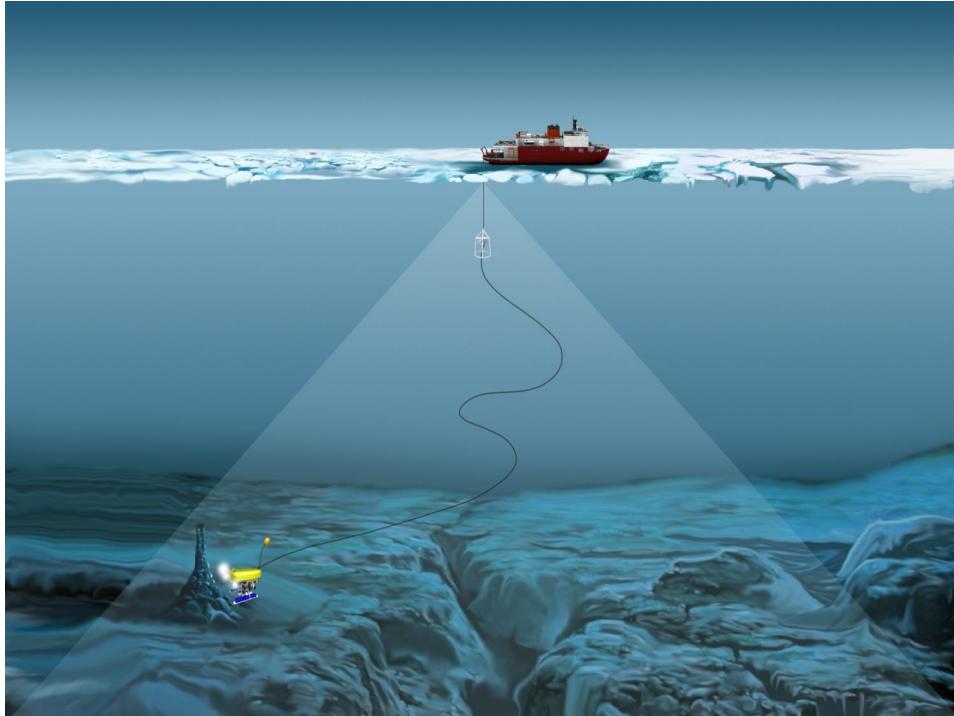
things

The Challenge of Coordinating Robotic Systems

- We are creating **increasingly autonomous** embedded systems that perform **critical functions** over **long** periods of time.
- These long-lived systems achieve **robustness** by coordinating a complex **network** of devices.
- Architecting and programming these systems robustly is becoming an increasingly daunting task.



“Autonomous” vehicles explore far away places .. but can end in disaster!

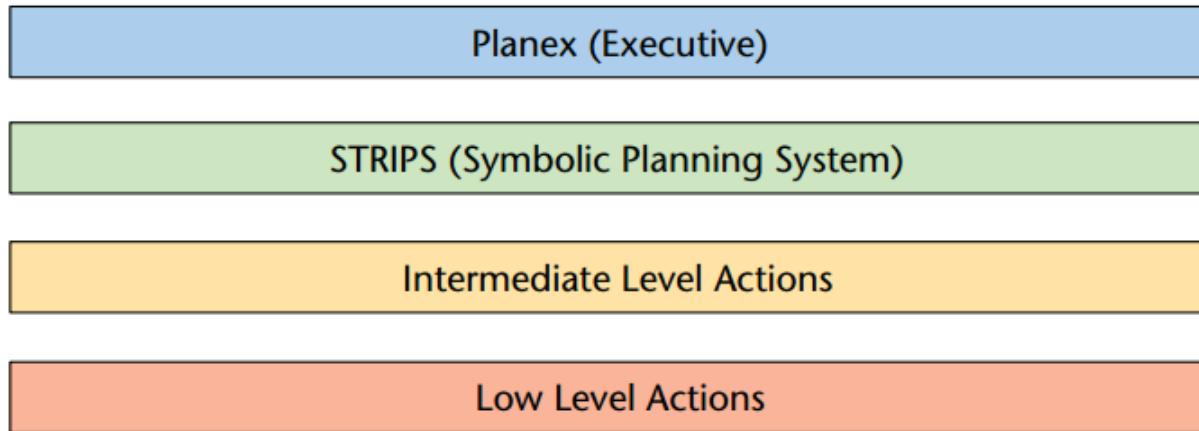


Most are just scripted!

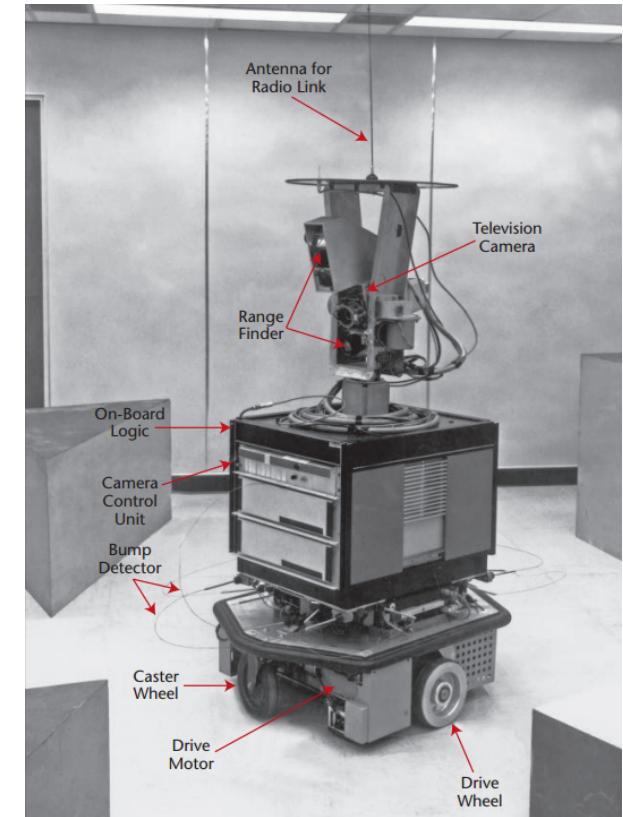
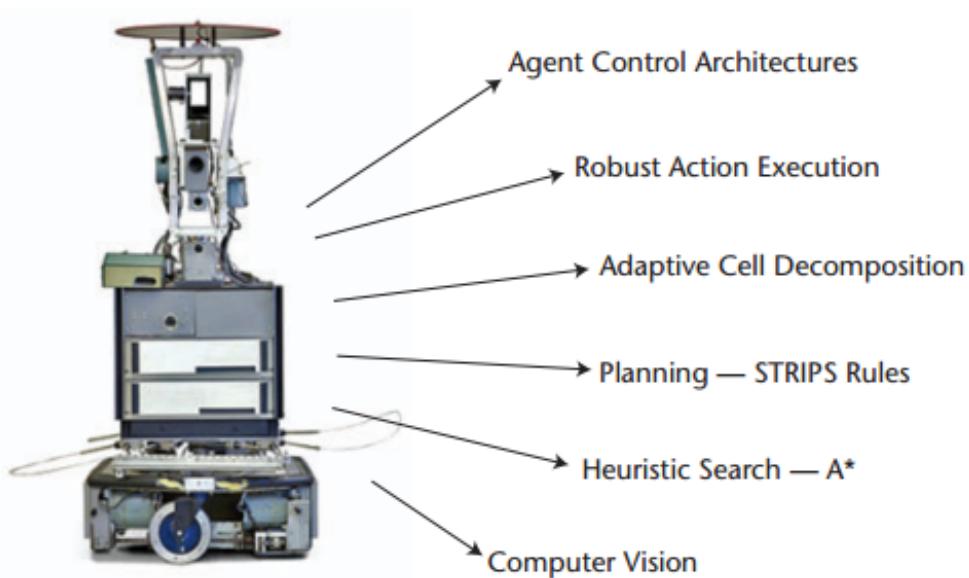
A Brief History of Architecting Robust Autonomous Systems

ANSERS

Shakey the Robot (SRI, 1966-72)



Layered Software Architecture



Shakey

Emphasizes Planning

- Activity
- Motion

Some similarities to RosPlan (Cashmore et al., 2015)

ANSRS

Layered Subsumption Architecture (Rod Brooks, 1986)



Behaviour-Based Control: Subsumption Architecture

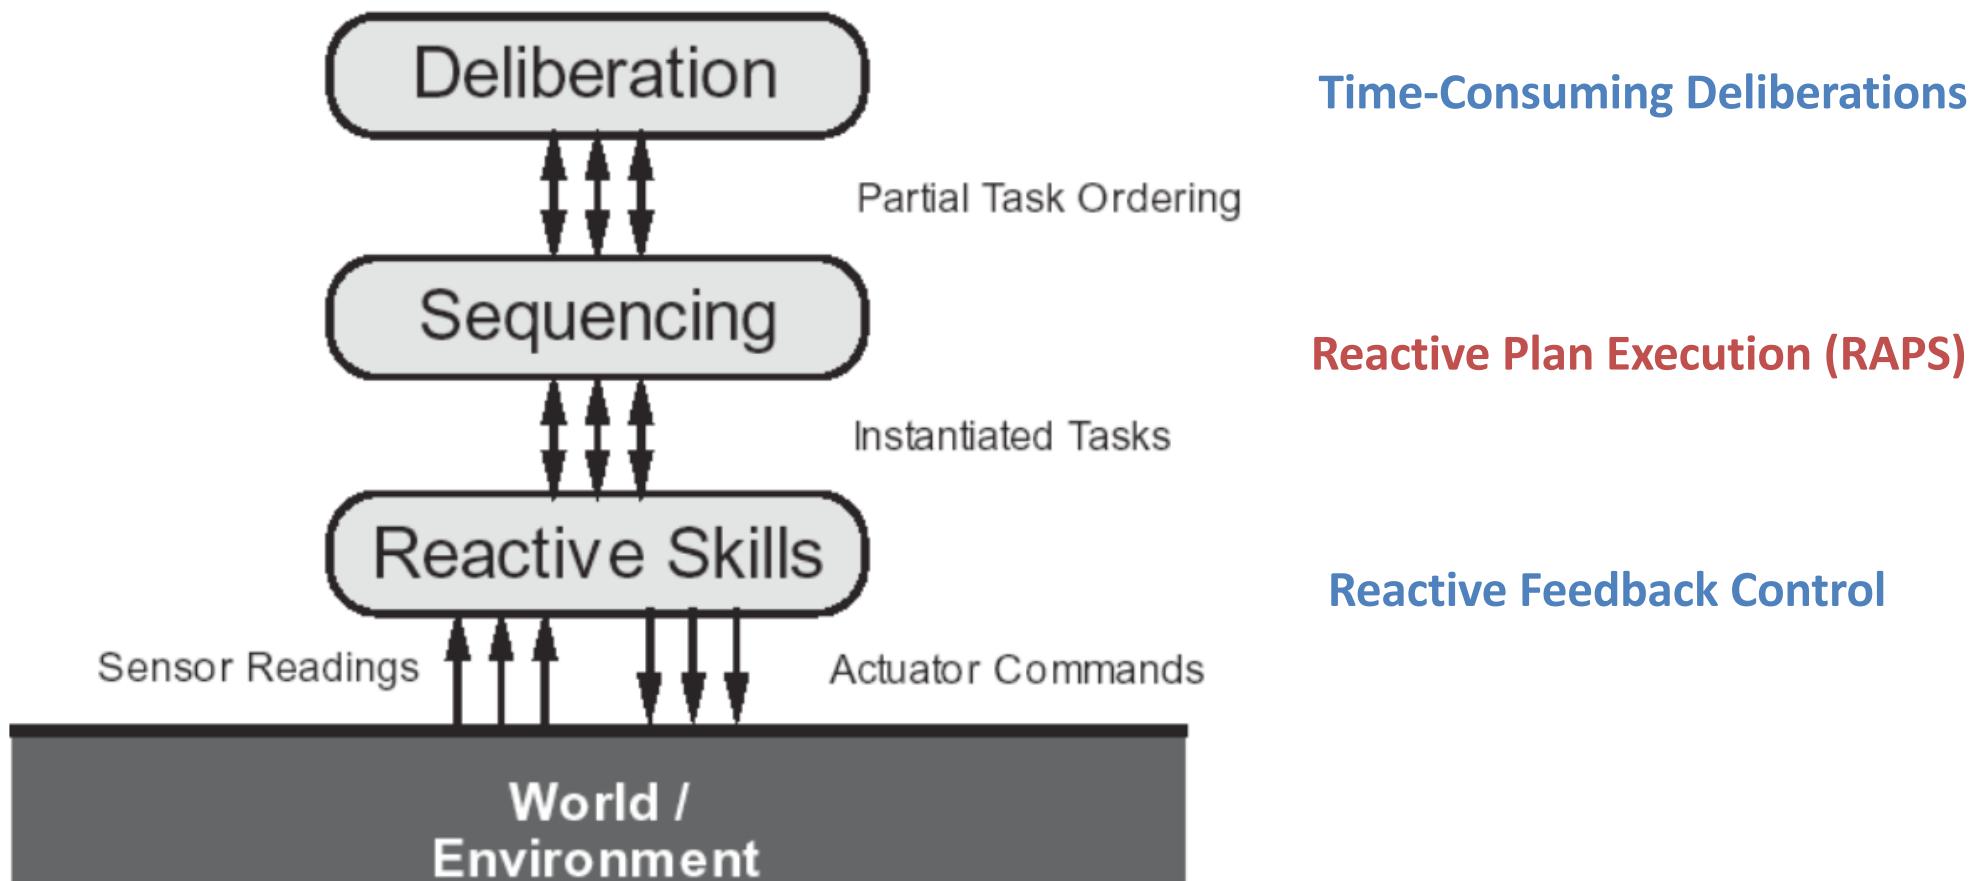
Modules act in parallel, in layers of increasing priority



Focus on layered reactive systems

Similarities to MOOS (Benjamin, Leonard et al.)

Atlantis: Three Layered Architecture (Jim Firby et al., 1990)



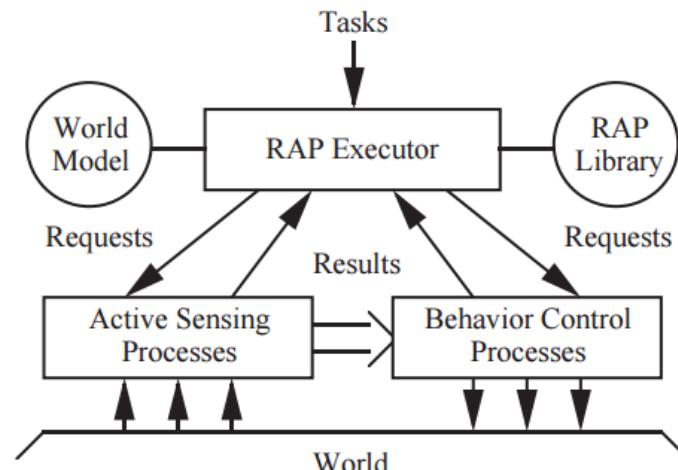
Similarities to 3T(Bonasso et al.) and many others

The RAPs Execution System (James Firby 1989)

```
(Define-rap (attach-at-site ?thesite)
  (succeed (docked ?thesite))
  (method
    (context (ferrous-hull ?thesite))
    (task-net
      (sequence
        (t1 (approach-site ?thesite))
        (t2 (magnetically-attach ?thesite)
          (wait-for (docked ?thesite))))))
  (method
    (context (not (ferrous-hull ?thesite)))
    (task-net
      (sequence
        (t1 (approach-site ?thesite))
        (t2 (grip-attach ?thesite)
          (wait-for (docked ?thesite)))))))
```

Repair Action Procedures (RAP):

- Goal-directed task-decomposition.
- Non-deterministic choice.
- Closed-loop monitoring on goals.



Similarities to Golog (Levesque et al), Alisp (Russel et al , and RMPL (Williams et al)

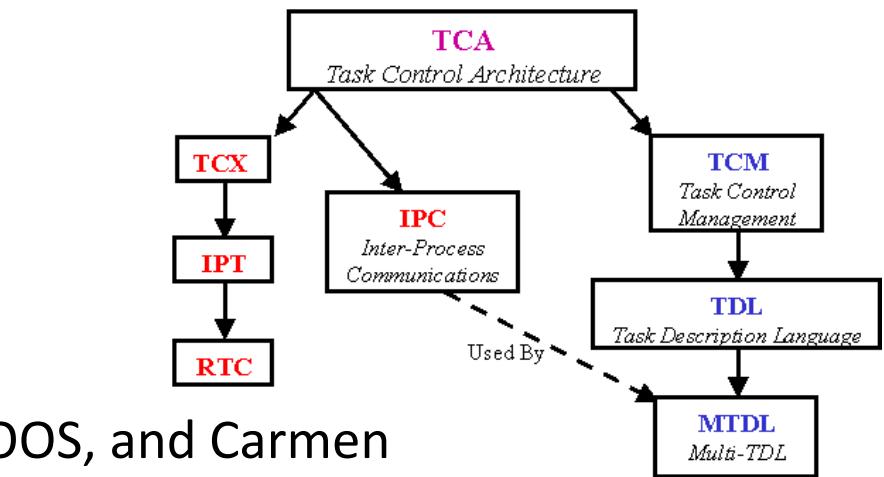
Task Control Architecture (Reid Simmons, 1994)

high-level robot operating system
that balances
goal-directed task decomposition
and reactive control.

Supports:

- distributed communications,
- task decomposition,
- resource management
- execution monitoring
- and error recovery.

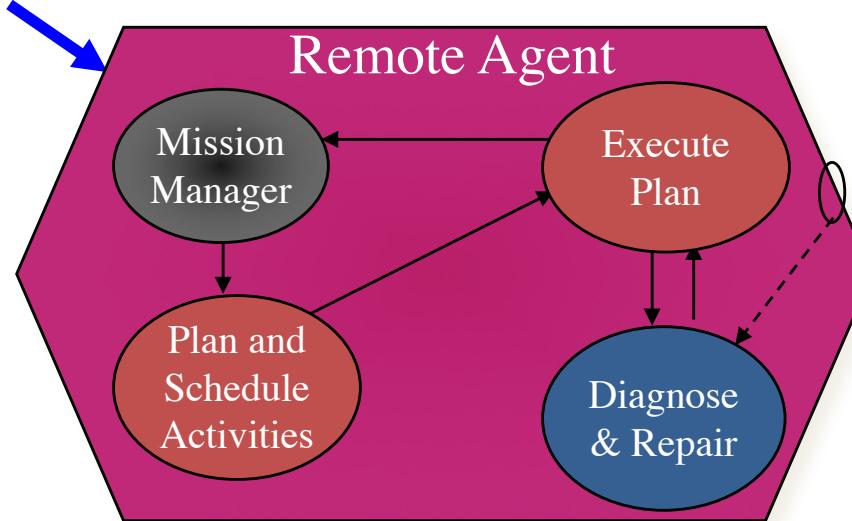
Similarities to Player Stage, Claraty, ROS, MOOS, and Carmen



Remote Agent on Deep Space One, Spring, 1999



Goals



1. Commanded by giving goals
2. Reasons extensively from models
3. Closes loop on goals
4. Diagnoses, repairs and continuously re-plans

Muscettola, Chien; Pell, Gat; Williams, Nayak, et. al.

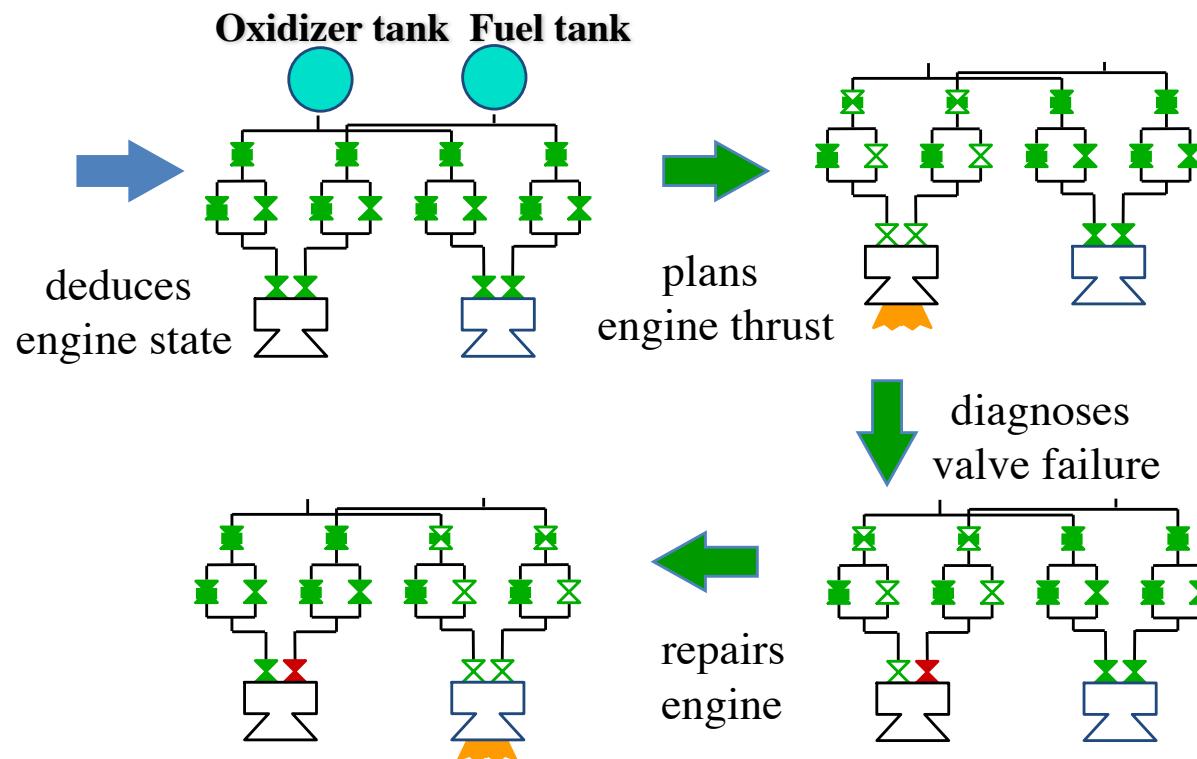
Similarities to IxTet, IDEA, Casper, MDS, TREX, Titan, Enterprise

Remote Agent Planners

(HSTS and Livingstone)



1. Plans concurrent processes (timelines).
2. Plans for time and resources.
3. Plans to control devices with indirect effects.
4. Plans and deduces at reactive times scales.



Programs on Hidden State

(Williams et al., 2003)

AAERS

Planners fully absorbed into execution language.

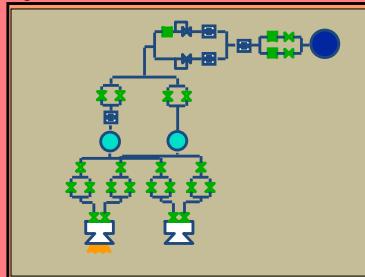
RMPL Model-based Program

Titan Model-based Executive

Control Program

- Executes concurrently
- Preempts
- Queries (hidden) states
- Asserts (hidden) state

System Model



Titan Model-based Executive

Generates target goal states
conditioned on state estimates

State estimates

State goals

Tracks
likely
plant states

Tracks least
cost goal states

Similarities to Esterel, RAPS, TCC, Golog,

Plant
Observations



Commands

And Many More . . .

- PRS (Georgeff, Lansky)
- Propel (Leveinson)
- IDEA (Pell, Muscettola)
- Casper (Chien, Rabideau et al)
- MDS (Rasmussen et al)
- Claraty (Volpe et al)
- IxTet
- TRex
- Moos (Benjamin et al)
- Carmen (Thrun et al)
- Player/Stage (Gerkey et al)
- RosPlan (Cashmore et al)

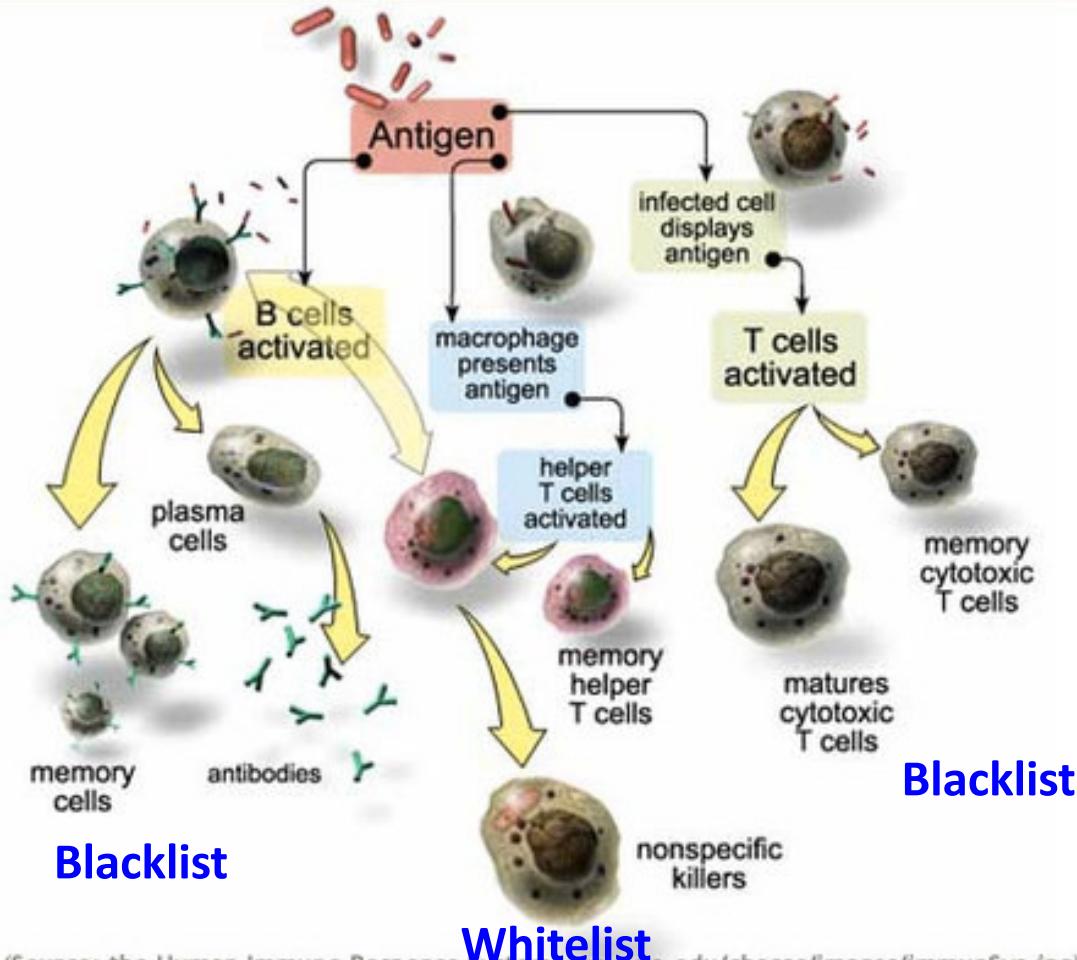


For autonomy to receive broad adoption,
managing risk is key.

What Helps to Make Humans Robust

ANSERS

Immune system cells



(Source: the Human Immune Response System www.uta.edu/chagas/images/immunSys.jpg)

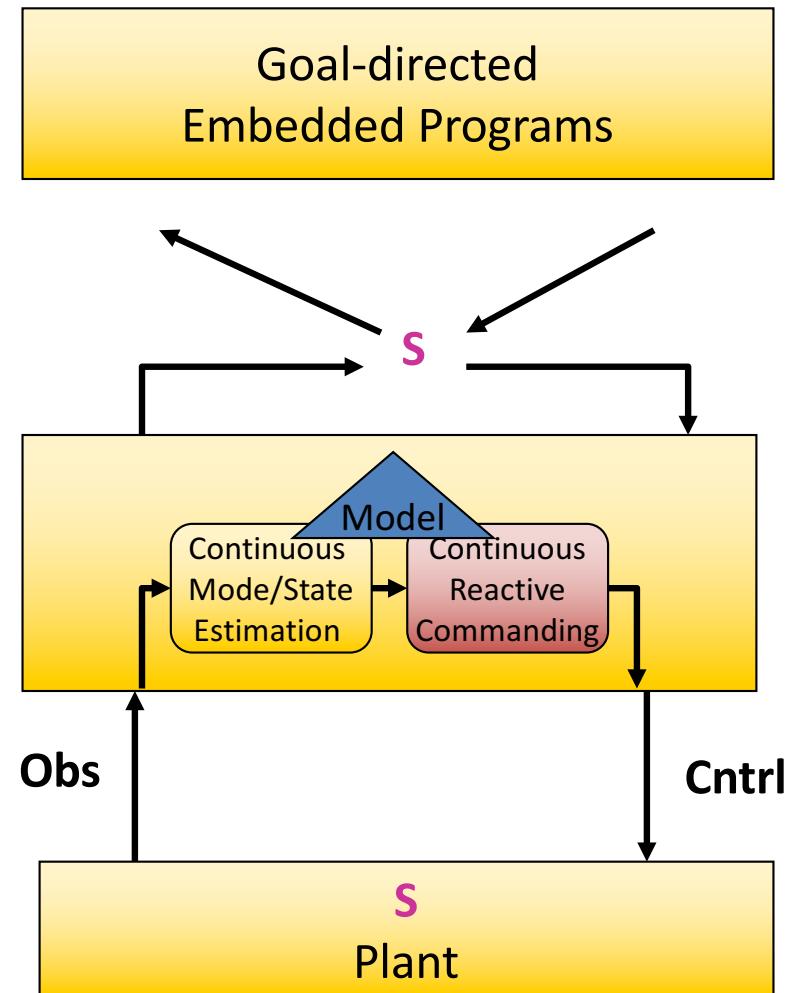
- Immune system
- Physical barriers
- inhospitable environments
- Tolerance
- Death

Complements of Jake Beal, BBN

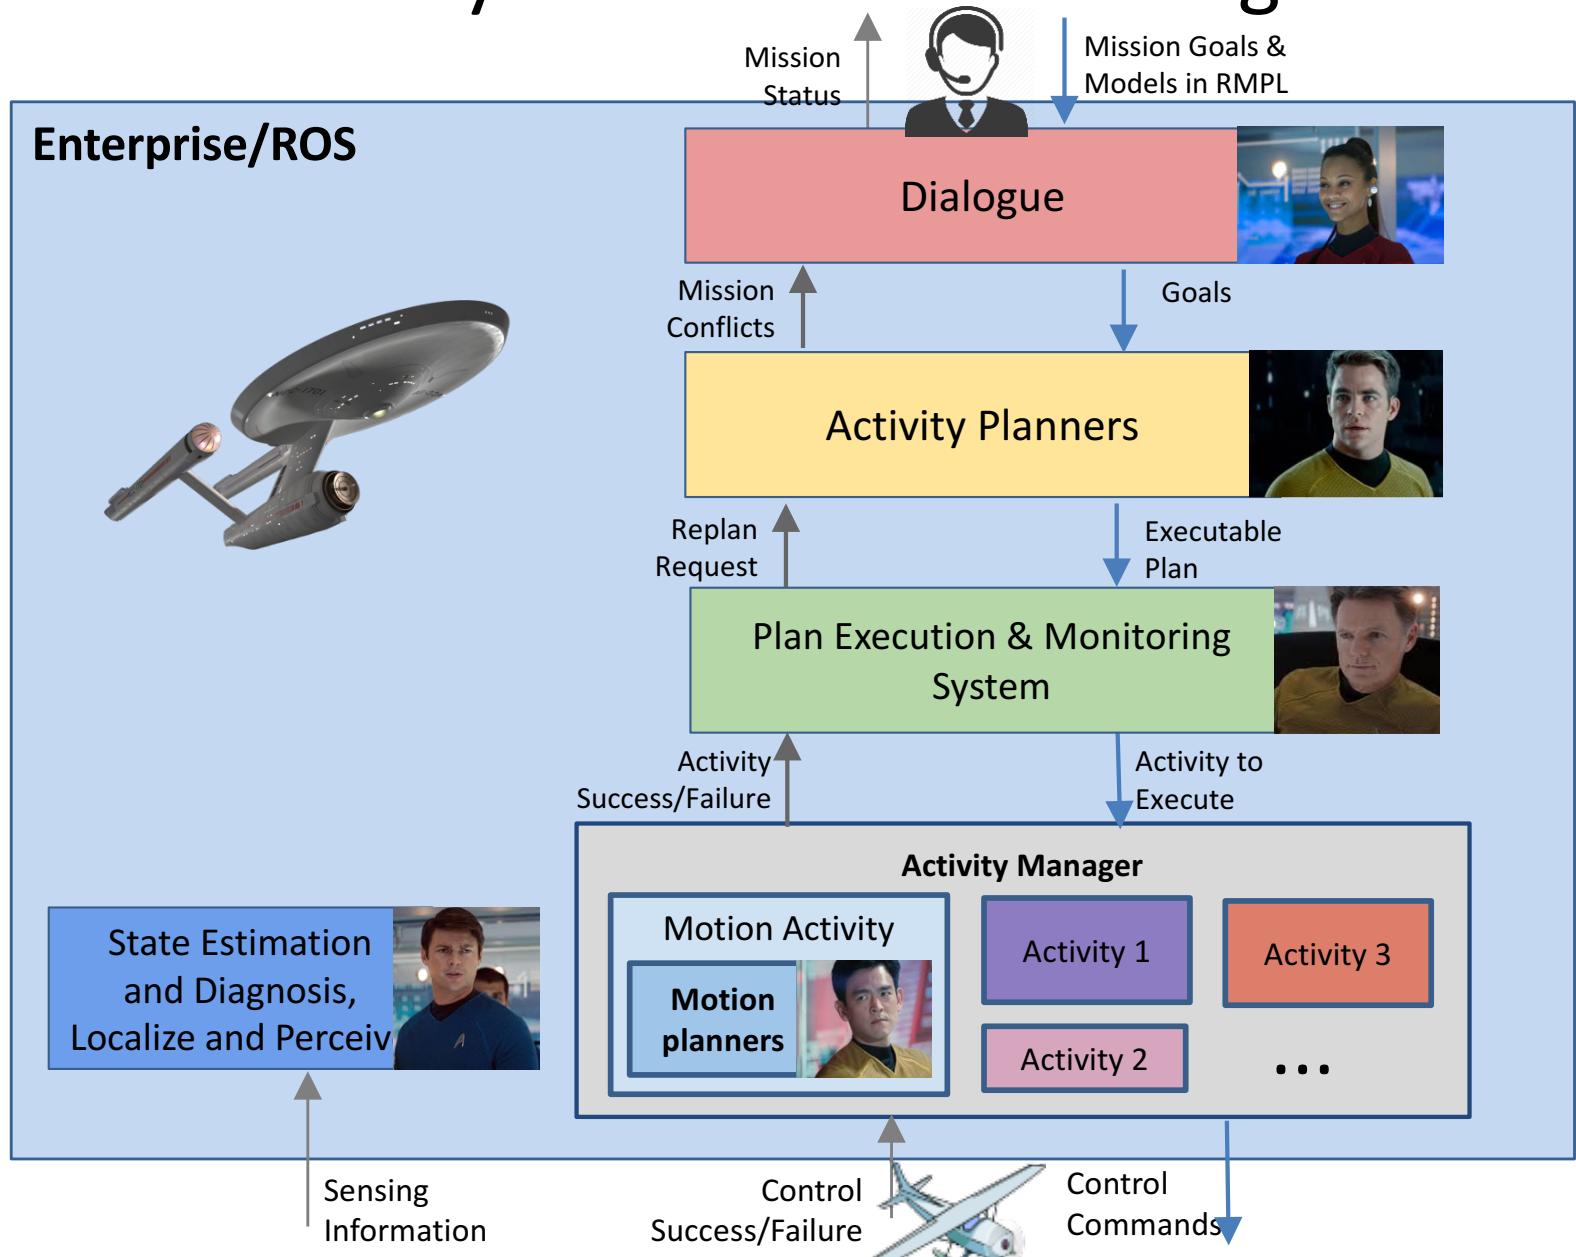
Goal-directed Systems Offer Layers of Defense

Architectures that achieve robustness through decision layers that are:

- **Suspicious (M)**
 - Monitor states and goals.
- **Adapt to disturbance (M)**
 - Adjust timing
 - Select contingencies
- **State Aware (W)**
 - Plan to achieve goals states.
- **Precise (T,R)**
 - Achieve continuous goal states.
- **Manage Risk (F)**
 - Executes plans in uncertain environments with bounded risk.
- **Collaborate (~M)**
 - Execute tasks, revise goals and plan with humans.

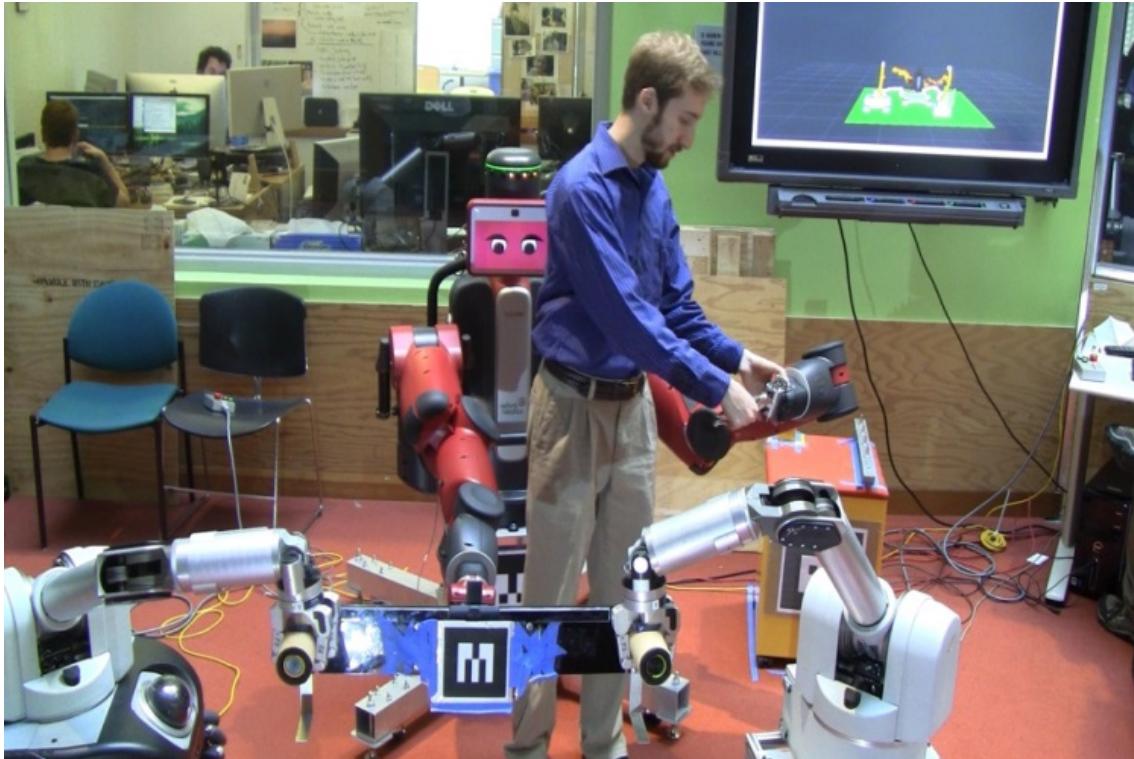


Enterprise/ROS: goal-directed model-based executive for mobility and decision-making



Components:

Planning as State Space Search



- Planning as HFS
- LPA*
- PRM
- RRT
- Belief update

Tutorials: T4, T5, T7, T8, T13, T14, . . .

Components:

Planning as Constraint Programming

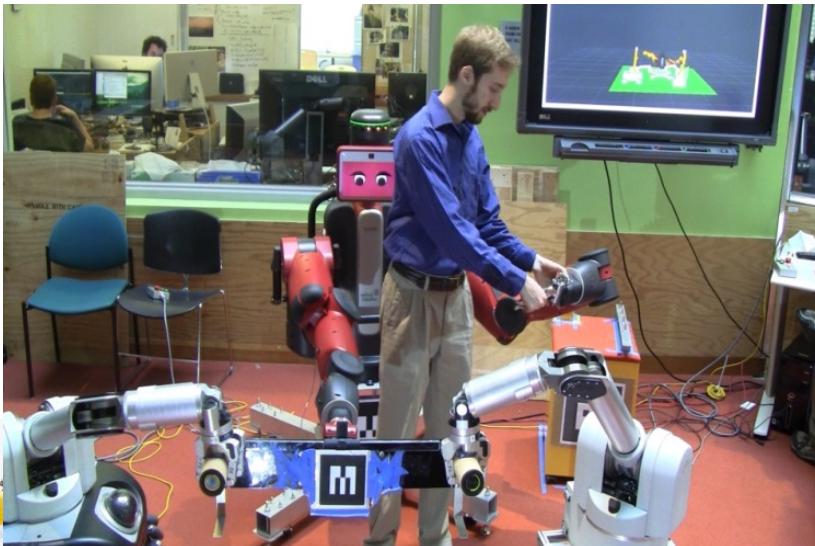
Constraint Types:

- Finite Domain
- Logical
- Temporal
- Global
- Linear
- Convex
- Mixed-integer



Tutorials: T2, T3, T6, T9, T10, T12, T15, . . .

Components: Perception

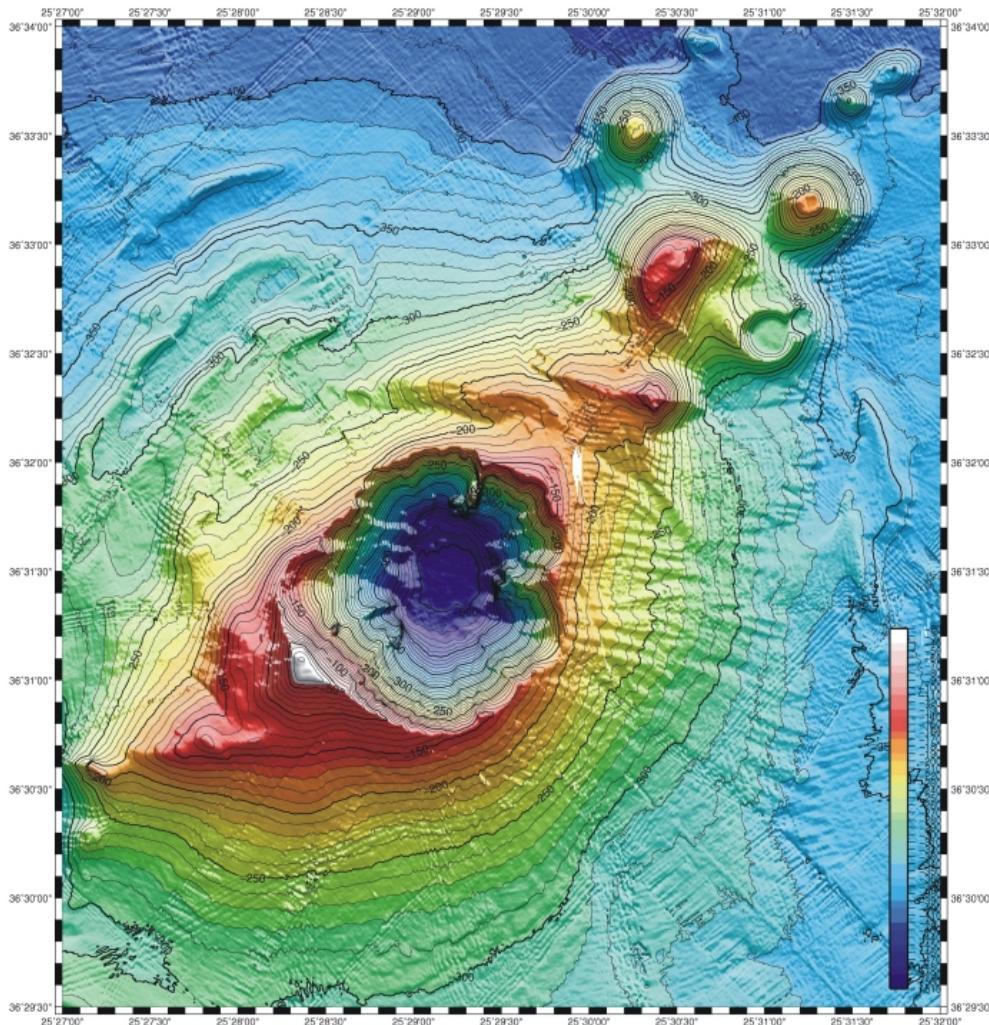


- Goal-monitoring
- Mode estimation
- Graphical models
- SLAM

Tutorials: T2, T11, T14

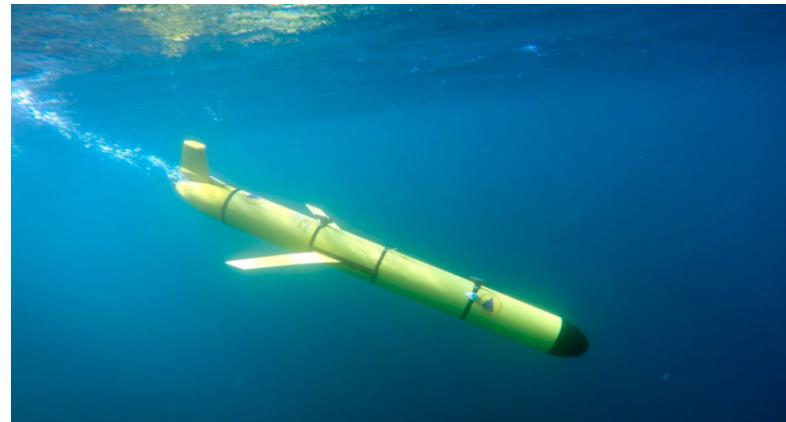
A Challenge that Requires Everything!

2019: Europa Analog Mission Demonstration

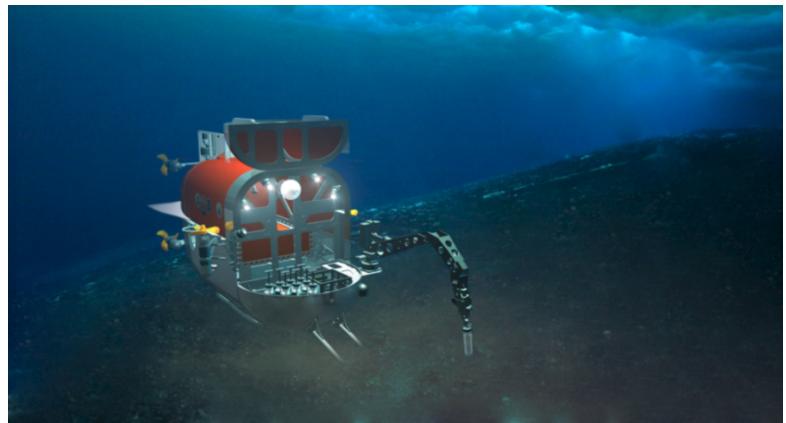


Mission: Look for evidence of “extreme” life at Kolumbo Deep-Sea Volcano near Santorini, Greece

Reconnaissance



Sample return



Sponsor: NASA PSTAR Program

Team: WHOI, MIT, ACFR, U. Michigan

Outline

- Summer School in a Nutshell
- Architectures for Autonomy
- **Programming Cognitive Robots**
 - Programs that make choices
 - Programs that monitor
 - Programs on state

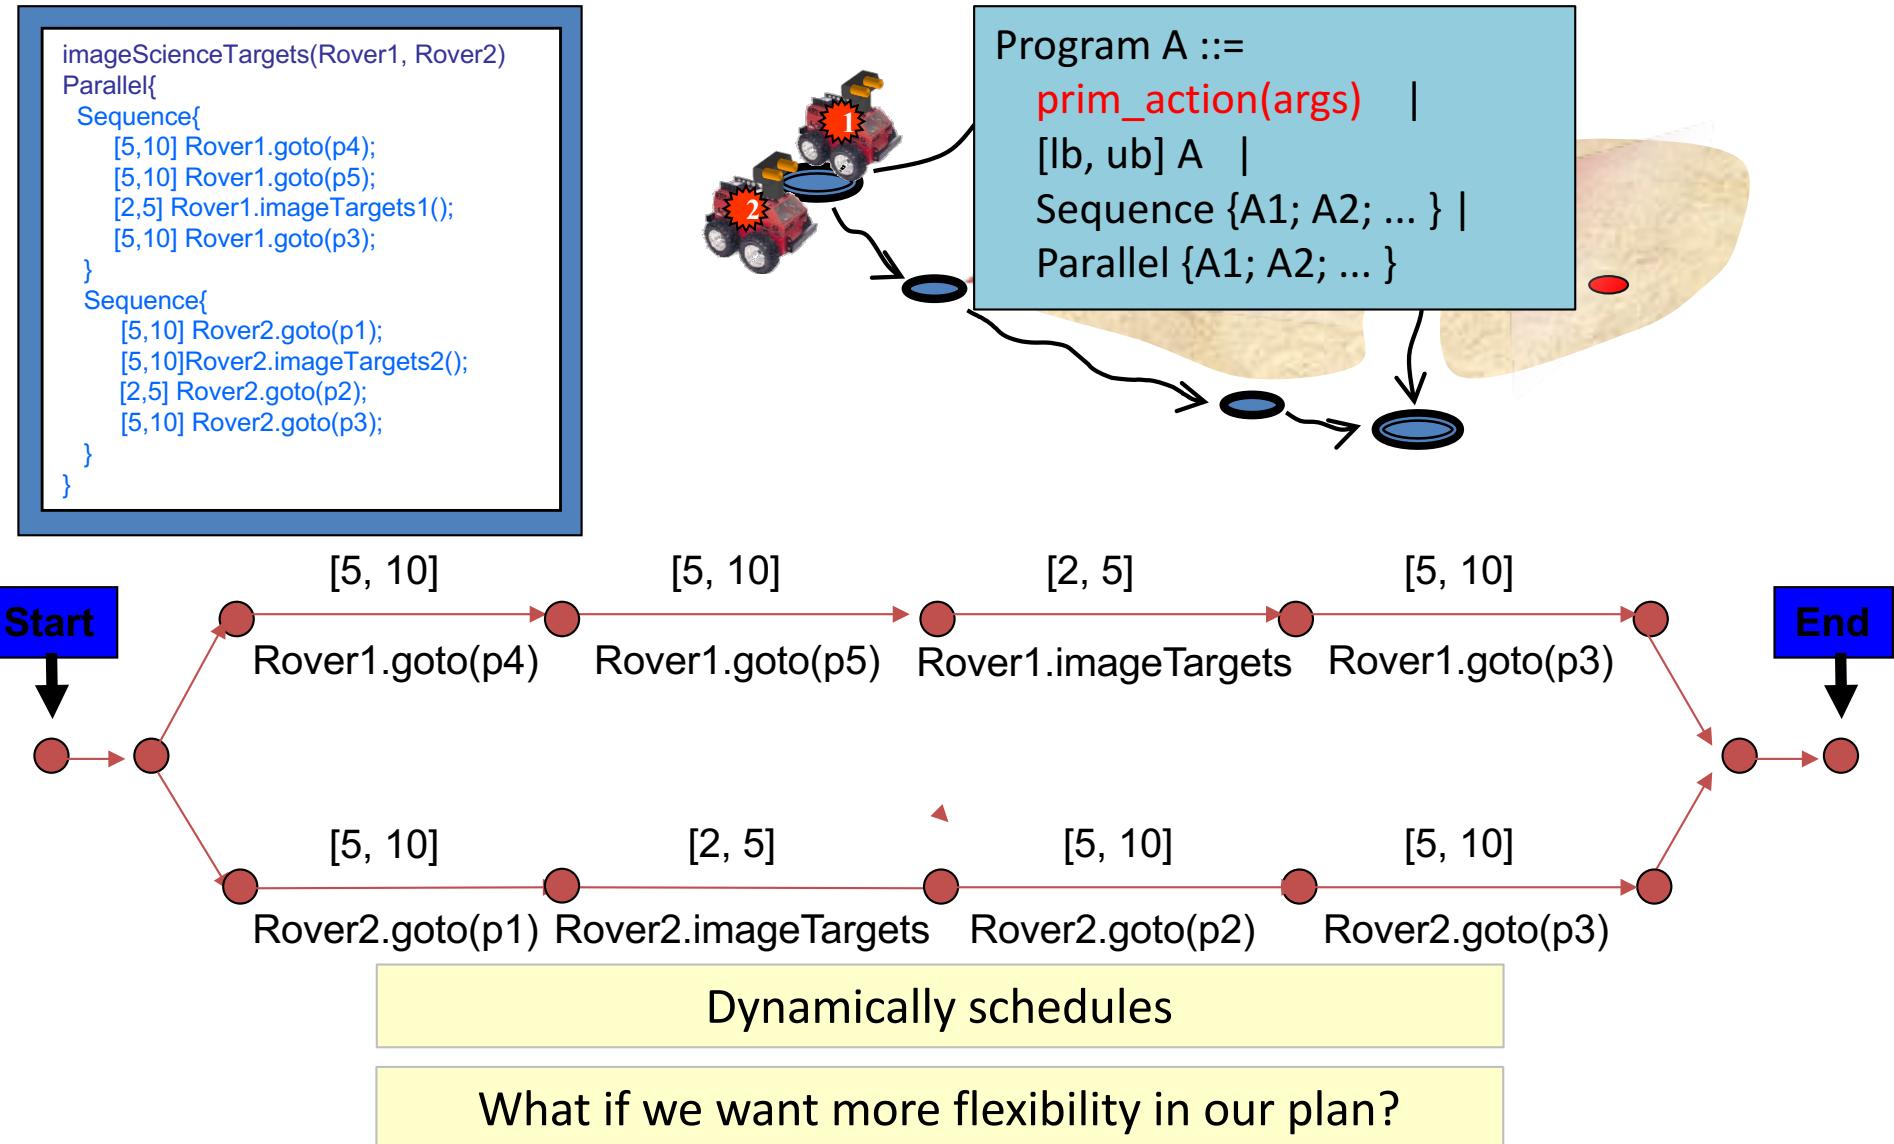
Non-deterministic Languages

```
(Define-rap (attach-at-site ?thesite)
  (succeed (docked ?thesite))
  (method
    (context (ferrous-hull ?thesite))
    (task-net
      (sequence
        (t1 (approach-site ?thesite))
        (t2 (magnetically-attach ?thesite)
          (wait-for (docked ?thesite)))))))
  (method
    (context (not (ferrous-hull ?thesite)))
    (task-net
      (sequence
        (t1 (approach-site ?thesite))
        (t2 (grip-attach ?thesite)
          (wait-for (docked ?thesite)))))))
```

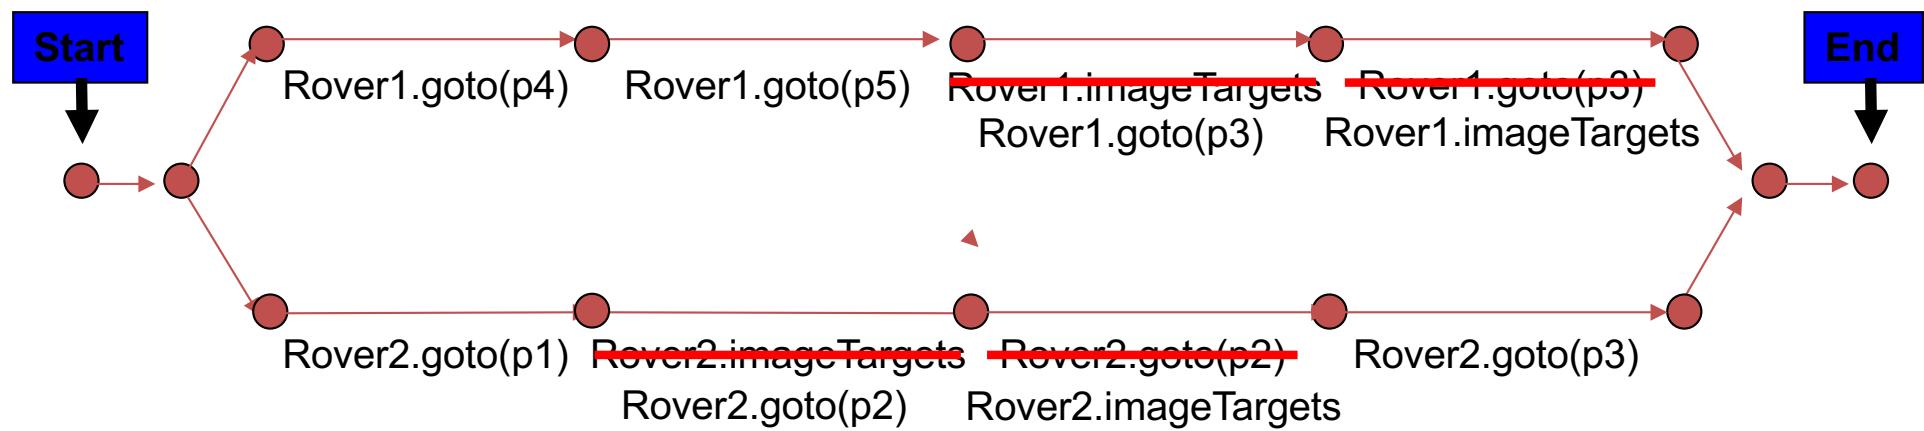
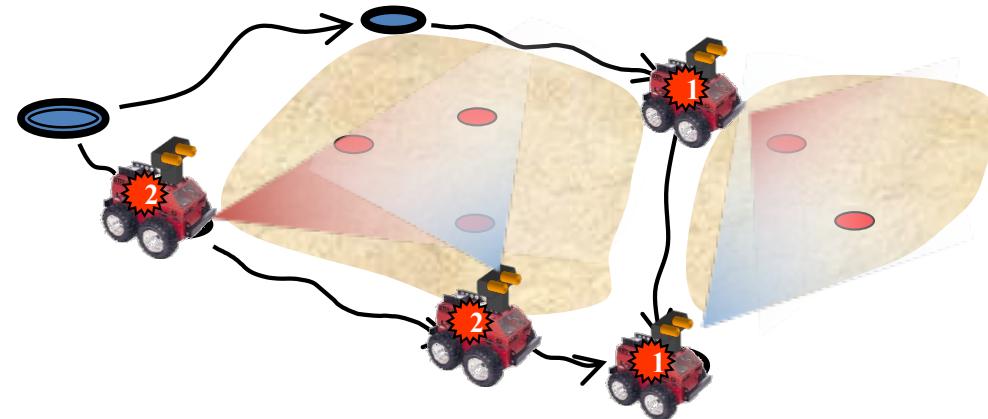
- RAPS (Firby)
- ESL (Gat)
- ALISP (Russel).
- Golog (Levesque)
- RMPL (Williams)
- ...

Robust Program and Plan Execution

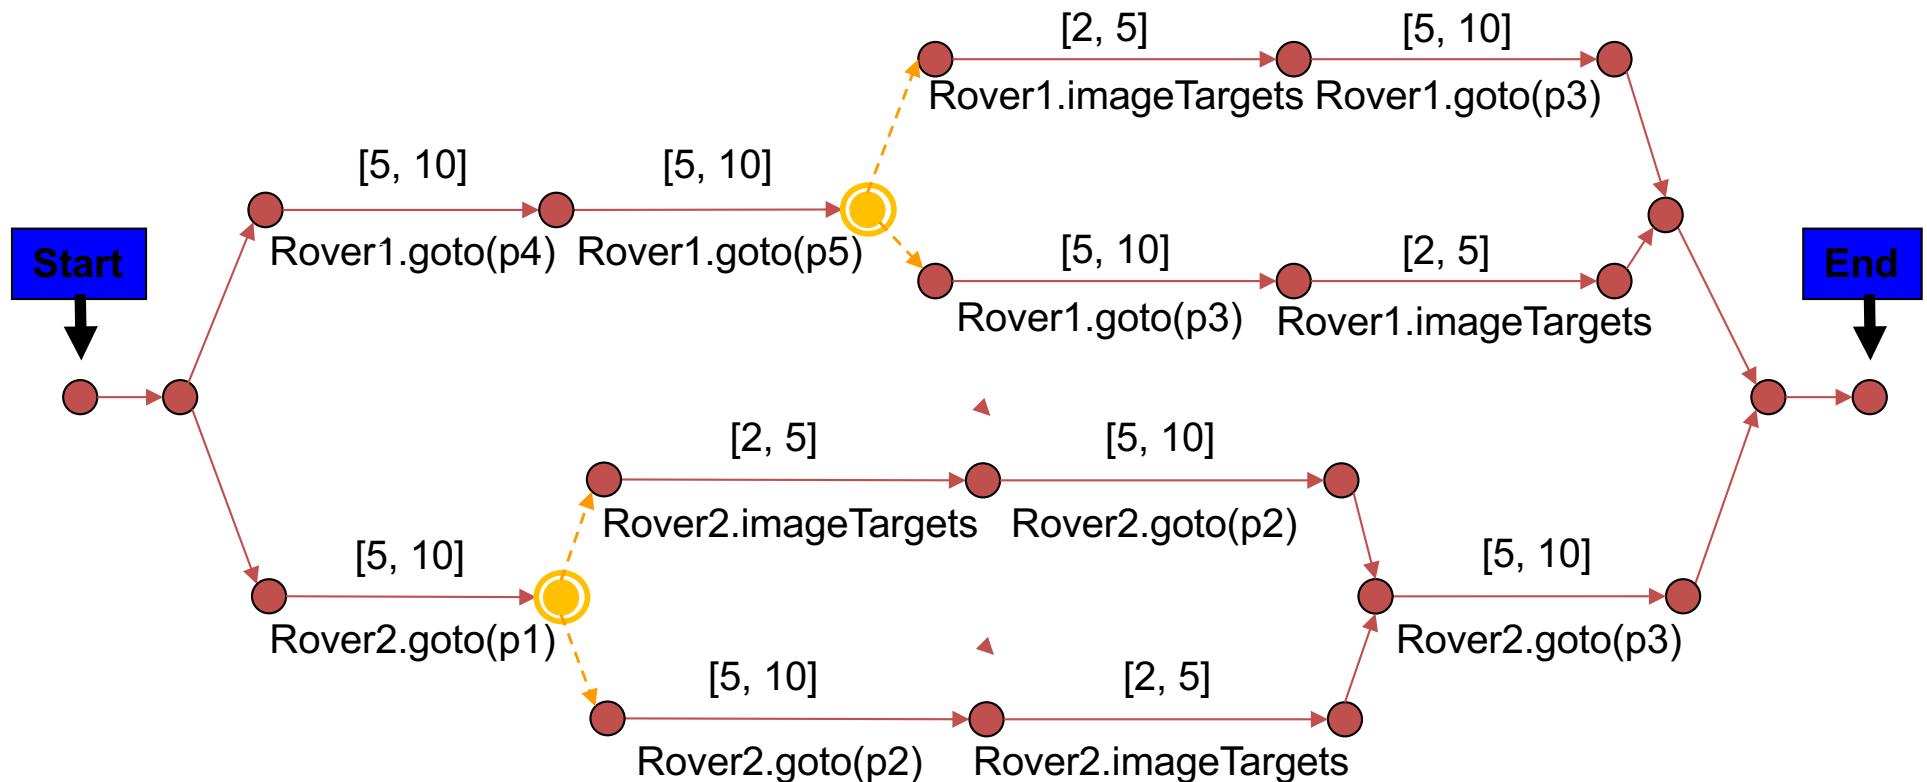
RMPL



A Different Choice



A Plan with Choice



Problem make highest utility choices that are schedulable

Assume for this plan, that edges without explicit temporal constraints are [0,0].

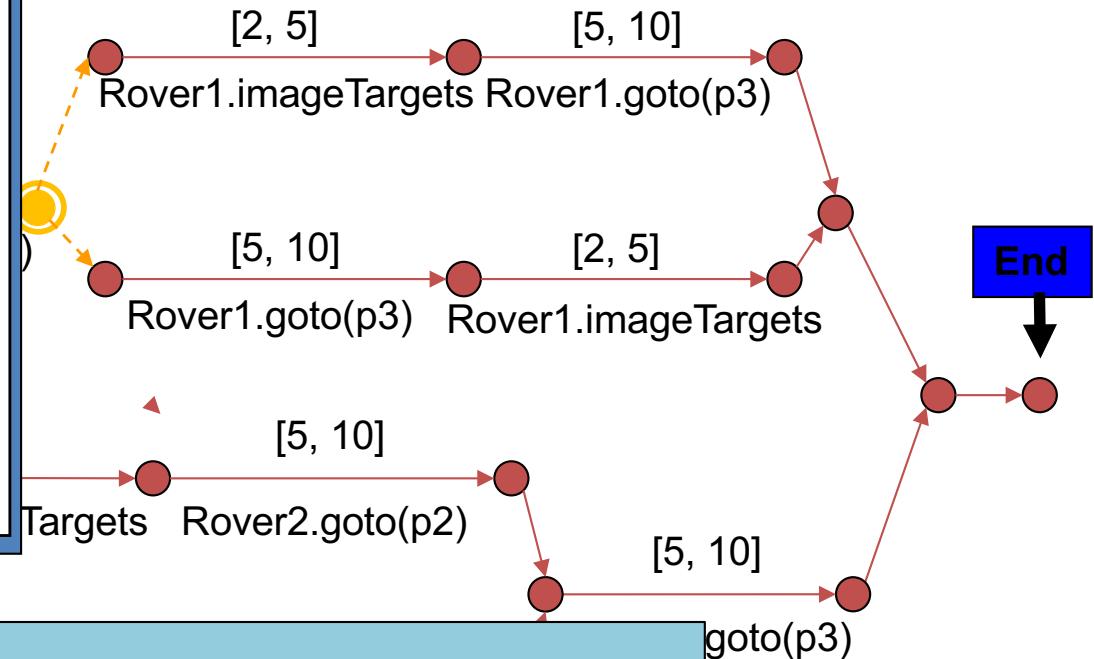
Programs with (Nondeterministic) Choice

RMPL

```

imageScienceTargets(Rover1, Rover2)
Parallel{
  Sequence{
    [5,10] Rover1.goto(p4);
    [5,10] Rover1.goto(p5);
    Choice{
      [2,5] Rover1.imageTargets1(); [5,10] Rover1.goto(p3);
      [5,10] Rover1.goto(p3); [2,5] Rover1.imageTargets1();
    }
  }
  Sequence{
    [5,10] Rover2.goto(p1);
    Choice{
      [5,10]Rover2.imageTargets2(); [2,5] Rover2.goto(p2);
      [2,5] Rover2.goto(p2); [5,10]Rover2.imageTargets2();
    }
    [5,10] Rover2.goto(p3);
  }
}

```



Program A ::=

```

prim_action(args)  |
[lb, ub] A  |
Sequence {A1; A2; ... } |
Parallel {A1; A2; ... }  |
Choose { [with reward R1] A1; [with reward] R2 A2; . . . }

```

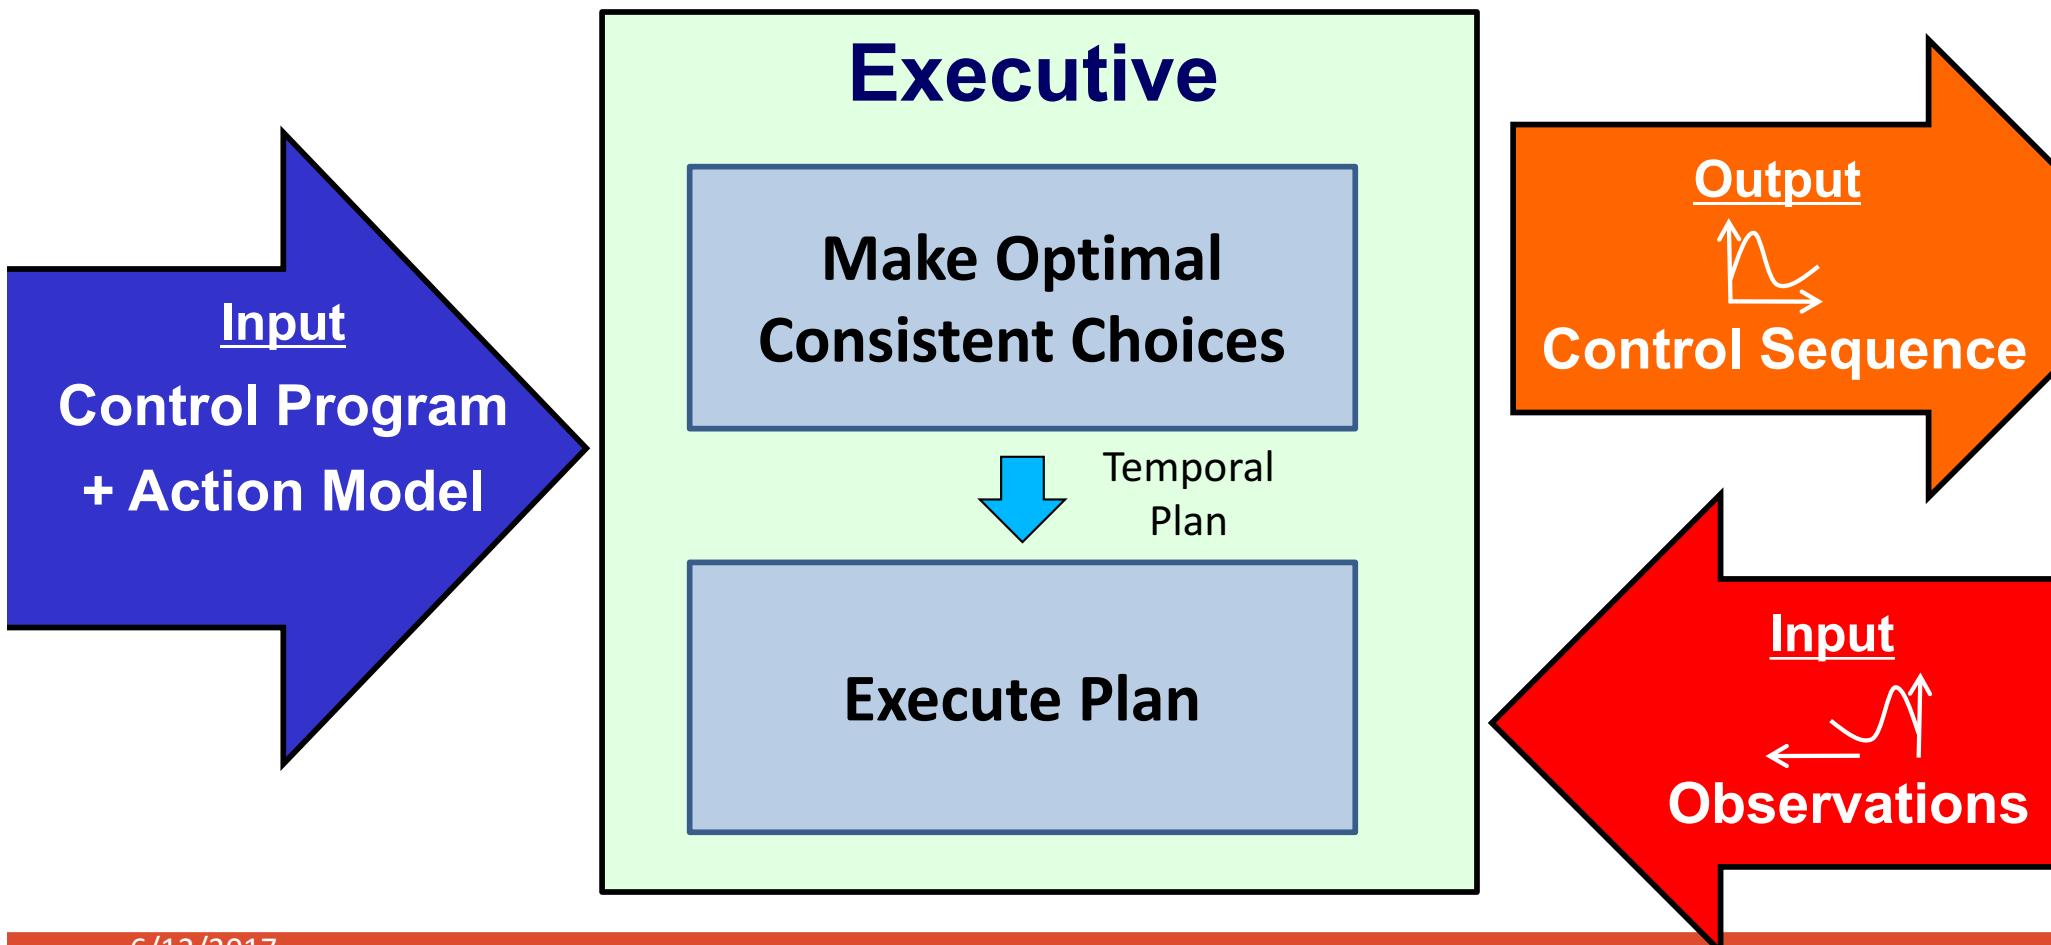
Non-deterministic programs
allow us to represent...



- **When:**
Alternative orderings of actions
- **With What:**
Alternative resource assignments
- **Who:**
Alternative task assignments
- **How:**
Alternative methods for completing a task

Executing Programs with Choice

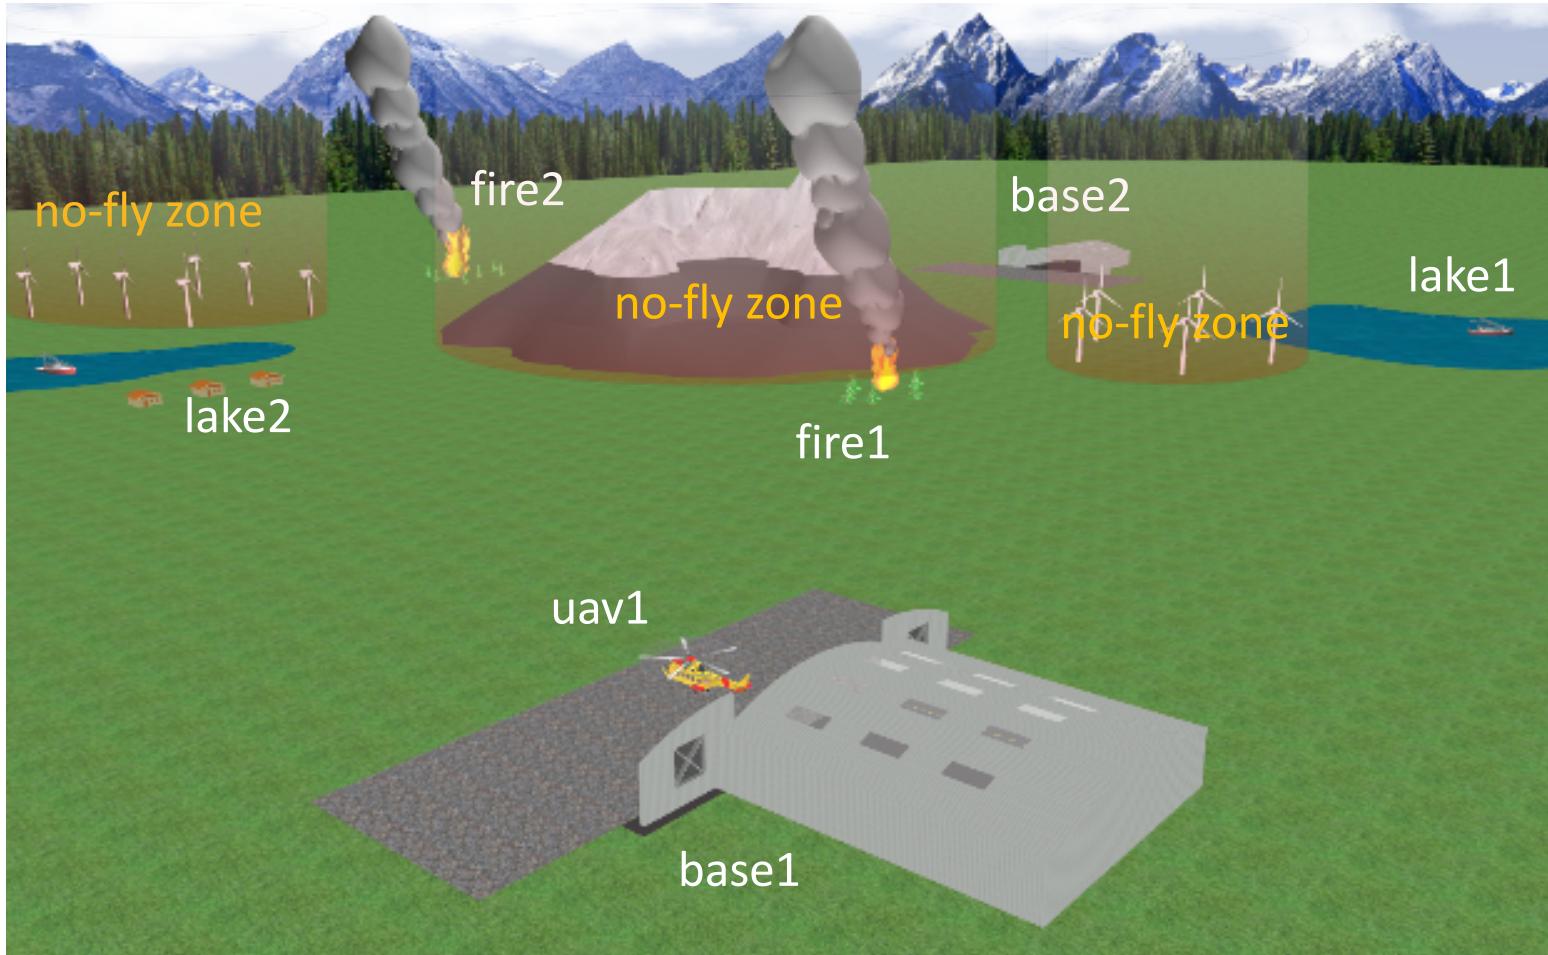
- **Plan**: Make optimal, consistent program choices.
- **Execute** : Dispatch and monitor resulting plan.



Outline

- Summer School in a Nutshell
- Architectures for Autonomy
- **Programming Cognitive Robots**
 - Programs that make choices
 - Programs on state
 - Programs that monitor

Programs on State - Firefighting Scenario



Firefighting in RMPL: Traditional Imperative program



```
class Main{
    UAV uav1;
    Lake lake1;
    Lake lake2;
    Fire fire1;
    Fire fire2;
    ...

method run() {
    sequence{
        uav1.takeoff();
        uav1.fly_base1_to_lake1();
        uav1.load_water(lake1);
        uav1.fly_lake1_to_fire1();
        uav1.drop_water_high_altitude(fire1);
        uav1.fly_fire1_to_lake1();
        uav1.load_water(lake1);
        uav1.fly_lake1_to_fire1();
        uav1.drop_water_low_altitude(fire1);
        uav1.fly_fire1_to_lake2();
        uav1.load_water(lake2);
        uav1.fly_lake2_to_fire2();
        uav1.drop_water_high_altitude(fire2);
        uav1.fly_fire2_to_lake2();
        uav1.load_water(lake2);
        uav1.fly_lake2_to_fire2();
        uav1.drop_water_low_altitude(fire2);
        uav1.fly_fire2_to_base1();
        uav1.land();
    }
}
```

Firefighting in RMPL: A Program on State

```
class Main{
    UAV uav1;
    Lake lake1;
    Lake lake2;
    Fire fire1;
    Fire fire2;
    ...
}

method run() {
    sequence{
        (fire1 == out);
        (fire2 == out);
        (uav1.flying == no &&
         uav1.location == base_1_location);
    }
}
```

Firefighting in RMPL: Setup & Initial Conditions

```
class Main{
    UAV uav1;
    Lake lake1;
    Lake lake2;
    Fire fire1;
    Fire fire2;
    ...

method run() {
    sequence{
        (fire1 == out);
        (fire2 == out);
        (uav1.flying == no &&
         uav1.location == base_1_location);
    }
}
```

```
Main () {
    uav1 = new UAV();
    uav1.location= base_1_location;
    uav1.flying = no;
    uav1.loaded = no;

    lake1 = new Lake();
    lake1.location = lake_1_location;

    lake2 = new Lake();
    lake2.location = lake_2_location;

    fire1 = new Fire();
    fire1.location = fire_1_location;
    fire1 = high;

    fire2 = new Fire();
    fire2.location = fire_2_location;
    fire2 = high;
}
```

Firefighting in RMPL: Model of Actions

```

class UAV {
    Roadmap location;
    Boolean flying;
    Boolean loaded;

    primitive method takeoff()
        flying == no => flying == yes;

    primitive method land()
        flying == yes => flying == no;

    primitive method load_water(Lake lakespot)
        ((flying == yes) && (loaded == no) && (lakespot.location == location)) => loaded == yes;

    primitive method drop_water_high_altitude(Fire firespot)
        ((flying == yes) && (loaded == yes) && (firespot.location == location) && (firespot == high))
        => ((loaded == no) && (firespot == medium));

    primitive method drop_water_low_altitude(Fire firespot)
        ((flying == yes) && (loaded == yes) && (firespot.location == location) && (firespot == medium))
        => ((loaded == no) && (firespot == out));

    #MOTION_PRIMITIVES(location, fly, flying==yes)
}

```

```

class Lake {
    Roadmap location;
}

class Fire{
    initial value high;
    value medium;
    value out;
    Roadmap location;
}

```

Firefighting in RMPL: Model of Actions

```

class UAV {
    Roadmap location;
    Boolean flying;
    Boolean loaded;

    primitive method takeoff()
        flying == no => flying == yes;

    primitive method land()
        flying == yes => flying == no;

    primitive method load_water(Lake lakespot)
        ((flying == yes) && (loaded == no) && (lakespot.location == location)) => loaded == yes;

    primitive method drop_water_high_altitude(Fire firespot)
        ((flying == yes) && (loaded == yes) && (firespot.location == location) && (firespot == high))
        => ((loaded == no) && (firespot == medium));

    primitive method drop_water_low_altitude(Fire firespot)
        ((flying == yes) && (loaded == yes) && (firespot.location == location) && (firespot == medium))
        => ((loaded == no) && (firespot == out));

    #MOTION_PRIMITIVES(location, fly, flying==yes)
}

```

```

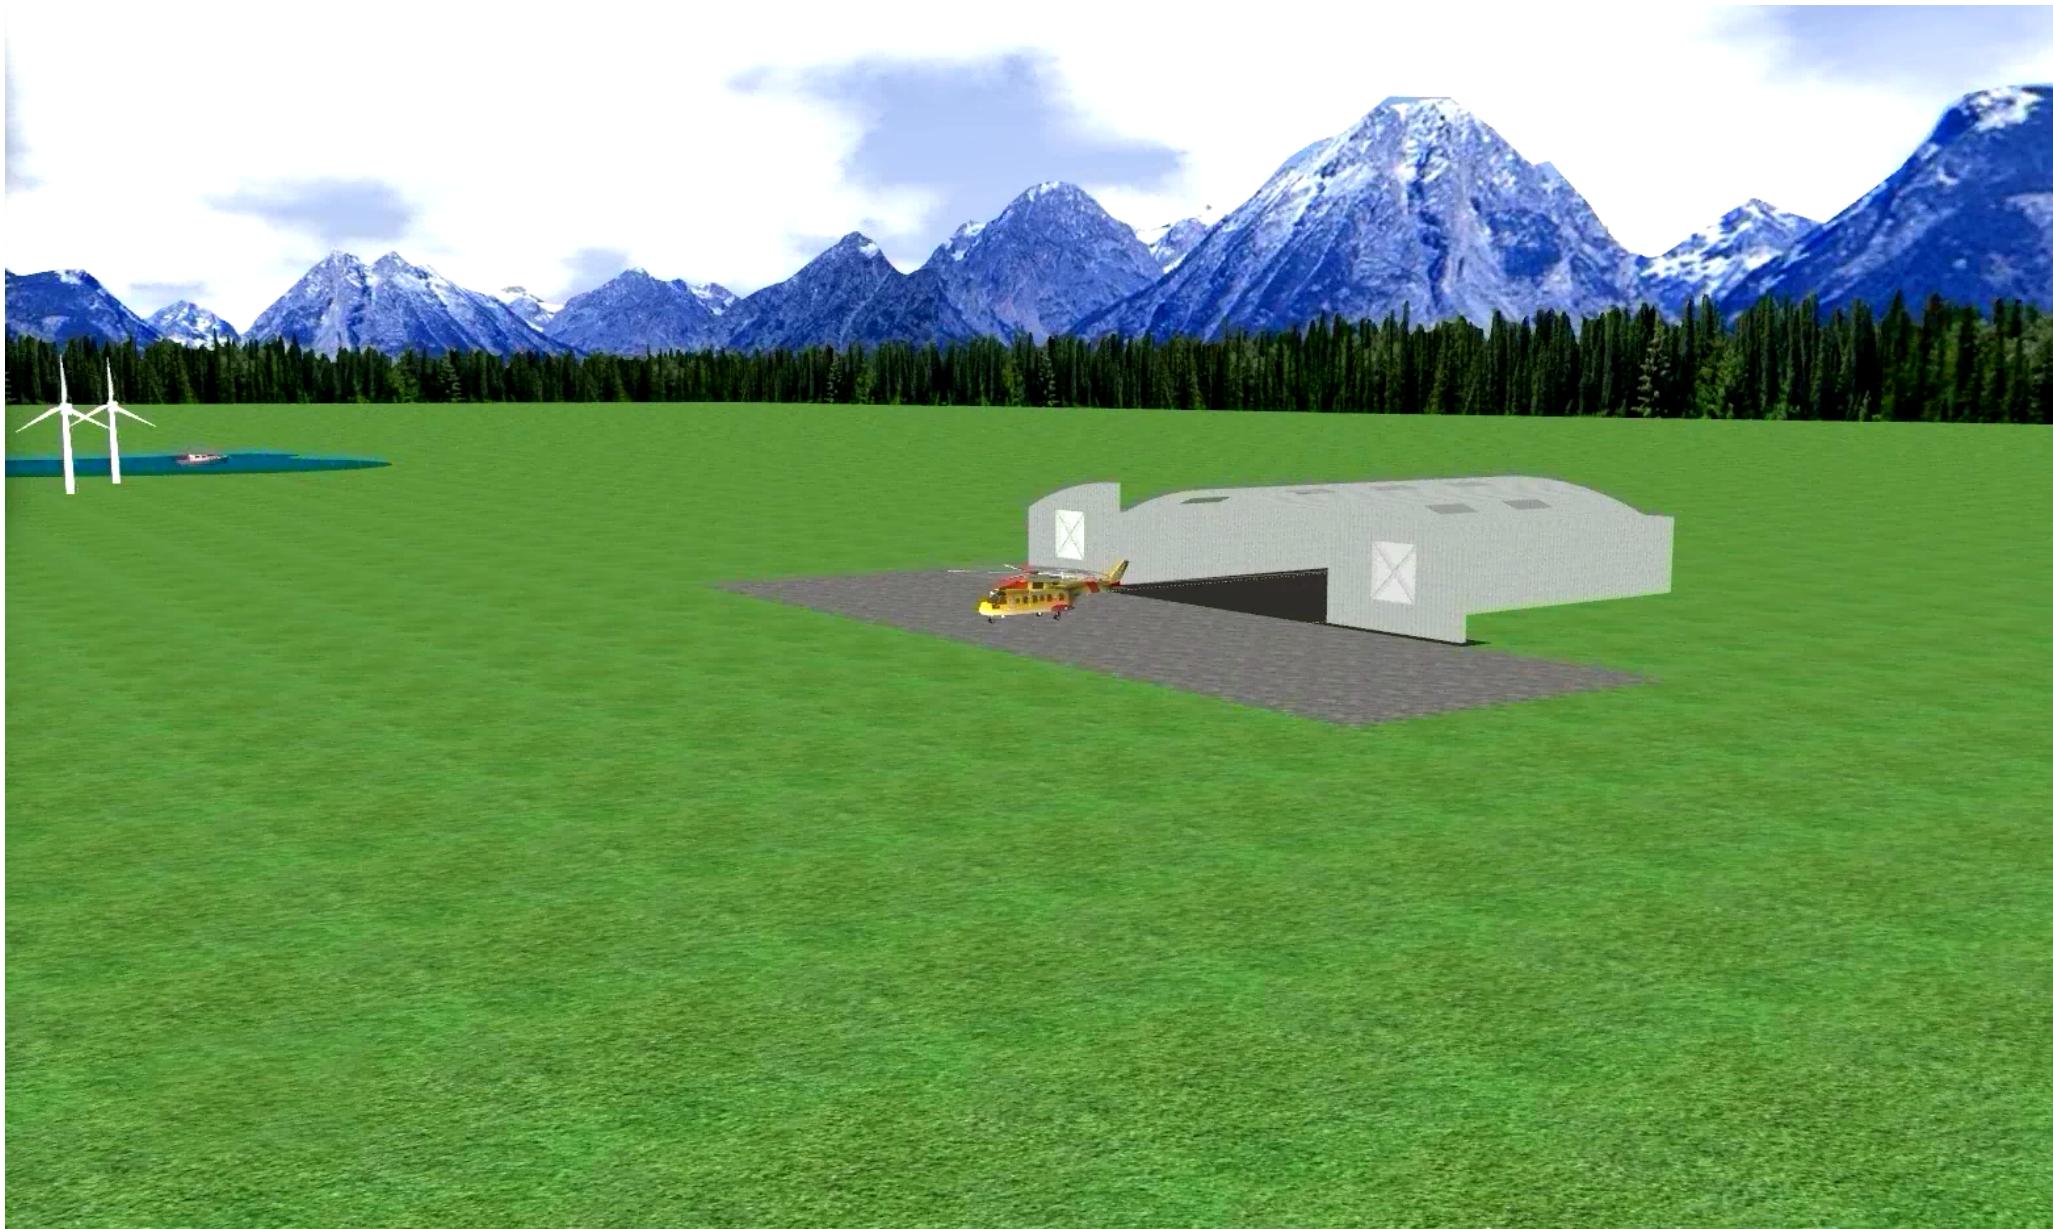
class Lake {
    Roadmap location;
}

class Fire{
    initial value high;
    value medium;
    value out;
    Roadmap location;
}

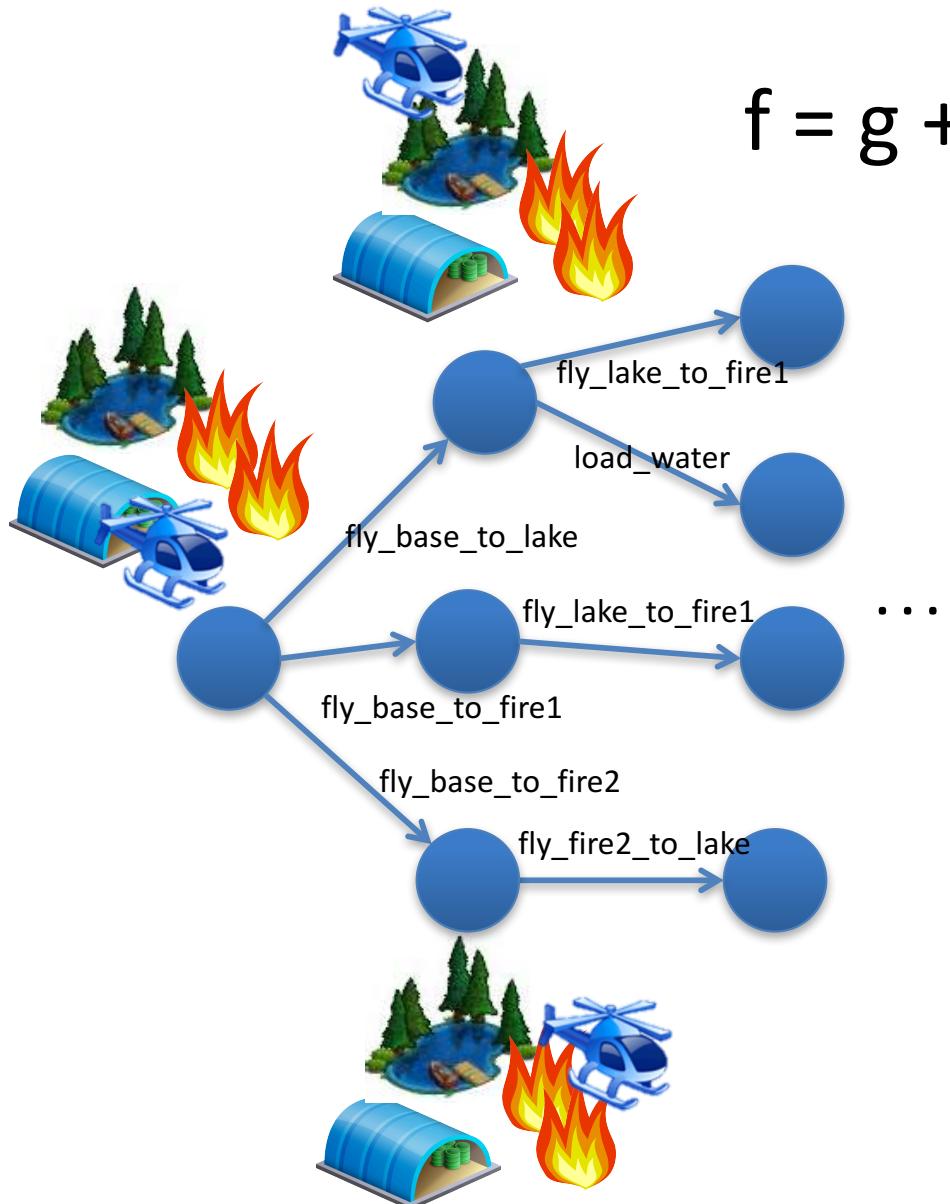
```

Simulation Testing: Firefighting Scenario

ANSRS



Decision-making algorithms: Activity Planning (W), Path Planning (T)



Solve Relaxed
Planning
Problem



Action Models are Instances of PDDL (Plan Domain Description Language)

(:durative-action pick-up-block

:parameters (?r - robot ?t ?b - block)

:duration (and (>= ?duration 10) (<= ?duration 30))

:condition (and (at start (clear-above ?t))
(at start (empty-gripper ?r))
(over all (can-reach ?r ?t))
(at start (on ?t ?b))))

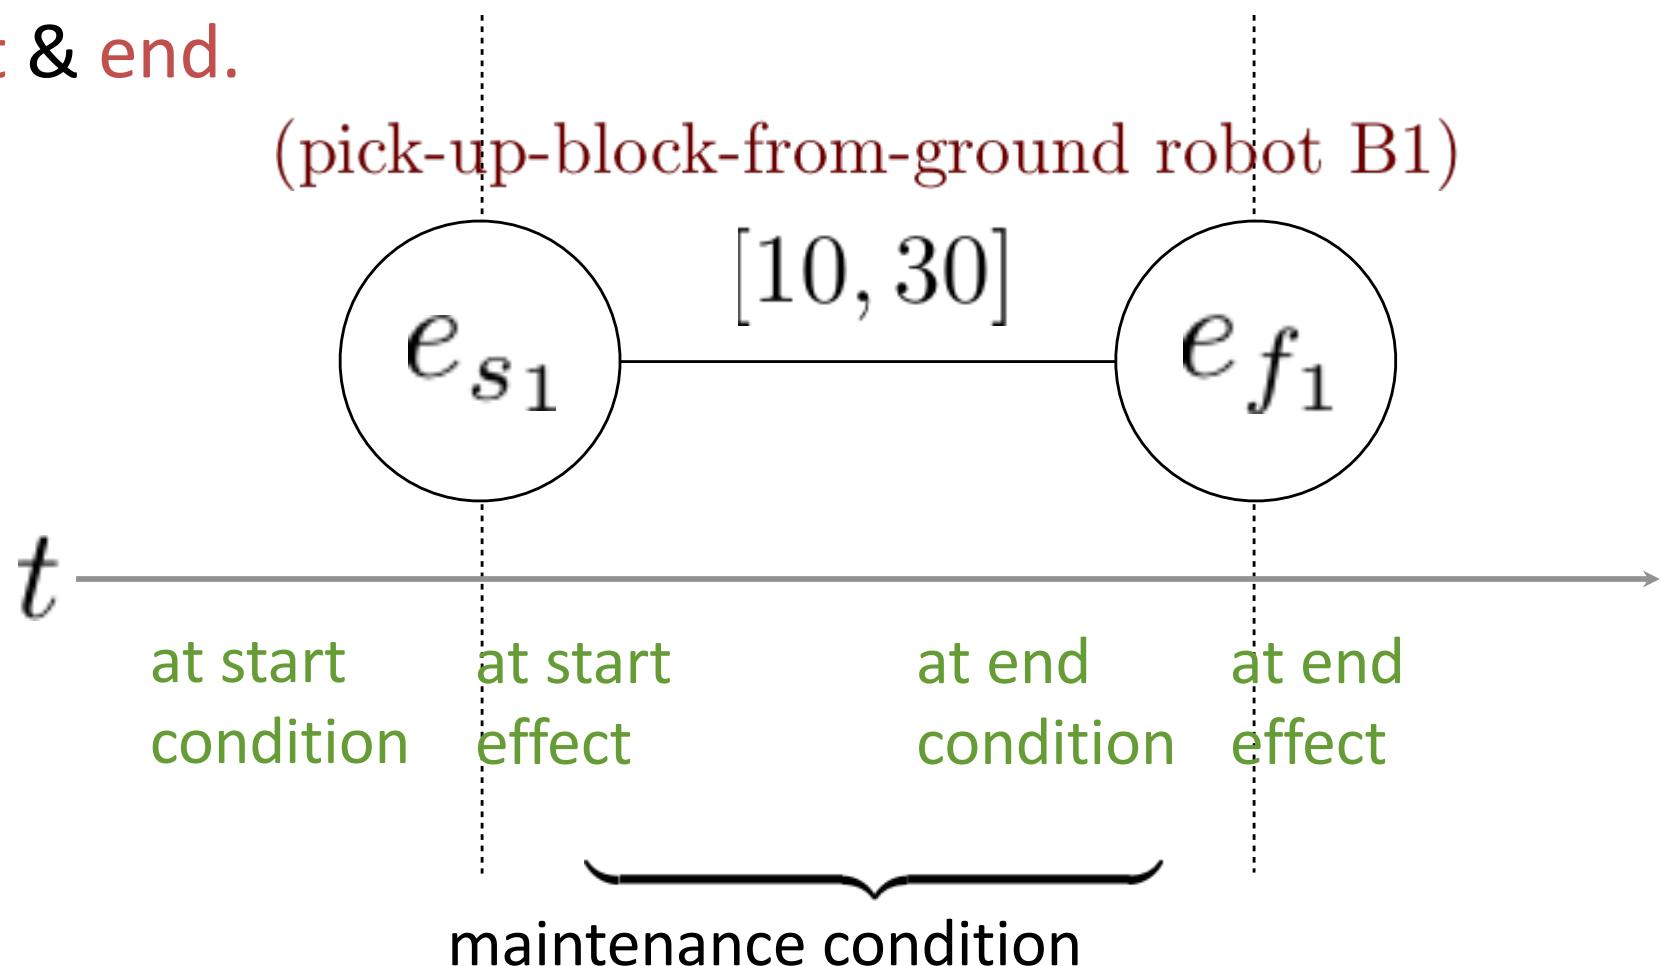
:effect (and (at end (not (empty-gripper ?r)))
(at end (not (on ?t ?b))))
(at end (holding ?r ?t))
(at end (clear-above ?b))
(at end (not (clear-above ?t))))

)

Primitive Action Model

Temporal actions modeled

- with **conditions & effects**,
- at **start & end**.



Outline

- Summer School in a Nutshell
- Architectures for Autonomy
- **Programming Cognitive Robots**
 - Programs that make choices
 - Programs on state
 - **Programs that monitor**

In an Uncertain World Task Executives need to be Suspicious



“something unexpected happens”
= not modeled in plan!

How should we react?

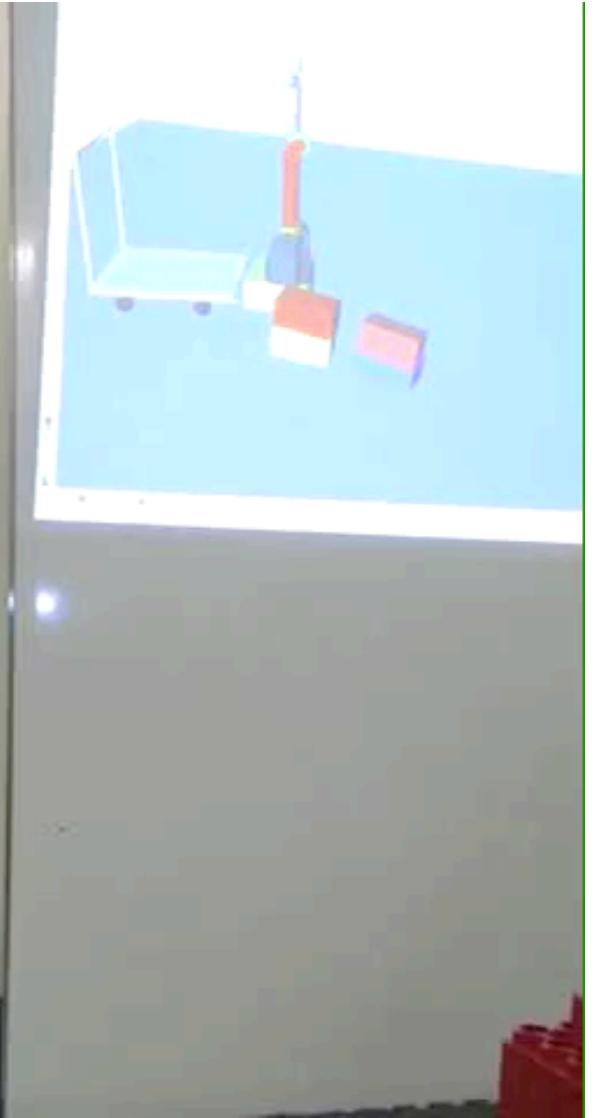


Languages that support Self-Monitoring

```
(Define-rap (attach-at-site ?thesite)
  (succeed (docked ?thesite))
  (method
    (context (ferrous-hull ?thesite))
    (task-net
      (sequence
        (t1 (approach-site ?thesite))
        (t2 (magnetically-attach ?thesite)
            (wait-for (docked ?thesite))))))
  (method
    (context (not (ferrous-hull ?thesite)))
    (task-net
      (sequence
        (t1 (approach-site ?thesite))
        (t2 (grip-attach ?thesite)
            (wait-for (docked ?thesite)))))))
```

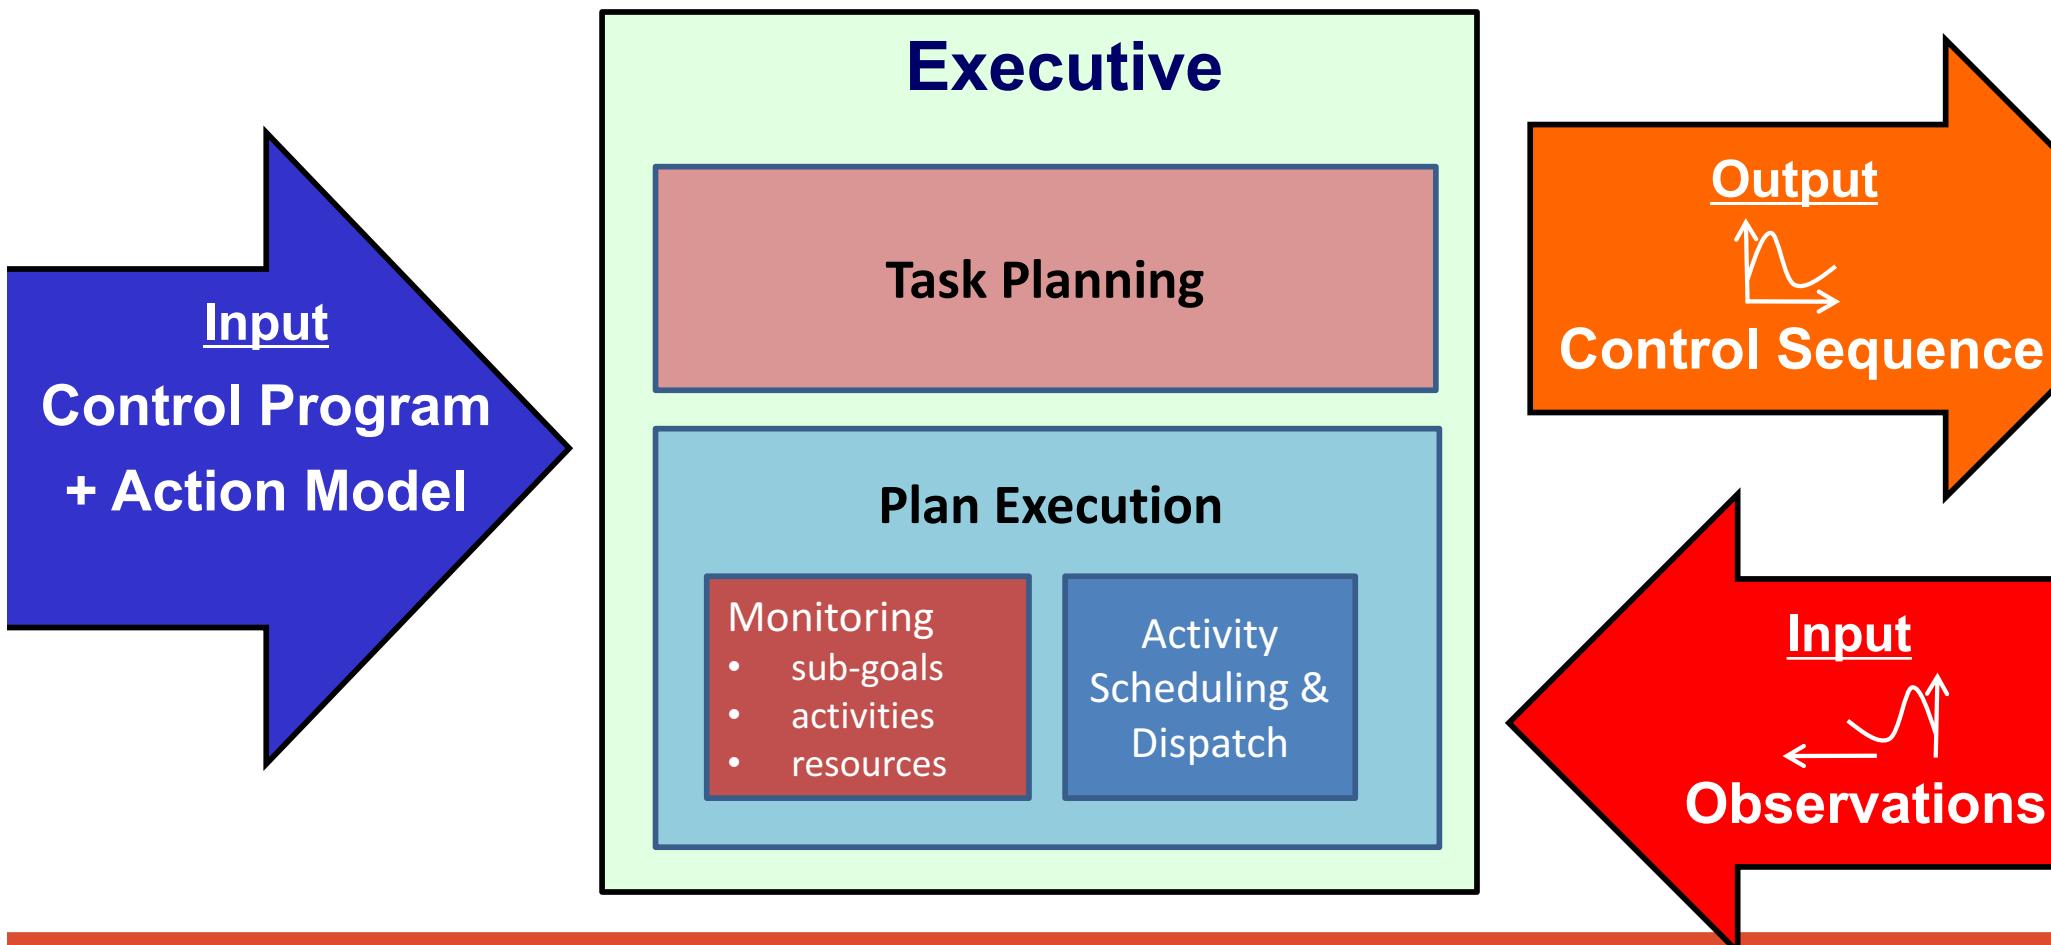
- RAPS (Firby)
- ESL (Gat)
- RMPL (Williams)
- ...

Self-Monitoring Systems

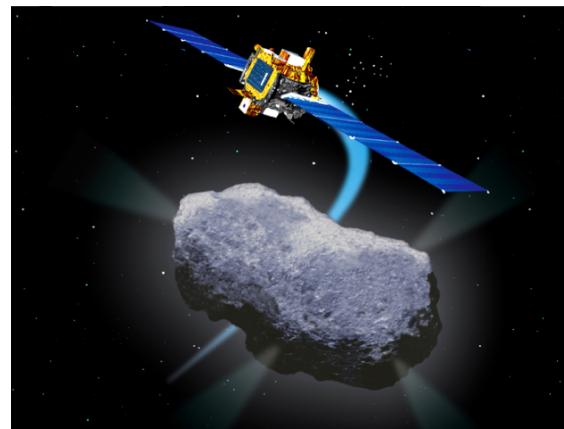


A Suspicious Task Executive Closes the Loop on Goals

- **Monitors**: detects as soon as sub-goals go wrong.
- **Adapts**: fixes the problem.



Summer School is about Creating Goal-directed Systems



1. Decision-making components for generating, refining, executing and monitoring plans on line ([15 Tutorials](#)).
2. Architectures and principles for creating goal-directed systems from decision-making components
(3T, Remote Agent, Claraty, MDS, Idea, [Enterprise/ROS](#), ROSplan).
3. Dynamic languages, elevated to the goal-level through partial specification and operation on state
(PRS, RAPS, TCC, Golog, [RMPL](#)).

QUESTIONS?