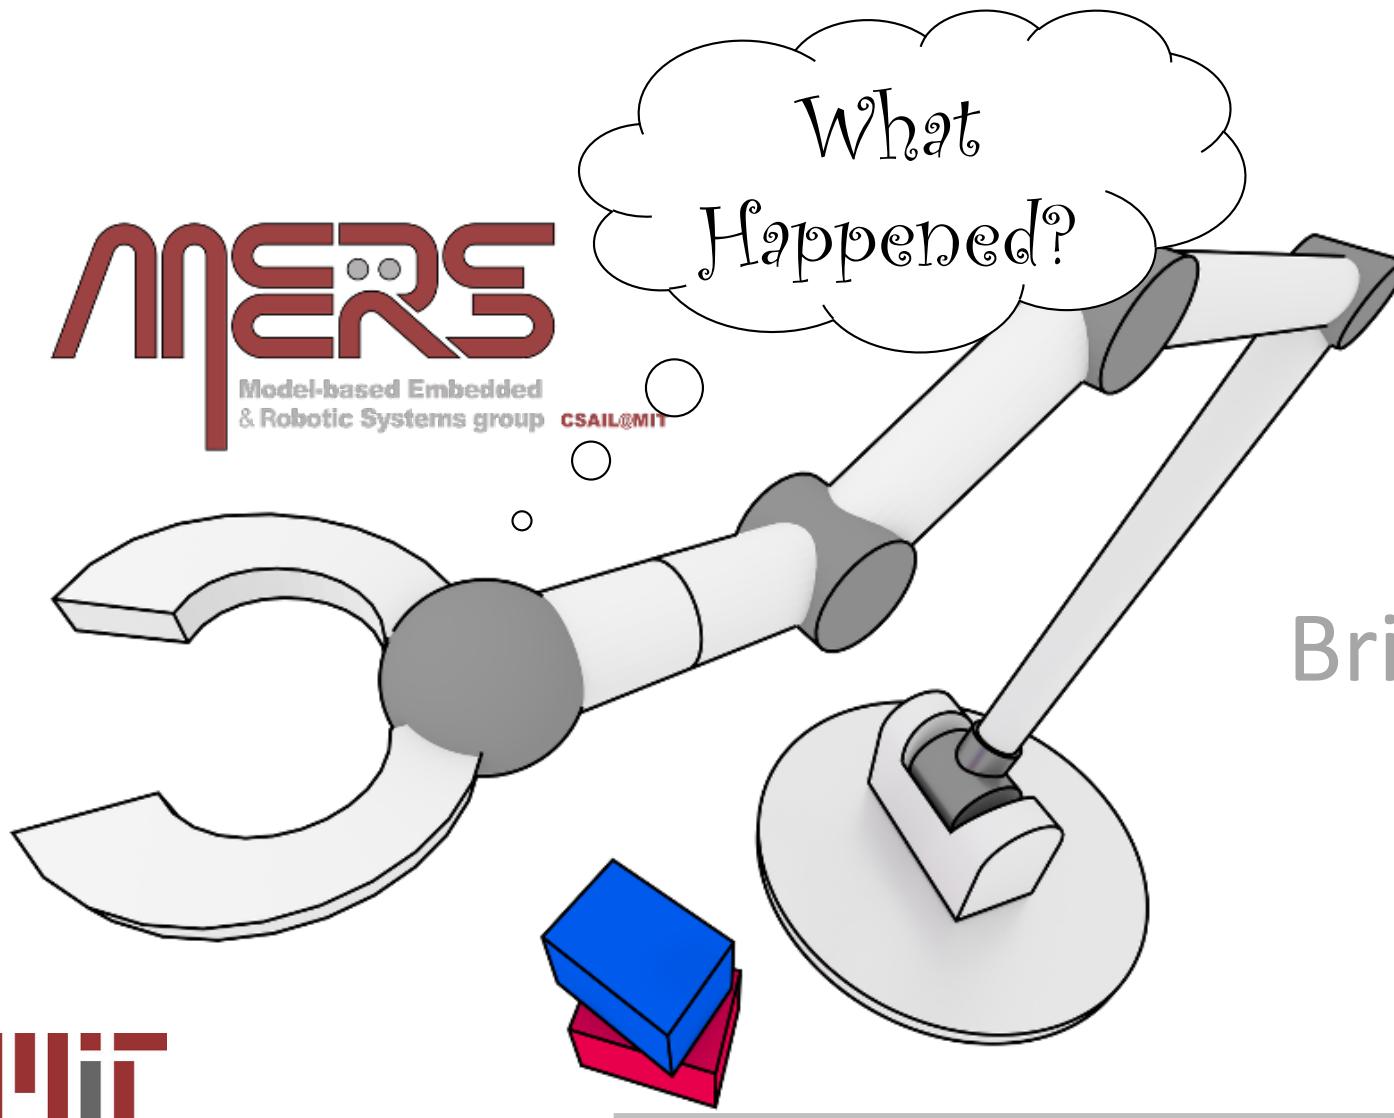


Self-Monitoring, Self-Diagnosing Systems



Brian Williams
June 12th, 2017

Slide Contributions:
Steve Levine
David Wang
Brian Williams

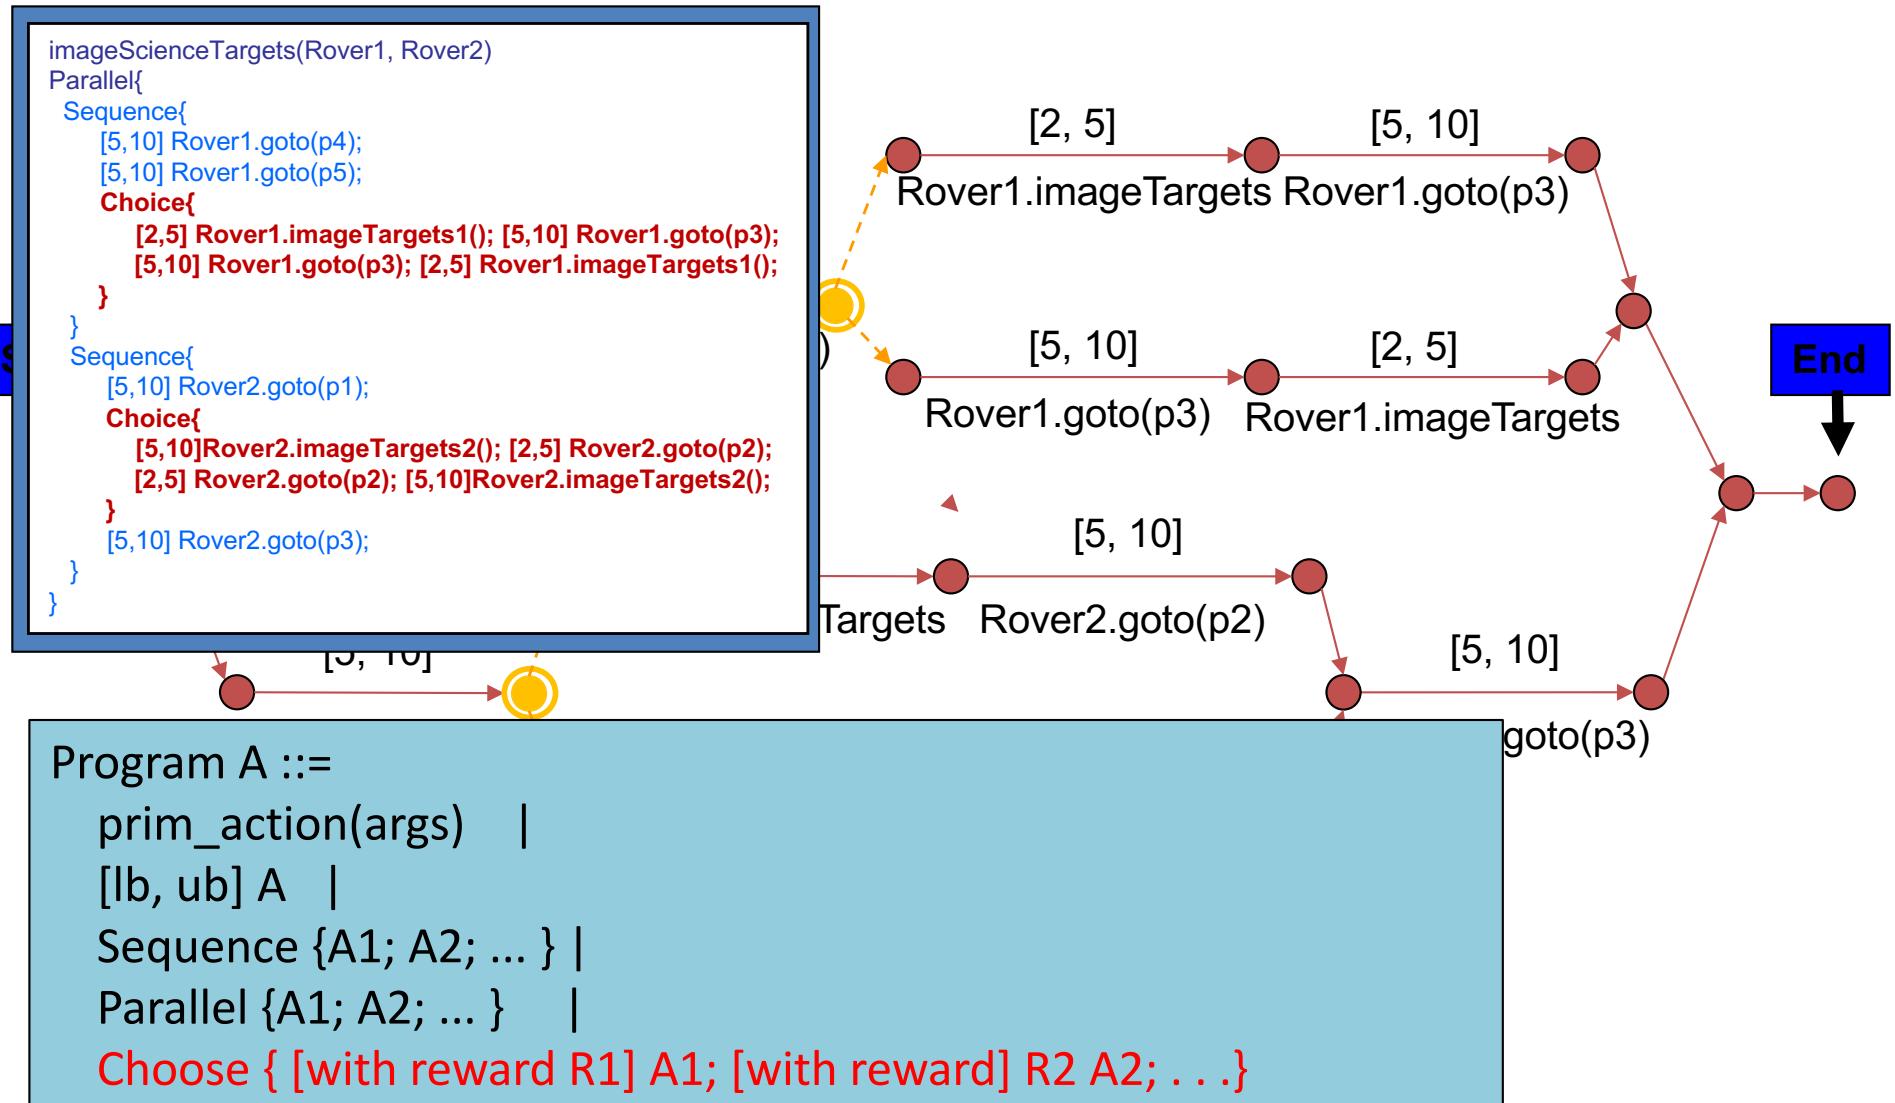
Some Related Papers

- B. C. Williams, and R. Ragno, "Conflict-directed A* and its Role in Model-based Embedded Systems," Special Issue on Theory and Applications of Satisfiability Testing, Journal of Discrete Applied Math, January 2003.
- B. C. Williams and P. P. Nayak, "A Model-based Approach to Reactive Self-Configuring Systems," AAAI, 1996.
- O. Martin, B. Williams and M. Ingham, "Diagnosis as Approximate Belief State Enumeration for Probabilistic Concurrent Constraint Automata," AAAI, 2005.
- M. Hofbaur and B. Williams, "Hybrid Estimation of Complex Systems," IEEE Trans SMC, v. 34, Oct., 2004.
- L. Blackmore, S. Funiak and B. Williams, "A Combined Stochastic and Greedy Hybrid Estimation Capability for Concurrent Hybrid Models with Autonomous Mode Transitions," J. Robotics and Autonomous Systems, V. 56, 2007.
- P. Yu and B. Williams, "Continuously Relaxing Over-constrained Conditional Temporal Problems through Generalized Conflict Learning and Resolution," IJCAI, Beijing, China, August 2013.
- P. Yu and C. Fang and B. Williams, "Resolving Uncontrollable Conditional Temporal Problems using Continuous Relaxations," ICAPS, Portsmouth, 2014.
- Kirk: Kim, P., Williams, B., Abramson, M., "Executing Reactive, Model-based Programs through Graph-based Temporal Planning." IJCAI, 2001.
- Pike: Levine, S., Williams, B., "Concurrent Plan Recognition and Execution for Human-Robot Teams." ICAPS, 2014.
- ITC: Shu, I., Effinger, R., Williams, B., "Enabling Fast Flexible Planning through Incremental Temporal Reasoning with Conflict Extraction." AAAI, 2005.
- DBT: Effinger, R., Williams, B., "Extending Dynamic Backtracking to Solve Weighted Conditional CSPs." AAAI, 2006.

Recall: Dynamic Languages

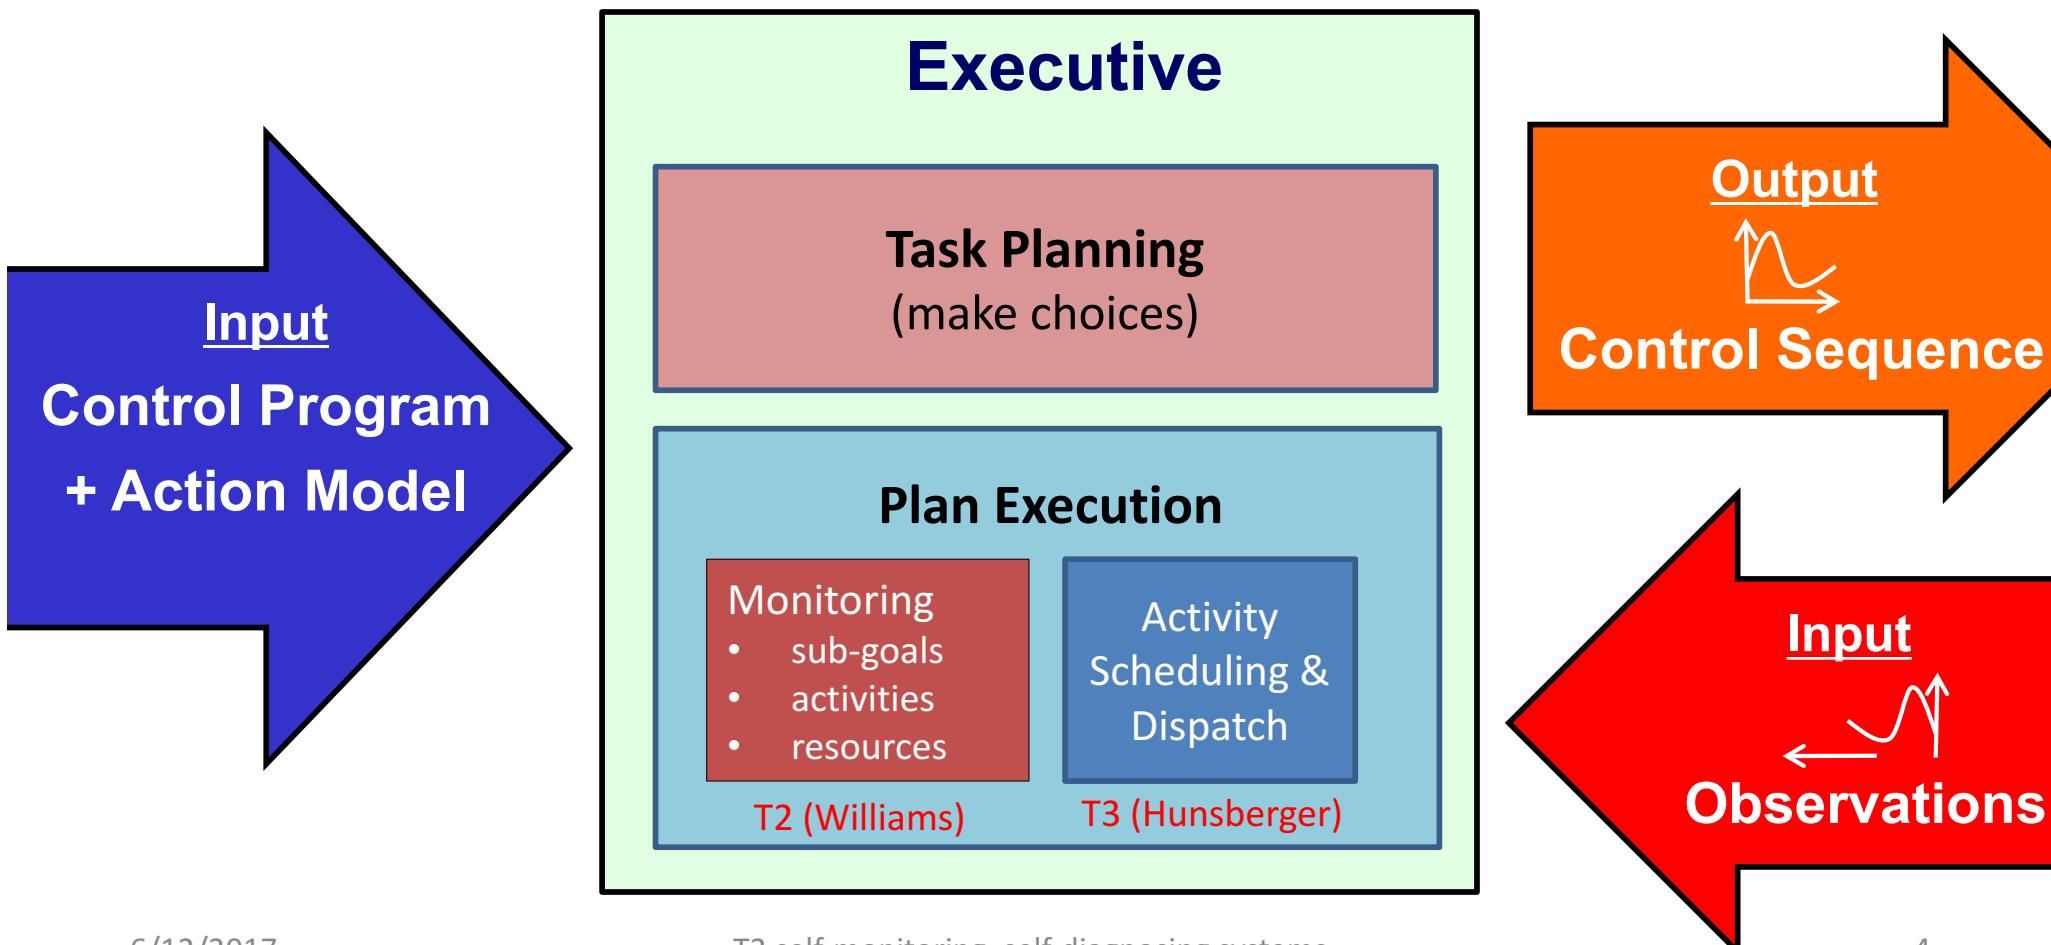
Leave Choices Open

RMPL



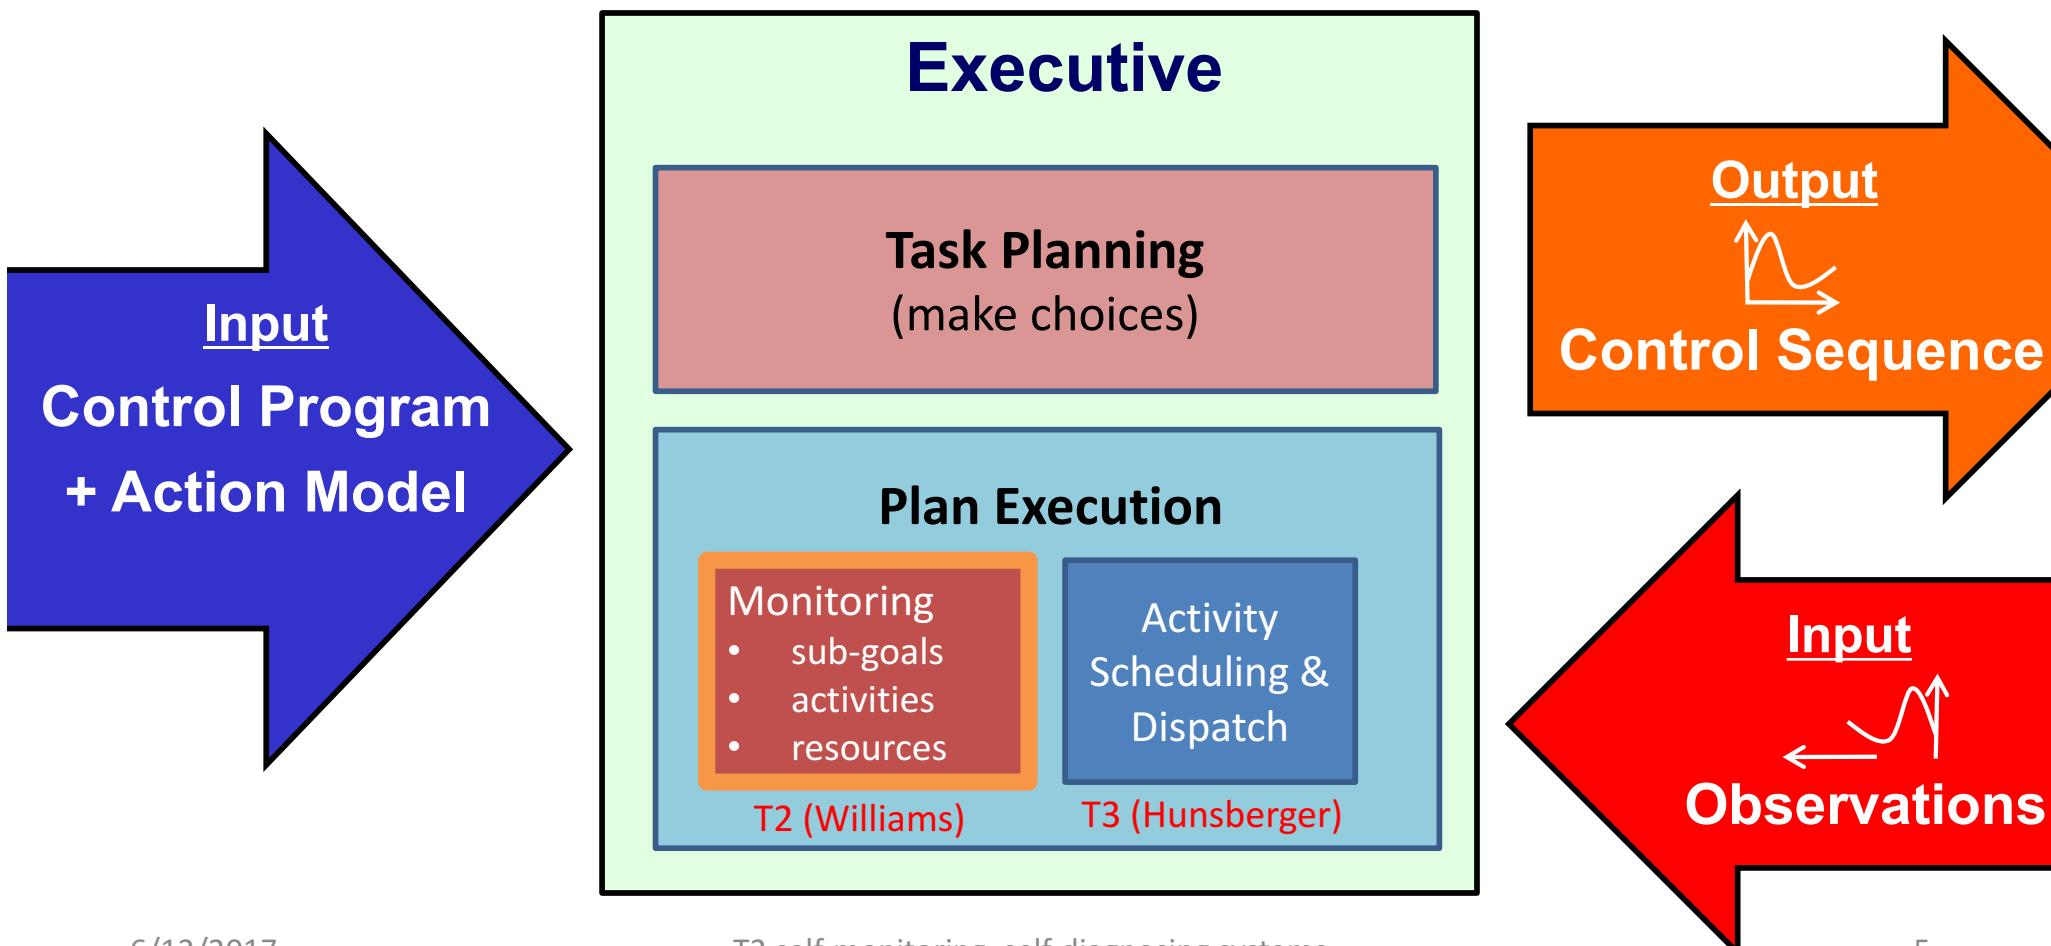
Recall: Program Executives Monitor Goals and Make Online Choices

- **Plan**: Make optimal, consistent **program choices**.
- **Execute** : Dispatch and monitor resulting **plan**.



Recall: A “Suspicious” Executive Closes the Loop on Goals

- **Monitors**: detects as soon as sub-goals fail, actions break.
- **Adapts**: fixes the problem.





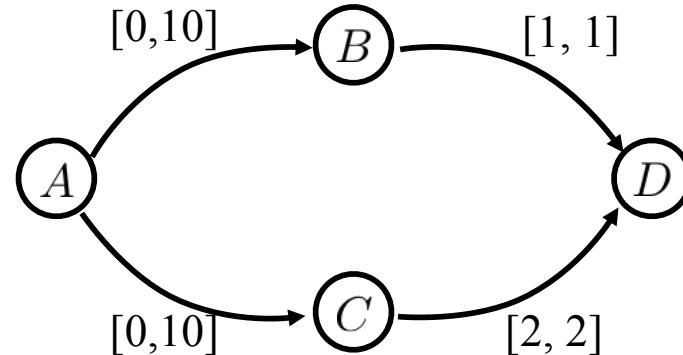
Recall: Self-Monitoring Systems



Detects failing actions and plans, due to:

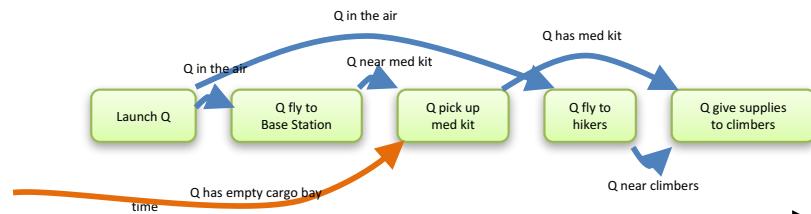
- Anomalies in relevant state (sub-goals).
- Faulty hardware.

We can create different **programming models**
by combining different methods



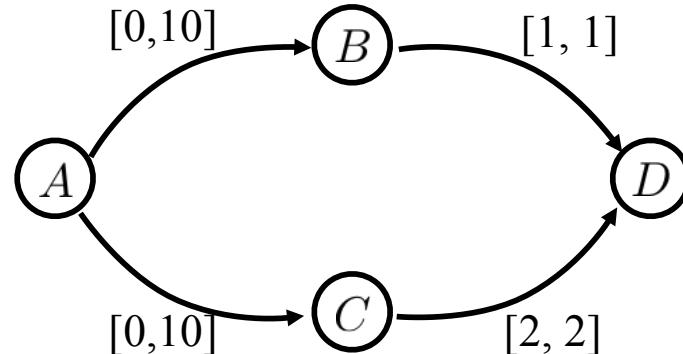
Flexible Time

+

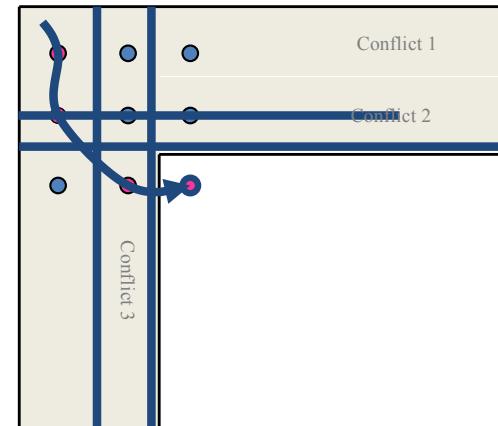


Execution monitoring

and



+



Diagnosis and Reconfiguration
Discrete Choice and Search

Outline

- Programs that Monitor State
- Programs that Self-Diagnose (opt)

Outline

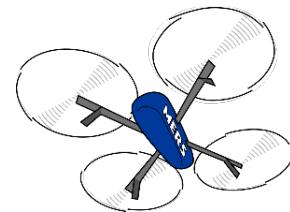
- Programs that Monitor State
 - Sub-goal monitoring
 - Model-based diagnosis and mode estimation
 - Programs that Self-Diagnose (opt)

Volcano eruption!

```
method run() {  
    sequence {  
        uav.launch();  
        uav.fly(base_station);  
        uav.pick_up(med_kit);  
        uav.fly(hikers);  
        uav.drop_off(med_kit);  
    }  
}
```



Base Station



Actions have specified
preconditions & effects



Program (Plan) is a sequence of actions in RMPL

Volcano Rescue in RMPL

```
class Main{
    QuadCopter uav;
    Location base-station;
    People climbers;
    MovableObject med-kit;
    ...

method run() {
    sequence{
        uav.launch();
        uav.fly(base_station);
        uav.pick_up(med_kit);
        uav.fly(climbers);
        uav.drop_off(med_kit) }}
```

Program A ::=
 prim_action(args) |
 [lb, ub] A |
 Sequence {A1; A2; ... } |
 Parallel {A1; A2; ... }

QuadCopter Action Model in RMPL

```
class QuadCopter {  
    Roadmap location;  
    Boolean flying;  
    Boolean loaded;
```

```
primitive method Launch()  
    flying == no => flying == yes;
```

```
primitive method land()  
    flying == yes => flying == no;
```

```
primitive method pickup(MoveableObject object)  
    ((flying == yes) && (loaded == no) && (object.location == location))  
    => loaded == yes && (object.on == self);
```

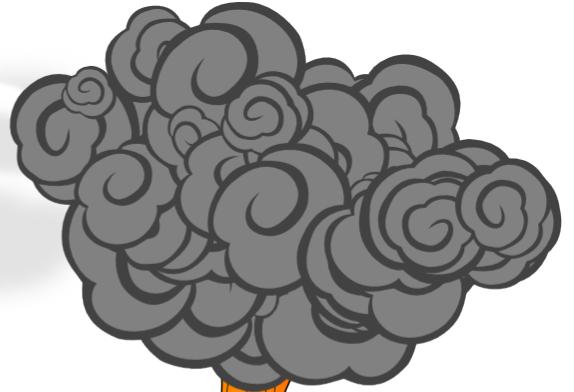
```
primitive method drop_off(MoveableObject object)  
    ((flying == yes) && (loaded == yes) && object.on == self)  
    => ((loaded == no) && (object.location == location) && object.on == none);
```

```
#MOTION_PRIMITIVES(location, fly, flying==yes)  
}
```



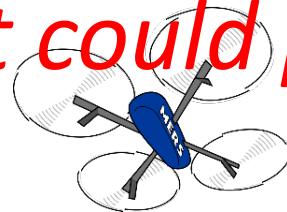
(PDDL over a rich action model)

Volcano



Base Station

Launch Q



Q fly to
Base Station

Q pick up
med kit

Q fly to
hikers

Q give med kit
to climbers

What could possibly go wrong??



Can set expectations from action preconditions & effects

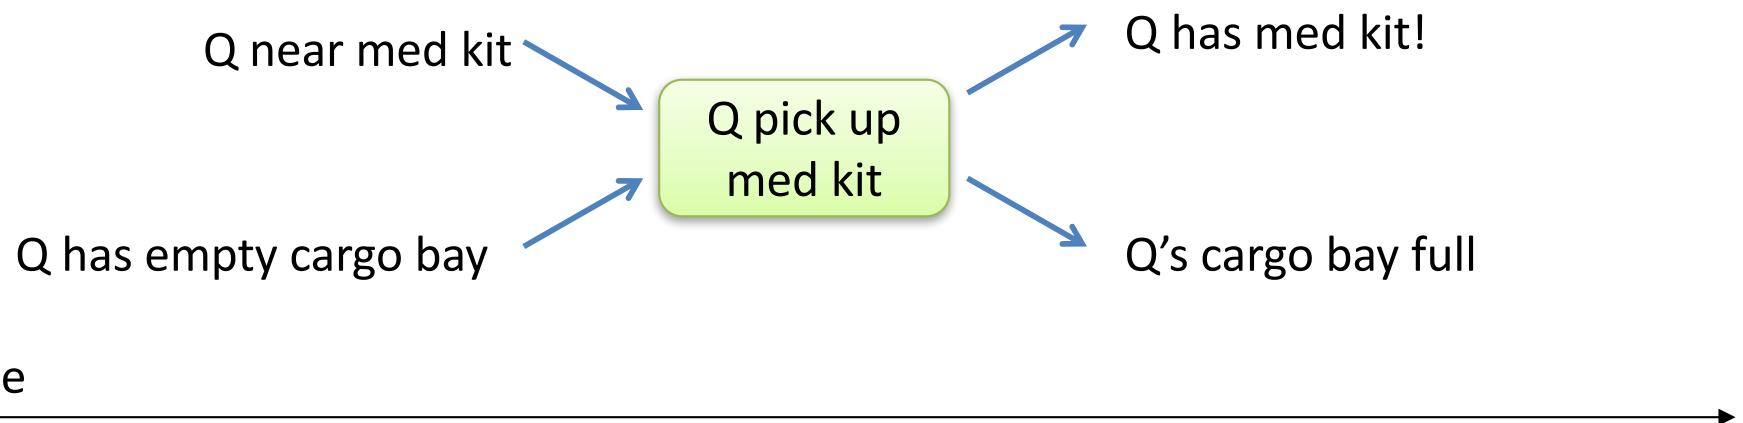


Can set expectations from action preconditions & effects

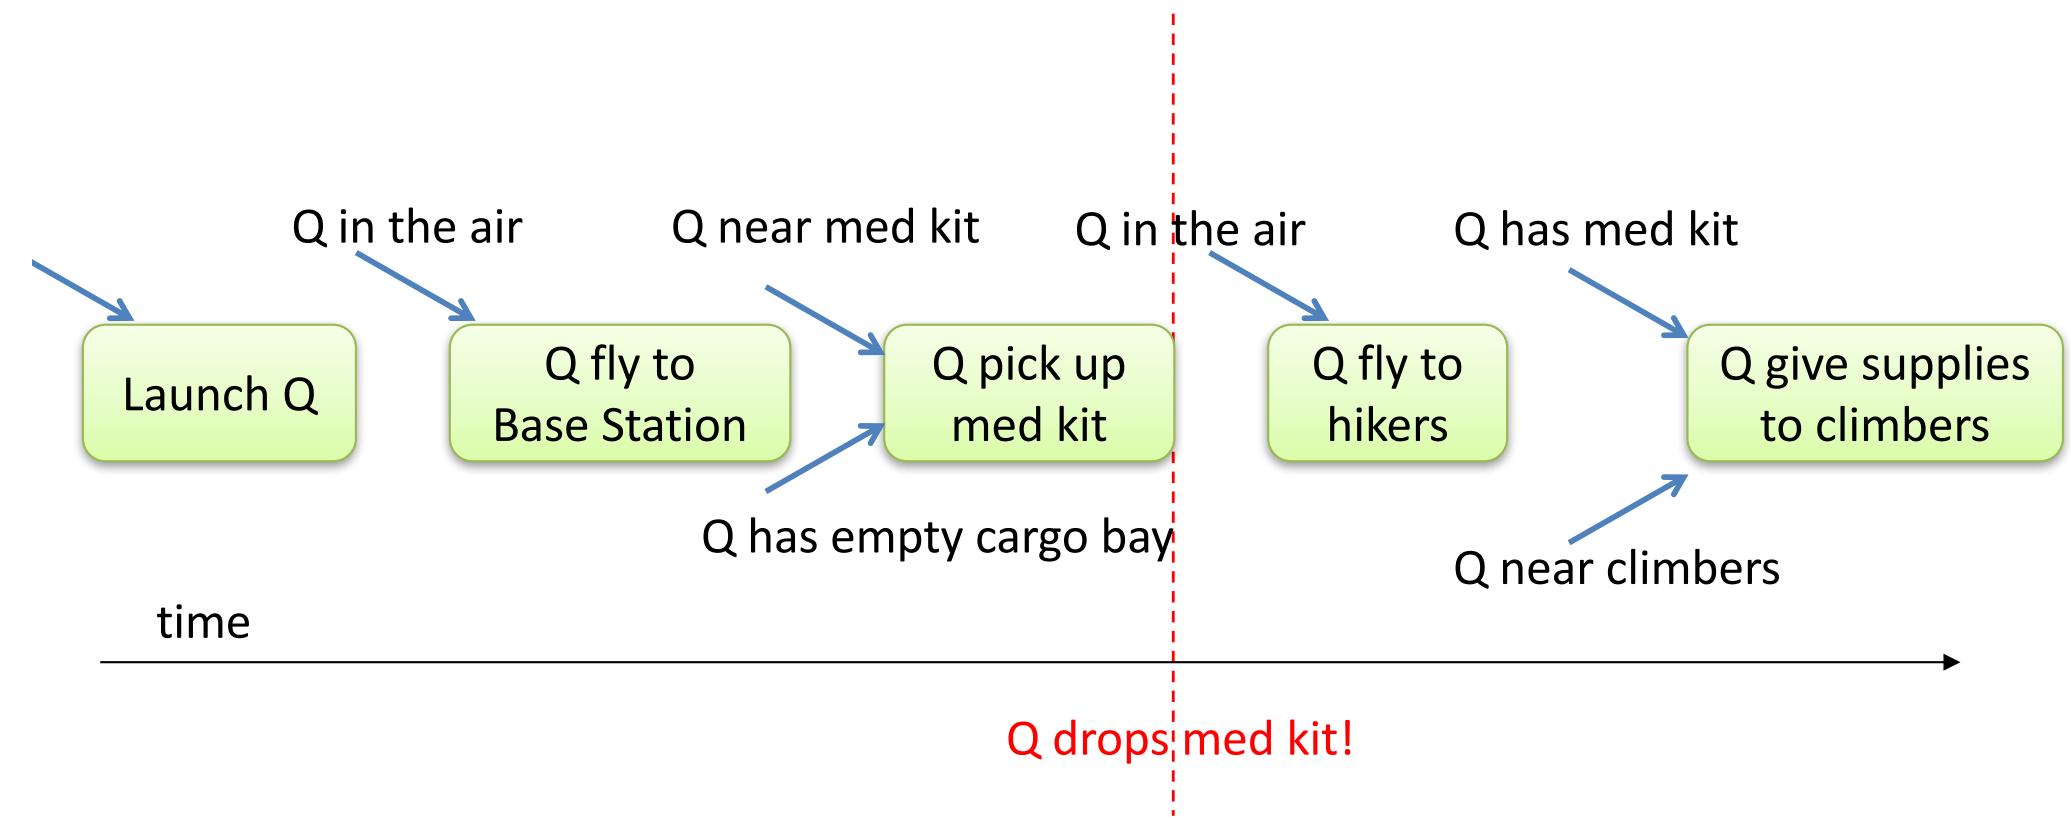
Q pick up
med kit

time

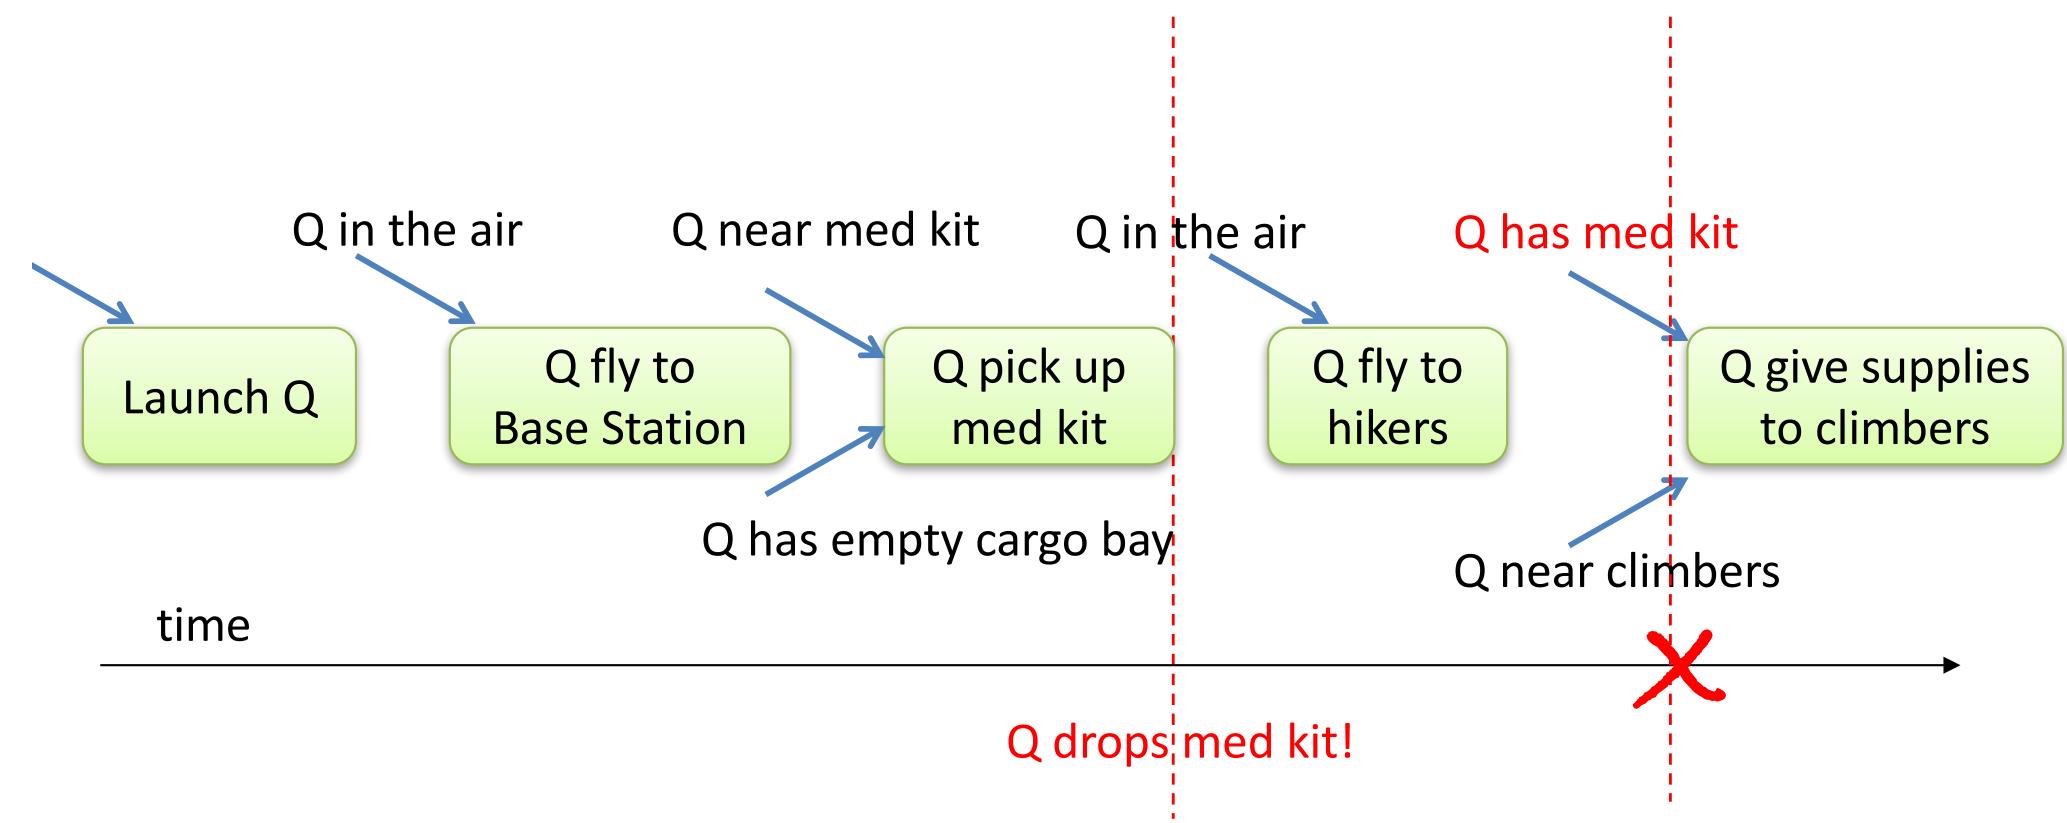
Can set expectations from action preconditions & effects



Preconditions tell executive what's relevant

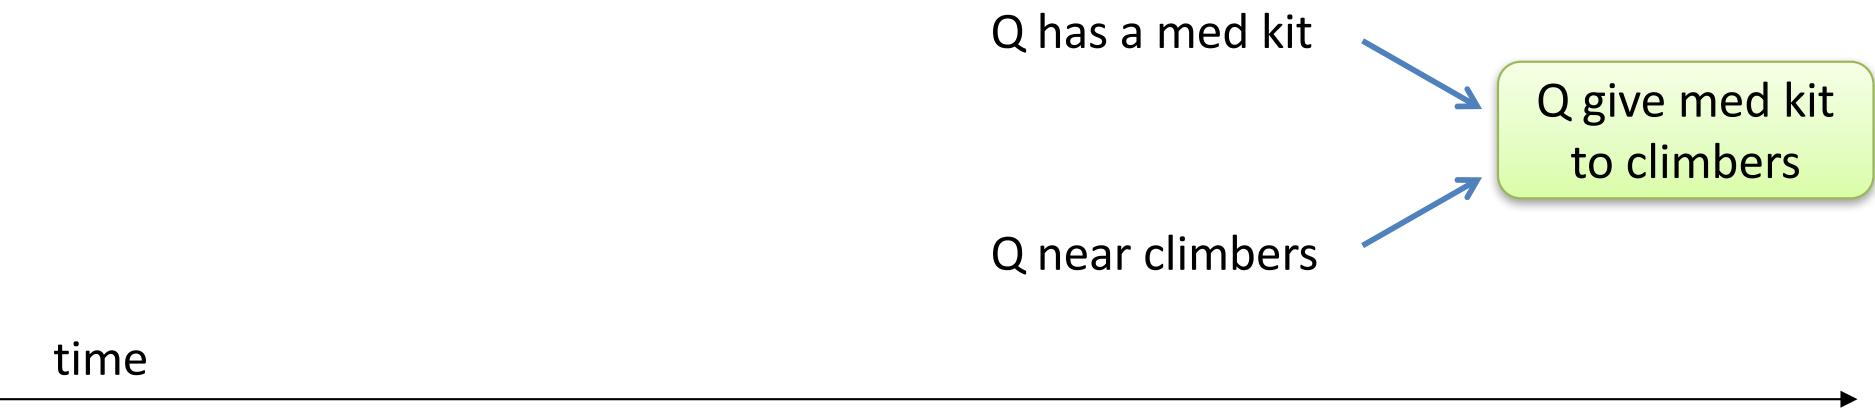


Preconditions tell executive what's relevant



Action monitoring checks preconditions just before action executed.
(e.g., RosPlan)

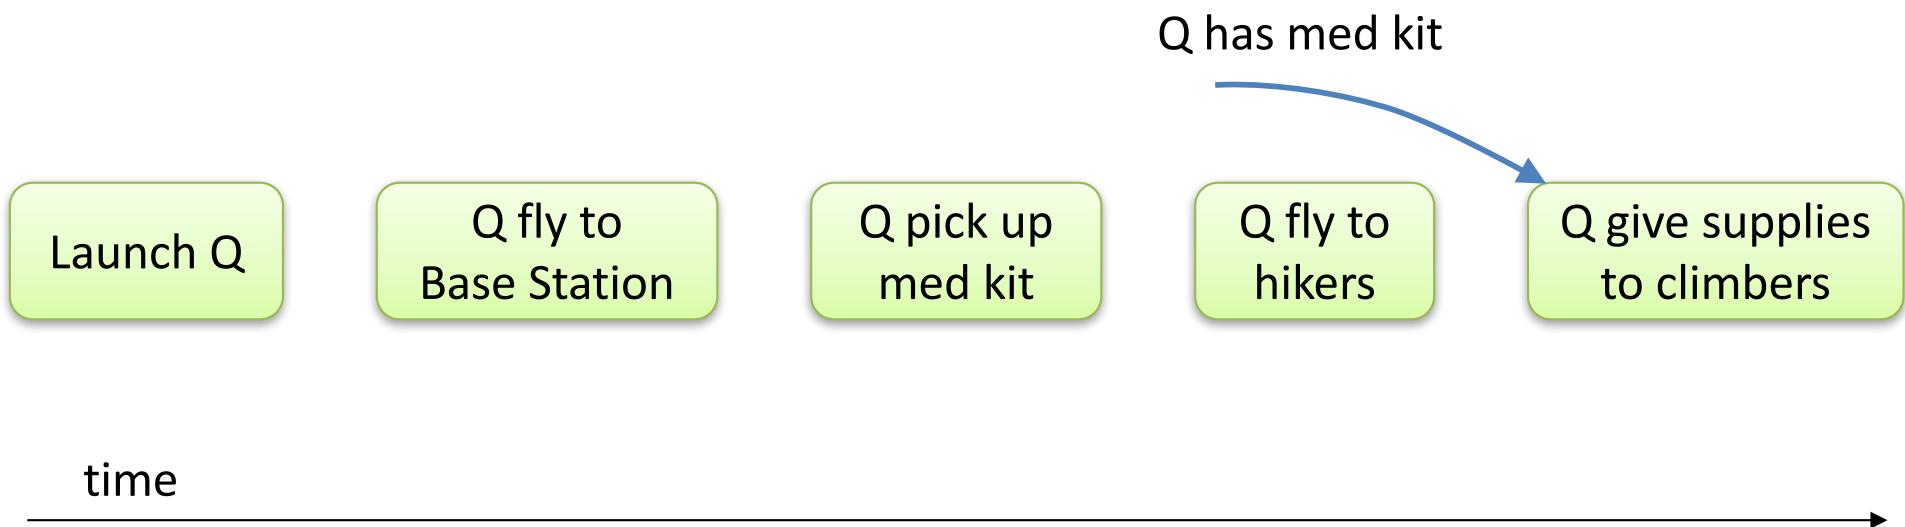
How do we anticipate failures?



Causal-link monitoring checks preconditions
as soon as desired effect is first established.

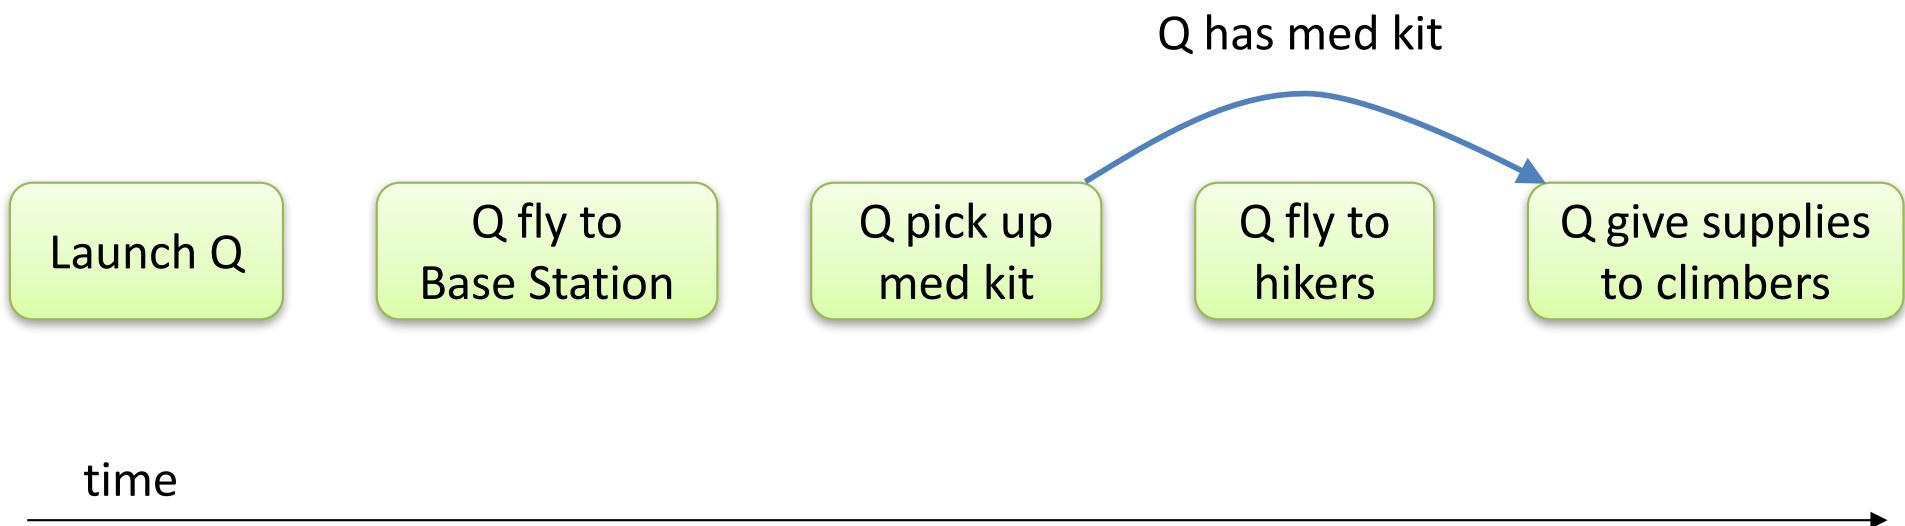
- But for how long should each precondition hold?

When are preconditions established?



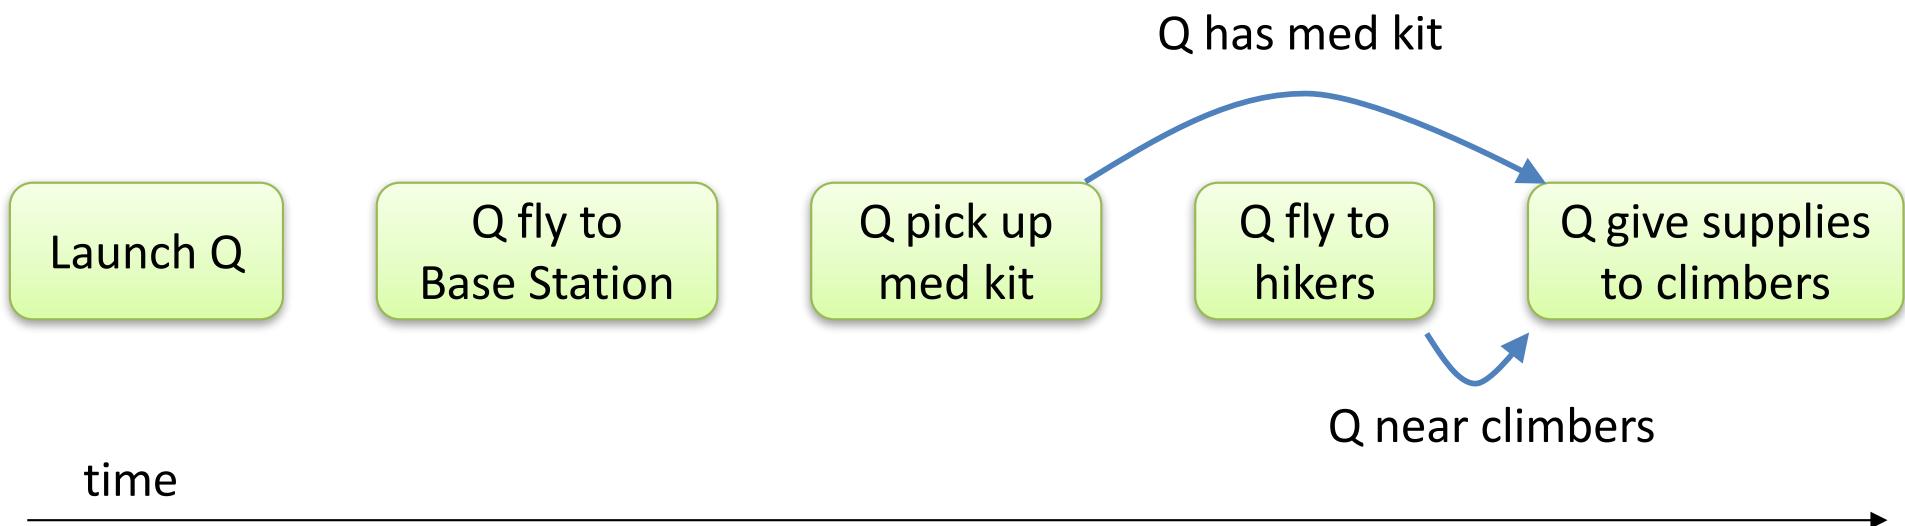
Causal-link monitoring checks preconditions
as soon as desired effect is first established.

When are preconditions established?



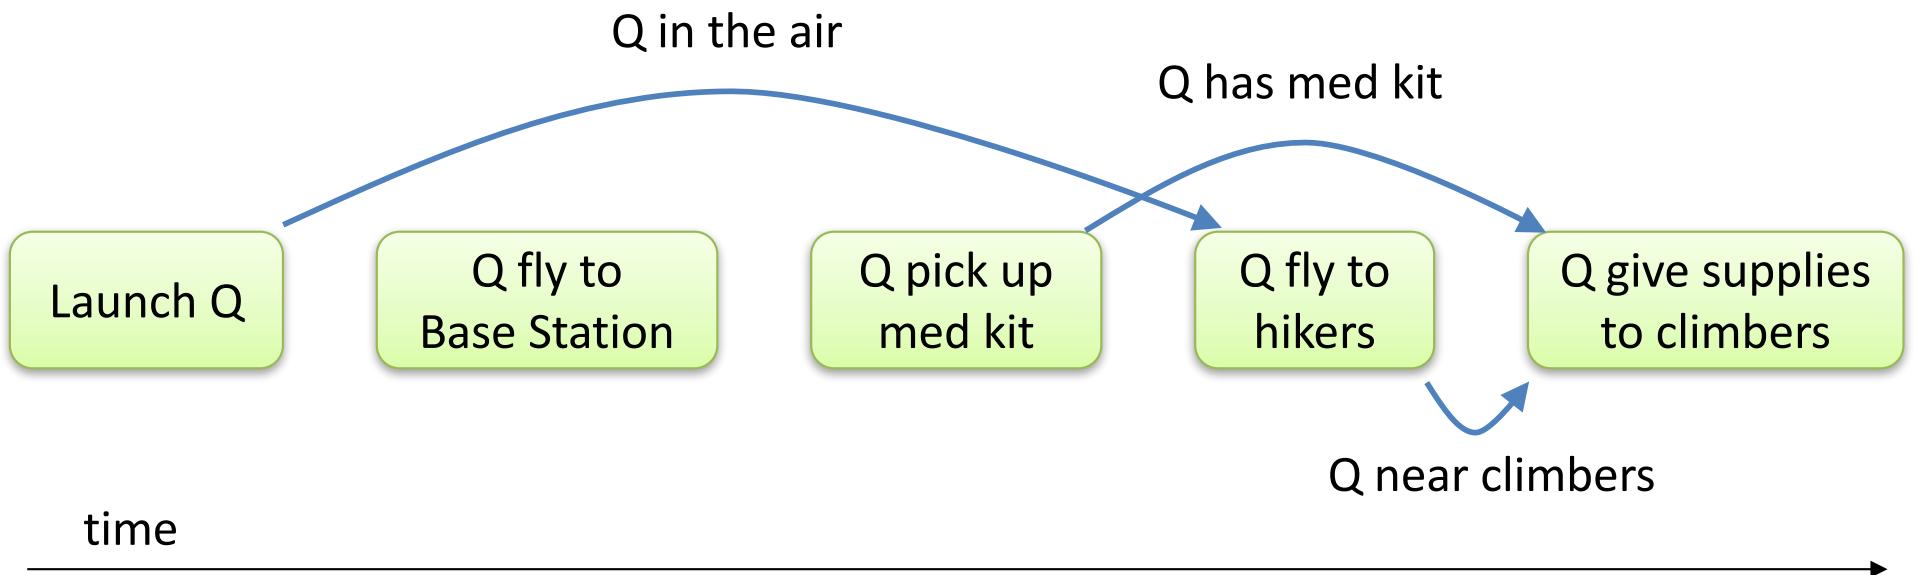
Causal-link monitoring checks preconditions
as soon as desired effect is first established.

When are preconditions established?



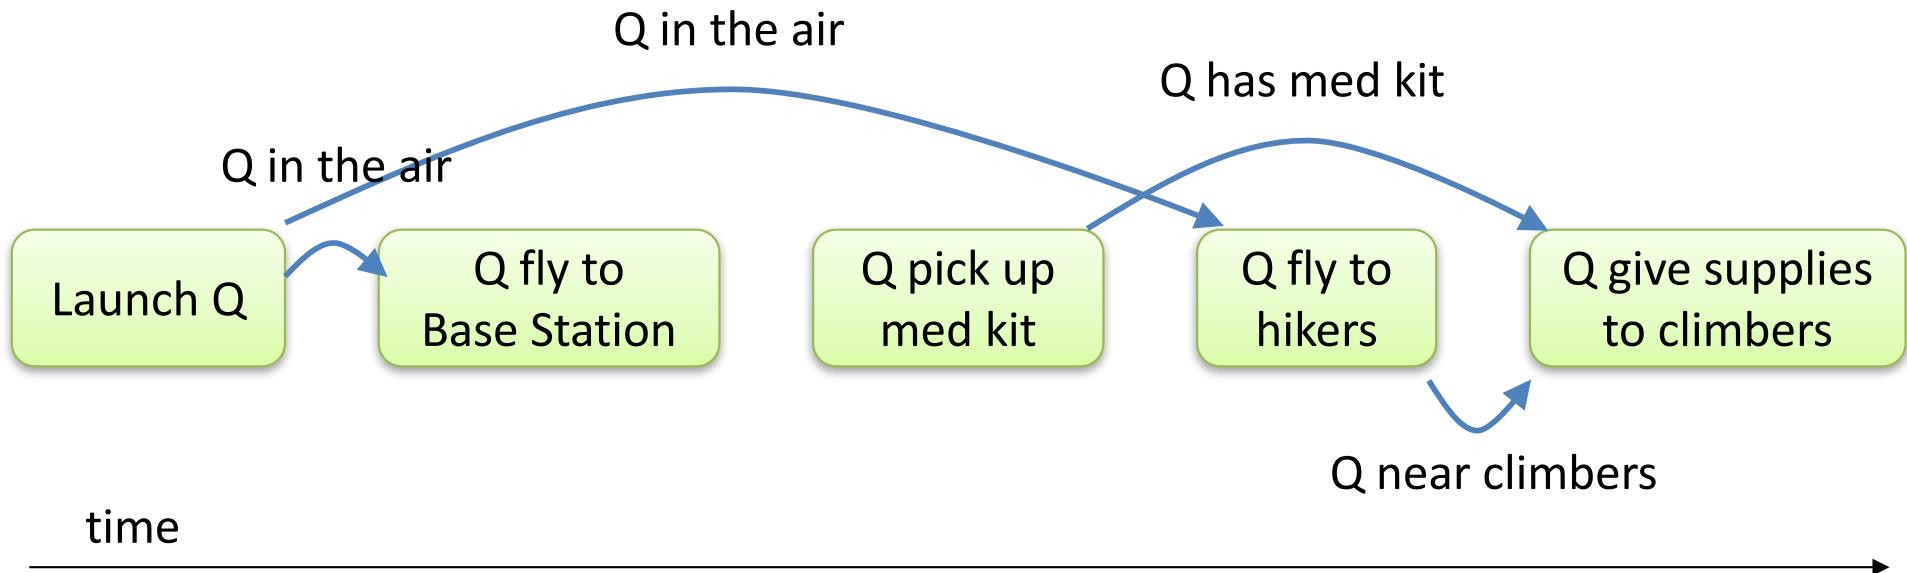
Causal-link monitoring checks preconditions
as soon as desired effect is first established.

When are preconditions established?



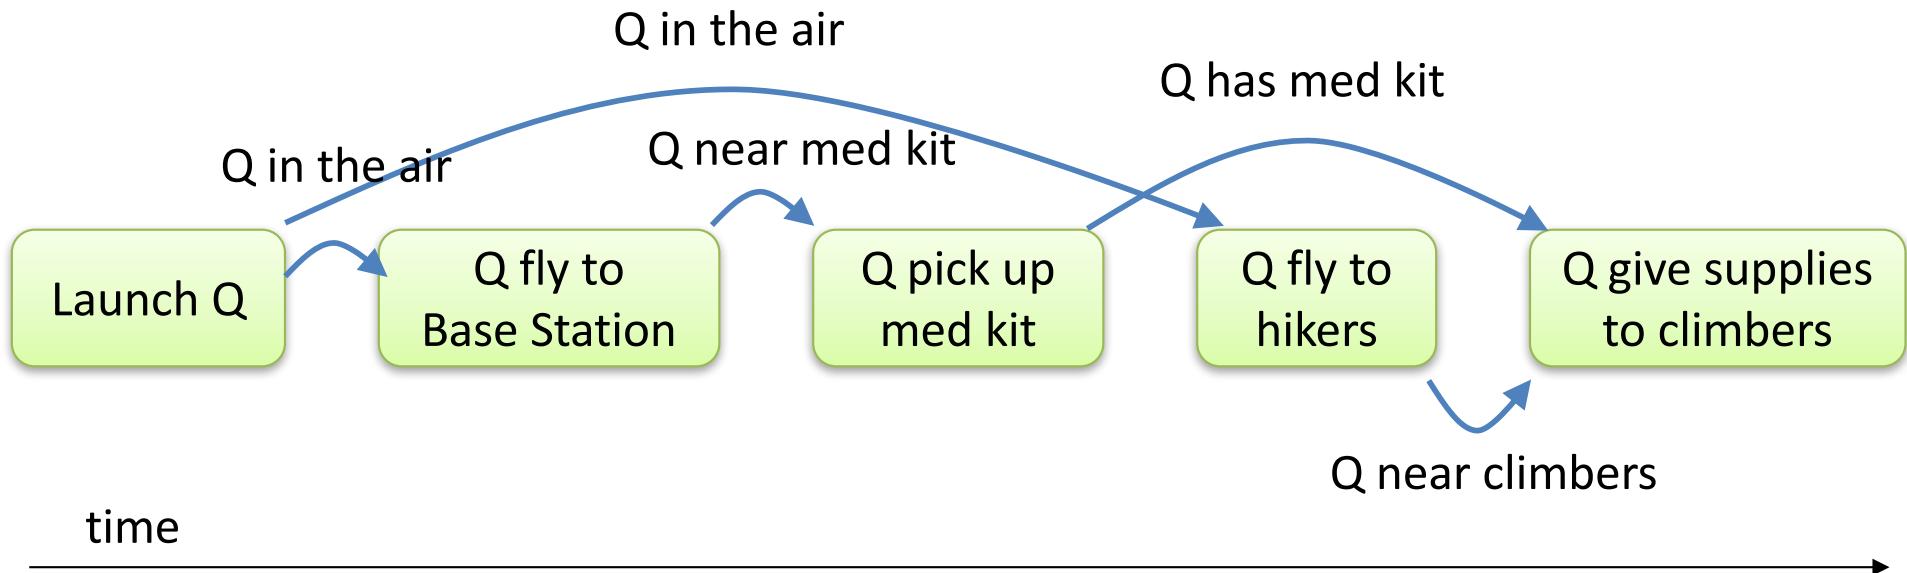
Causal-link monitoring checks preconditions
as soon as desired effect is first established.

When are preconditions established?



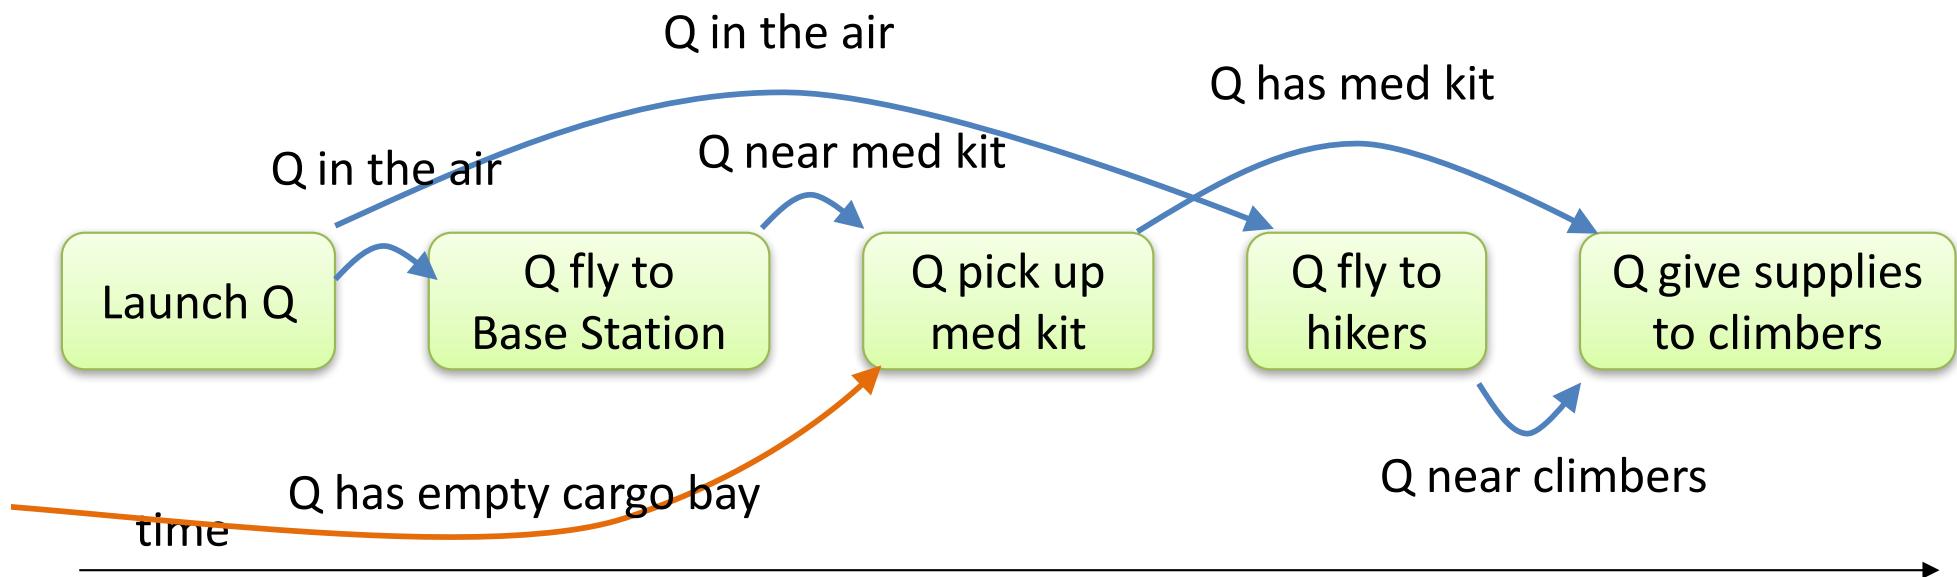
Causal-link monitoring checks preconditions
as soon as desired effect is first established.

When are preconditions established?



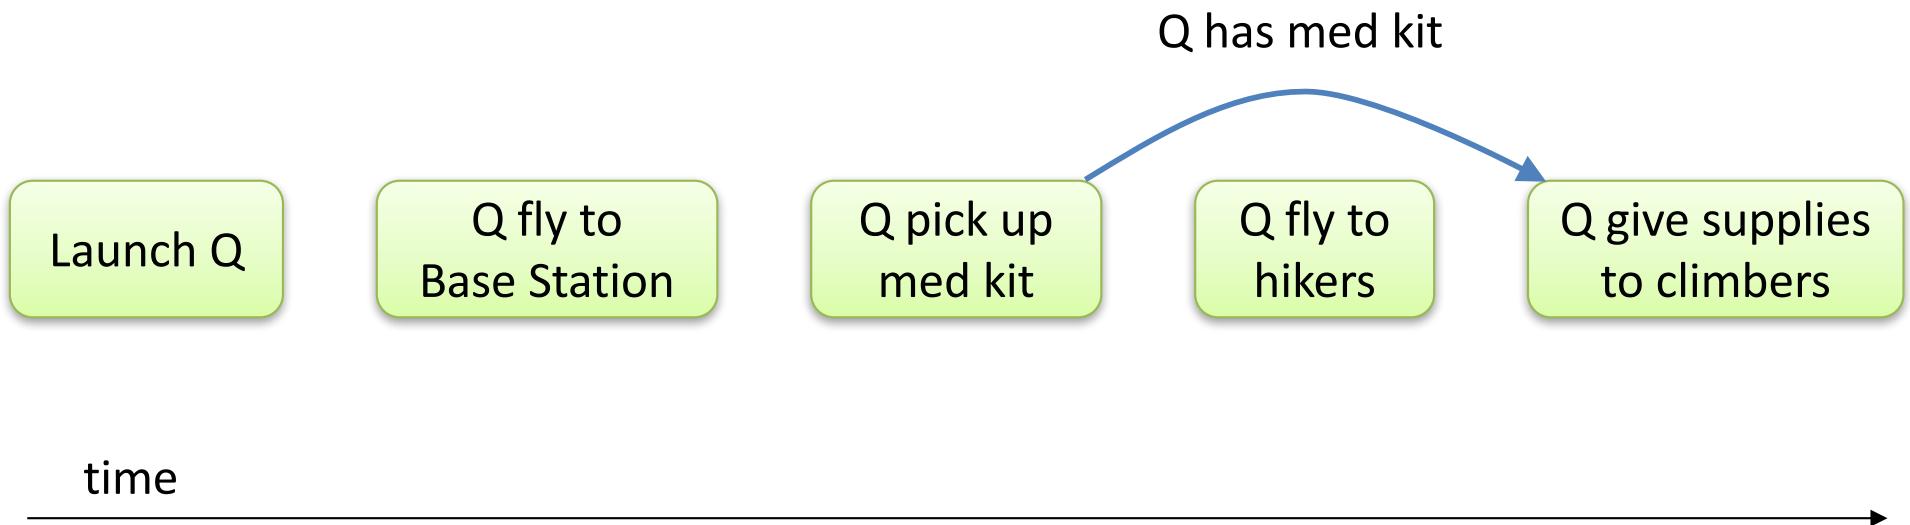
Causal-link monitoring checks preconditions
as soon as desired effect is first established.

When are preconditions established?



Causal-link monitoring checks preconditions
as soon as desired effect is first established.

These are called “causal links”



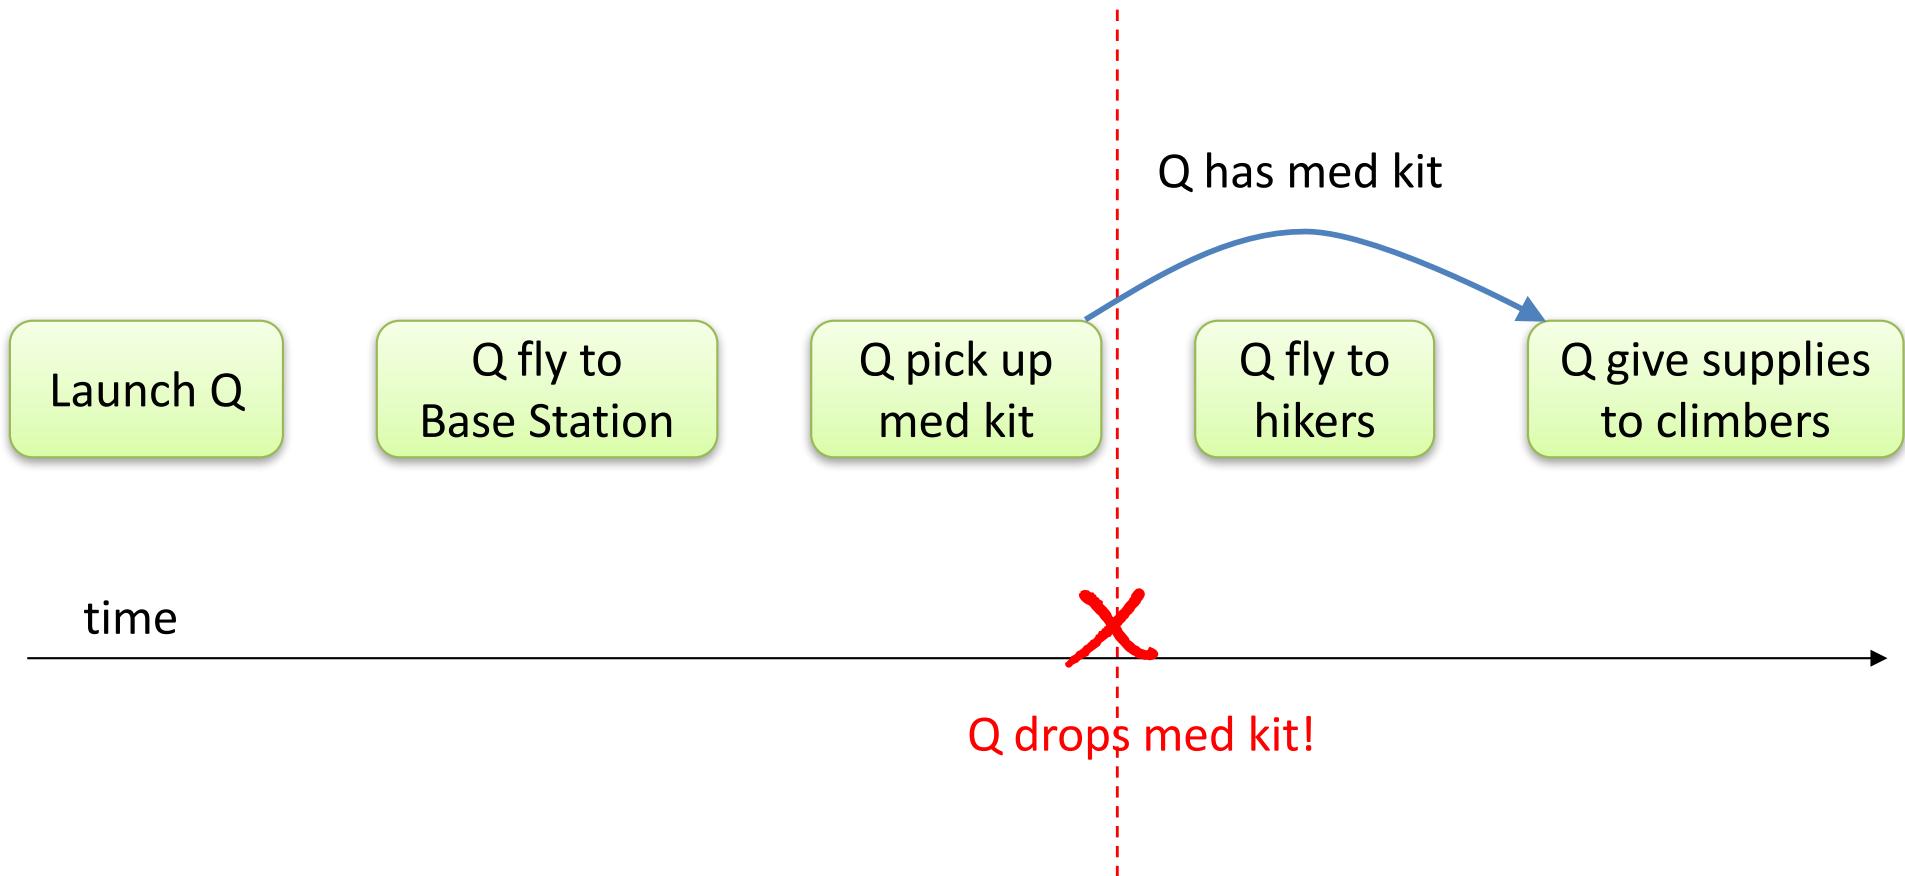
Causal link: identifies how an action's effect is **consumed** by a later action.

Causal Link

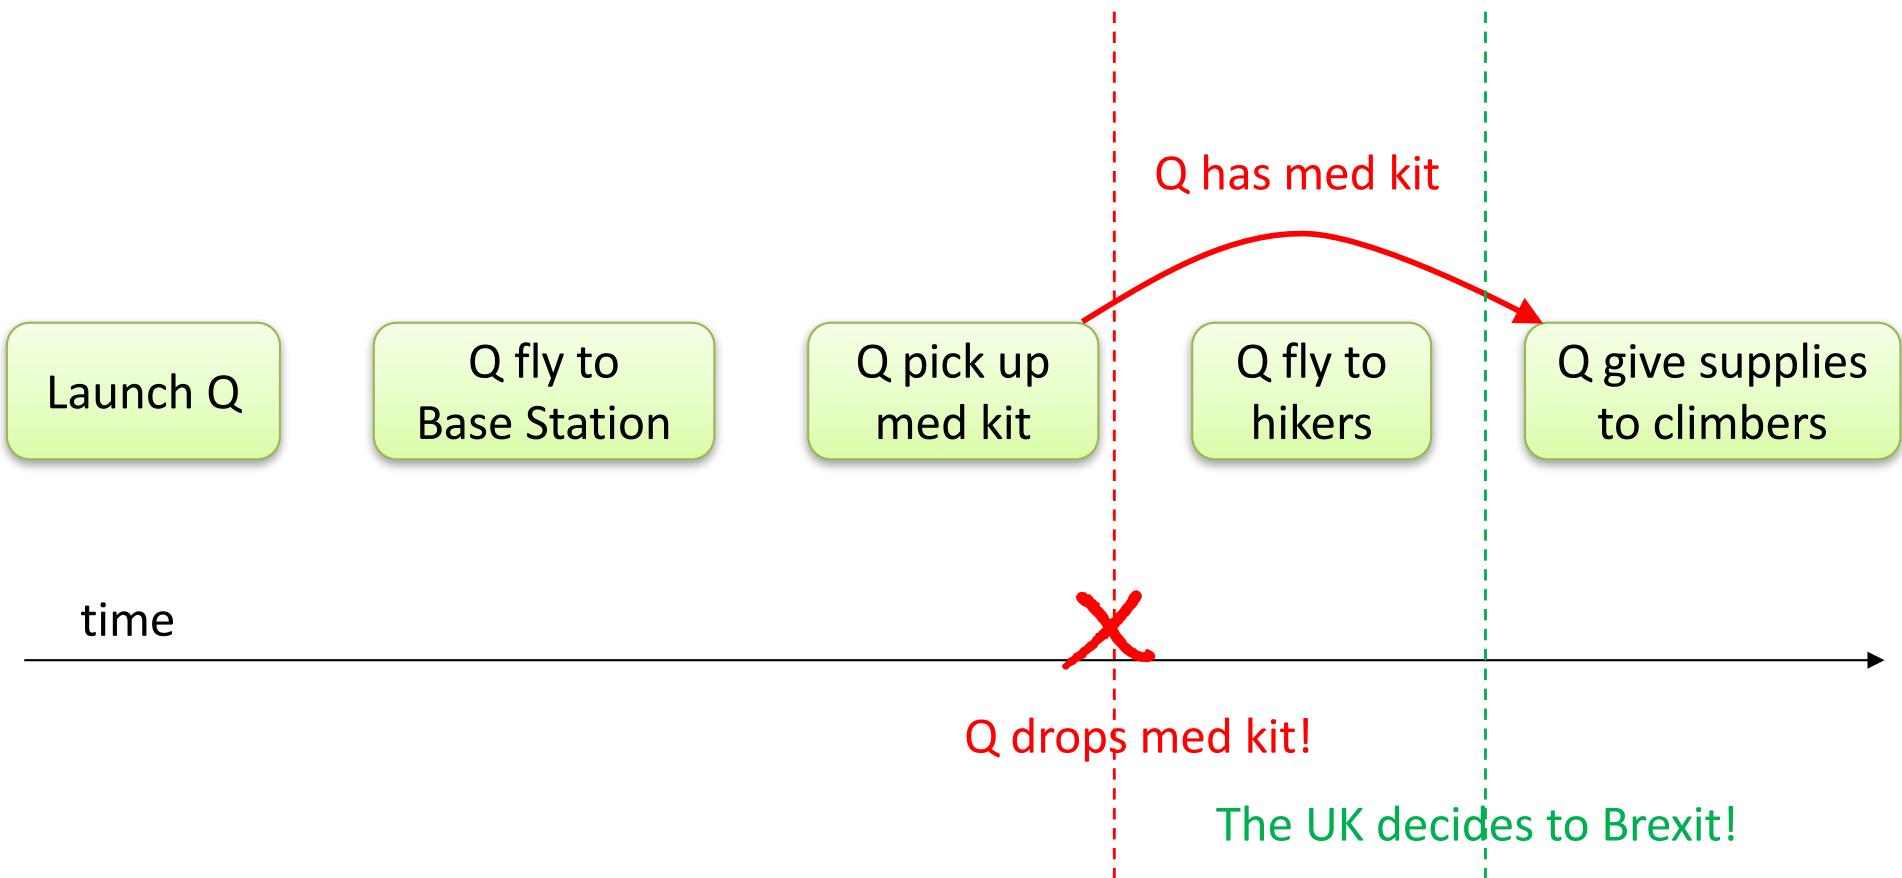
$\langle a_{producer}, \ p, \ a_{consumer} \rangle$

- Proposition p
- Producer: action with p as effect
- Consumer: action with p as precondition
- Producer must precede consumer
- No action interferes in between

Links tell executive what's *relevant* to plan success at any time.



Links tell executive what's *relevant* to plan success at any time.



How do we do perform causal link execution monitoring?

Offline:

- *Extract* causal links from plan & action model

Online:

- Continuously *monitor “active”* causal links against sensor measurements

Active links:

< ... , Q near med kit,

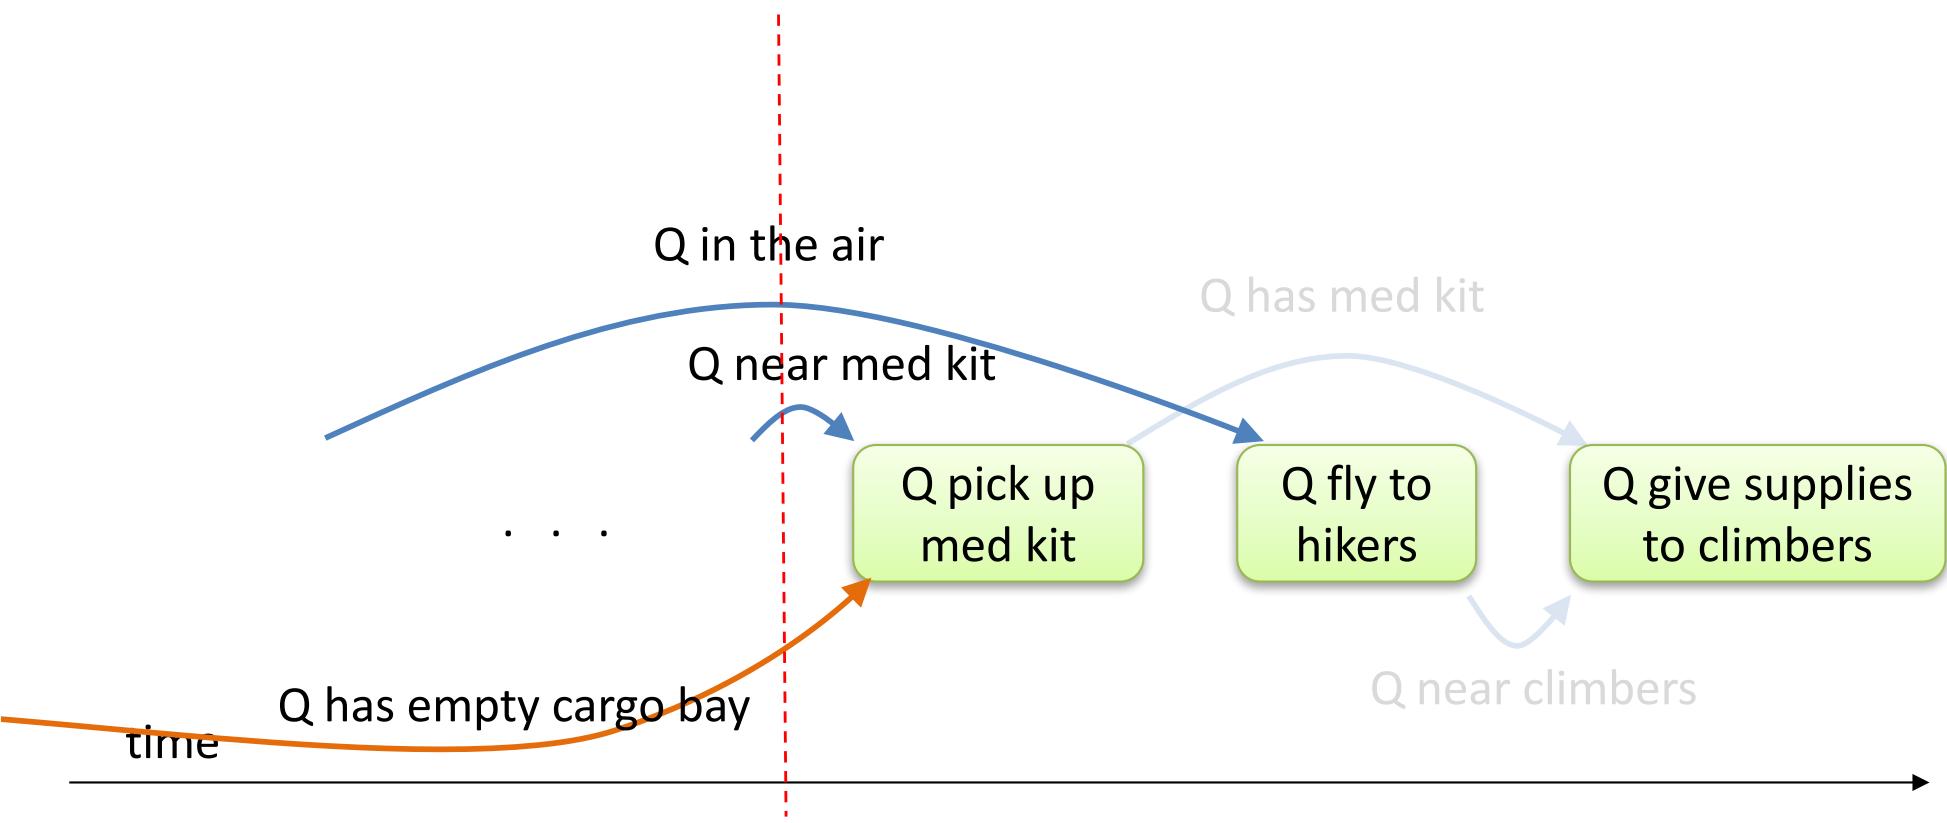
< ... , Q in the air,

< ... , Q has empty cargo bay,

Q pick up med kit>

Q pick up med kit>

Q pick up med kit>



Active links:

< ... , Q near med kit,

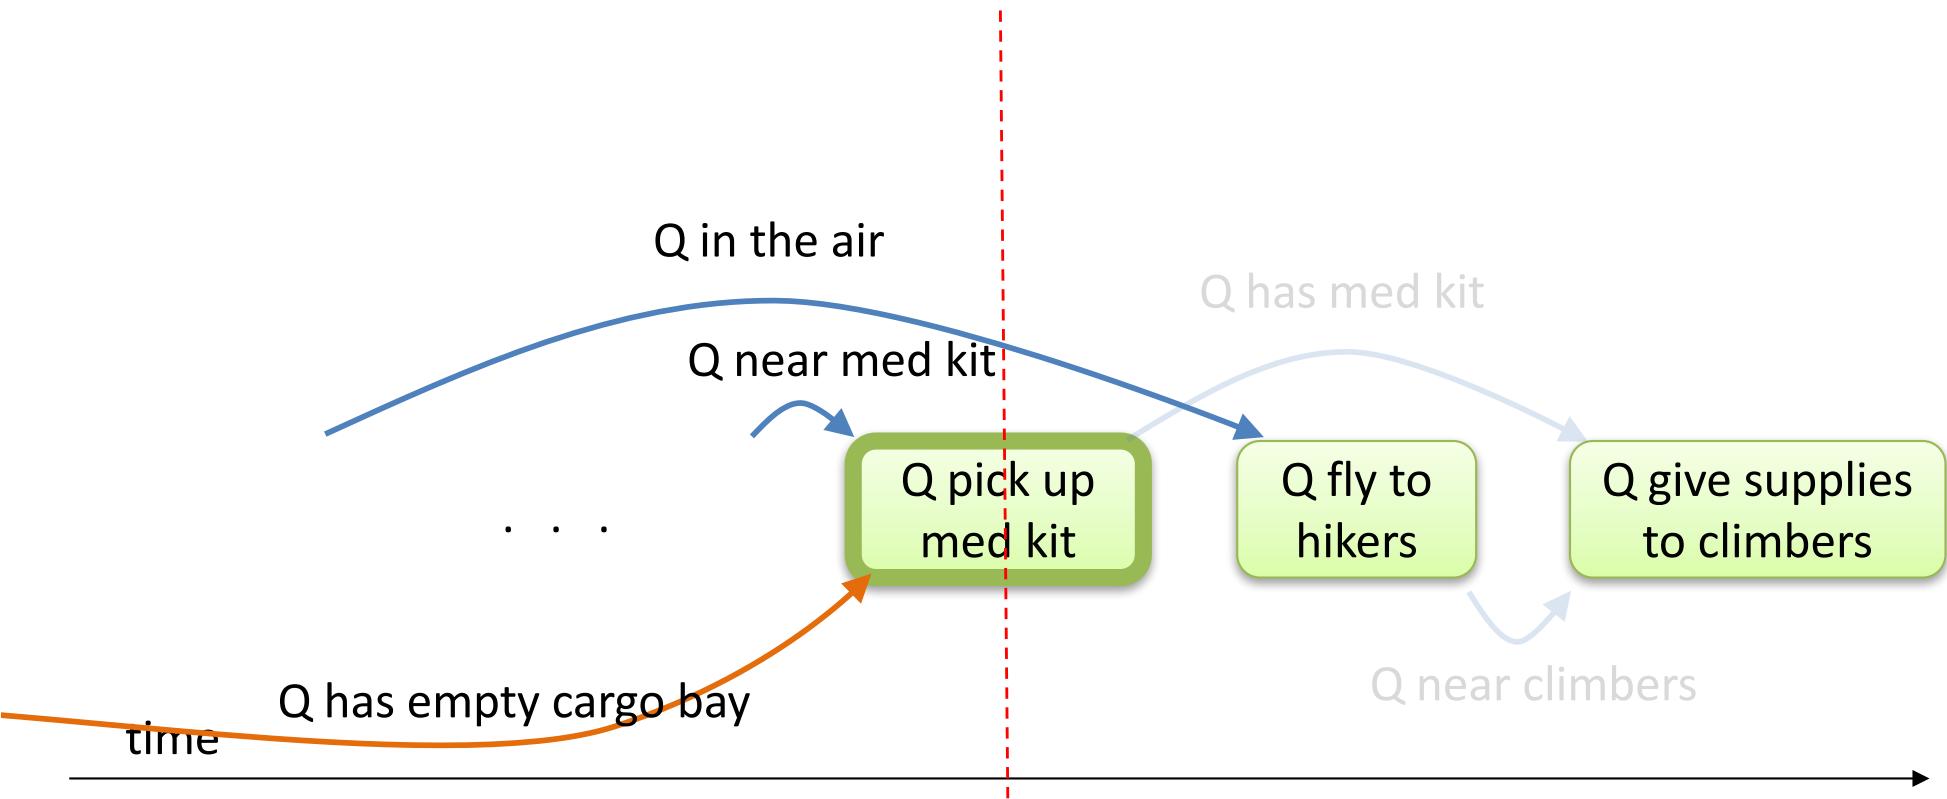
< ... , Q in the air,

< ... , Q has empty cargo bay,

Q pick up med kit>

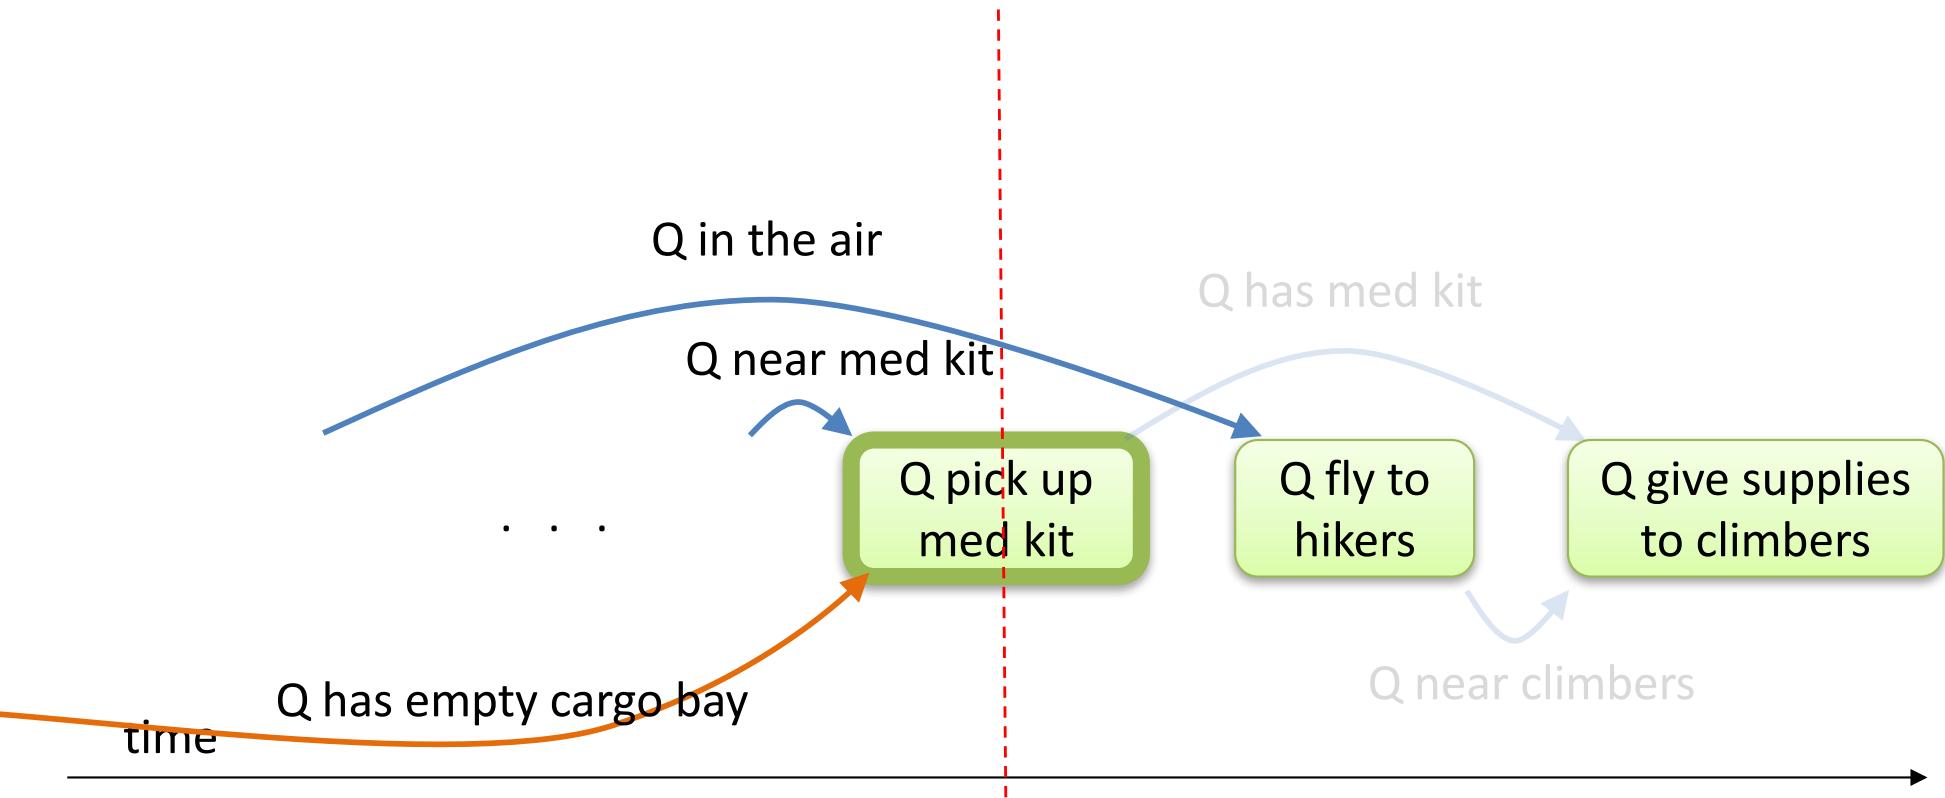
Q pick up med kit>

Q pick up med kit>



Active links:

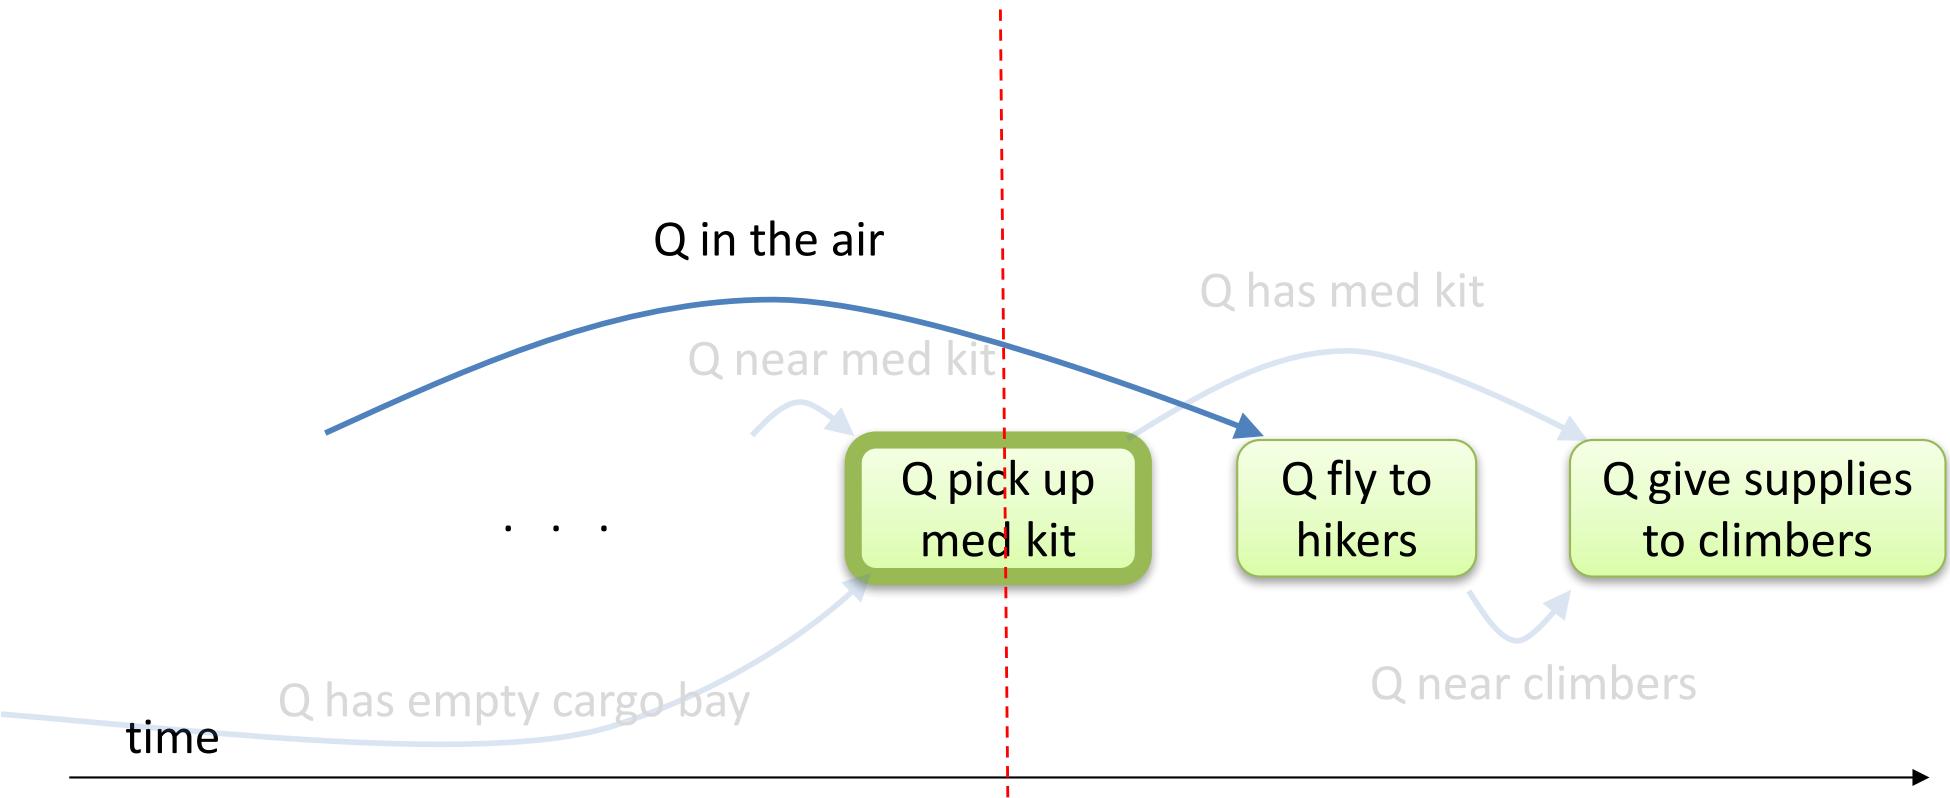
< ... , Q near med kit,	Q pick up med kit>
< ... , Q in the air,	Q pick up med kit>
< ... , Q has empty cargo bay,	Q pick up med kit>



Active links:

< ... , Q in the air,

Q pick up med kit>

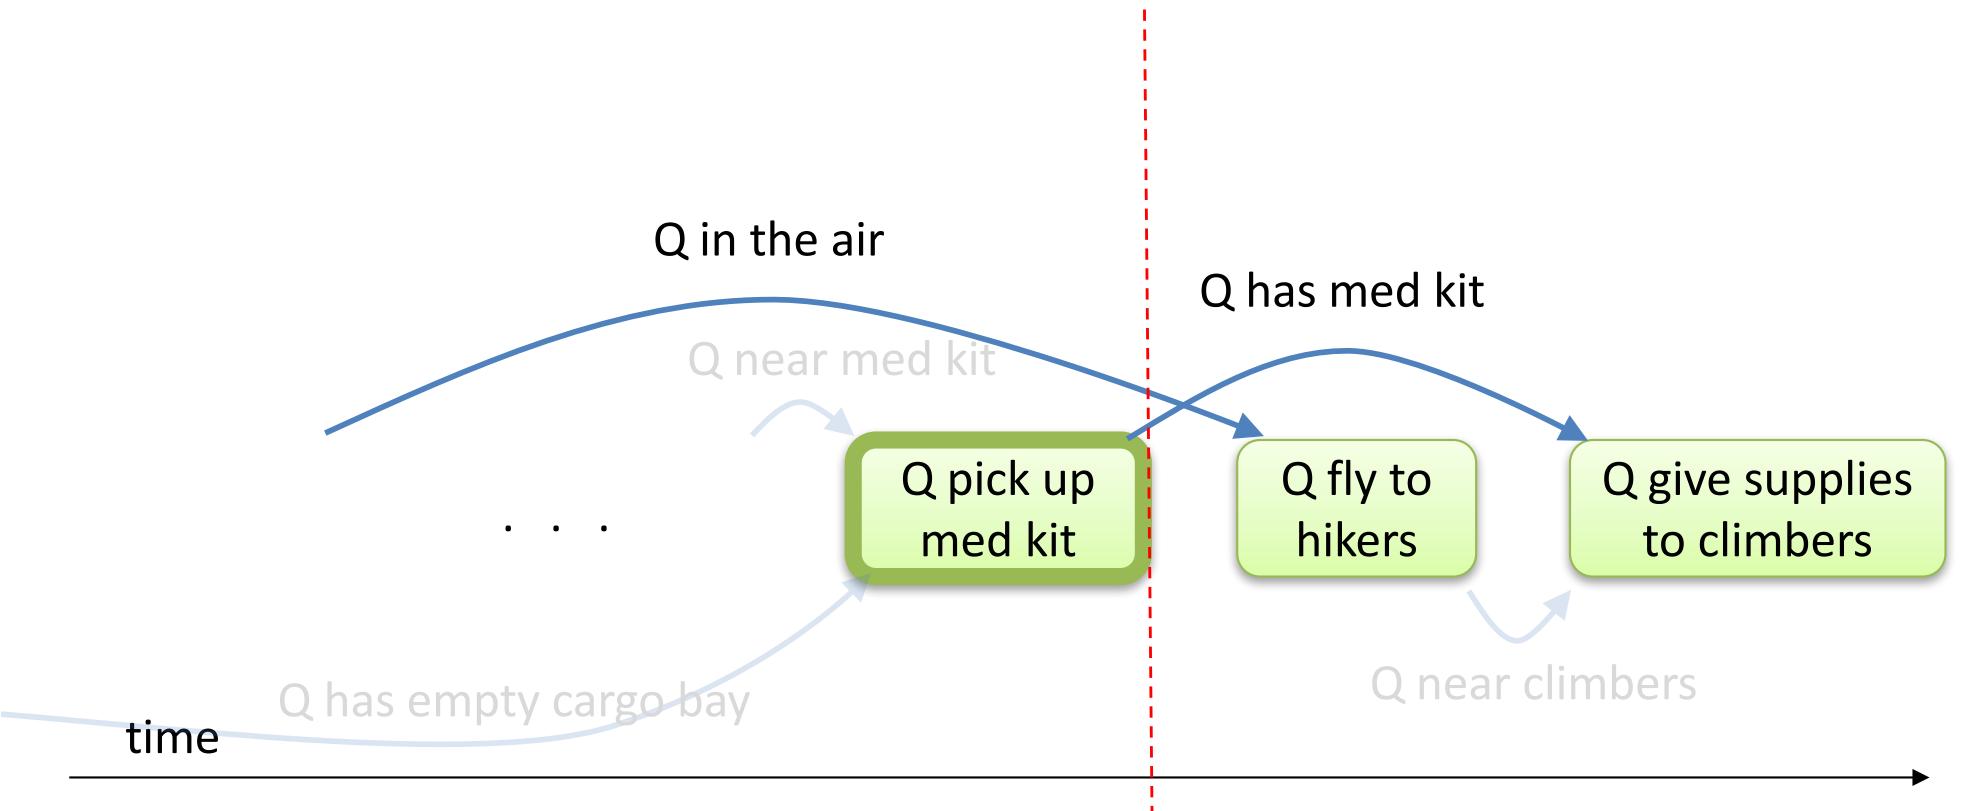


Active links:

<... , Q in the air,

Q pick up med kit>

<Q pick up med kit, Q has med kit, Q give supplies>

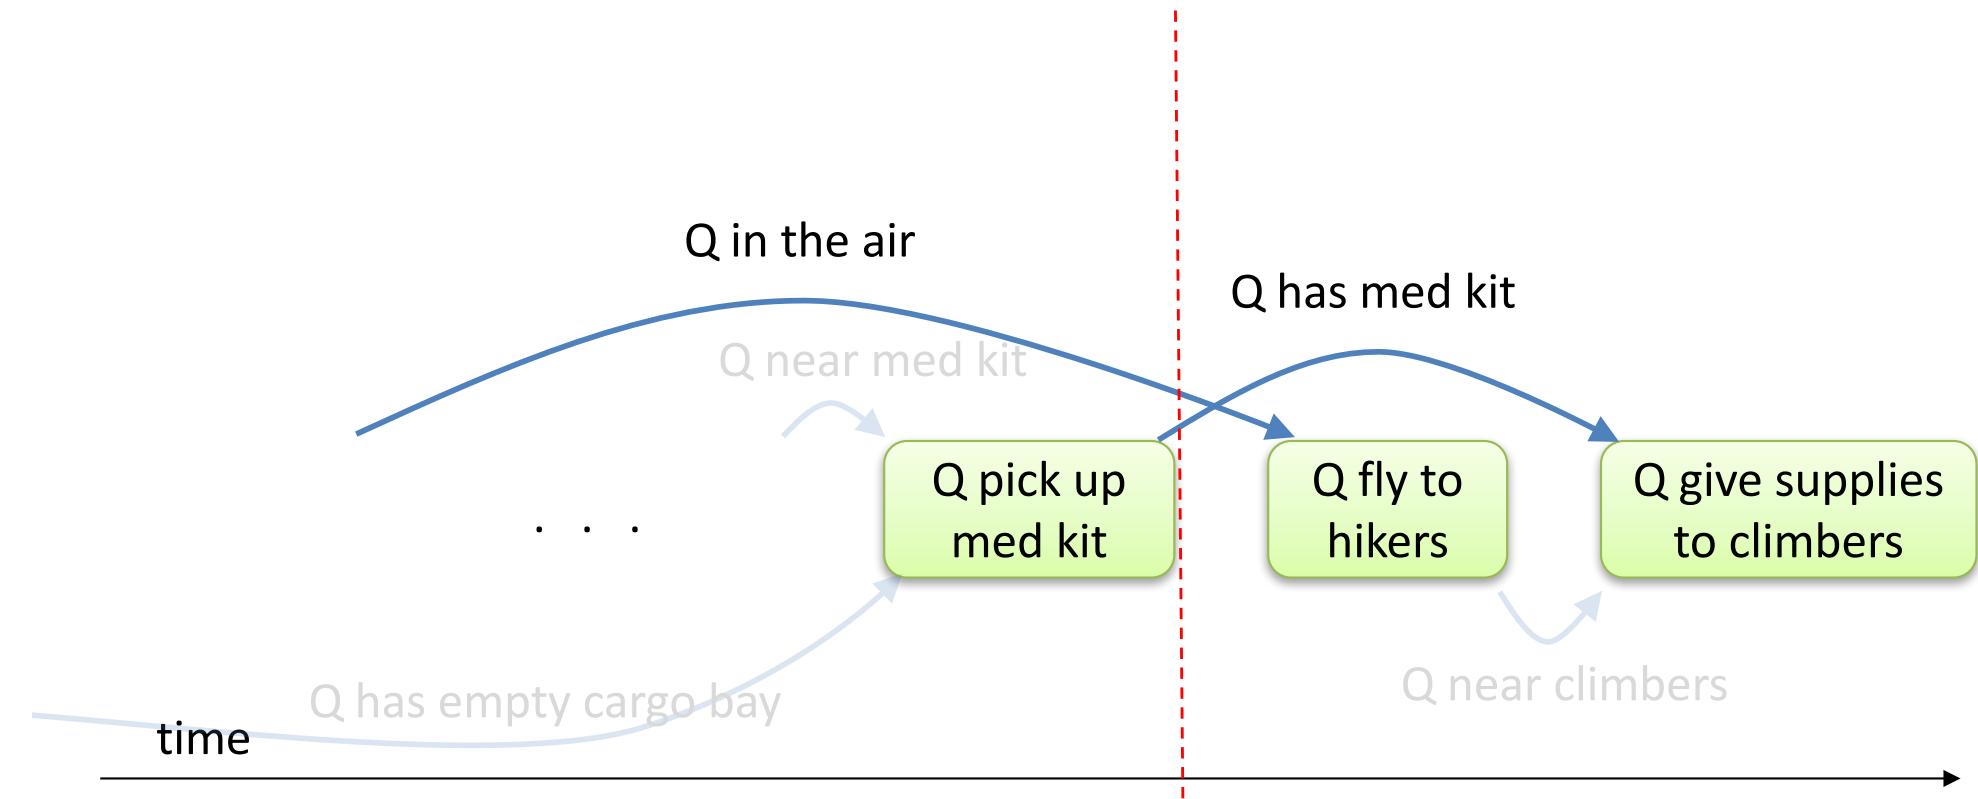


Active links:

<... , Q in the air,

Q pick up med kit>

<Q pick up med kit, Q has med kit, Q give supplies>



Active links:

<... , Q in the air,

Q pick up med kit>

<Q pick up med kit, Q has med kit, Q give supplies>

S:

Q has med kit,

Q in the air,

...

Q in the air

Q near med kit

Q pick up
med kit

Q has med kit

Q fly to
hikers

Q give supplies
to climbers

Q has empty cargo bay

time

Q near climbers



Active links:

<... , Q in the air,

Q pick up med kit>

<Q pick up med kit, Q has med kit, Q give supplies>

S:

Not (Q has med kit),

Q in the air,

...

Q in the air

Q near med kit

Q has med kit

Q pick up
med kit

Q fly to
hikers

Q give supplies
to climbers

Q has empty cargo bay

time

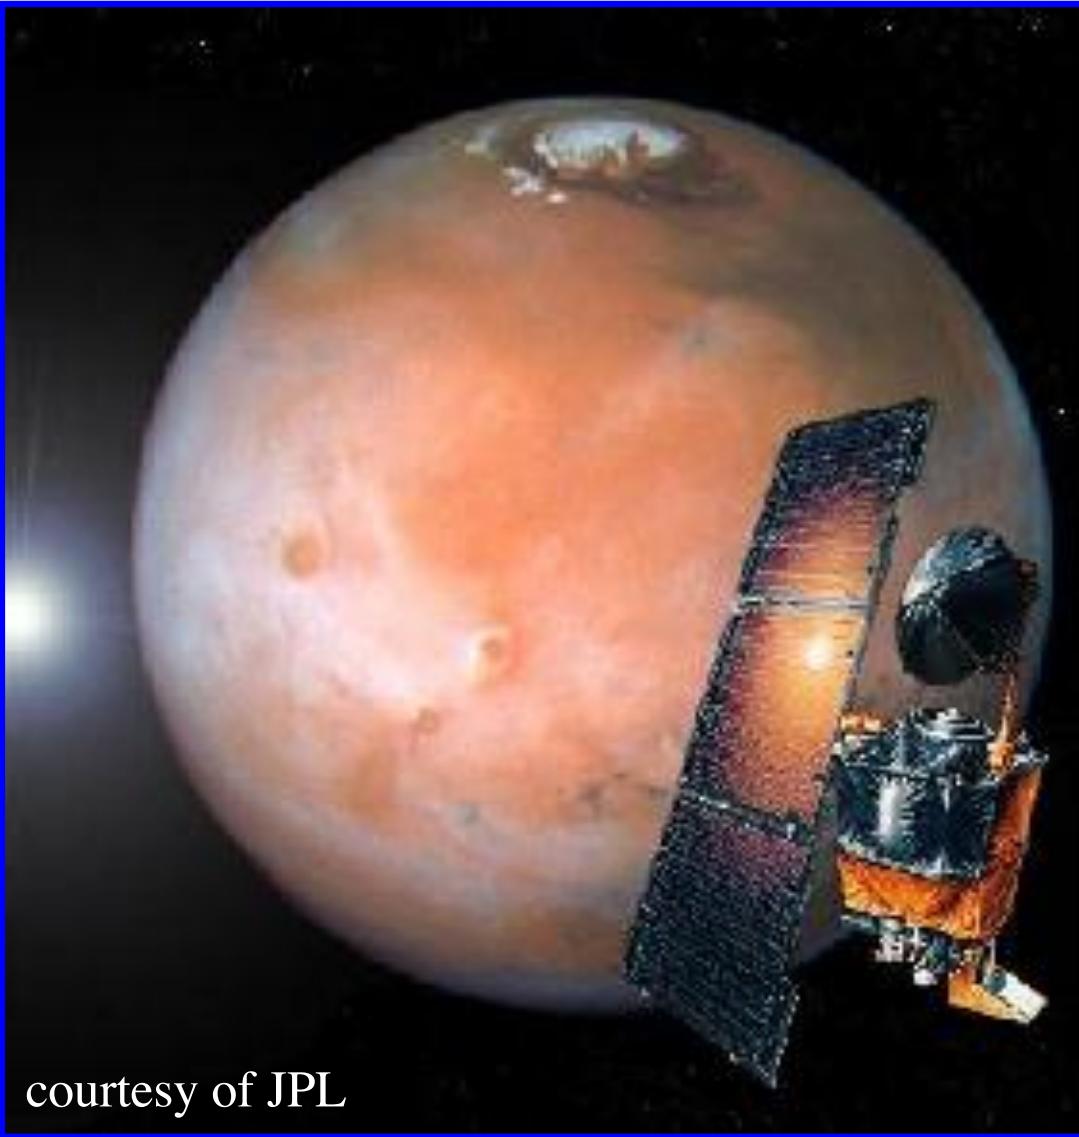
Key takeaways

- Need to monitor plan sub-goals during execution, to **anticipate failures**.
- Causal links:
 - Encode **what** needs to be true, and **when**
 - Offline (check completeness) and online (**monitor**)
- See Appendix for details.

Outline

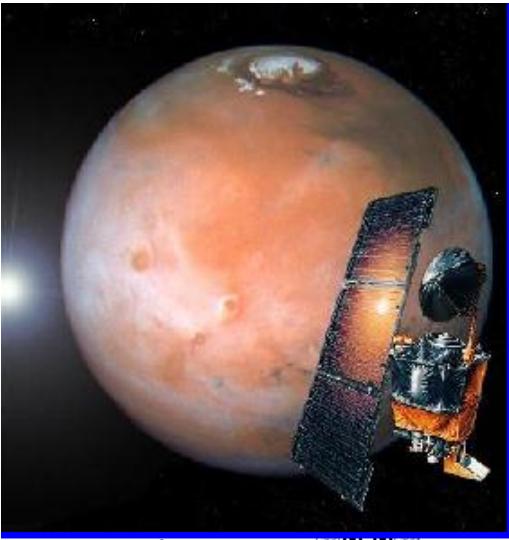
- Programs that Monitor State
 - Sub-goal monitoring
 - Model-based diagnosis and mode estimation
- Programs that Self-Diagnose (opt)

Motivation: Mission Loss Highlights the Challenge of Robustness



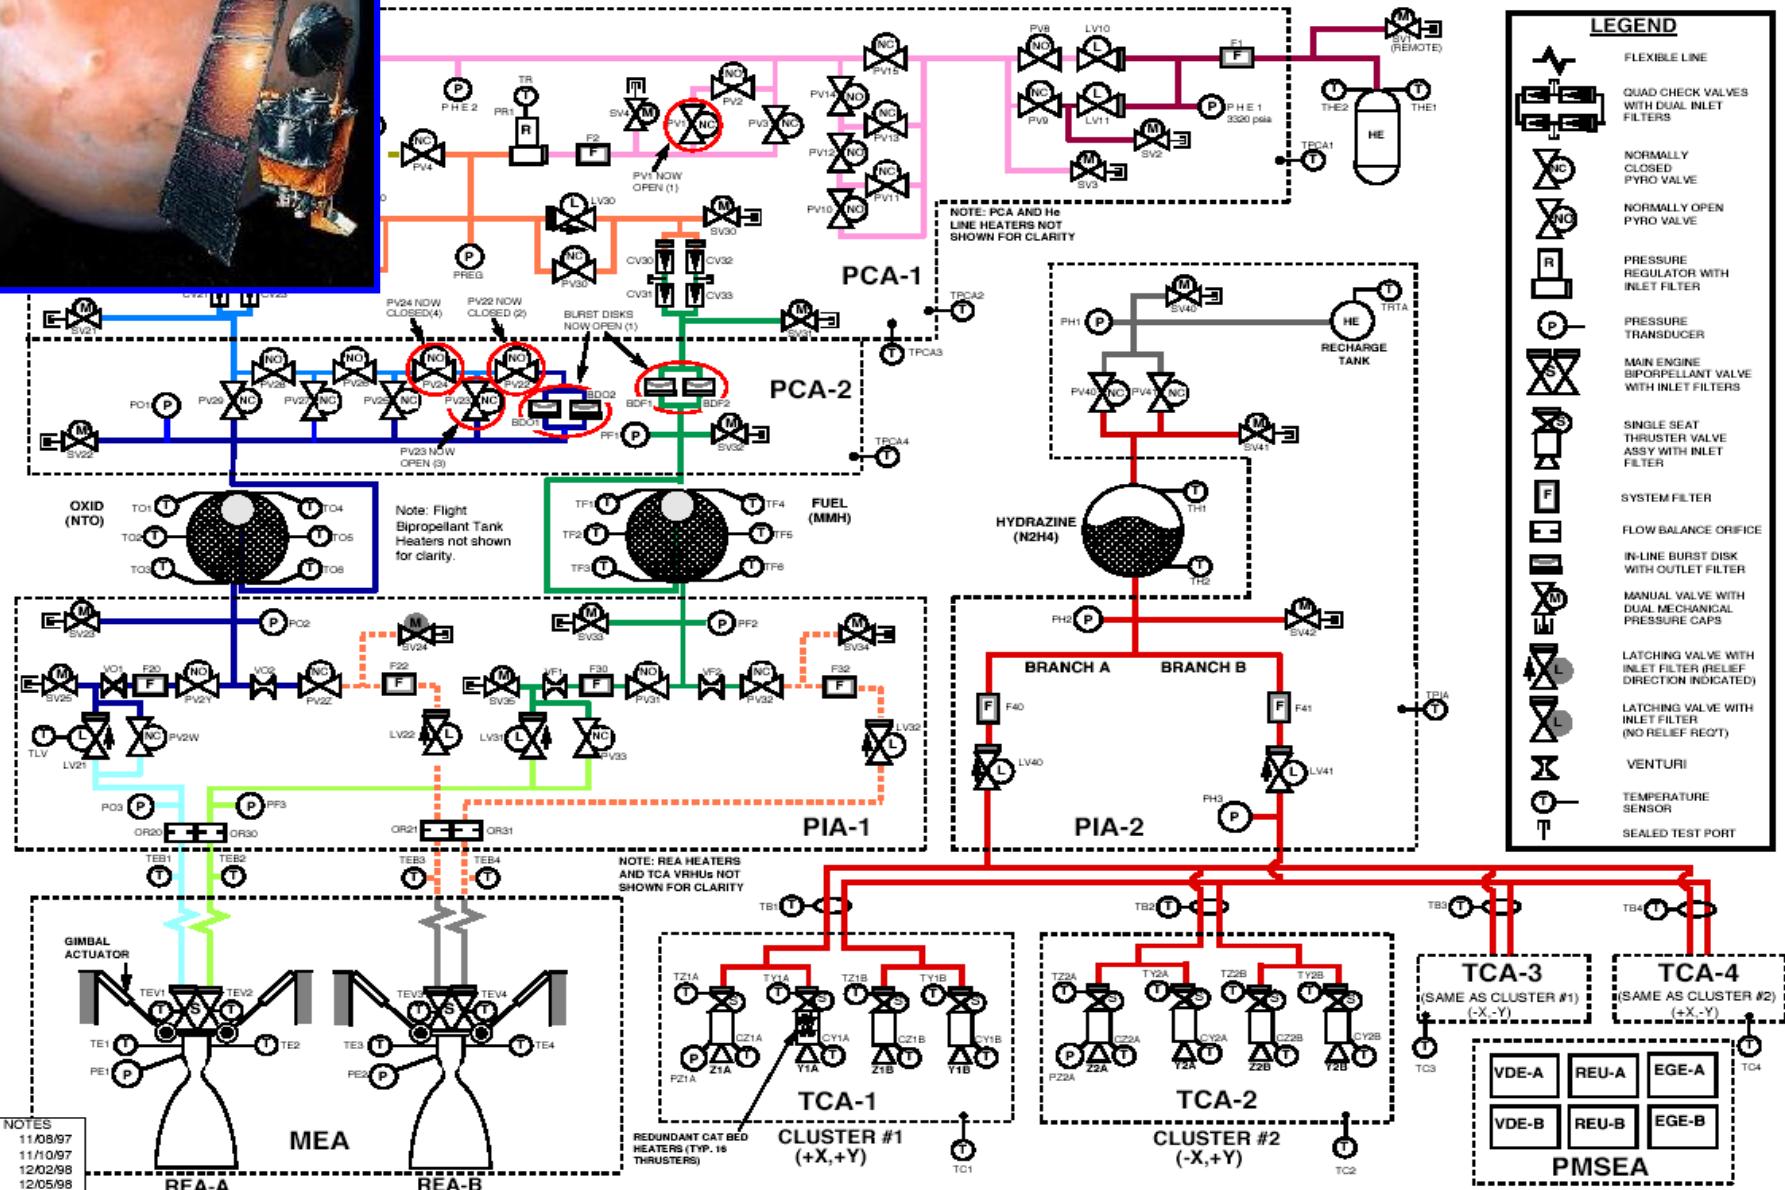
courtesy of JPL

- Mars Observer
- Clementine
- Mars Climate Orbiter
- Mars Polar Lander



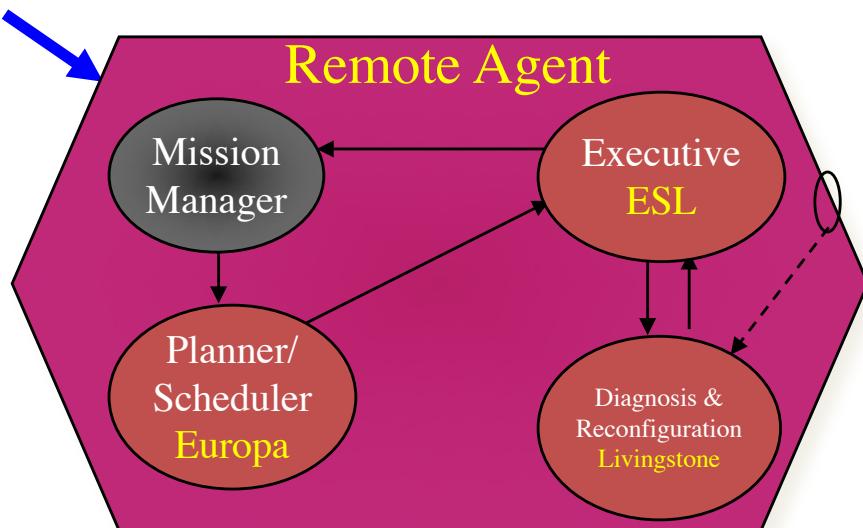
Mars Observer Failure

High Pressure 2416 psia	Ox (N ₂ O ₄) Tank 229 psia	Fuel (MMH) Line 275 psia	Vented REA-B Line 0 psia
LV10-REG 225 psia	Ox (N ₂ O ₄) Line 311 psia	RTA Pressure 2380 psia	Pad Pressure 50 psia
REG-Chk Valve 241 psia	Fuel (MMH) Tank 239 psia	Hydrazine N ₂ H ₄ 340 psia	



Remote Agent on Deep Space One - May, 1999

Goals



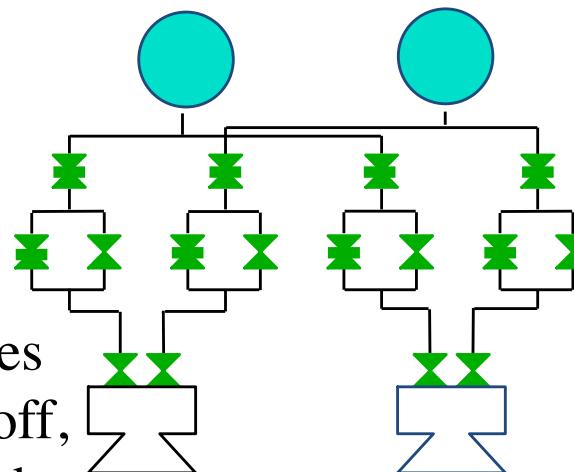
1. Commanded by giving goals
2. Reasons from commonsense models
3. Closes loop on goals
4. Diagnoses, repairs and re-plans.

[Williams & Nayak, AAAI 95;
Muscettola et al, AIJ 00]

Executive assigns EngineA = Thrusting, and Livingstone. . . .

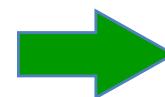
Estimate Modes

Oxidizer tank Fuel tank

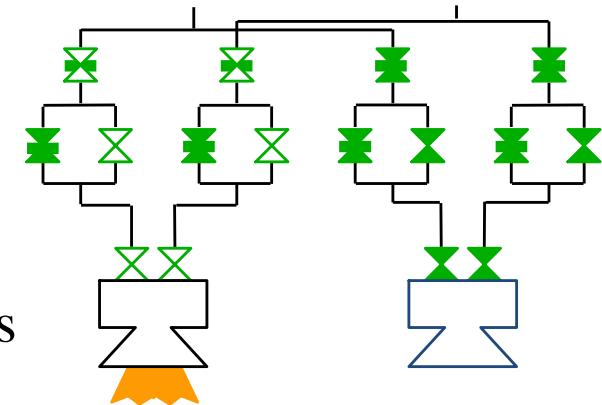


Deduces that valves
are closed, thrust off,
and engine is healthy

Reconfigure Modes



Selects valves
To open, and
plans actions
to open valves



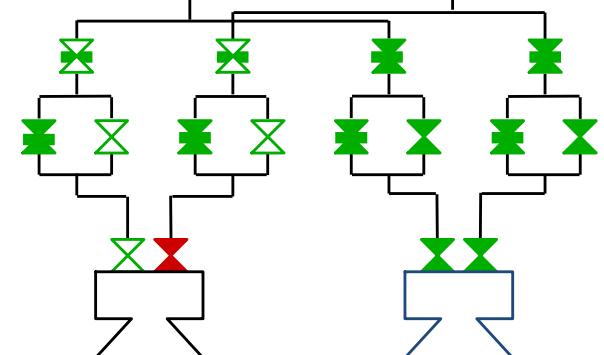
Deduces that a valve
failed - stuck closed

Reconfigure Modes

Prog: EngineB = Thrusting



Selects valves in the
backup engine B that
will achieve thrust, and
plans needed actions.



Model: Probabilistic Constraint Automata (PCA)

component modes...

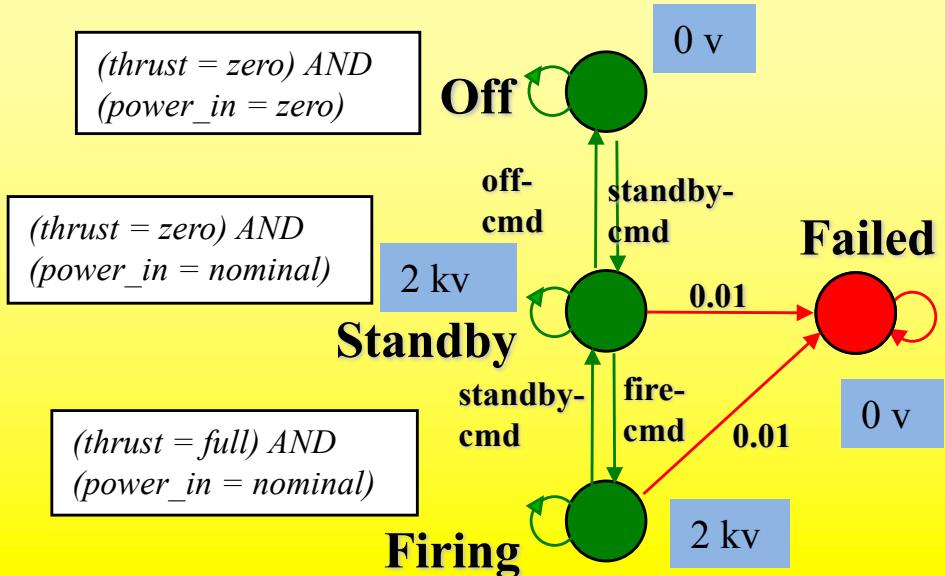
described by finite domain constraints on variables...

guarded deterministic and probabilistic transitions

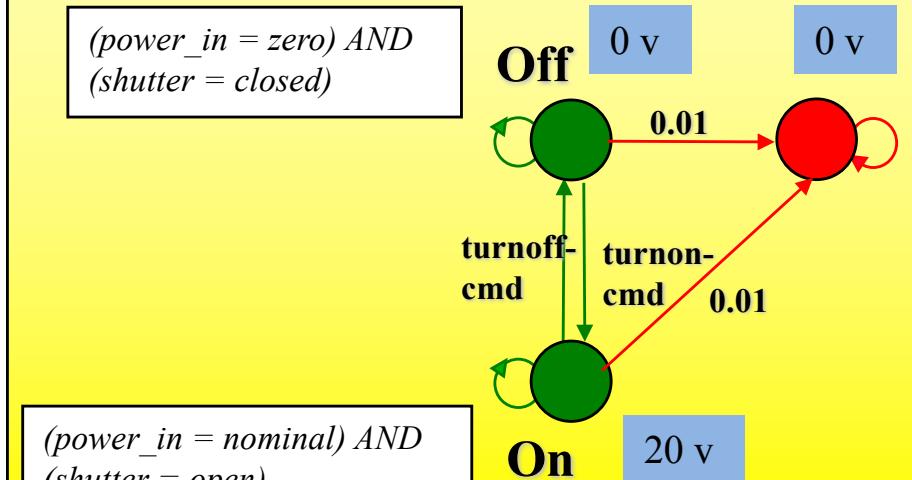
cost / reward & prior distribution

[Williams & Nayak AAAI 95,
Williams et al. IEEE Proc. 01]

Engine Model



Camera Model



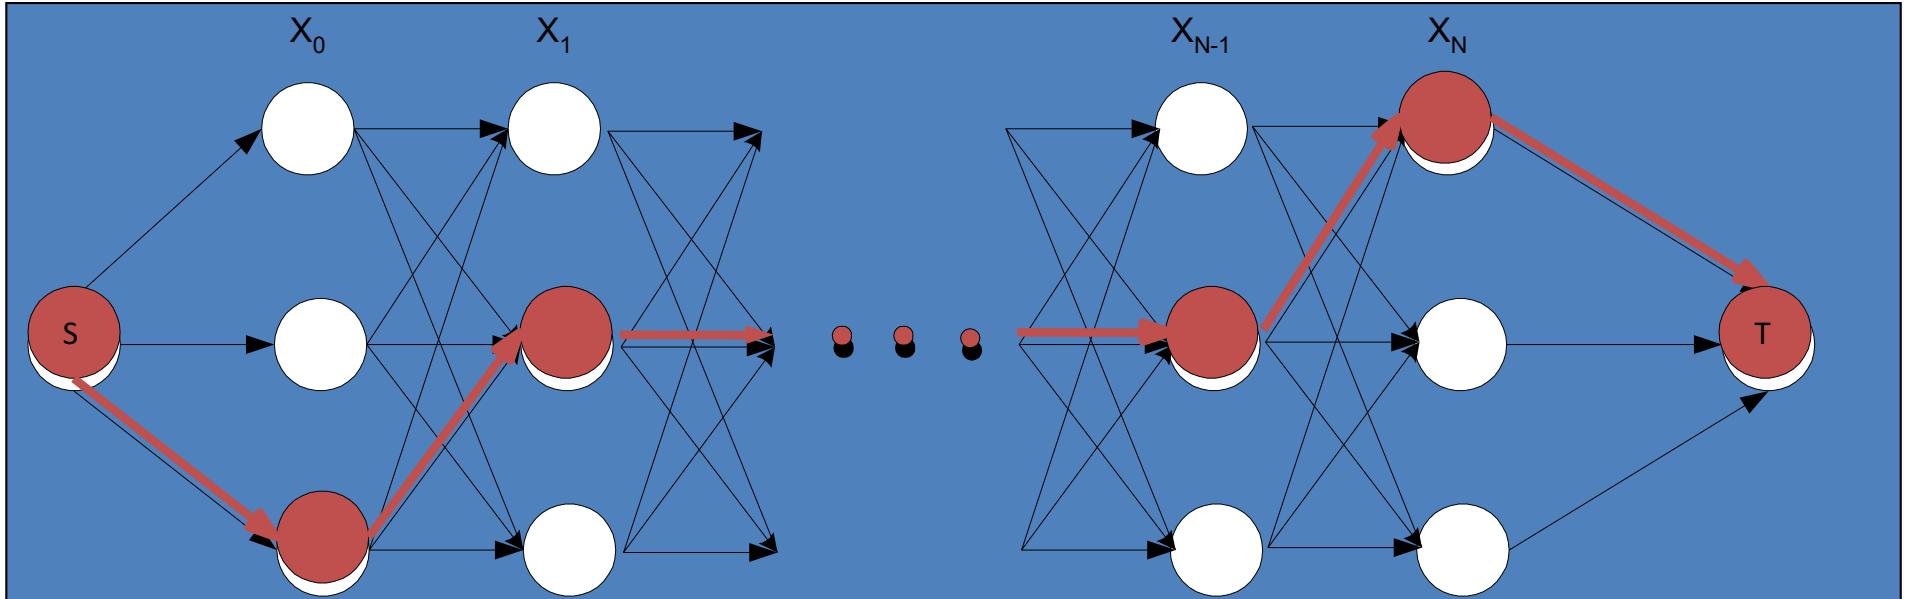
one per component ... operating concurrently

6/12/2017

T2 self-monitoring, self-diagnosing systems

46

Mode Estimation = Belief State Update



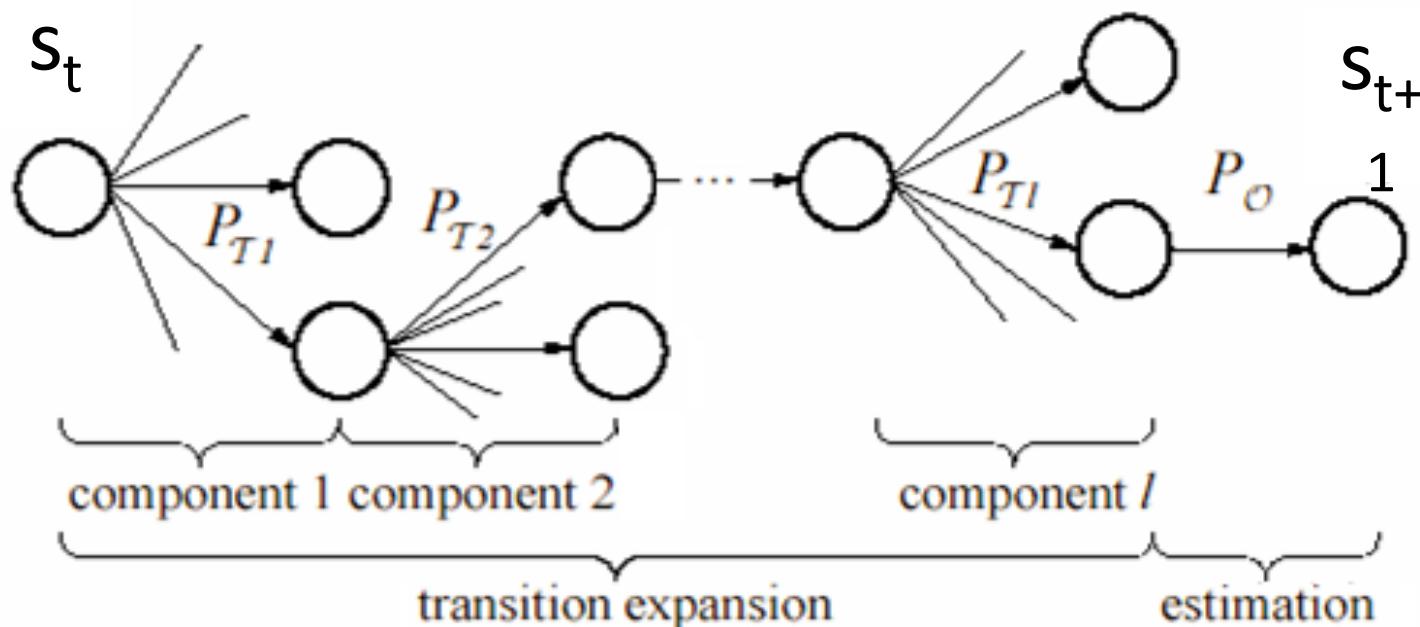
Given prior commands and observations:

1. **What?** Infer distribution on **states** (modes).
2. **How?** Infer most likely state (*mode*) **trajectories**.

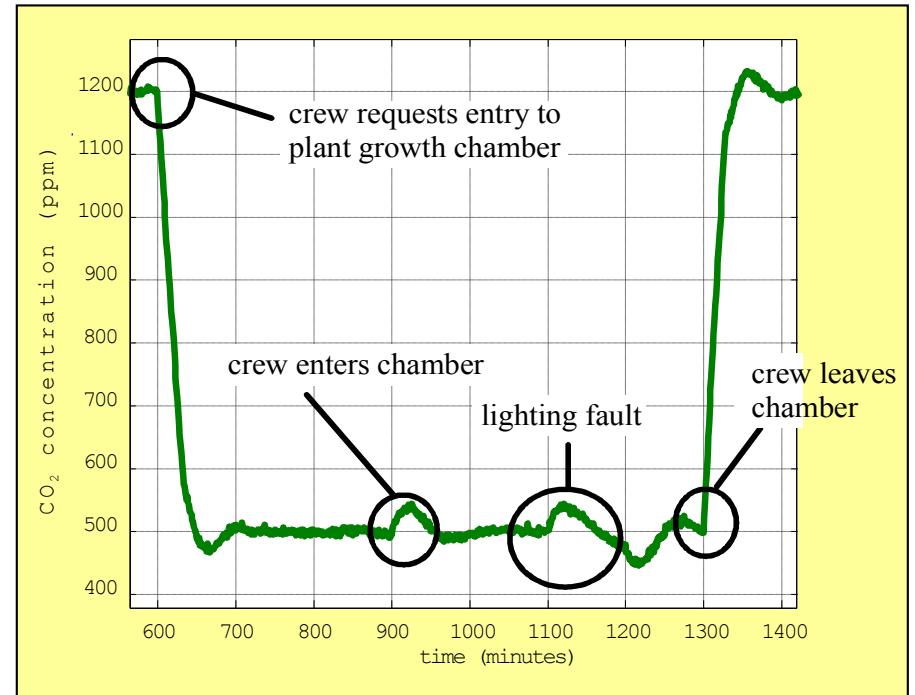
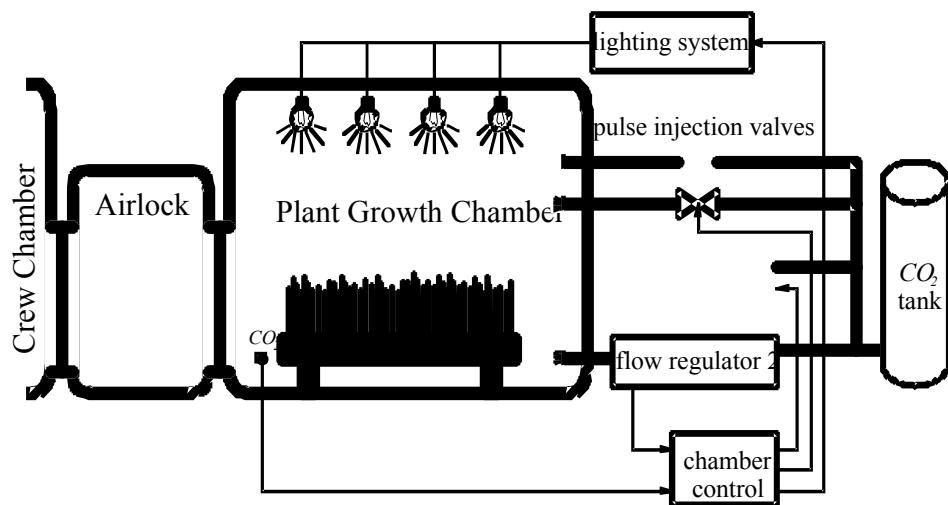
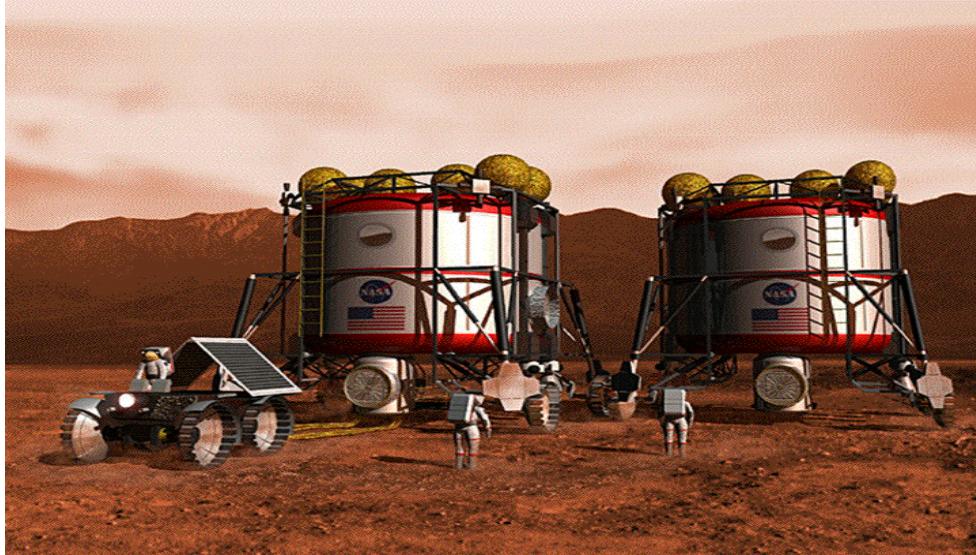
Approximate by enumerating most **likely states** or **trajectories**.

Best-first Tree Search at Each Time Step

- Path search through component mode transitions, while checking constraints.
 - E.g., Enumerate best using conflict-directed A*



Detecting Subtle Failures: Hybrid Mode Estimation

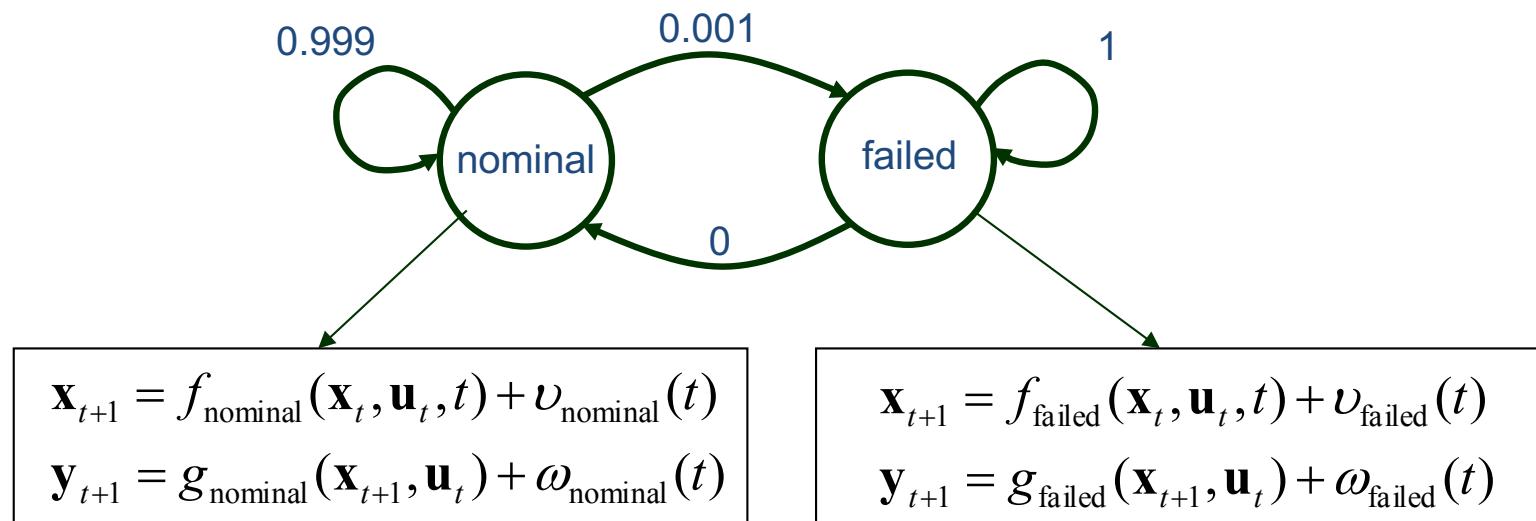


Hybrid:

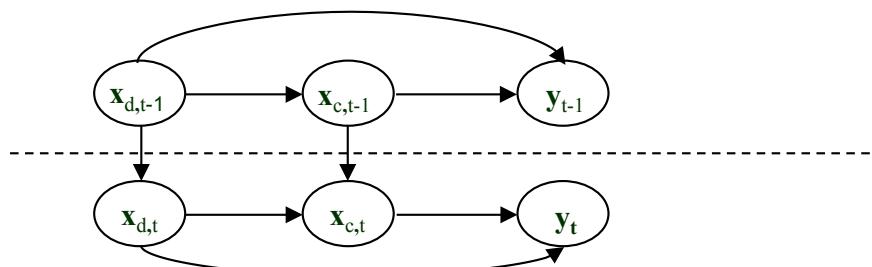
- Identify **discrete** failure mode.
- Interpret observations using **continuous** model.

Model – Hybrid HMM

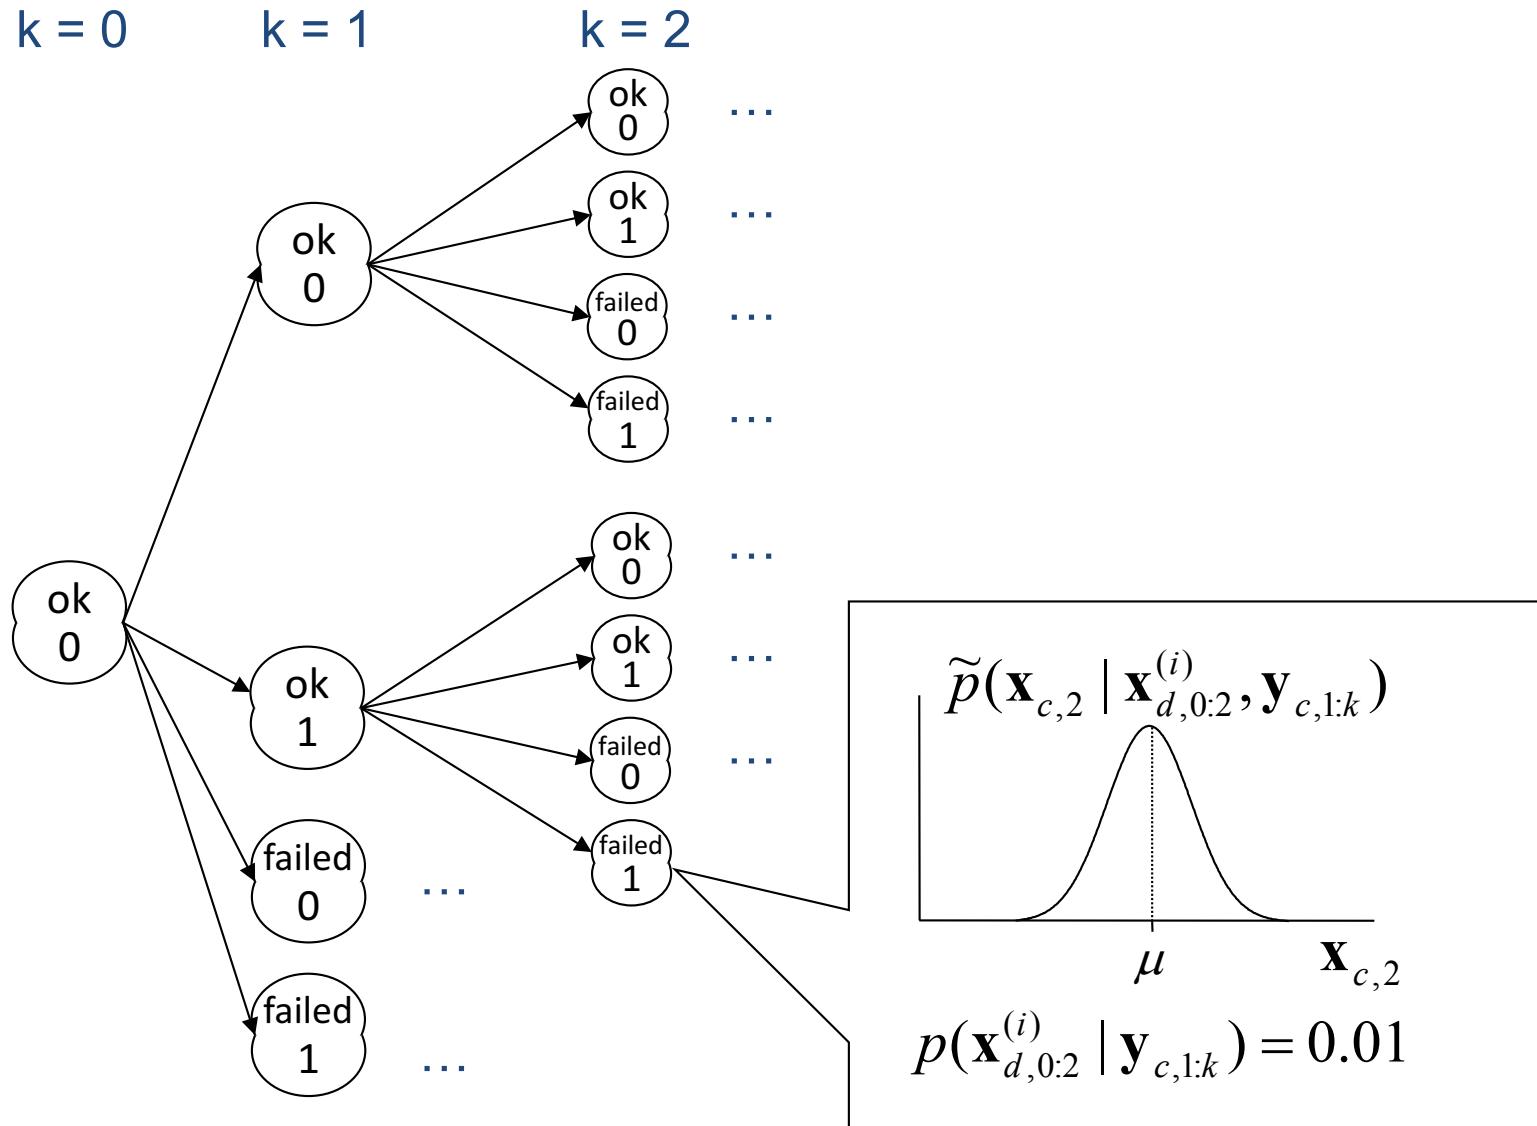
- **Discrete modes** and transitions between them



- **Continuous dynamics** corresponding to each mode



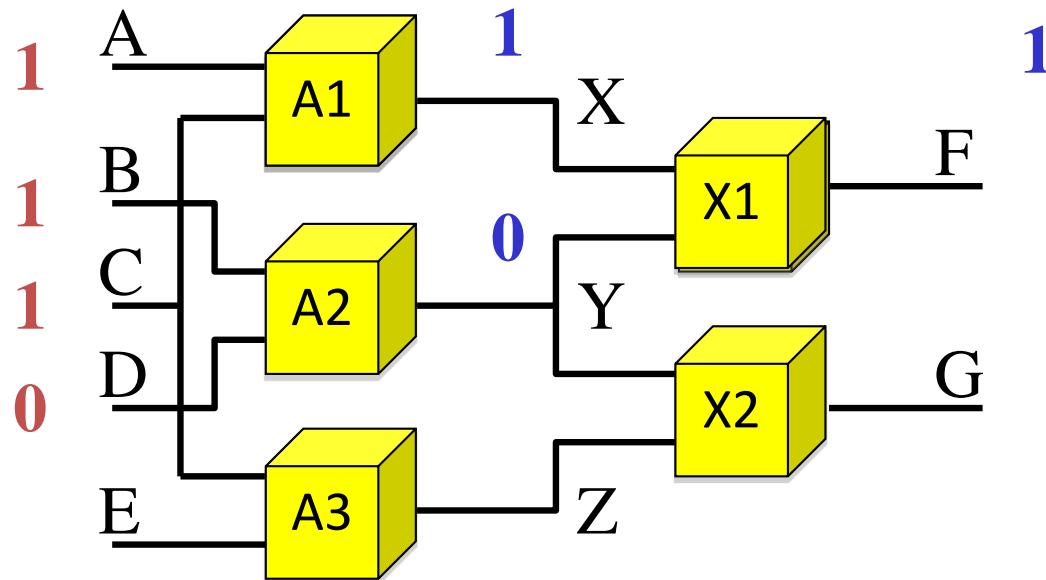
Kalman Filters Track Trajectories



Outline

- Programs that Monitor State
 - Sub-goal monitoring
 - Model-based diagnosis and mode estimation
 - Static mode estimation
 - Optimal CSPs and conflict-directed A*
- Programs that Self-Diagnose (opt)

Pedagogical Example: Polycell



OR	0	1
0	0	1
1	1	1

AND	0	1
0	0	0
1	0	1

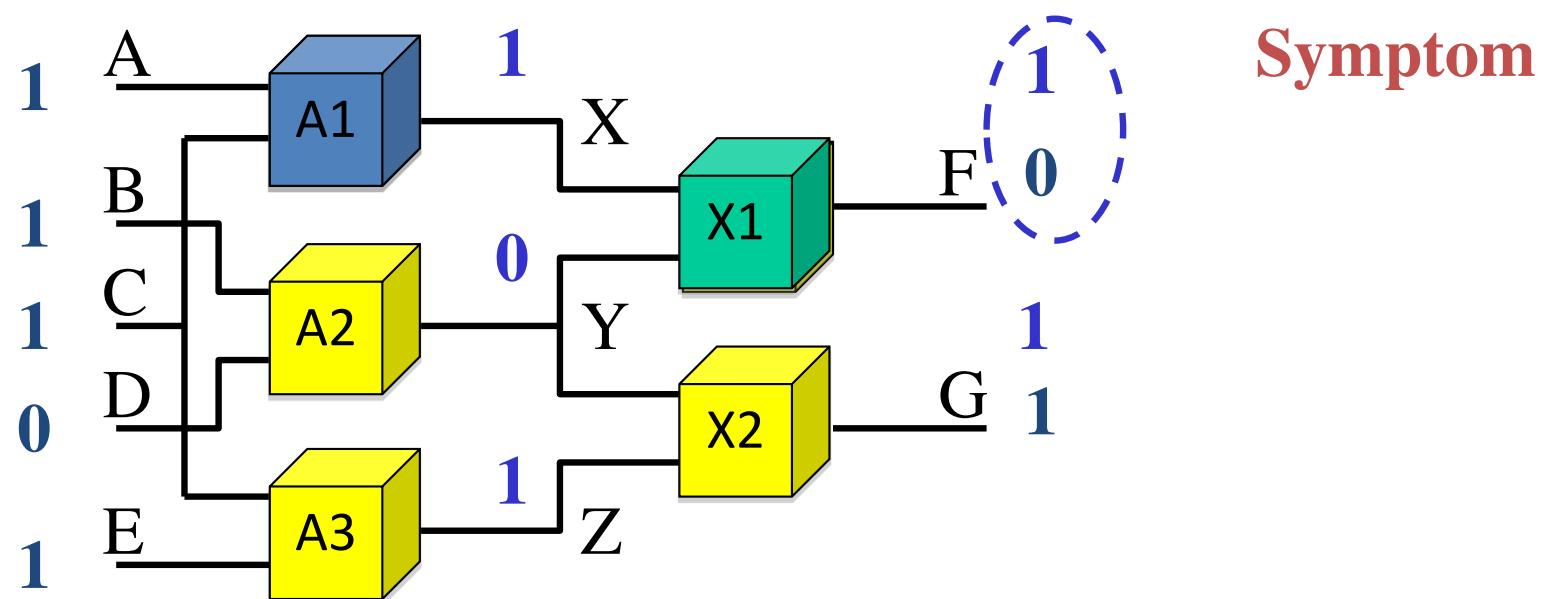
XOR	0	1
0	0	1
1	1	0

NOT	
0	1
1	0

Model-based Diagnosis

Input: **Observations** of a **system** with symptomatic behavior, and a **model Φ** of the system.

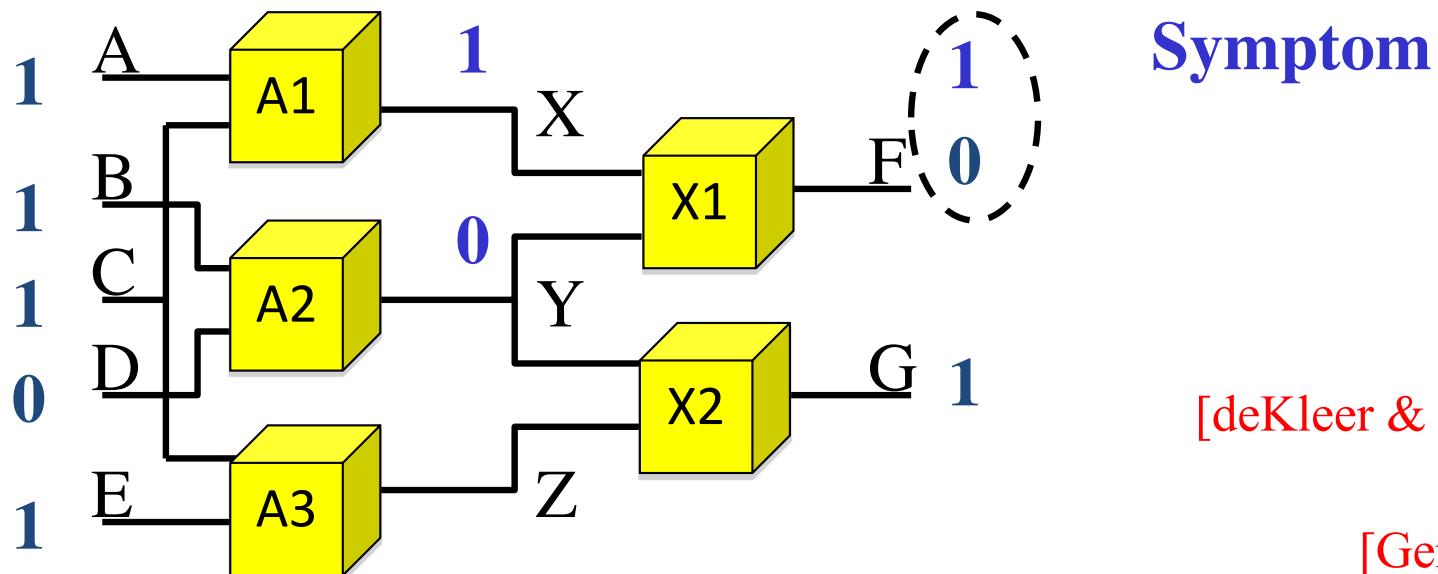
Output: **Diagnoses** that **account for the symptoms**.



How Should Diagnoses Account for Novel Symptoms?

Consistency-based Diagnosis: Given symptoms, find diagnoses that are **consistent** with symptoms.

Suspending Constraints: For **novel** faults, make **no presumption** about faulty component behavior.

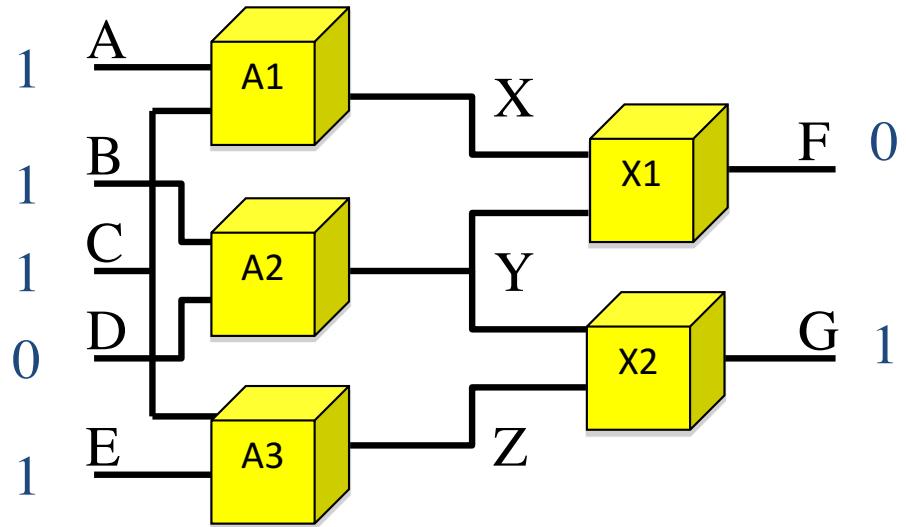


[deKleer & Brown, 83]
[Davis, 84]
[Genesereth, 84]

Multiple Faults: Identify all Combinations of Consistent “Unknown” Modes

And(I):

- I=G:
Out(i) = In1(i) AND In2(i)
- I=U:
True



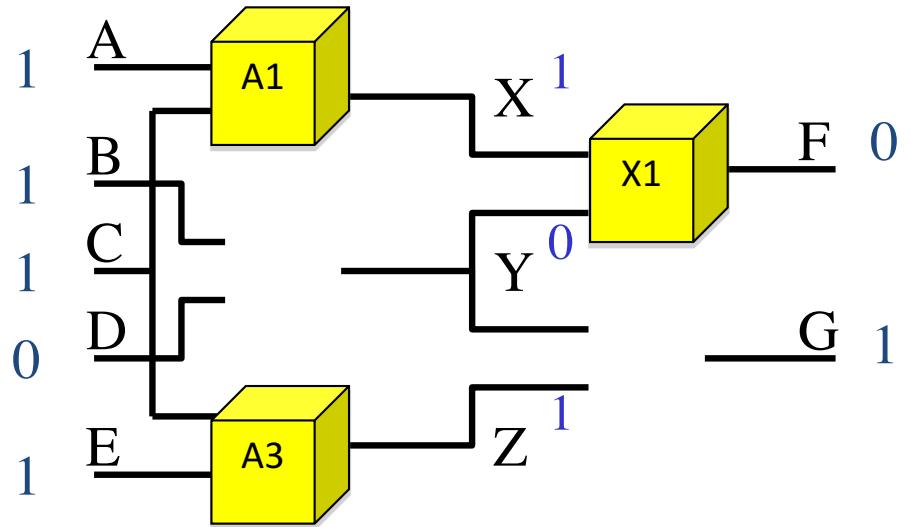
Candidate = {A1=G, A2=G, A3=G, X1=G, X2=G}

- Candidate: Assignment of **G** or **U** to each component.

Multiple Faults: Identify all Combinations of Consistent “Unknown” Modes

And(I):

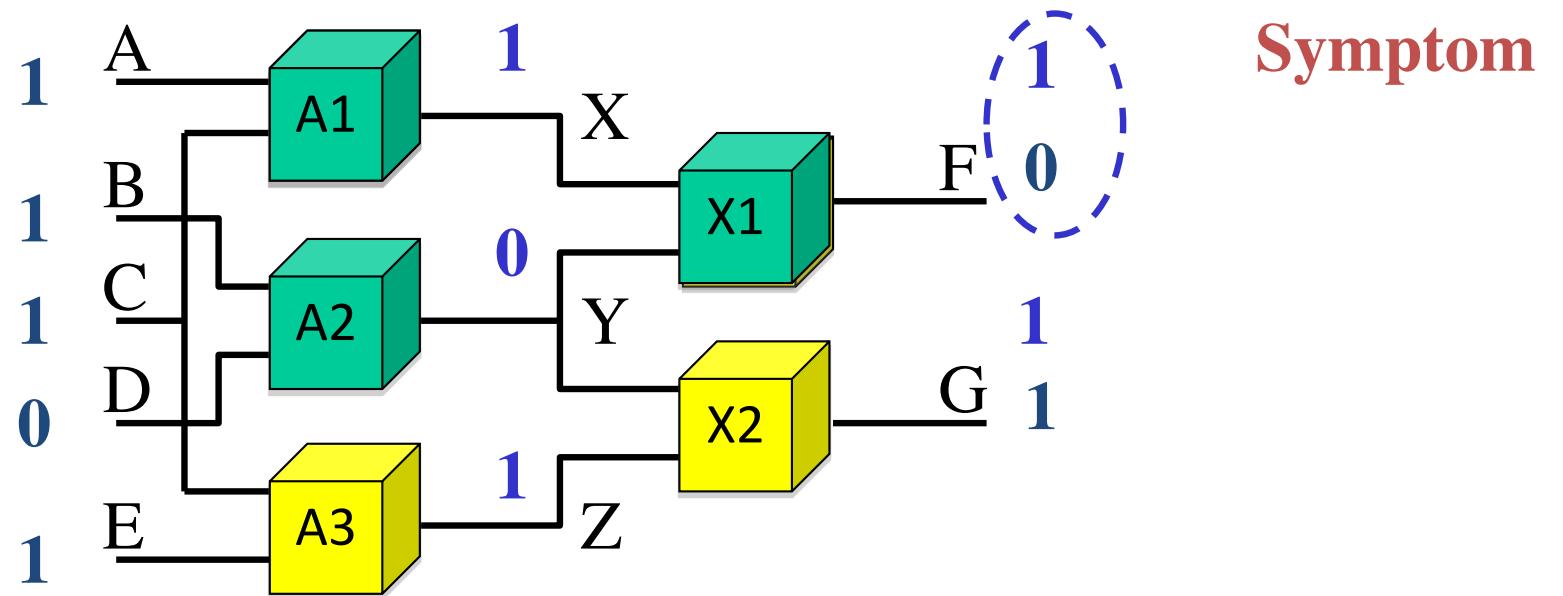
- I=G:
 $\text{Out}(i) = \text{In}_1(i) \text{ AND } \text{In}_2(i)$
- I=U:
True



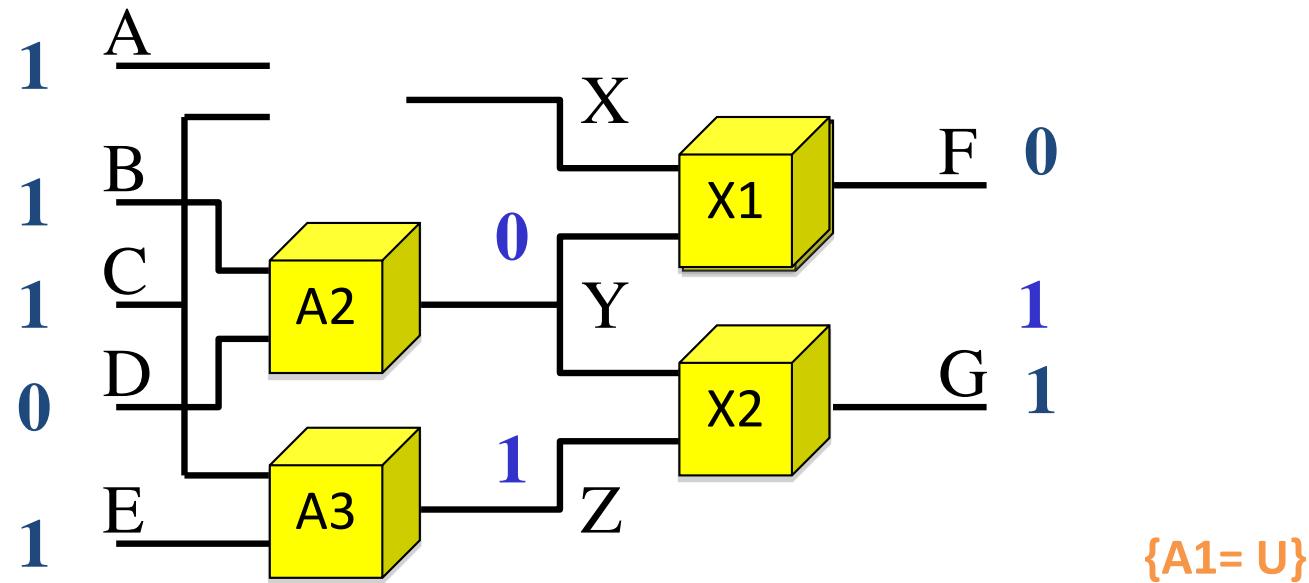
$$\text{Diagnosis} = \{A1=G, A2=U, A3=G, X1=G, X2=U\}$$

- Candidate:
Assignment of G or U to each component.
- Diagnosis:
Candidate **consistent** with **model** and **observations**.

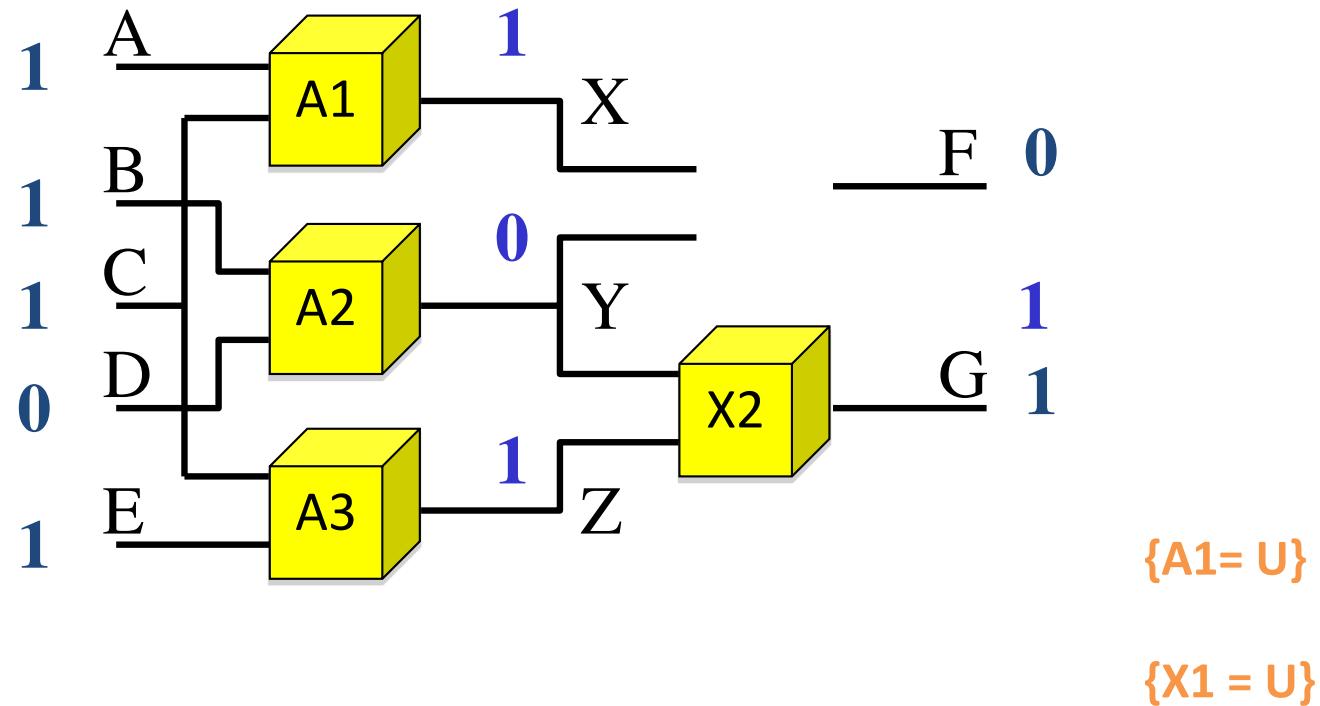
Consistency-based Diagnoses



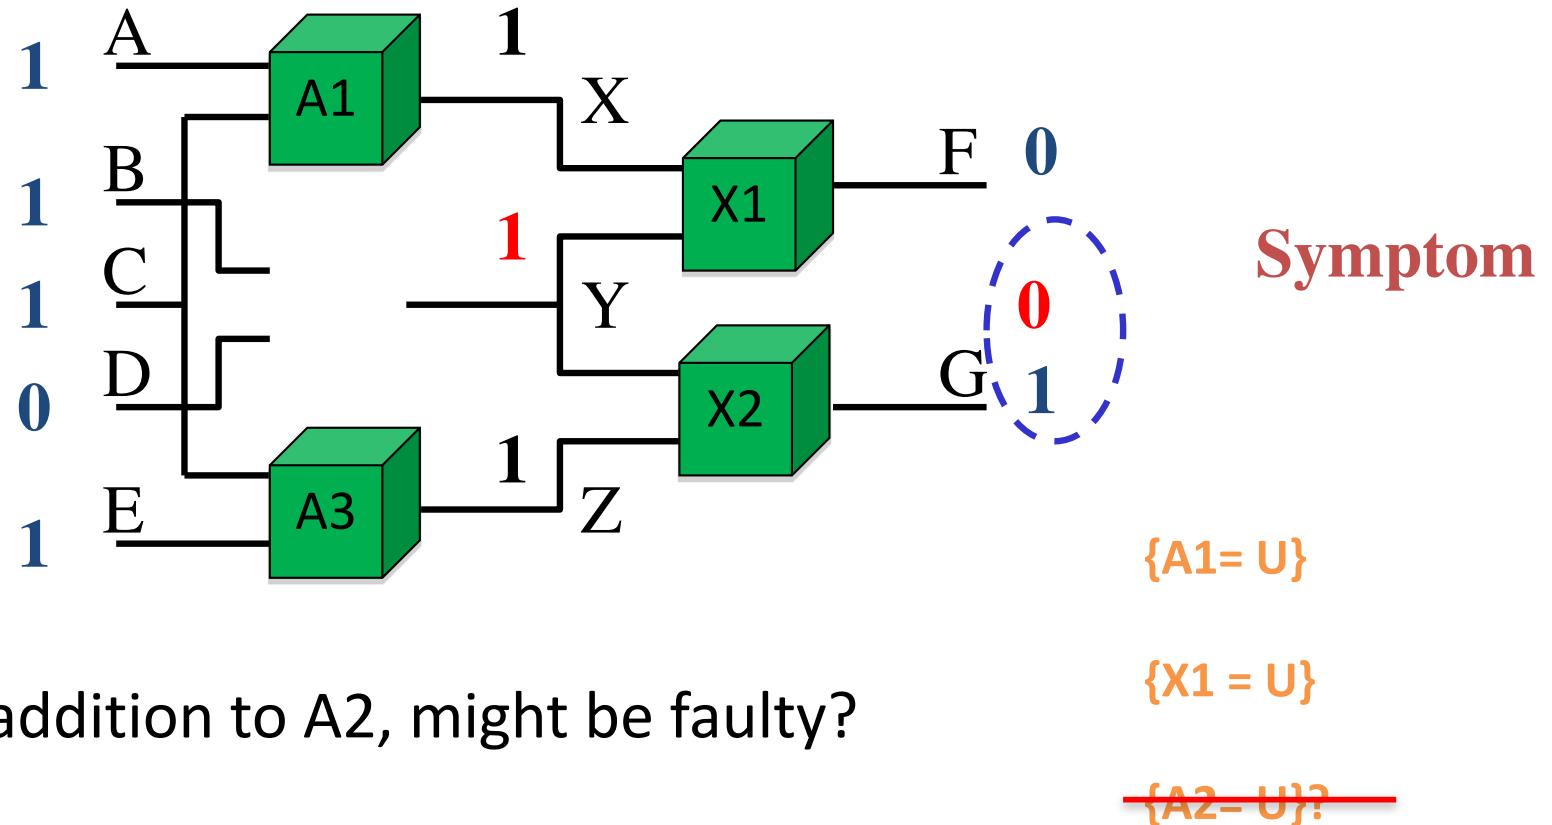
Consistency-based Diagnoses



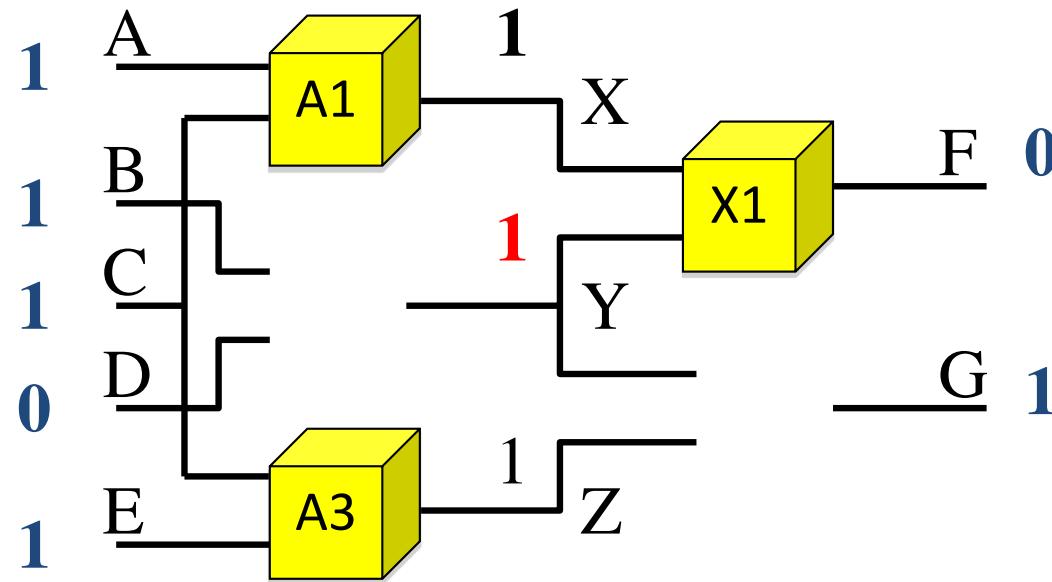
Consistency-based Diagnoses



Consistency-based Diagnoses



Consistency-based Diagnoses



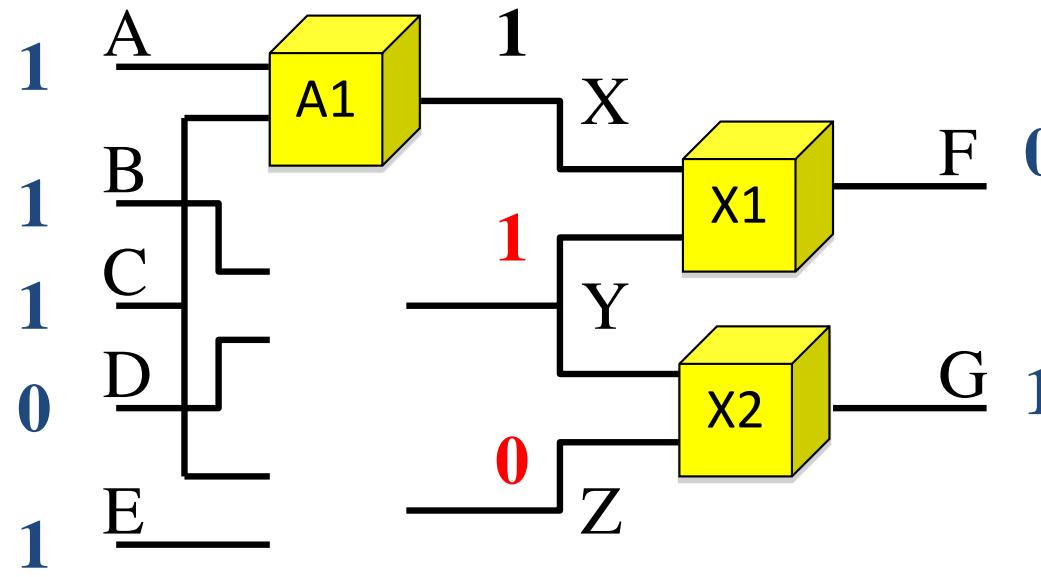
What, in addition to A2, might be faulty?

{A1 = U}

{X1 = U}

{A2 = U, X2 = U}

Consistency-based Diagnoses



{A1 = U}

{X1 = U}

{A2 = U, X2 = U}

{A2 = U, A3 = U}

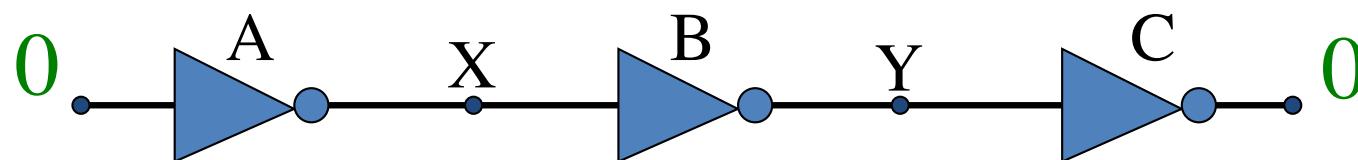
What, in addition to A2, might be faulty?

- These diagnoses are minimal.
- All extensions are also diagnoses.

⇒ Kernel Diagnoses

Incorporating Failure Modes: Mode Estimation

Sherlock
[de Kleer & Williams, IJCAI 89]



Inverter(I):

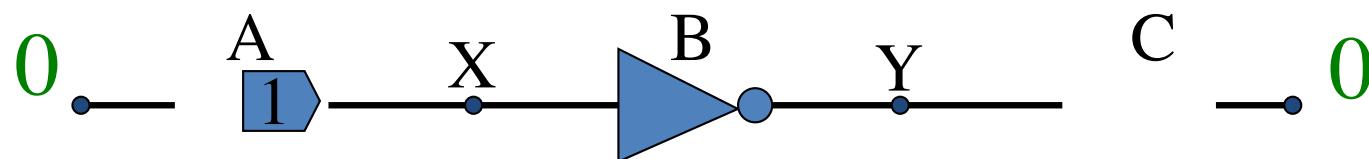
- $I=G$: $\text{Out}(i) = \text{not}(\text{In}(i))$
- $I=S1$: $\text{Out}(i) = 1$
- $I=S0$: $\text{Out}(i) = 0$
- $I=U$: True

- Isolates unknown.
- Explains.

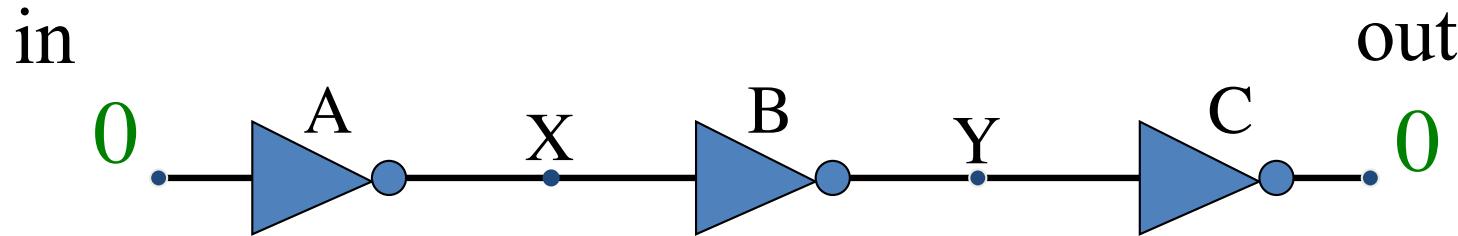
Good, Faulty and Unknown Modes

Example Mode Estimate

Sherlock
[de Kleer & Williams, IJCAI 89]



Diagnosis: [A=S1, B=G, C=U]



Diagnoses: (42 of 64 candidates)

Sherlock
[de Kleer & Williams, IJCAI 89]

Fully Explained Failures

- [A=G, B=G, **C=S0**]
- [A=G, **B=S1**, **C=S0**]
- [**A=S0**, B=G, C=G]

...

Partially Explained

- [A=G, **B=U**, **C=S0**]
- [**A=U**, **B=S1**, C=G]
- [**A=S0**, **B=U**, C=G]

...

Fault Isolated, But Unexplained

- [A=G, B=G, **C=U**]
- [A=G, **B=U**, C=G]
- [**A=U**, B=G, C=G]

Mode Estimation Problem

Given:

- Mode, State, Observable Variables:
- Model:

And(I):

G(I):

$$\text{Out}(i) = \text{In}1(i) \text{ AND } \text{In}2(i)$$

U(I):

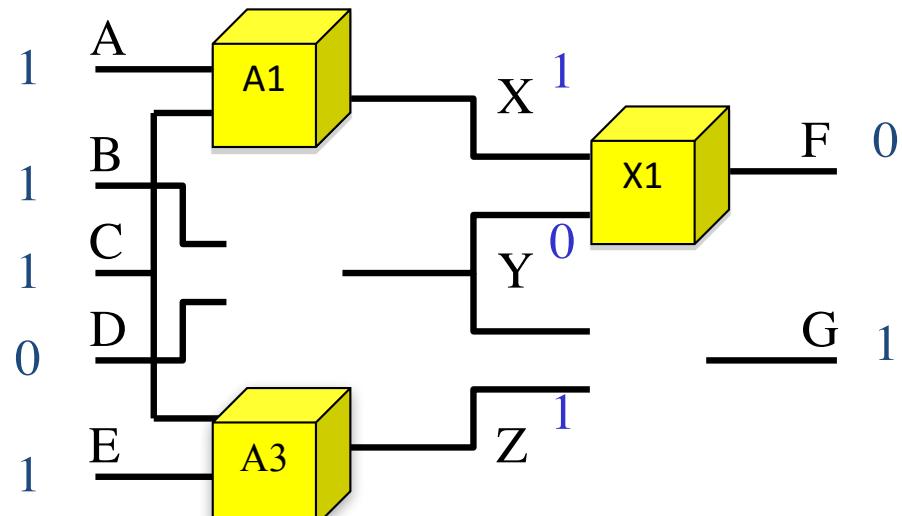
$X, Y, O \subseteq Y$

$$\Phi(X, Y) = \text{components} + \text{structure}$$

- All behaviors associated with modes.
- All components have “unknown mode” U, which enables no constraint.

Sherlock

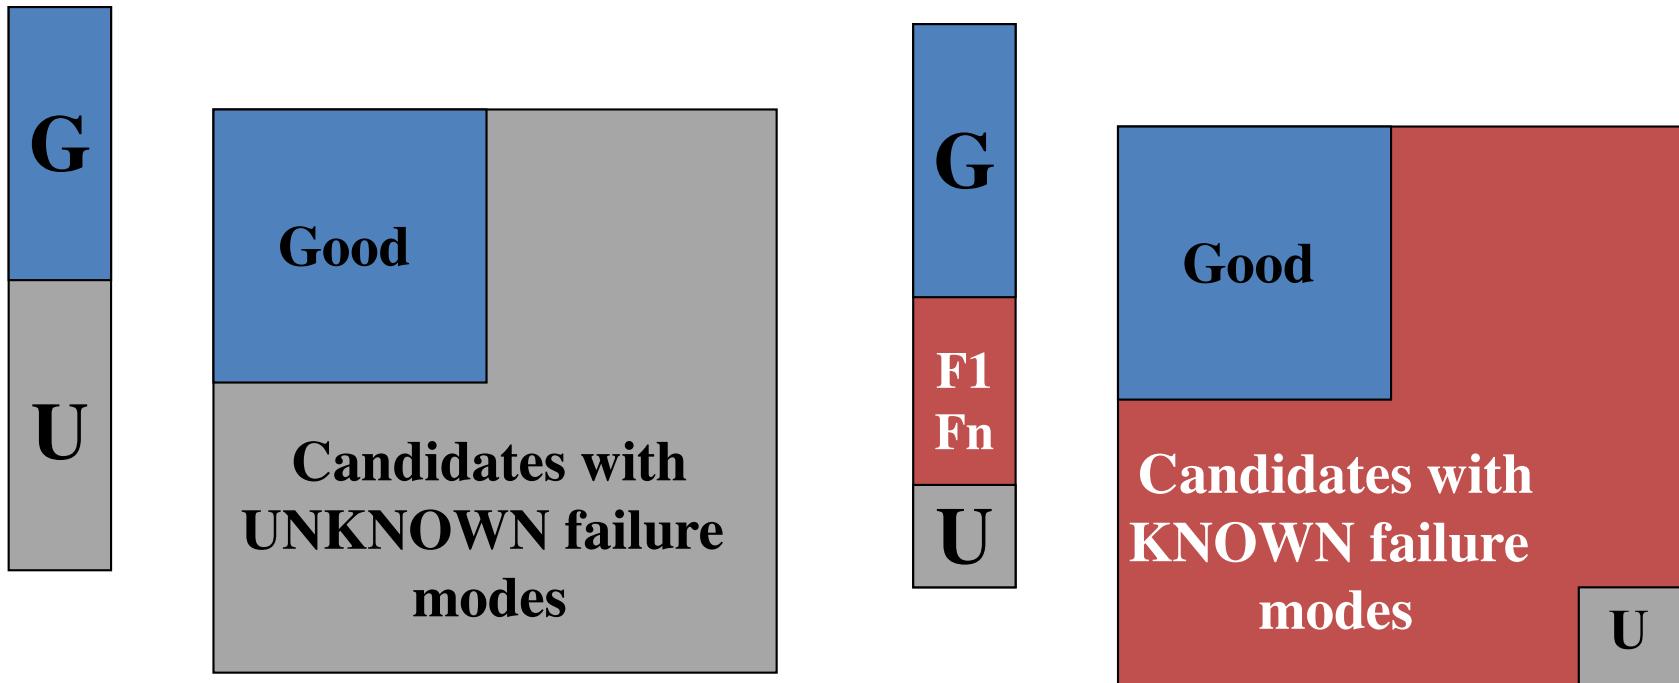
[de Kleer & Williams, IJCAI 89]



Return: All diagnoses

$$D_{\Phi, \text{obs}} \equiv \{X \in D_X \mid \exists Y \in D_Y \text{ st Obs} \wedge \Phi(X, Y)\}$$

Problem: Due to the **unknown mode**, there tends to be an **exponential number** of mode estimates.



Ideas:

1. Compute probability of each mode estimate.
2. Enumerate most likely mode estimates X_i , based on probability.

Prefix (k) (Sort $\{X_i\}$ by decreasing $P(X_i | O)$)

Probabilistic, Mode Estimation

Inputs:

- Mode X, State Y, and Observation (O subset Y) variables, with finite domains.
- Model $\Phi(X;Y)$.
- Observations $O = o$.
- **Prior distribution $P(X)$.**
- **Observation function $P(o | X, \Phi(X;Y)) = P(o | X)$.**

Outputs:

- Exact: $P(X | o)$ Posterior, given observations.
- Approximate: k most likely mode estimates X_i ,
[Prefix (k) (Sort $\{X_i\}$ by decreasing $P(X_i | o)$)].

A Simple Approximation for Probabilistic, Mode Estimation

$$P(X|o) = \frac{P(o|X)P(X)}{P(o)} = \alpha P(o|X)P(X)$$

$$P(o) = \frac{1}{\alpha} = \sum_{c_i \in C} P(o|c_i)P(c_i) \quad \text{for candidates } C \text{ consistent with } o$$

1. Assume modes are *a priori* independent:

$$P(X) = \prod_{X_i \in X} P(X_i)$$

2. Assume consistent observations are equally likely for a given mode assignment:

$$P(o|X) = \begin{cases} 0 & \text{if } \Phi \wedge o \wedge X \text{ is inconsistent} \\ 1/n & \text{else } n = |\{o_i \mid \Phi \wedge o_i \wedge X \text{ is consistent}\}| \end{cases}$$

A Simple Approximation for Probabilistic, Mode Estimation

$$P(X|o) = \frac{P(o|X)P(X)}{P(o)} = \alpha P(o|X)P(X)$$

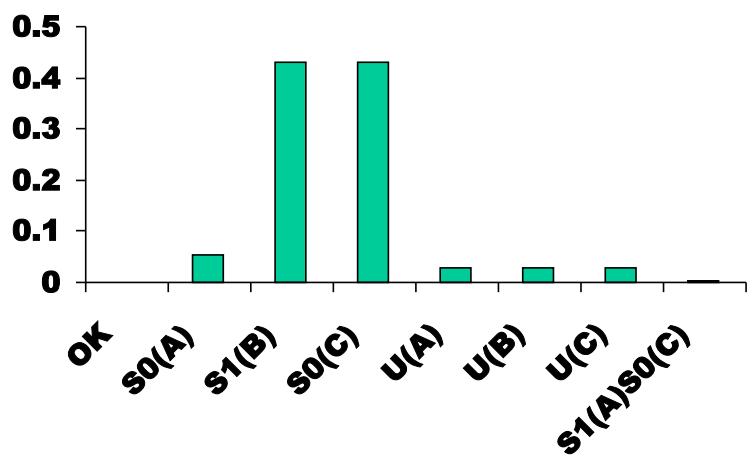
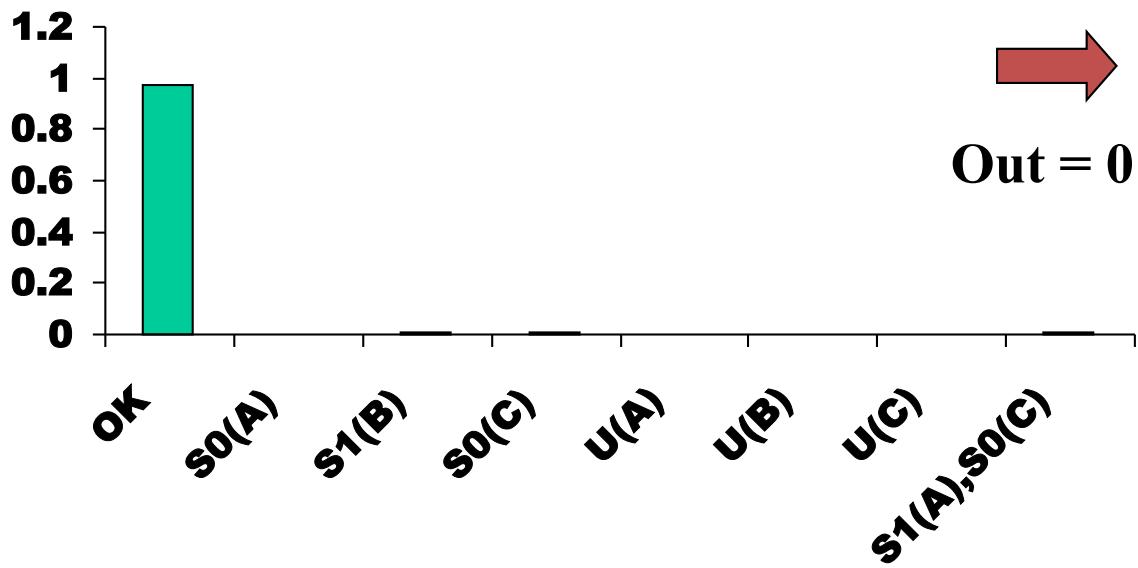
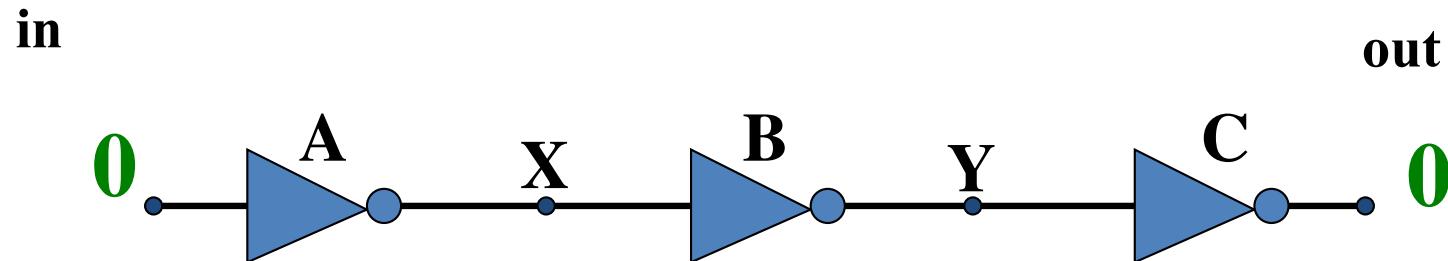
$$P(o) = \frac{1}{\alpha} = \sum_{c_i \in C} P(o|c_i)P(c_i) \quad \text{for candidates } C \text{ consistent with } o$$

1. Assume modes are *a priori* independent:

$$P(X) = \prod_{X_i \in X} P(X_i)$$

2. Assume consistent *interpretations* are equally likely, given observations, model and mode assignment.

⇒ Model Counting: Use exhaustive DPLL to count models.



Top 6 of 64 = 98.6% of P

Leading mode estimates before and after output observed.

Outline

- Programs that Monitor State
 - Sub-goal monitoring
 - Model-based diagnosis and mode estimation
 - Static mode estimation
 - Optimal CSPs and conflict-directed A*
- Programs that Self-Diagnose (opt)

Optimal CSP

Mode Estimation: k most likely mode estimates X_i ,

Prefix (k) (Sort $\{X_i\}$ by decreasing $P(X_i | o)$)



Given OCSP= $\langle X, f, \text{CSP} \rangle$

- X are *decision variables*, with finite domain D_X .
- $f: D_X \rightarrow \mathbb{R}$ is a *utility function*.
- CSP is a set of *constraints* over variables $\langle X; Y \rangle$.

Find leading $\arg \max f(X)$

$$X \in D_X$$

$$\text{s.t. } \exists Y \in D_Y . C(X, Y)$$

Example: Encode Constraints in Propositional State Logic

$$Y \in \{1, 0\}$$

$$\begin{aligned} Y=1 \vee Y=0 \\ \neg[Y=1 \wedge Y=0] \end{aligned}$$

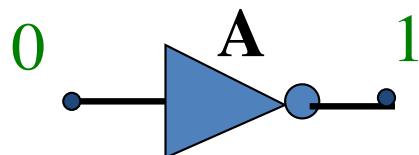
$$I \in \{G, S1, S0, U\}$$

...

Inverter(I):

- $I=G$: $\text{Out}(I) = \text{NOT In}(I)$
- $I=S1$: $\text{Out}(I) = 1$
- $I=S0$: $\text{Out}(I) = 0$

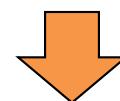
$$\begin{aligned} I=G &\rightarrow [\text{In}(I)=1 \text{ iff } \text{Out}(I)=0] \\ I=S1 &\rightarrow \text{Out}(I)=1 \\ I=S0 &\rightarrow \text{Out}(I)=0 \end{aligned}$$



Alternative constraints:

- Simple temporal networks
- Linear programs
- Global constraints

$$I=G \rightarrow [\text{In}(I)=1 \text{ iff } \text{Out}(I)=0]$$



$$\begin{aligned} [\neg(I=G) \vee \neg(\text{In}(I)=1) \vee \text{Out}(I)=0] \wedge \\ [\neg(I=G) \vee \neg(\text{Out}(I)=0) \vee \text{In}(I)=1] \end{aligned}$$

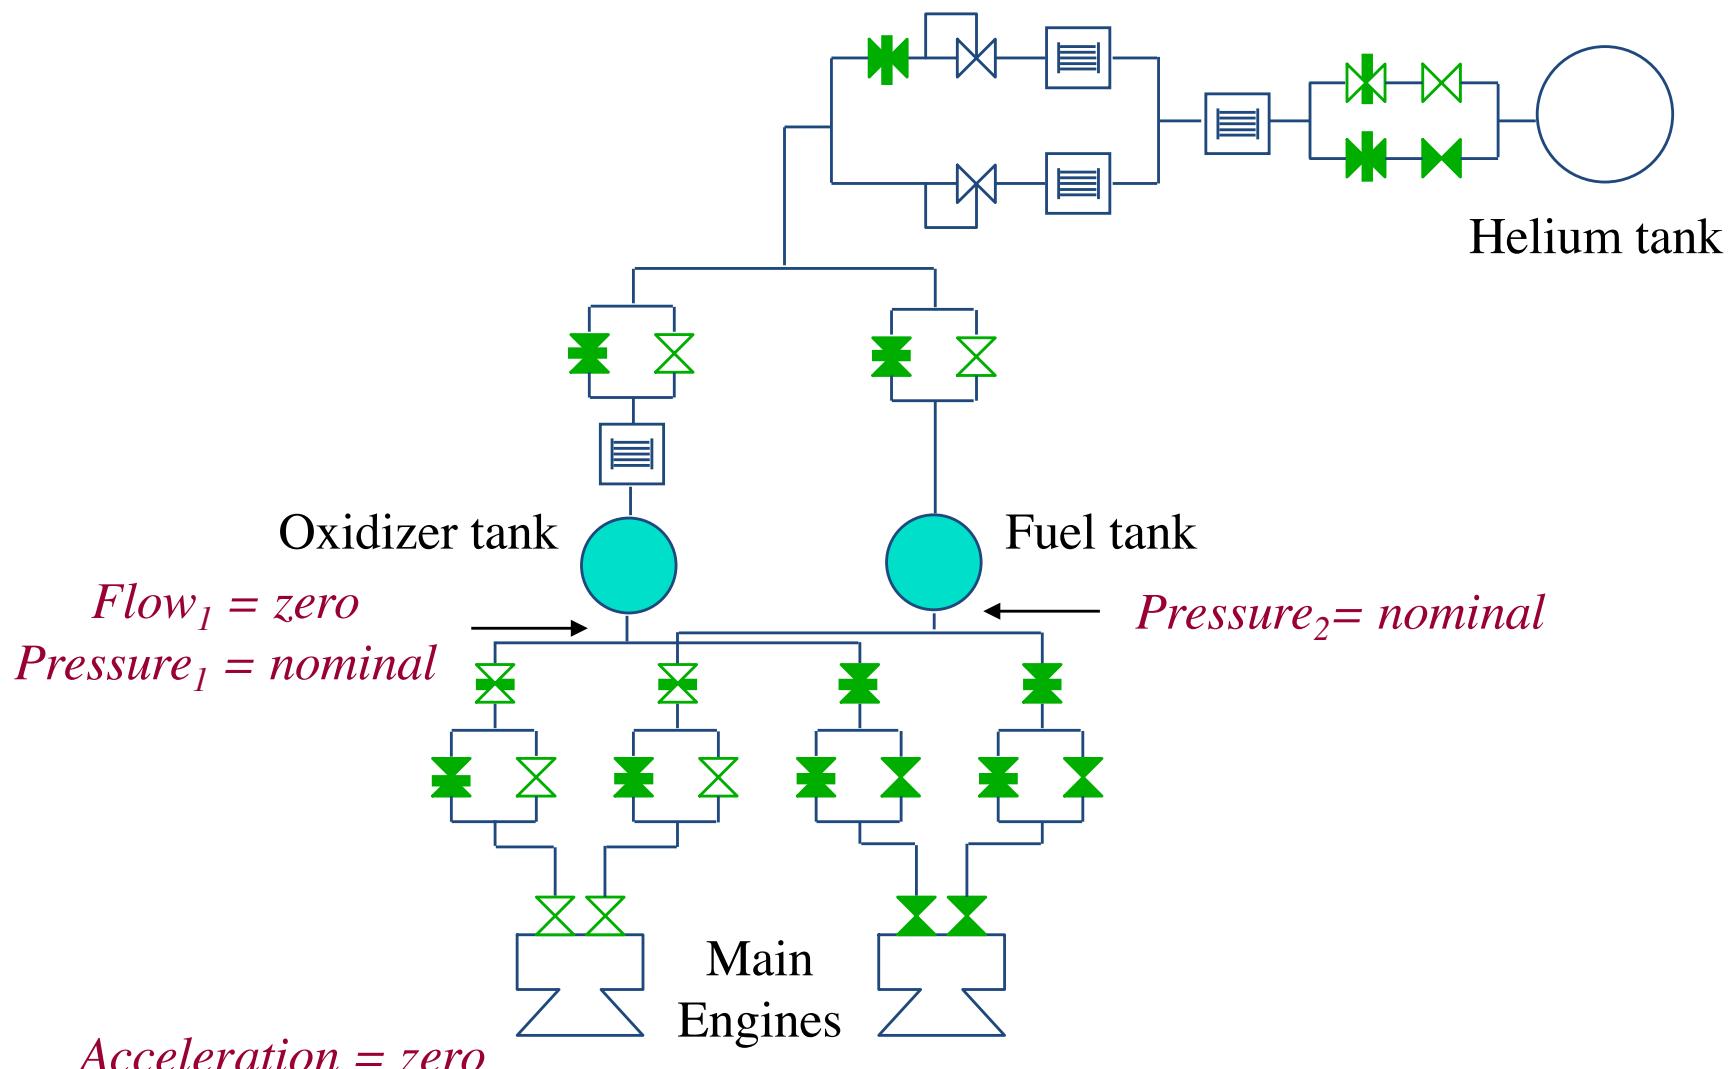
Model-based Diagnosis as Conflict-directed Best First Search

When you have eliminated the impossible,
whatever remains, however improbable,
must be the truth.

- Sherlock Holmes. The Sign of the Four.

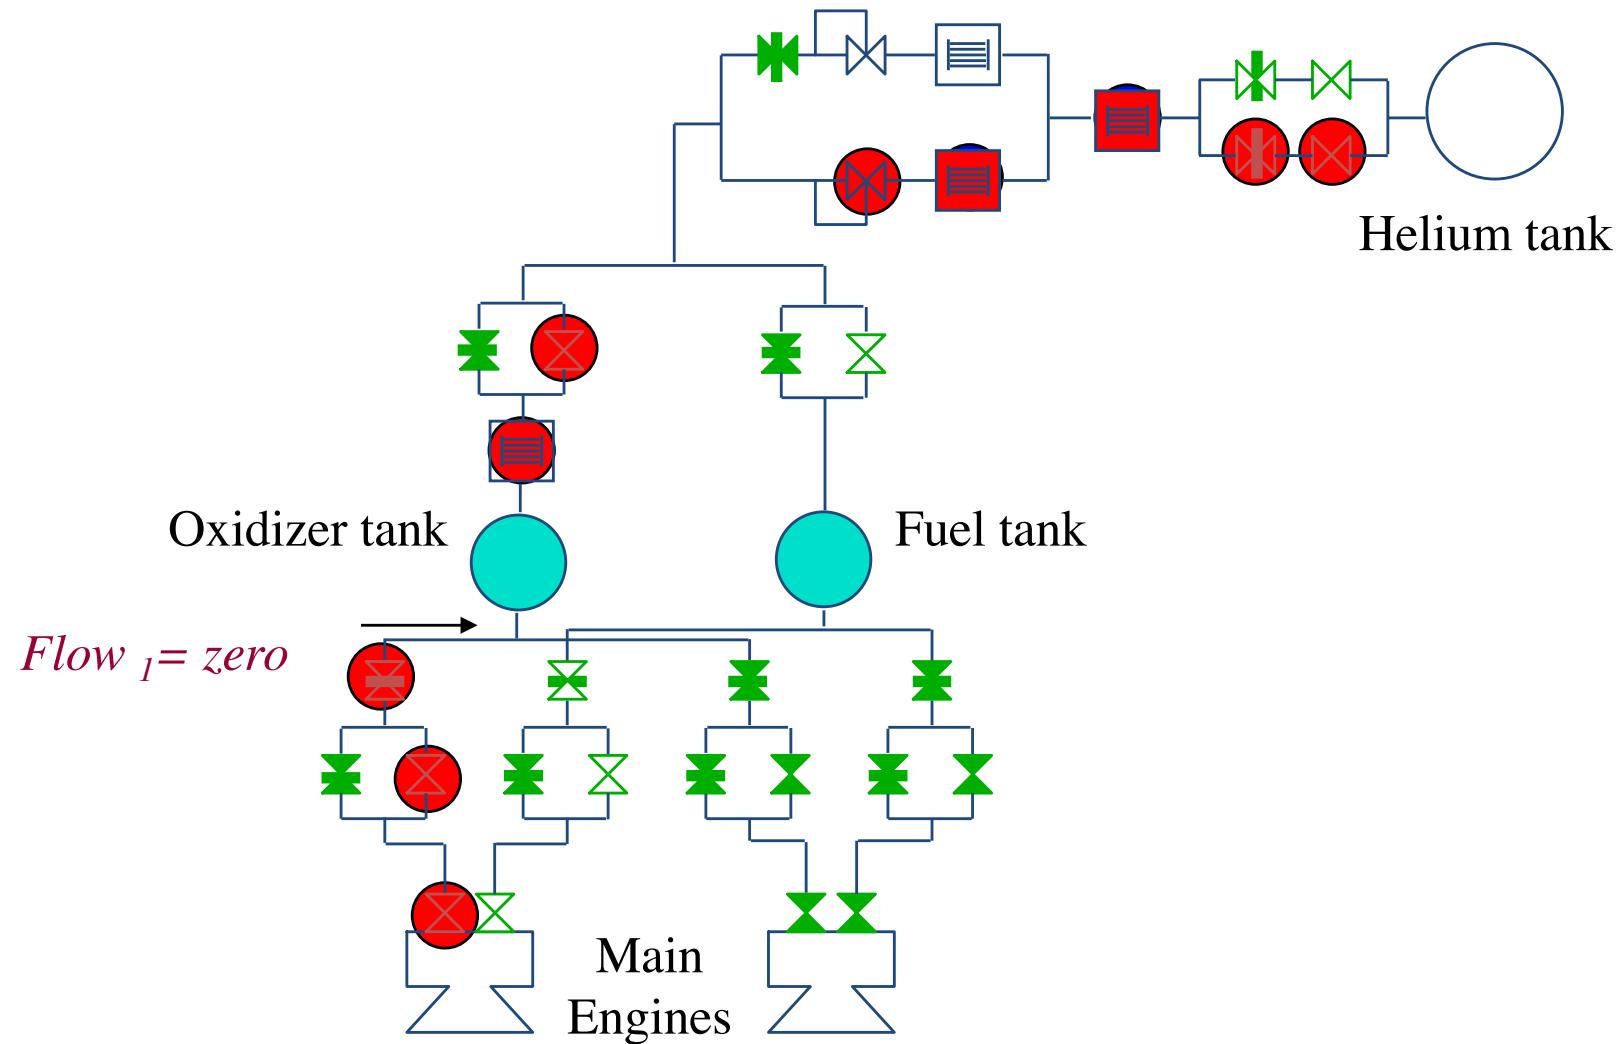
- 
1. Generate most likely candidate.
 2. Test candidate.
 3. If Inconsistent, learn reason for inconsistency
(a **conflict**).
 4. Use **conflicts** to leap over similarly infeasible options
to the next best candidate.

Compare Most Likely Candidate to Observations



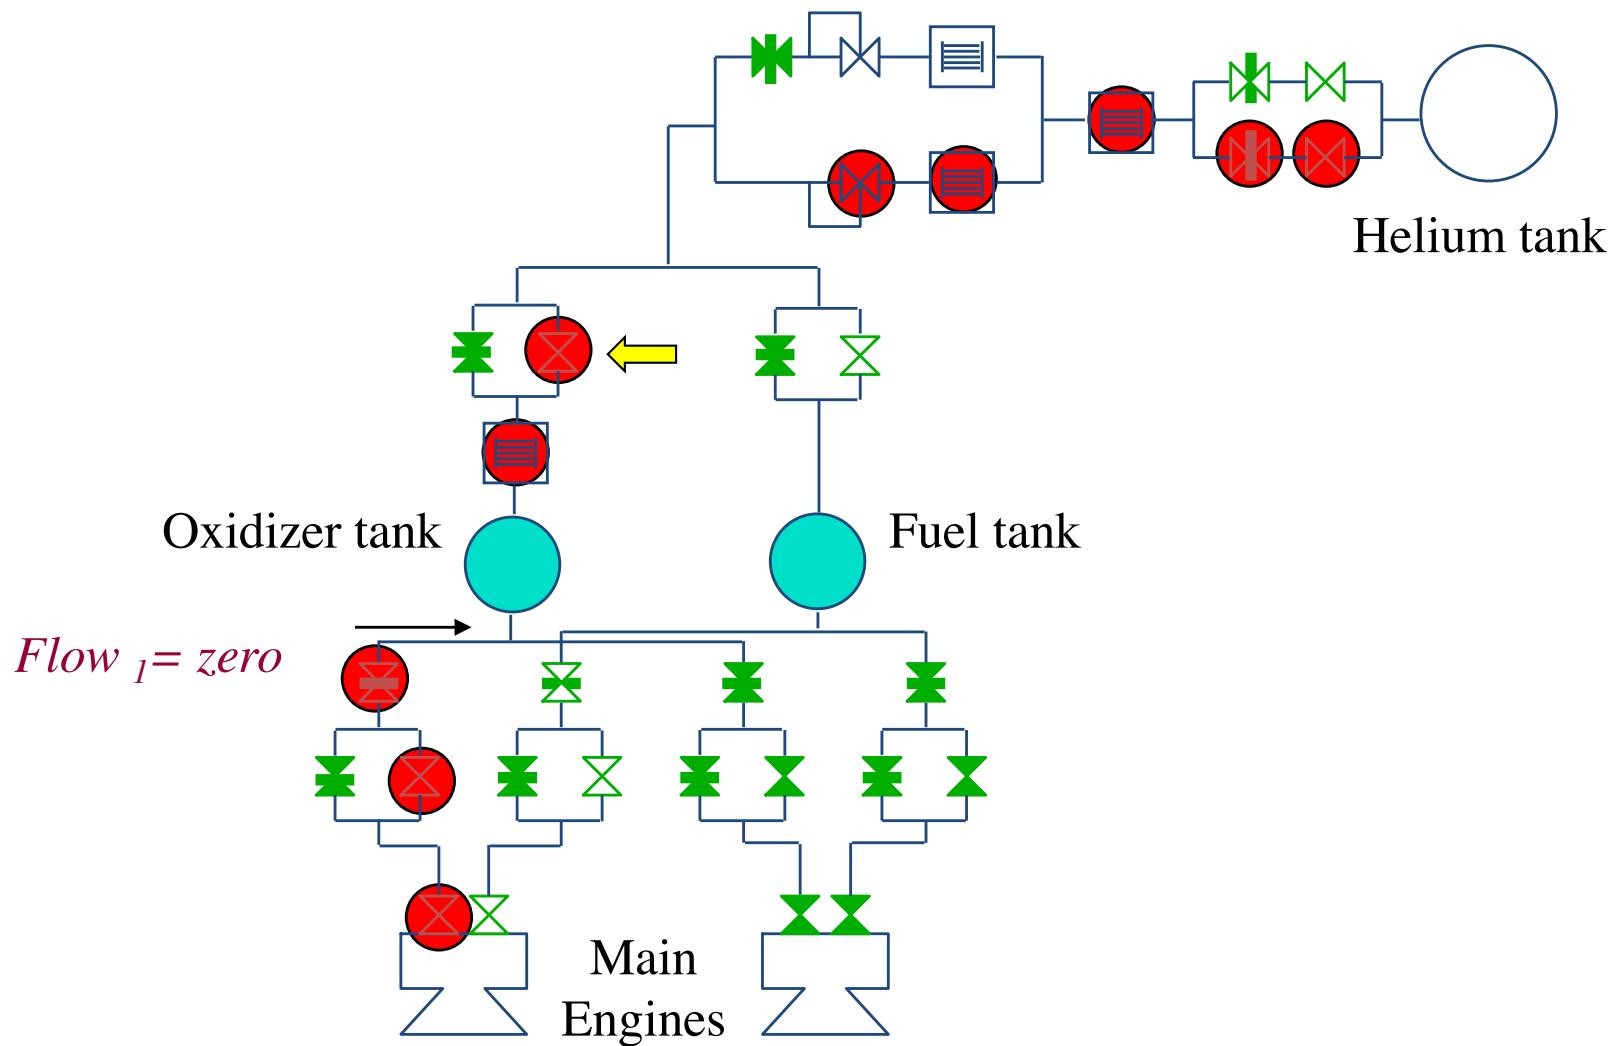
It is most likely that all components are okay.

Isolate Conflicting Information



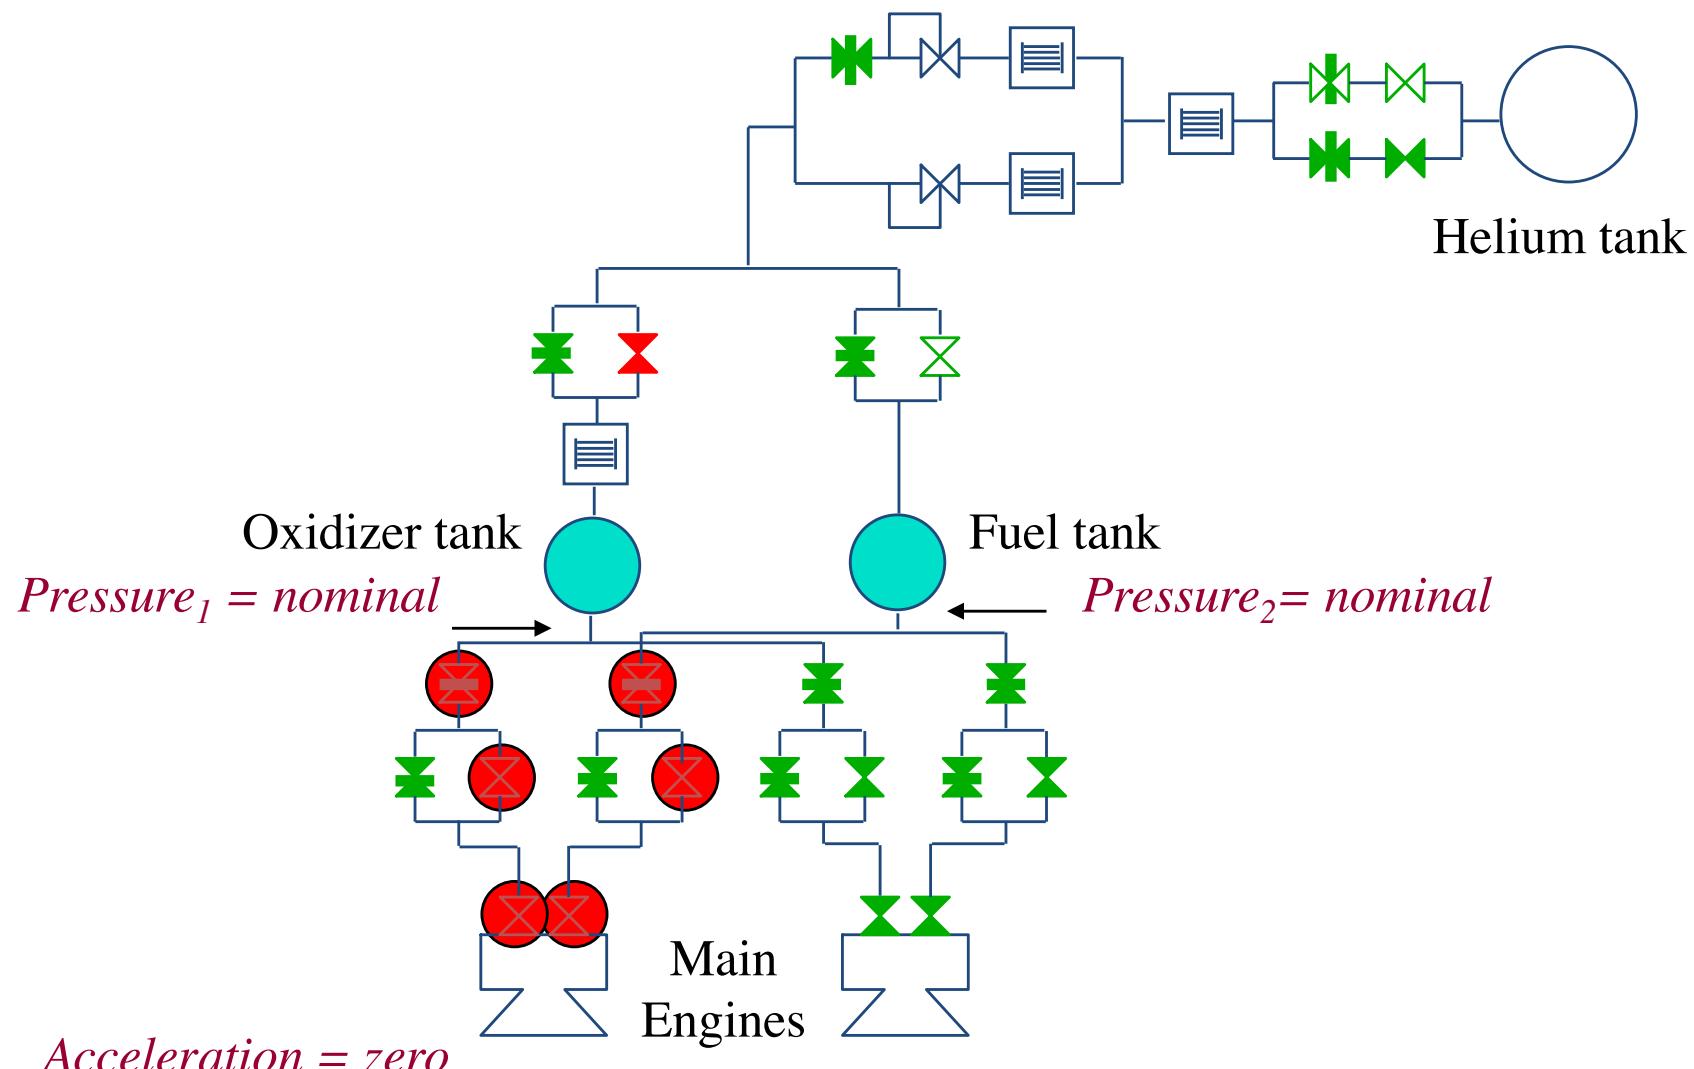
The red component modes *conflict* with the model and observations.

Leap to the Next Most Likely Candidate that Resolves the Conflict



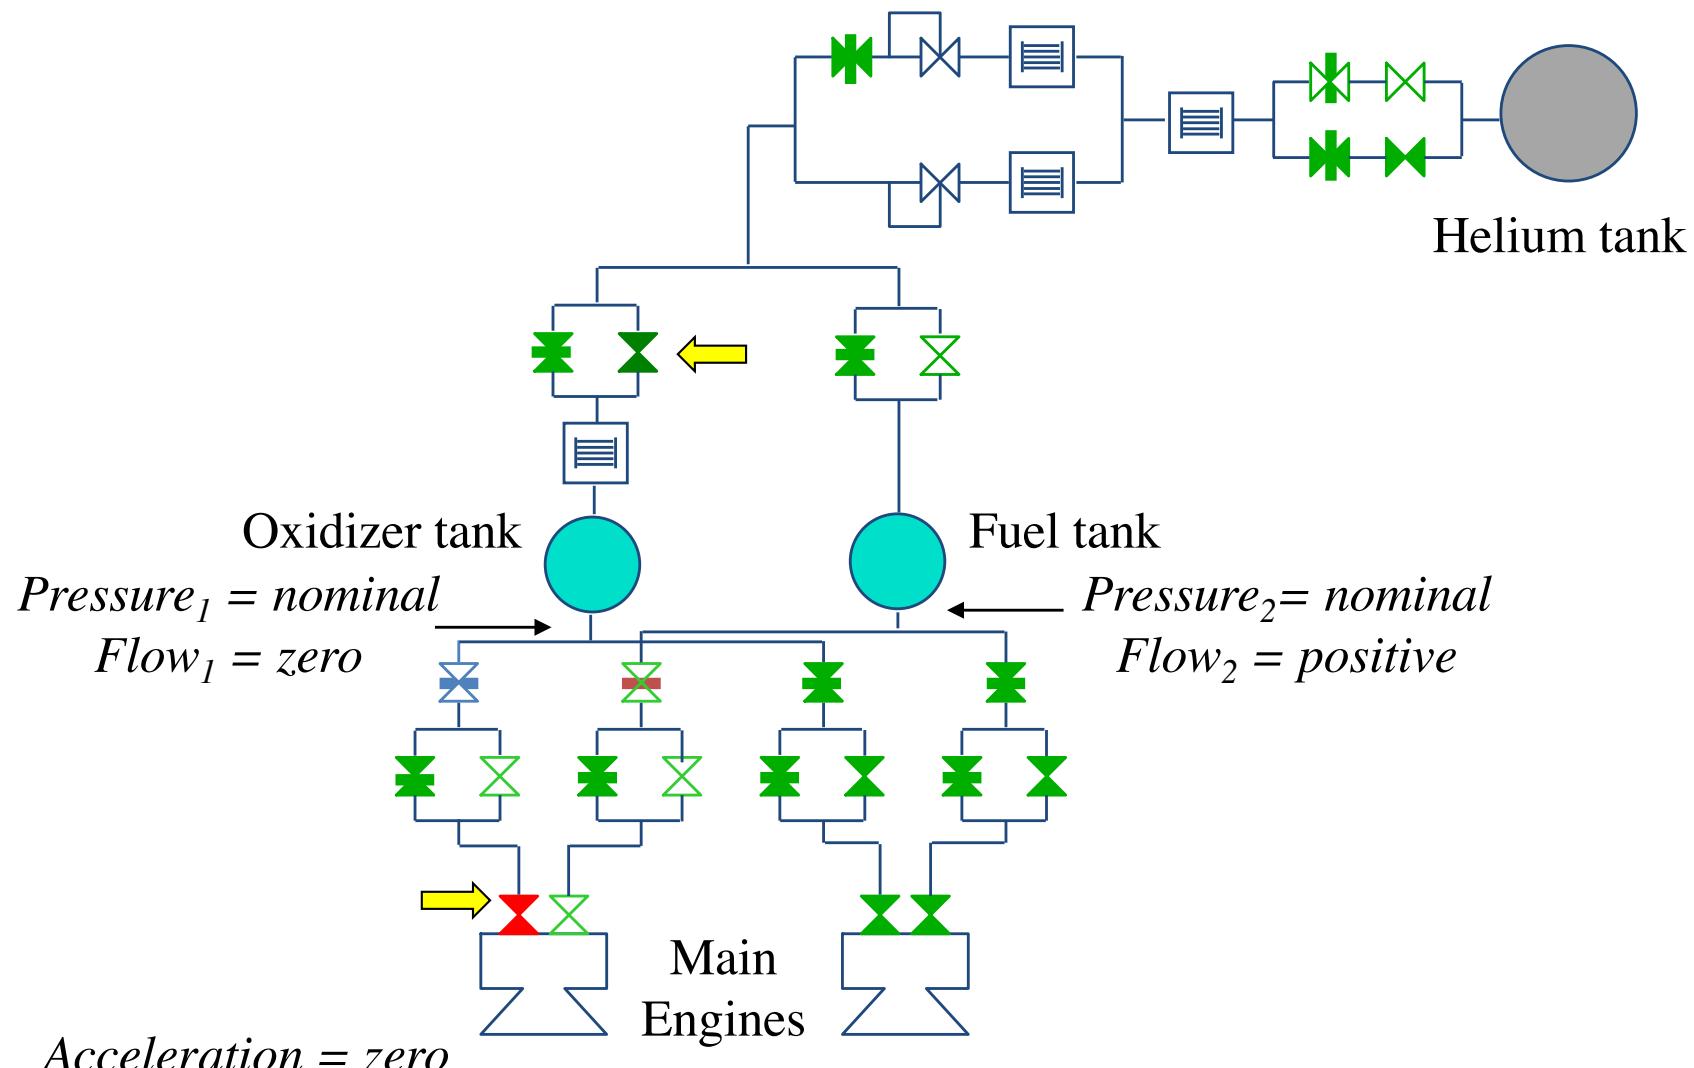
The next candidate **must remove** the **conflict**.

New Candidate Exposes Additional Conflicts



Another **conflict**, try **removing both**.

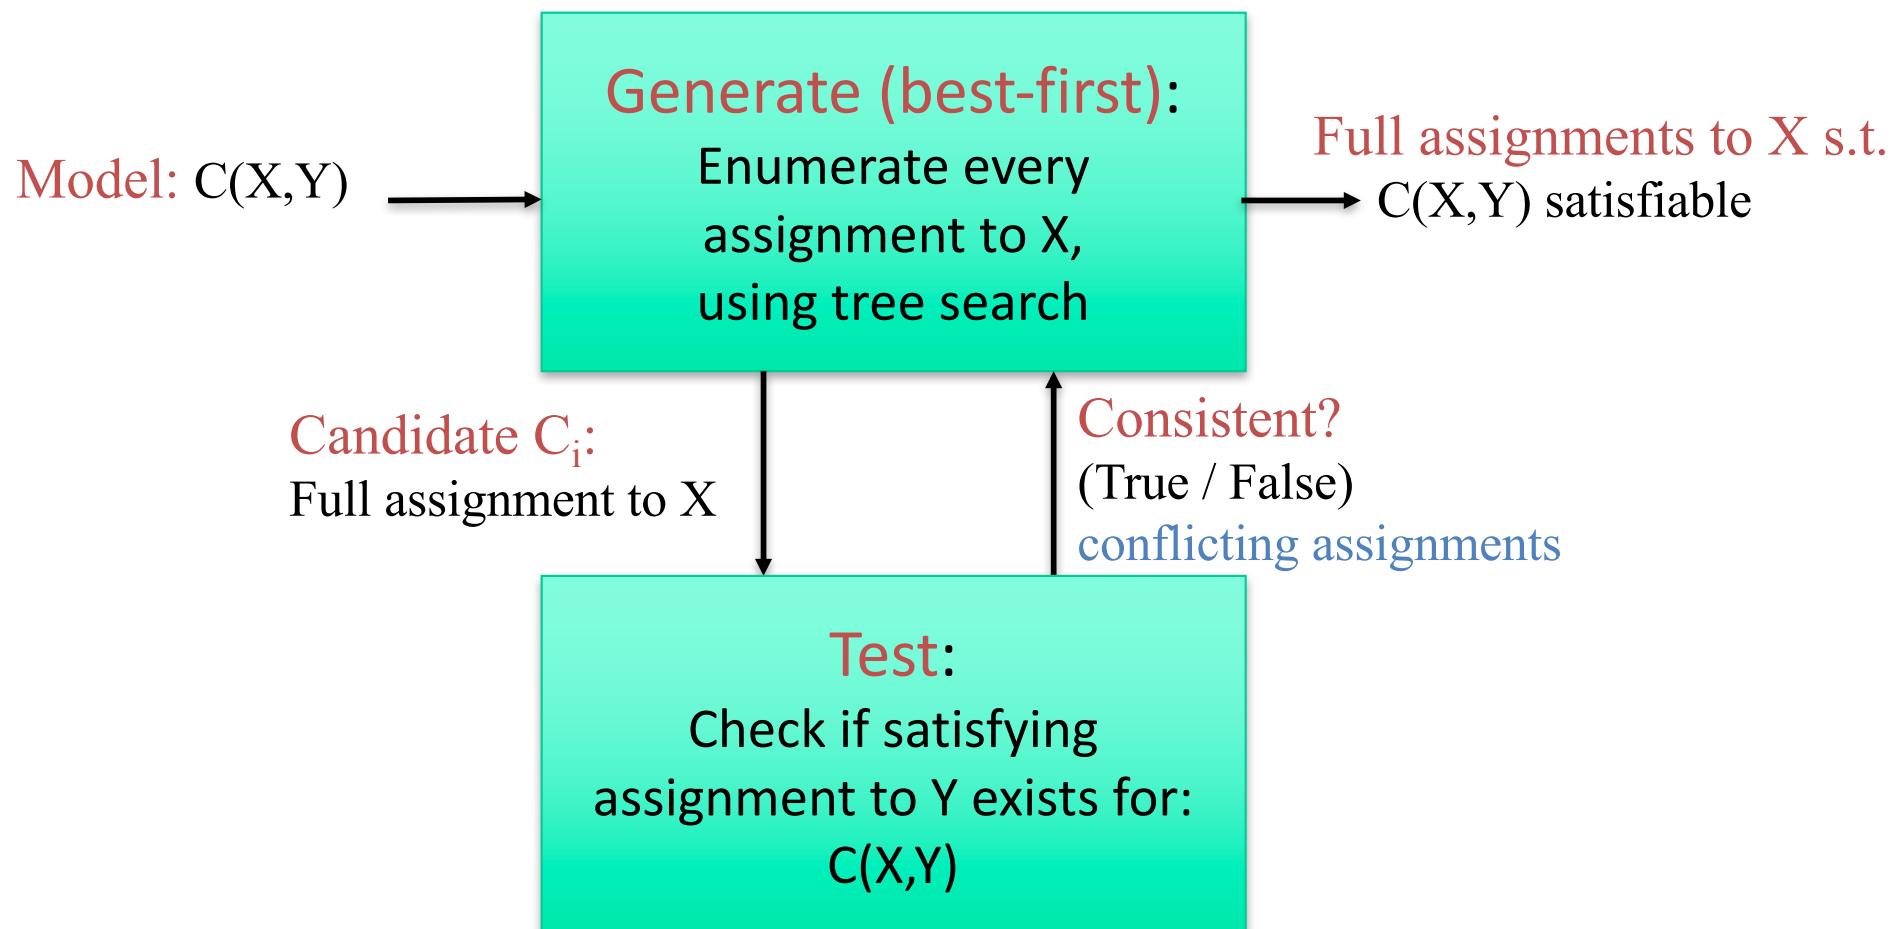
Final Candidate Resolves all Conflicts



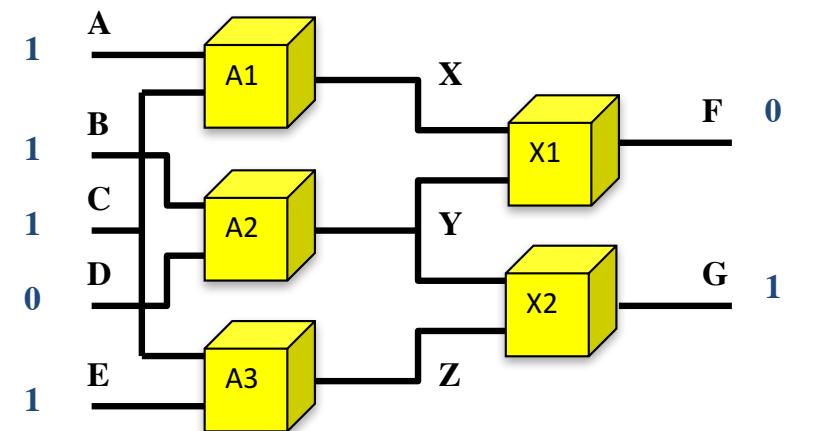
Implementation: Conflict-directed A* search.

Solver: OpSat

$$C_{\Phi} \equiv \{X \in D_X | C(X, Y) \text{ satisfiable}\}$$



Search Tree of Mode Assignments



A1

A2

A3

X1

X2

{ }

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

U

G

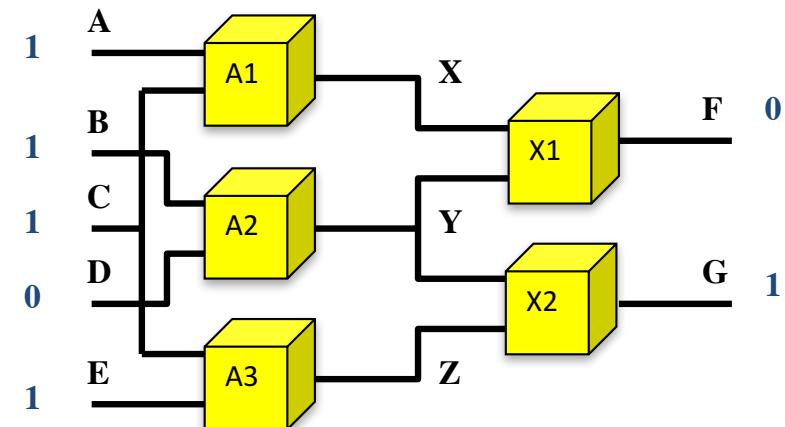
U

Constraint-based A*

$$P_{Xi=G} \gg P_{Xi=U}$$

$$P_{\text{single}} \gg P_{\text{double}}$$

$$P_{A2=U} > P_{A1=U} > P_{A3=U} > P_{X1=U} > P_{X2=U}$$



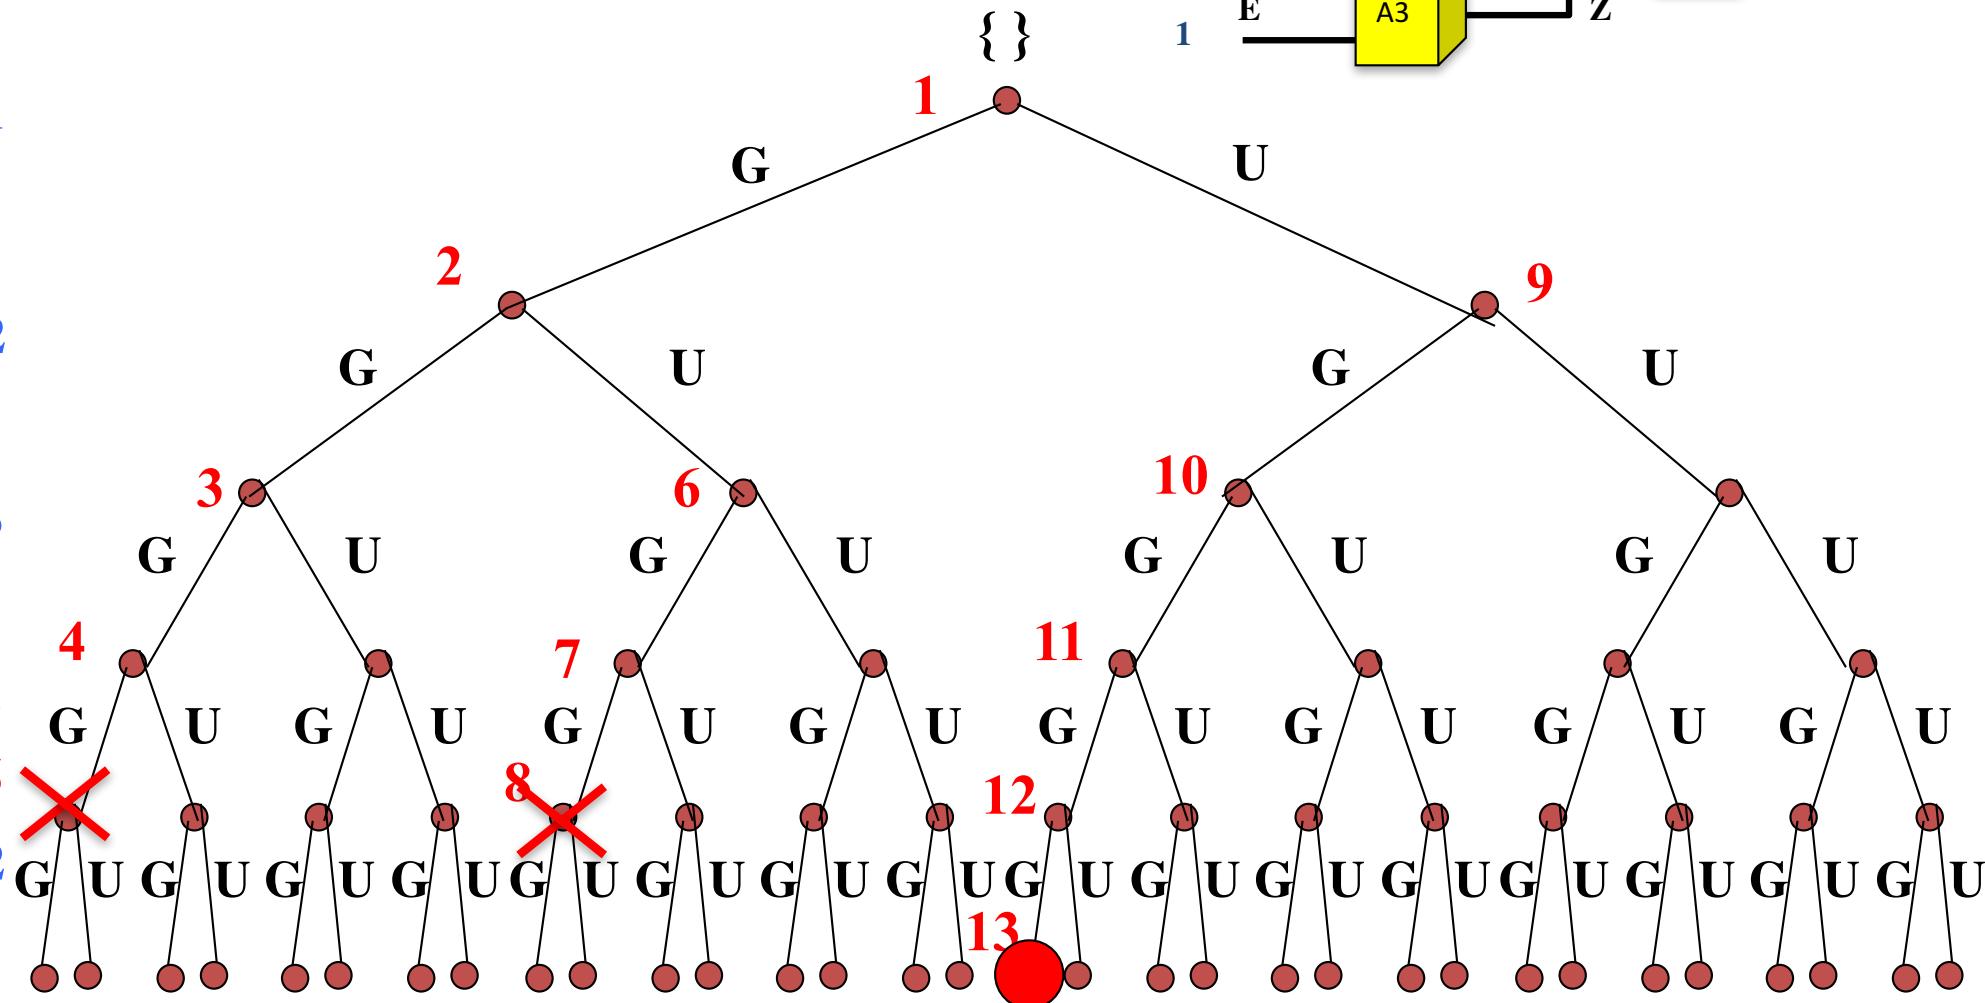
A1

A2

A3

X1

X2



Optimal CSP Candidate Generation using Constraint-based A*

Initial State:

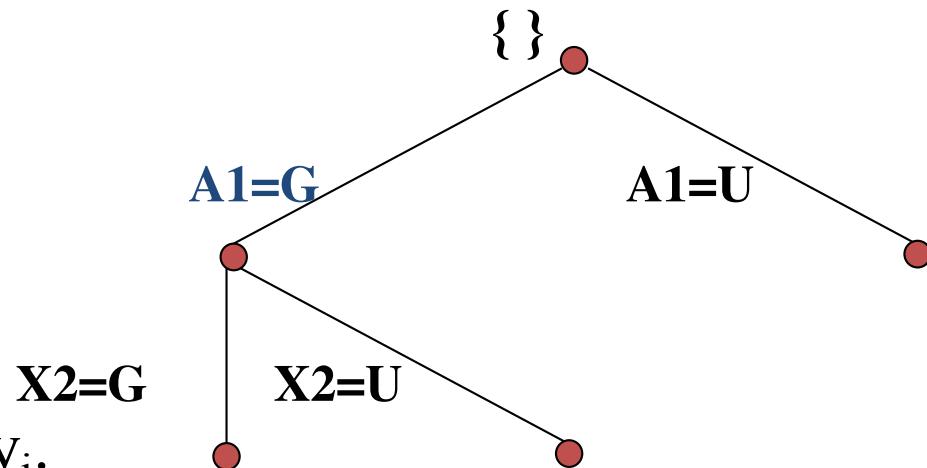
- Empty assignment $\{\}$.

Child Expansion:

- Select unassigned variable y_i .
- Each child adds an assignment from D_i .

Cutoff:

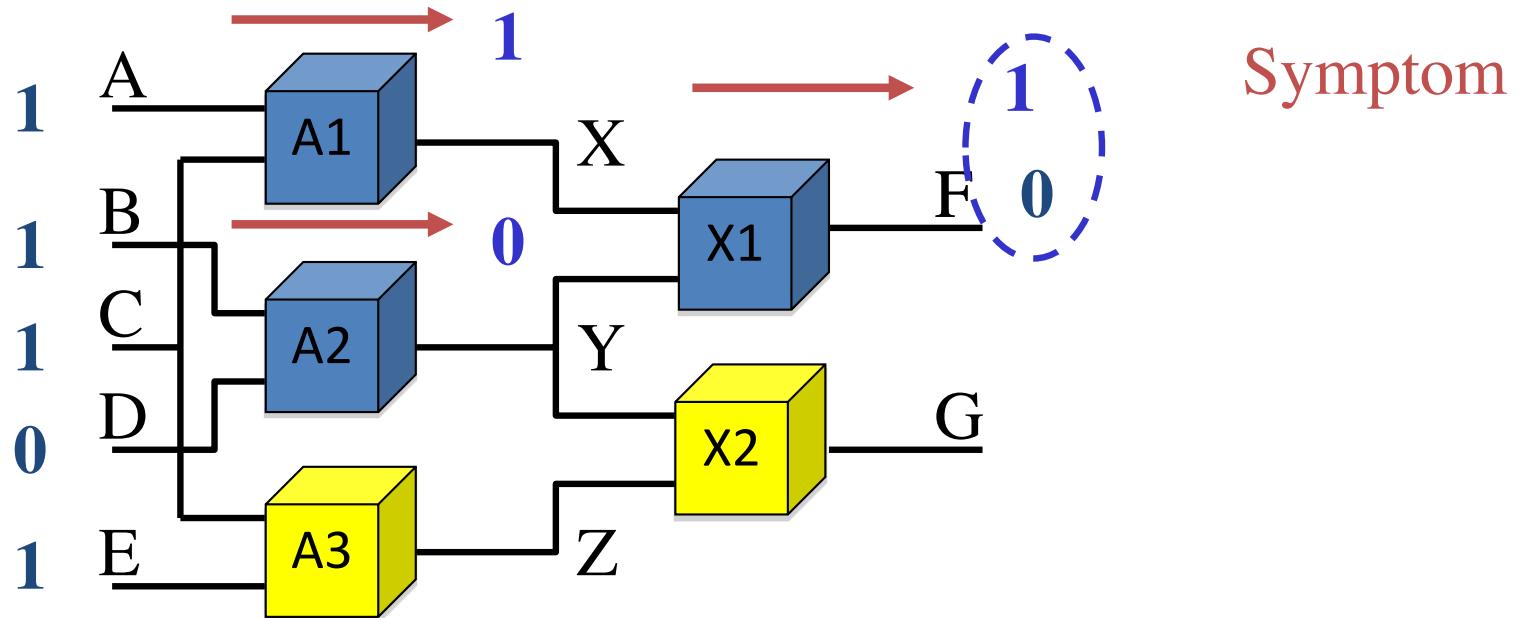
- $C(X, Y)$ is inconsistent.



Goal Test:

- Full assignment to X .
- $\exists Y \in D_Y . C(X, Y)$

Conflicts Indicate How to Remove Symptoms



Symptom:

F is observed 0, but predicted to be 1 if A1, A2 and X1 are okay.

Conflict 1: Learn $\{A1=G, A2=G, X1=G\}$ is inconsistent.

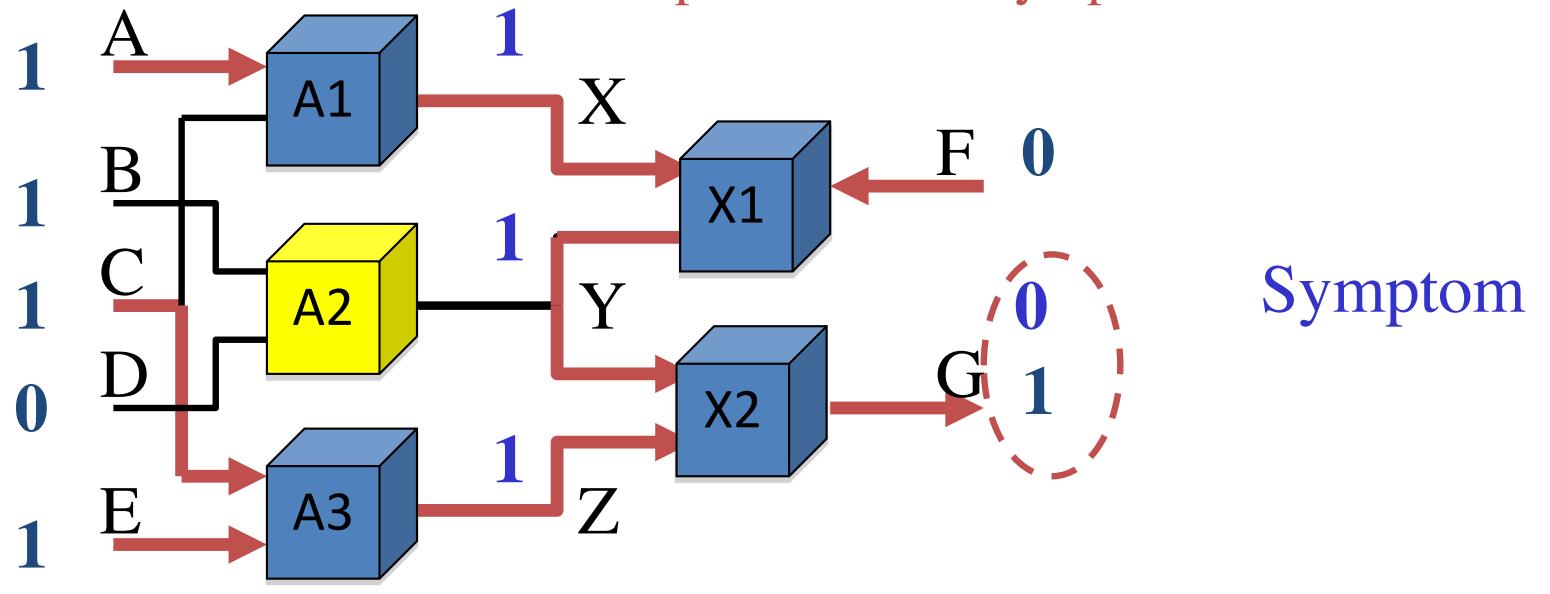
→ One of A1, A2 or X1 must be broken (*diagnosis of conflict*).

Conflict:

Partial assignment to mode variables X, inconsistent with $\Phi \wedge O=0$.

Second Conflict

Conflicting modes aren't always upstream from symptom.



Symptom: G is observed 1, but predicted 0.

Conflict 2: Learn $\{A1=G, A3=G, X1=G, X2=G\}$ is inconsistent.

→ One of A1, A3, X1 or X2 must be broken.

Generate Kernel Diagnoses by Set Covering

$\{A1=G, A2=G, X1=G\}$

Conflict 1.

$\{A1=G, A3=G, X1=G, X2=G\}$

Conflict 2.

$\{A1=U, A2=U, X1=U\}$

Diagnoses of Conflict 1.

$\{A1=U, A3=U, X1=U, X2=U\}$

Diagnoses of Conflict 2.

Kernel Diagnoses =

1. Compute cross product.
2. Remove supersets.
 - Old subset New.
 - New subset Old.

“Smallest” sets of modes that remove all conflicts.

Generate Kernel Diagnoses by Set Covering

$\{A1=G, A2=G, X1=G\}$

Conflict 1.

$\{A1=G, A3=G, X1=G, X2=G\}$

Conflict 2.

$\{A1=U, A2=U, X1=U\}$

Diagnoses of Conflict 1.

$\{A1=U, A3=U, X1=U, X2=U\}$

Diagnoses of Conflict 2.

Kernel Diagnoses = $\{A1=U\}$

1. Compute cross product.
2. Remove supersets.
 - Old subset New.
 - New subset Old.

“Smallest” sets of modes that remove all conflicts.

Generate Kernel Diagnoses by Set Covering

$\{A_1=G, A_2=G, X_1=G\}$

Conflict 1.

$\{A_1=G, A_3=G, X_1=G, X_2=G\}$

Conflict 2.

$\{A_1=U, A_2=U, X_1=U\}$

Diagnoses of Conflict 1.

$\{A_1=U, A_3=U, X_1=U, X_2=U\}$

Diagnoses of Conflict 2.

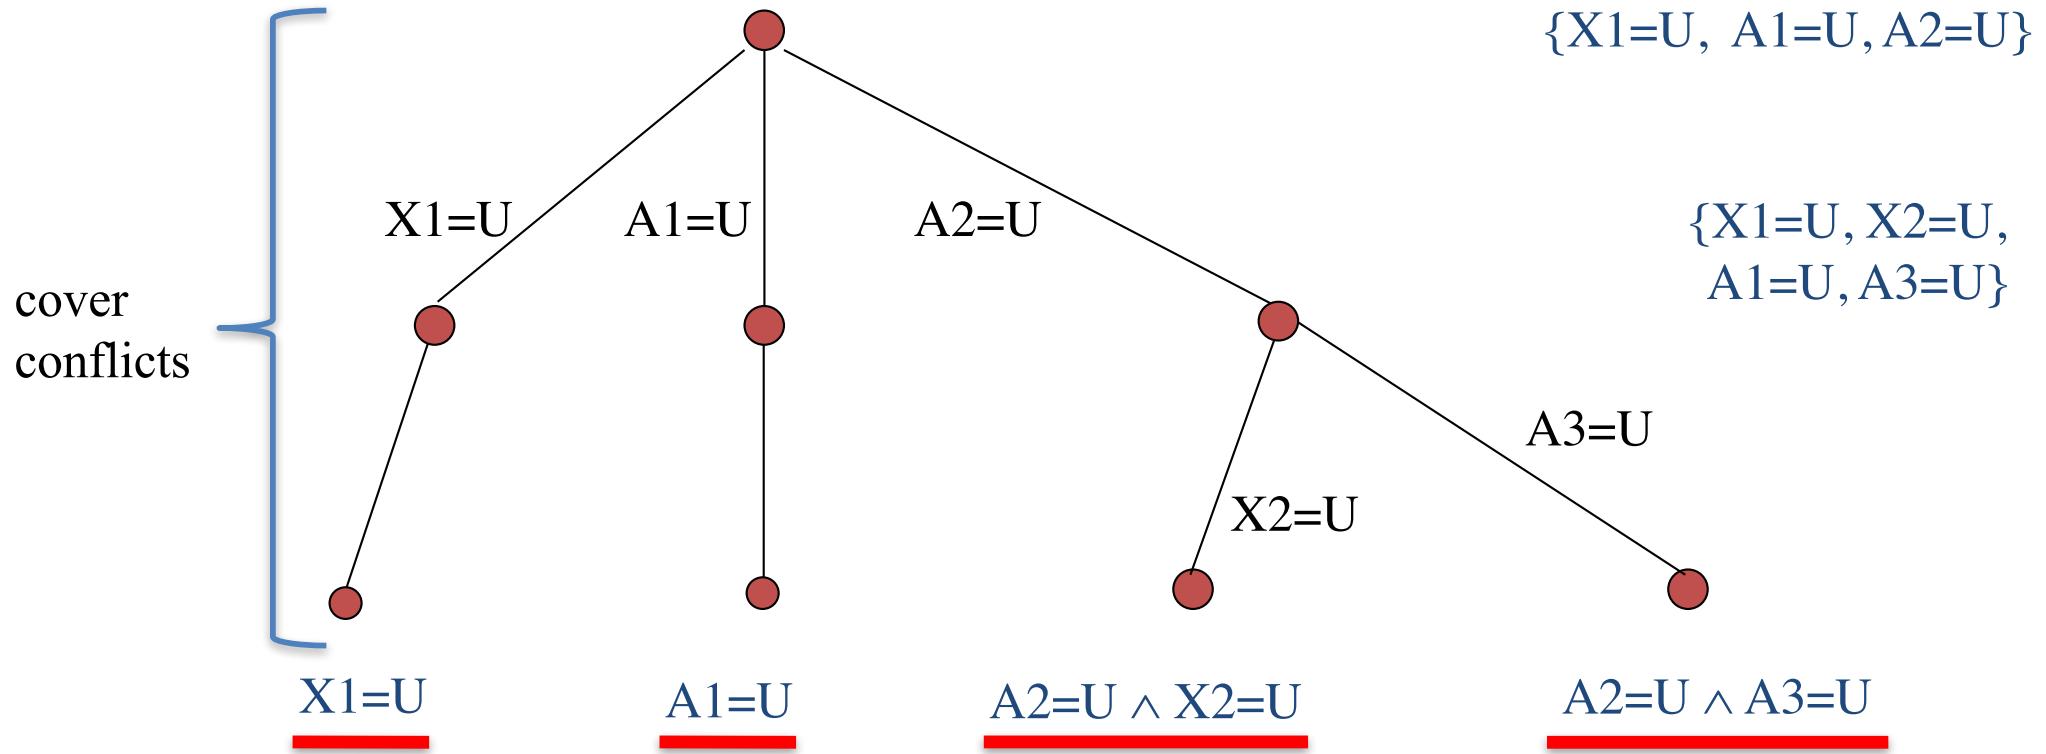
Kernel Diagnoses = $\{A_1=U, A_3=U\}$
 $\{A_1=U\}$

1. Compute cross product.
2. Remove supersets.
 - Old subset New.
 - New subset Old.

“Smallest” sets of modes that remove all conflicts.

Kernel Generation Tree

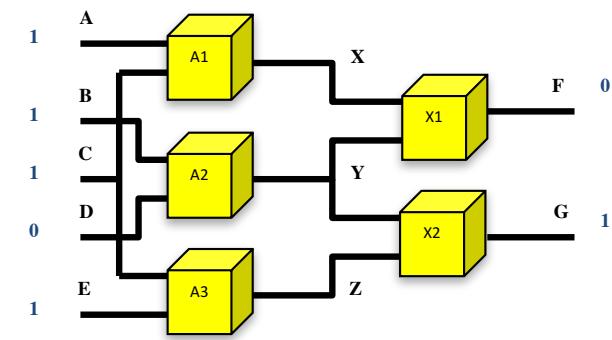
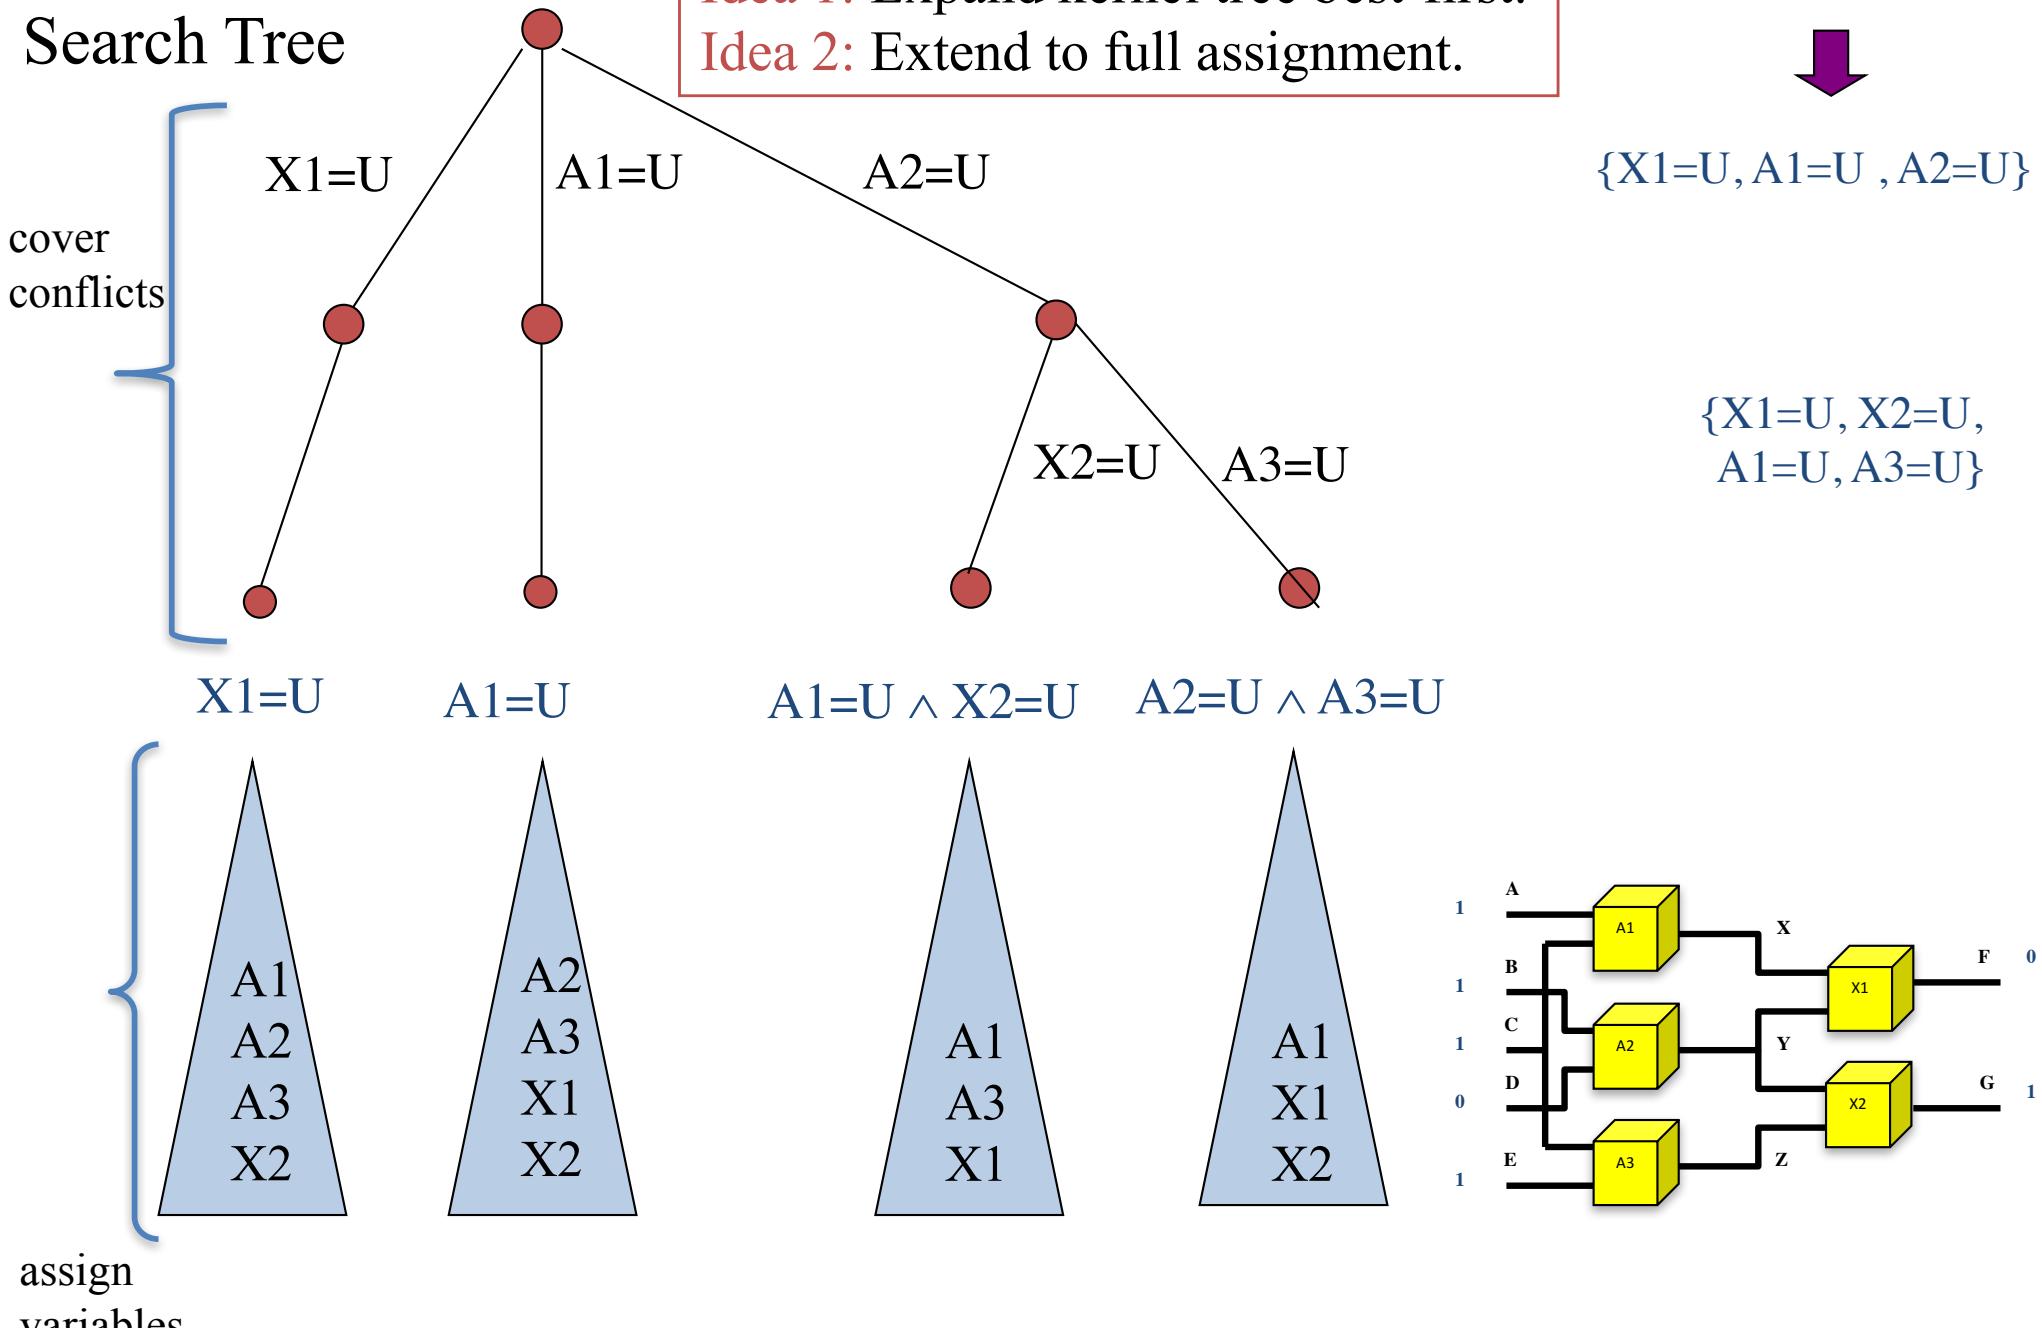
Constituent Kernels



Conflict-directed A* Search Tree

Idea 1: Expand kernel tree best-first.
 Idea 2: Extend to full assignment.

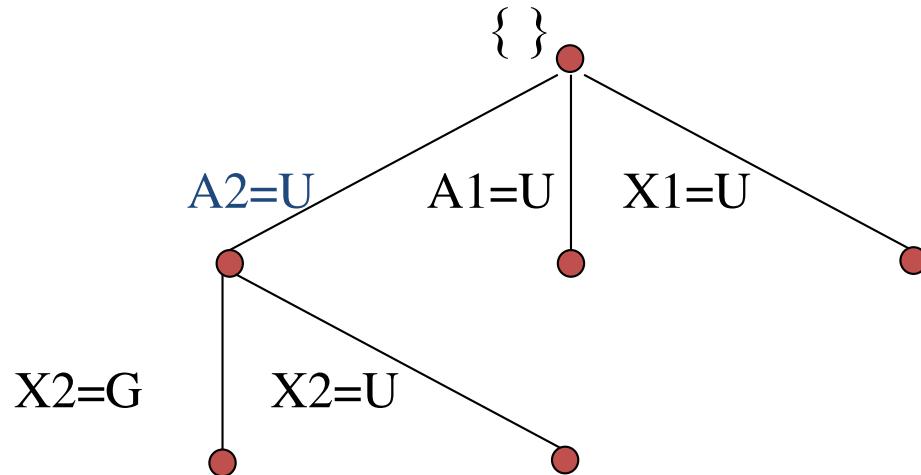
Constituent Kernels



Conflict-directed A* Child Expansion

Conflict

$$\neg (A2=G \wedge A1=G \wedge X1=G)$$



If Unresolved Conflicts:

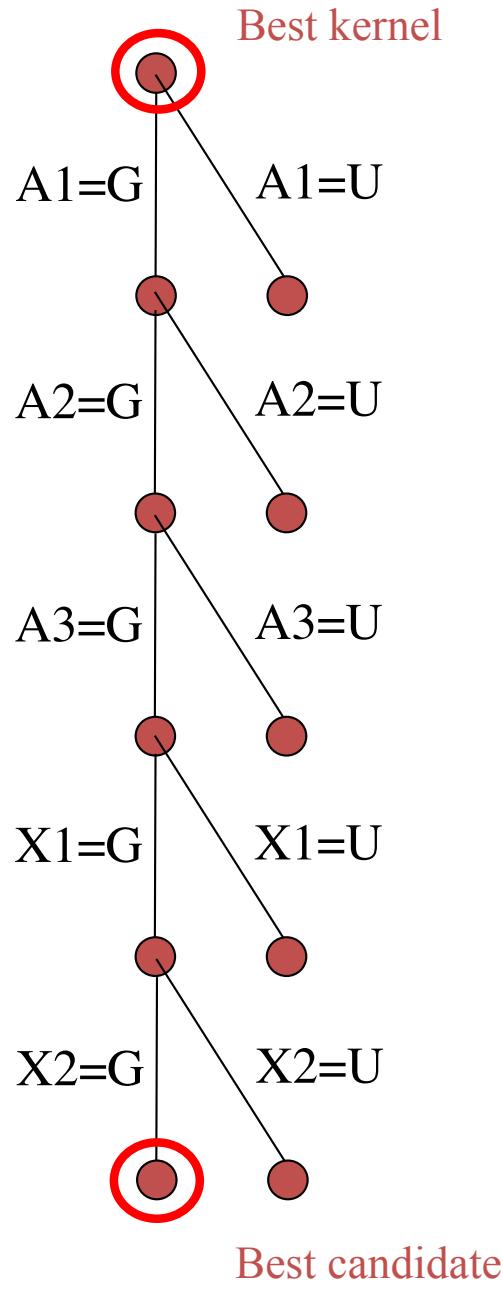
- Select unresolved conflict.
- Each child adds a conflict diagnosis (*constituent kernel*).

If All Conflicts Resolved:

- Select unassigned variable y_i .
- Each child adds an assignment from D_i .

1st Round:

cover
conflicts

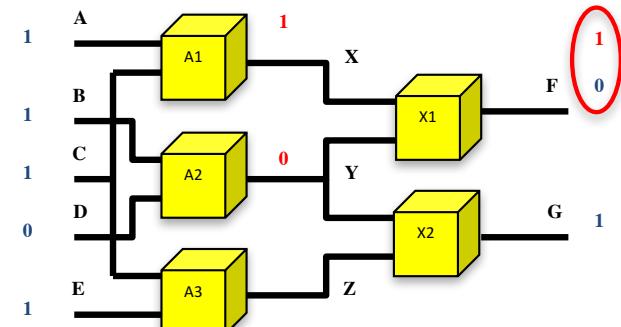


Conflict Diagnoses



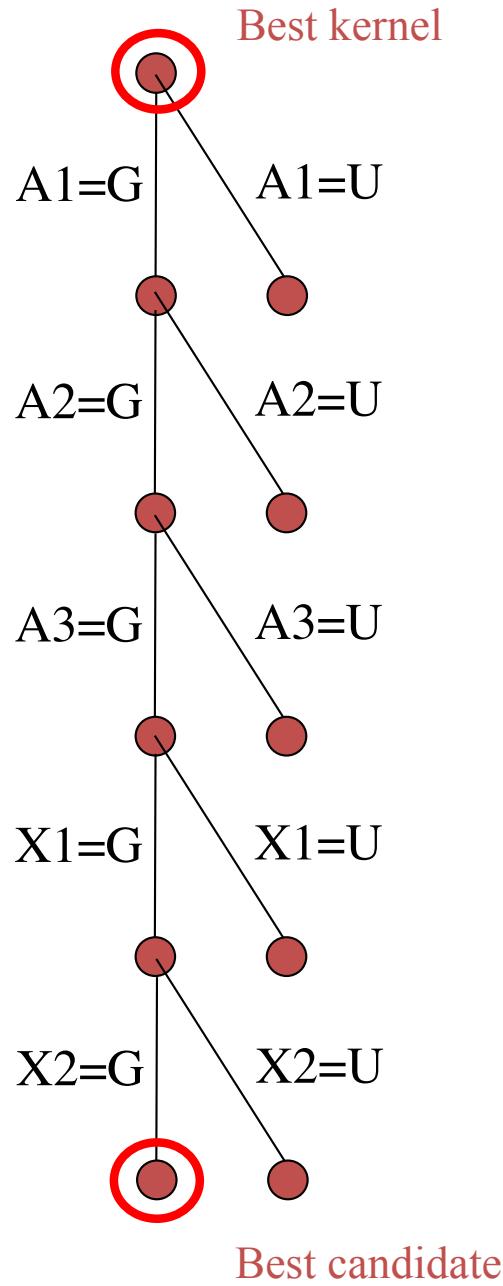
$P_{Xi=G} \gg P_{Xi=U}$
 $P_{\text{single}} \gg P_{\text{double}}$
 $P_{A2=U} > P_{A1=U}$
 $> P_{A3=U} > P_{X1=U}$
 $> P_{X2=U}$

Test Candidate 1:



1st Round:

cover conflicts }
assign variables

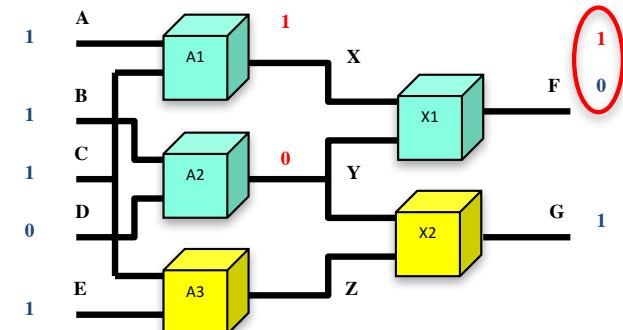


Conflict Diagnoses



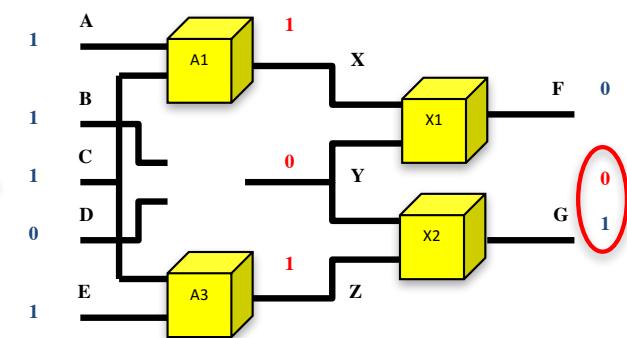
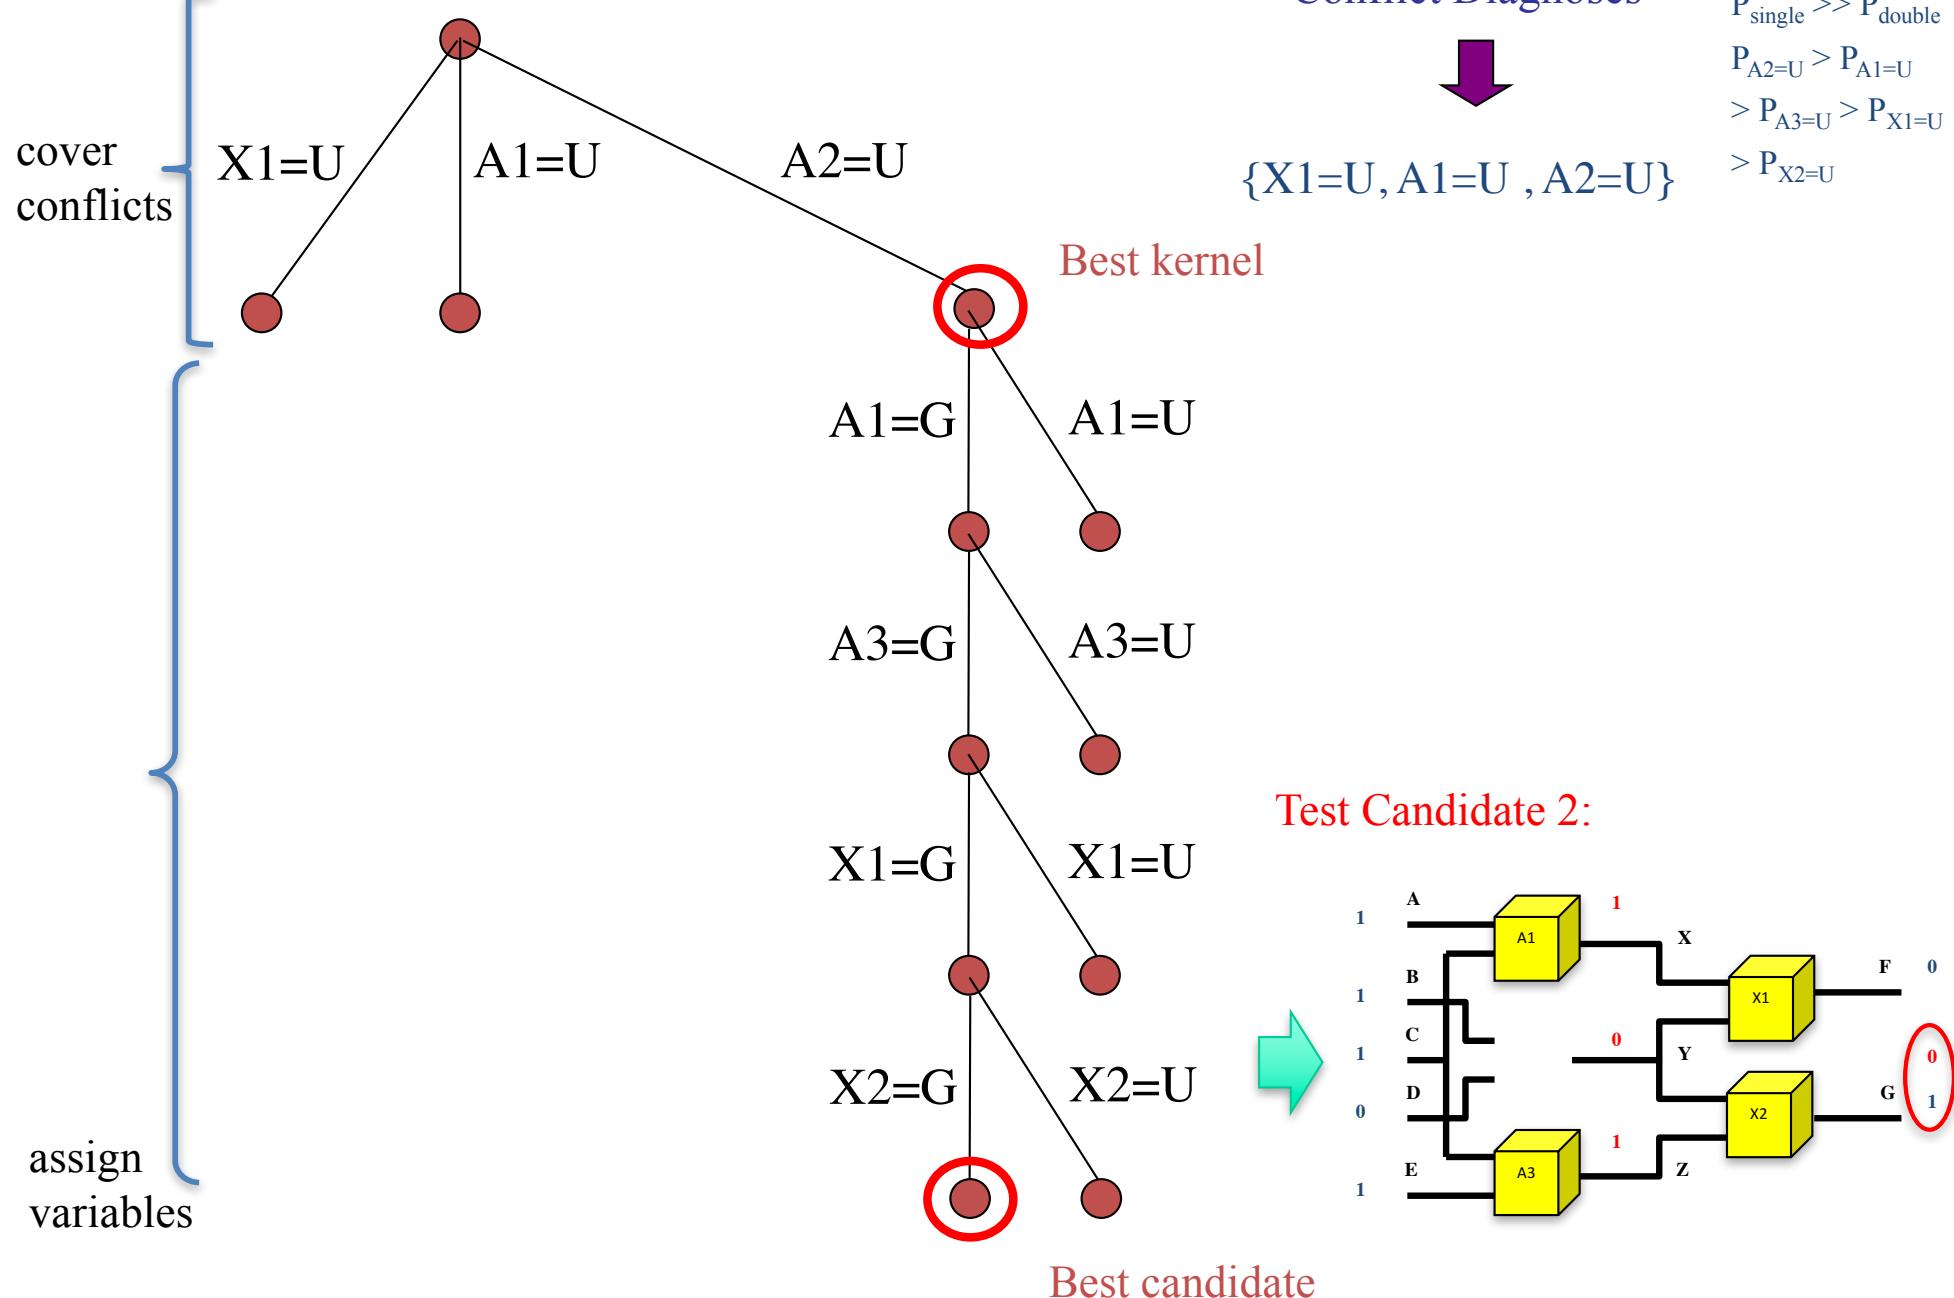
$P_{Xi=G} >> P_{Xi=U}$
 $P_{\text{single}} >> P_{\text{double}}$
 $P_{A2=U} > P_{A1=U}$
 $> P_{A3=U} > P_{X1=U}$
 $> P_{X2=U}$

Test Candidate 1:

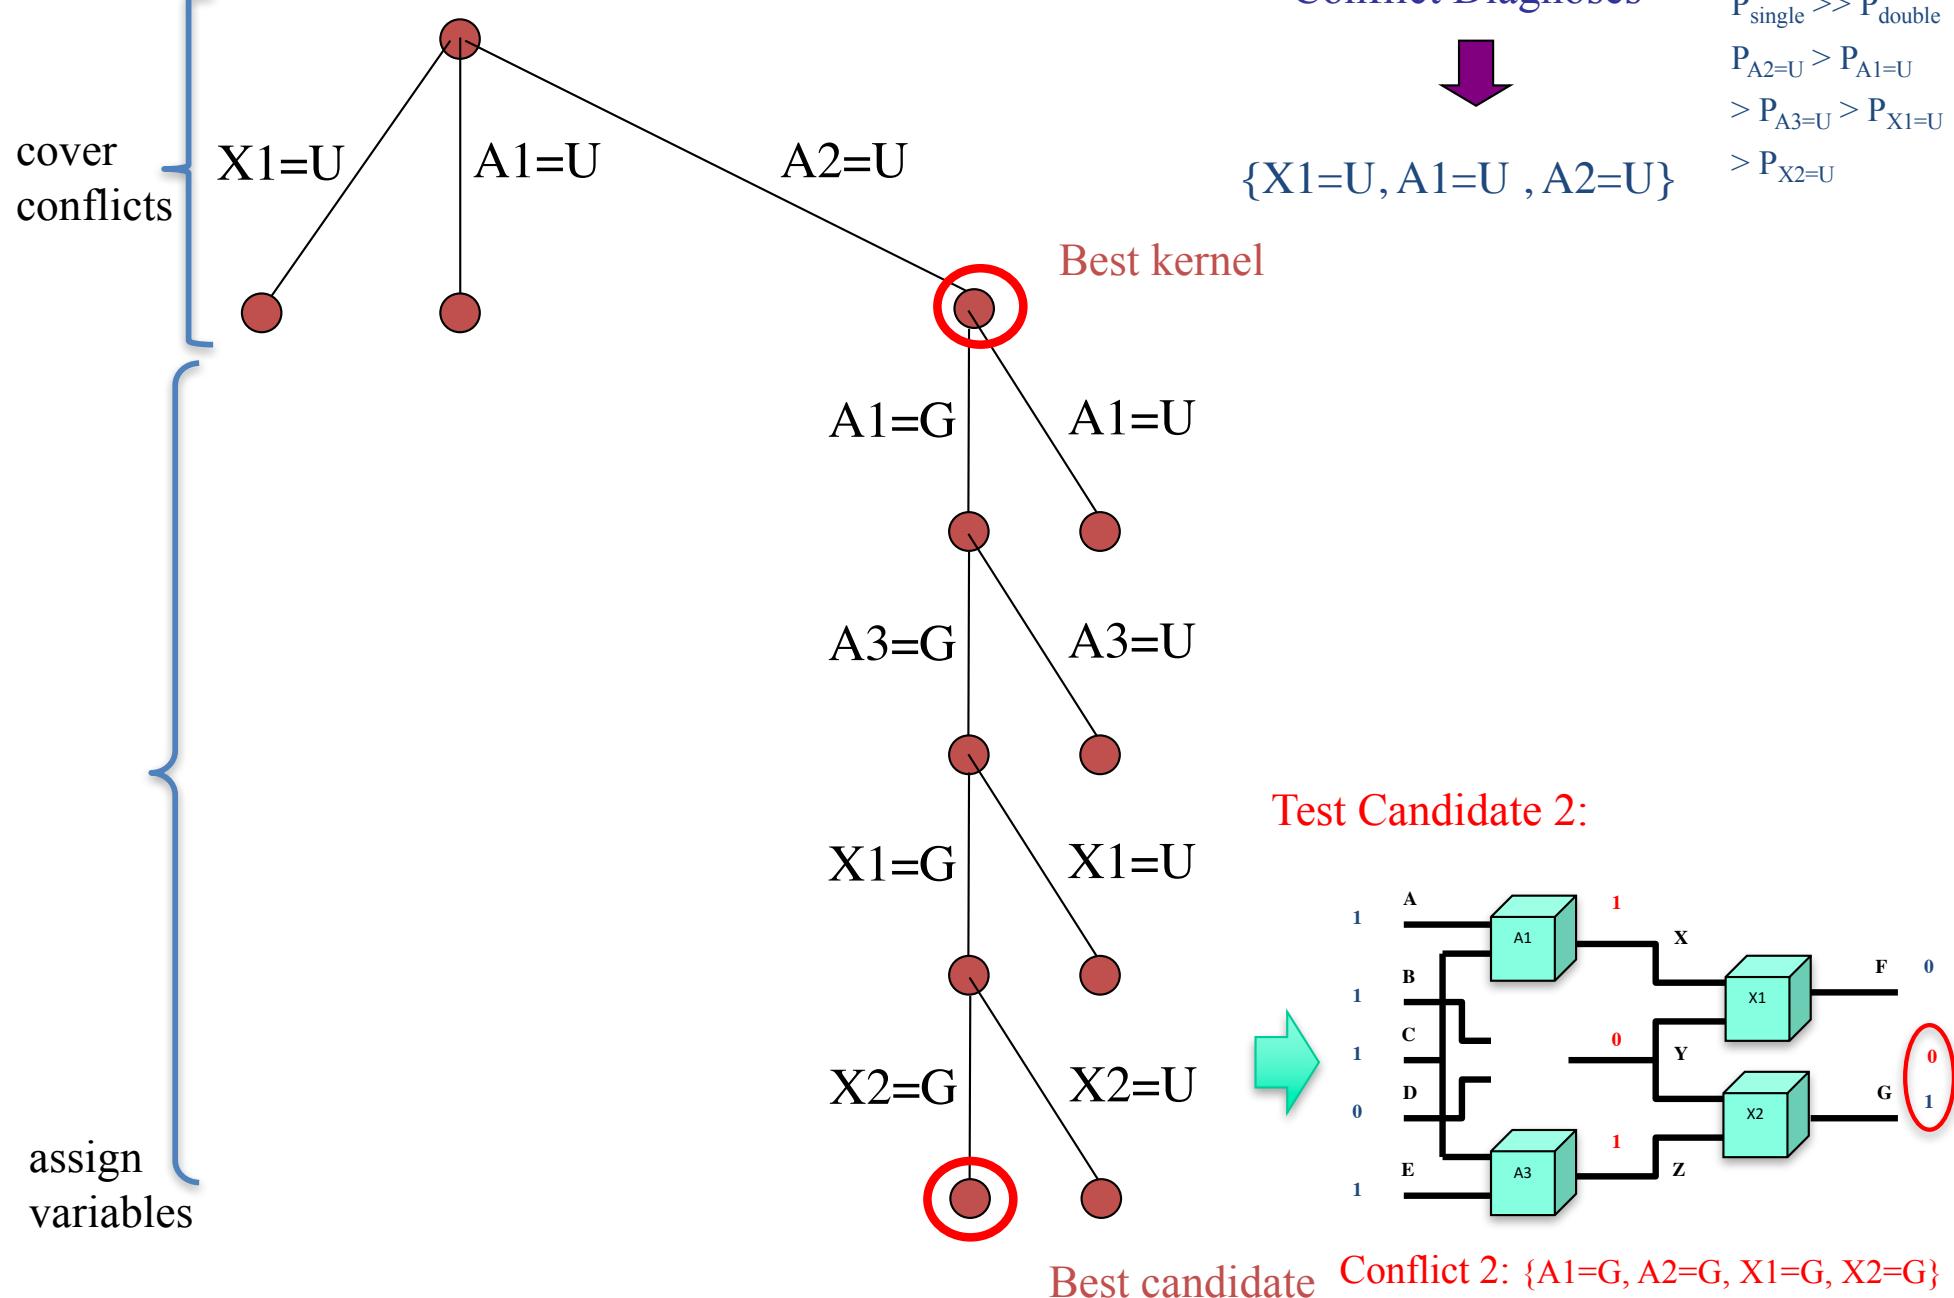


Conflict 1: {A1=G, A2=G, X1=G}

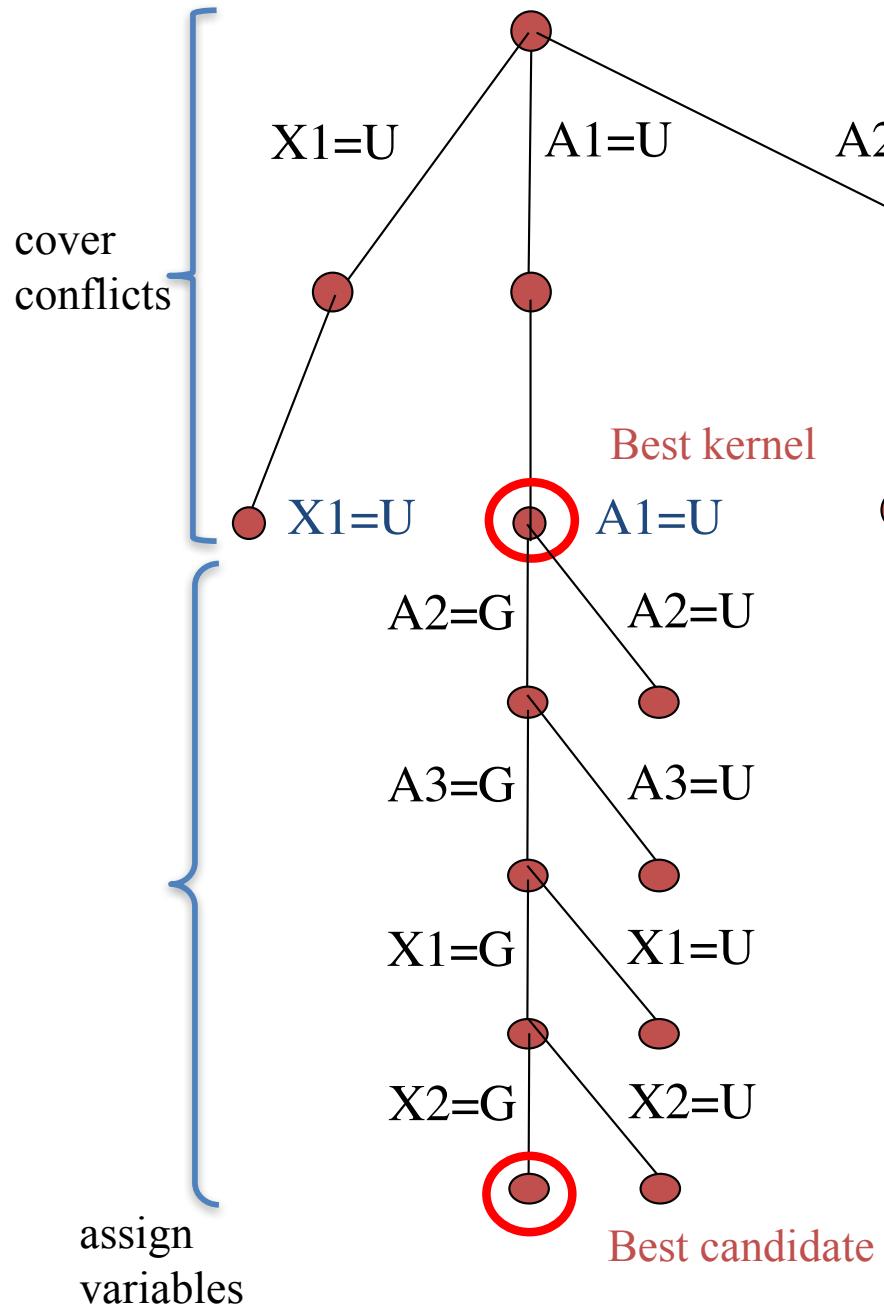
2nd Round:



2nd Round:



3rd Round:

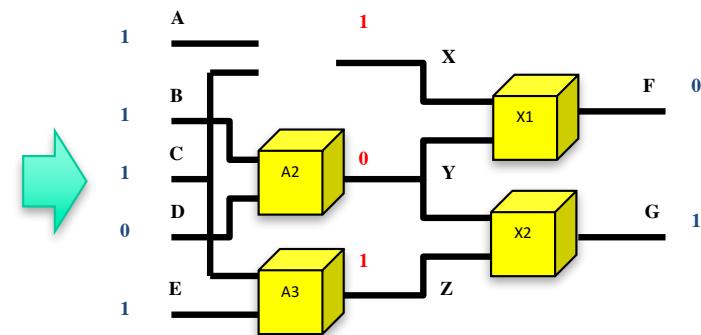


Conflict Diagnoses

$\{X1=U, A1=U, A2=U\}$

$\{X1=U, X2=U, A1=U, A3=U\}$

Test Candidate 3:



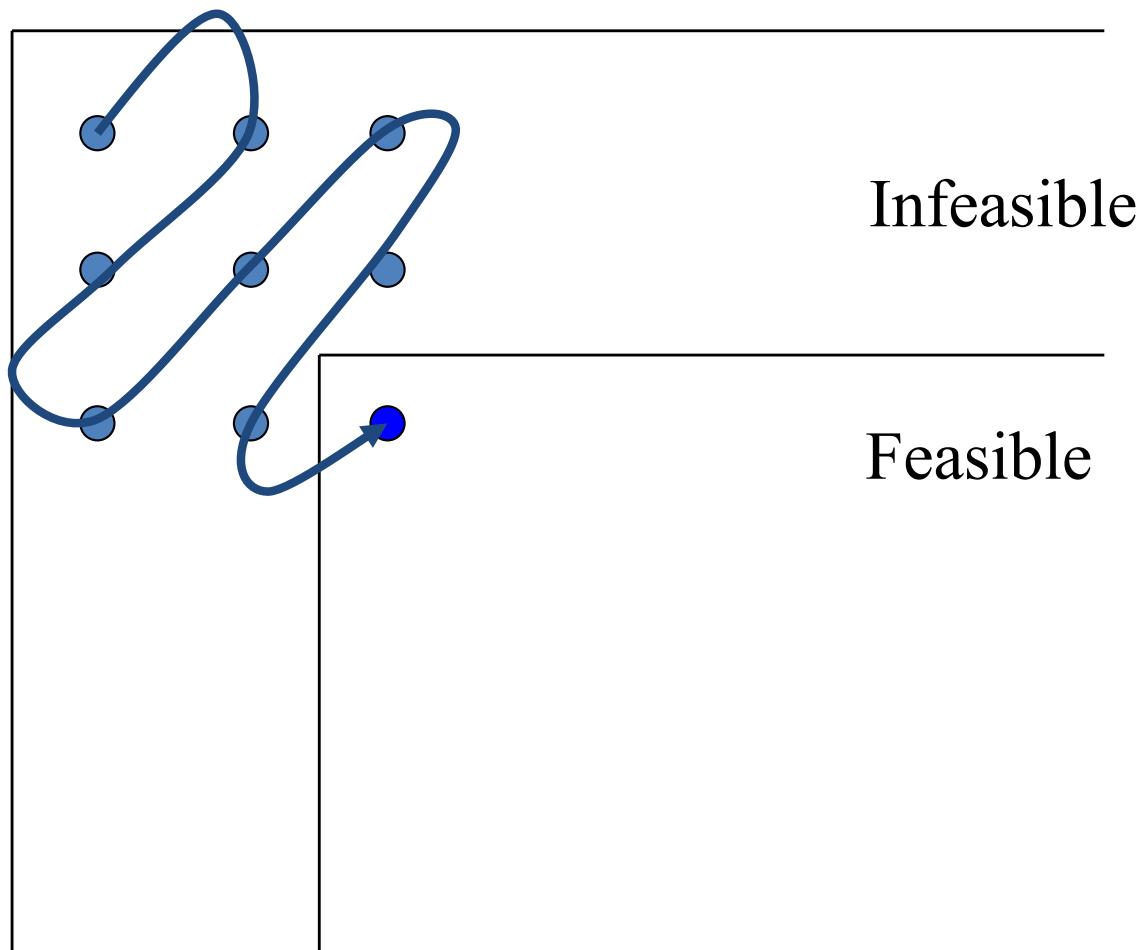
Consistent!

- All conflicts discovered.
- ⇒ All remaining candidates are diagnoses.

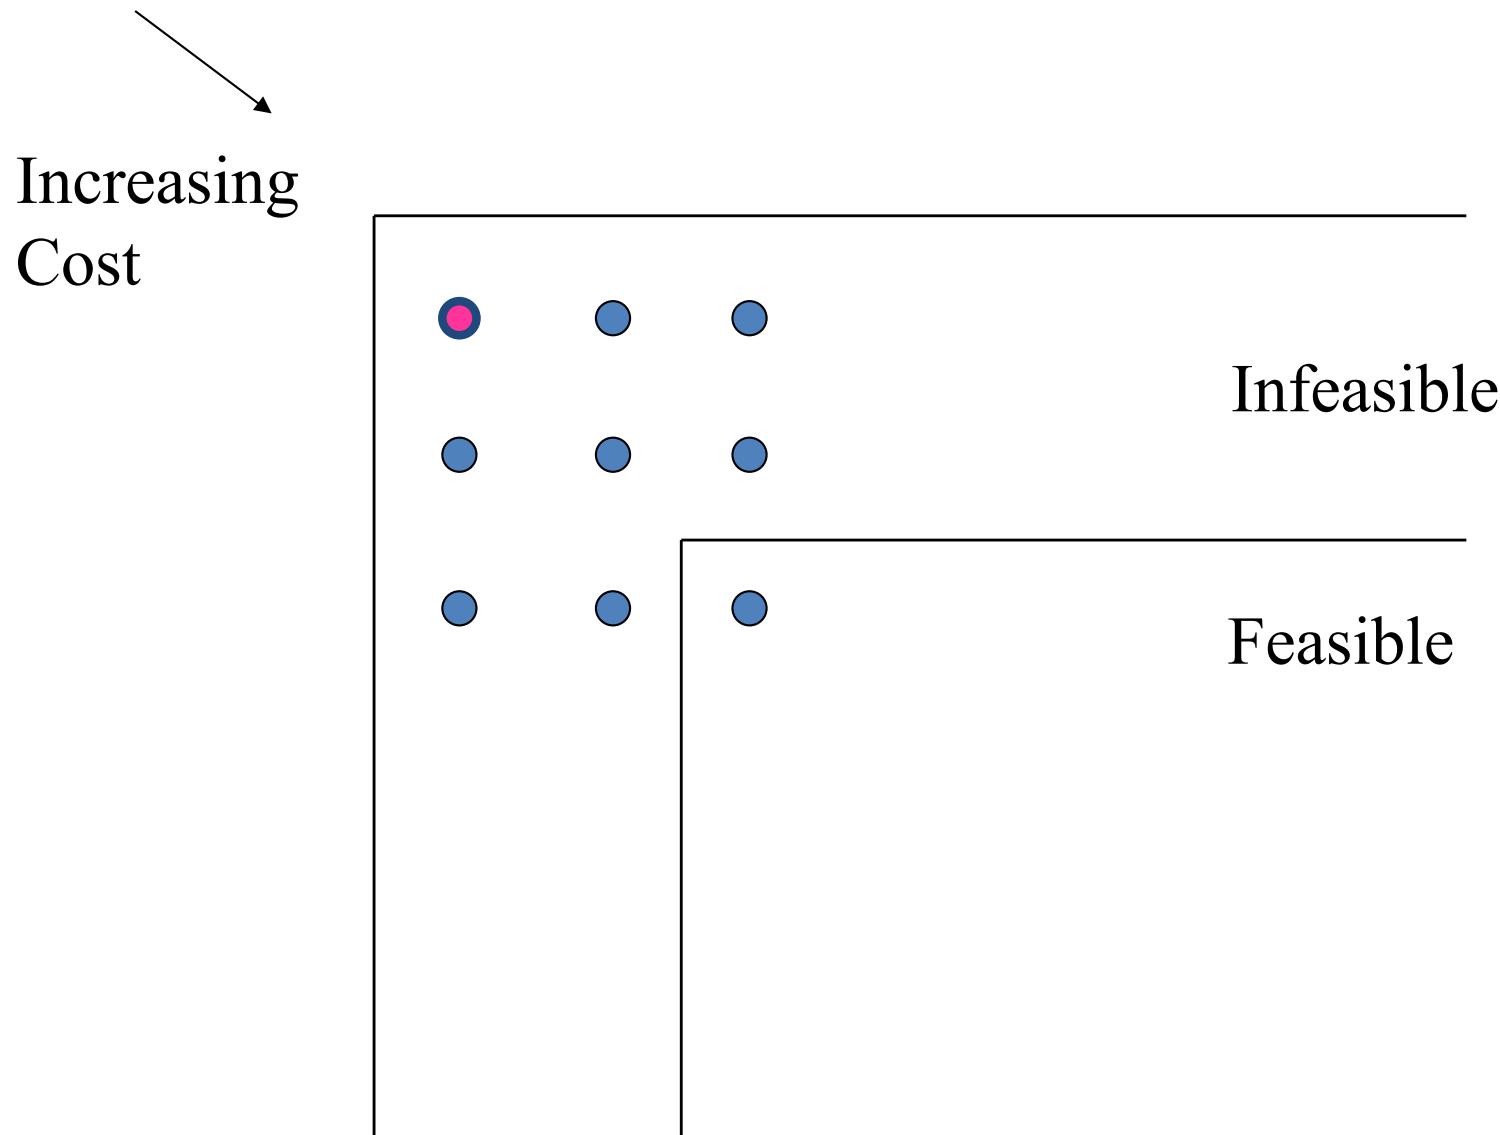
$$\begin{aligned} P_{Xi=G} &>> P_{Xi=U} \\ P_{\text{single}} &>> P_{\text{double}} \\ P_{A2=U} &> P_{A1=U} \\ &> P_{A3=U} > P_{X1=U} \\ &> P_{X2=U} \end{aligned}$$

A^*

Increasing
Cost

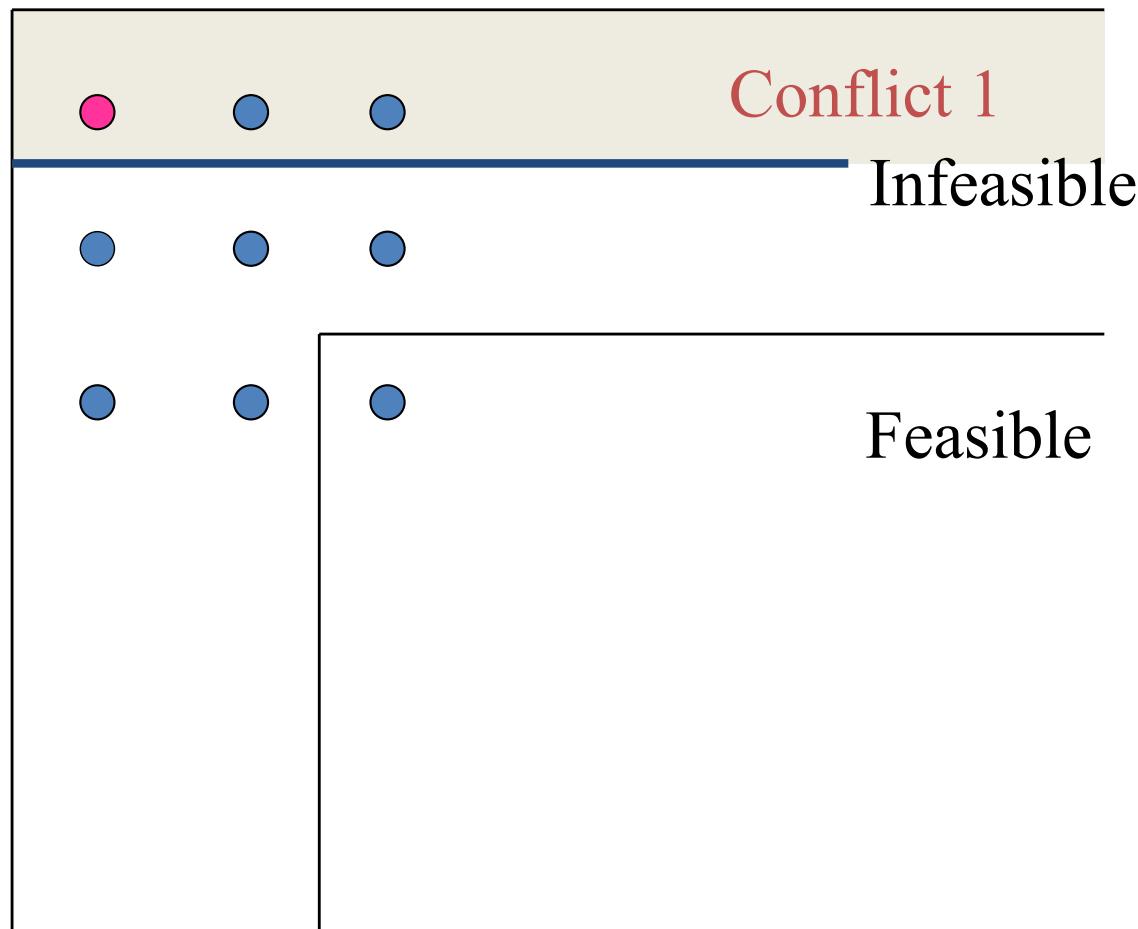


Conflict-directed A*



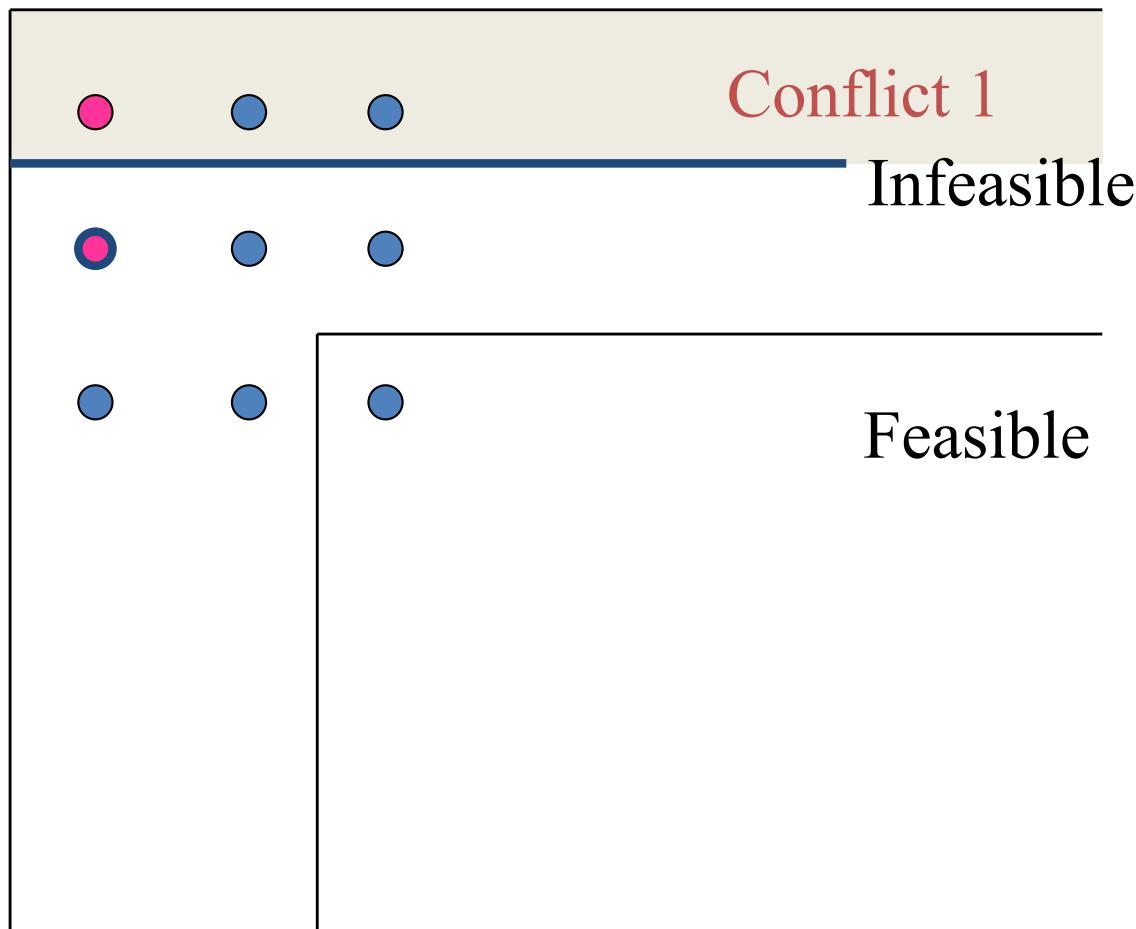
Conflict-directed A*

Increasing
Cost



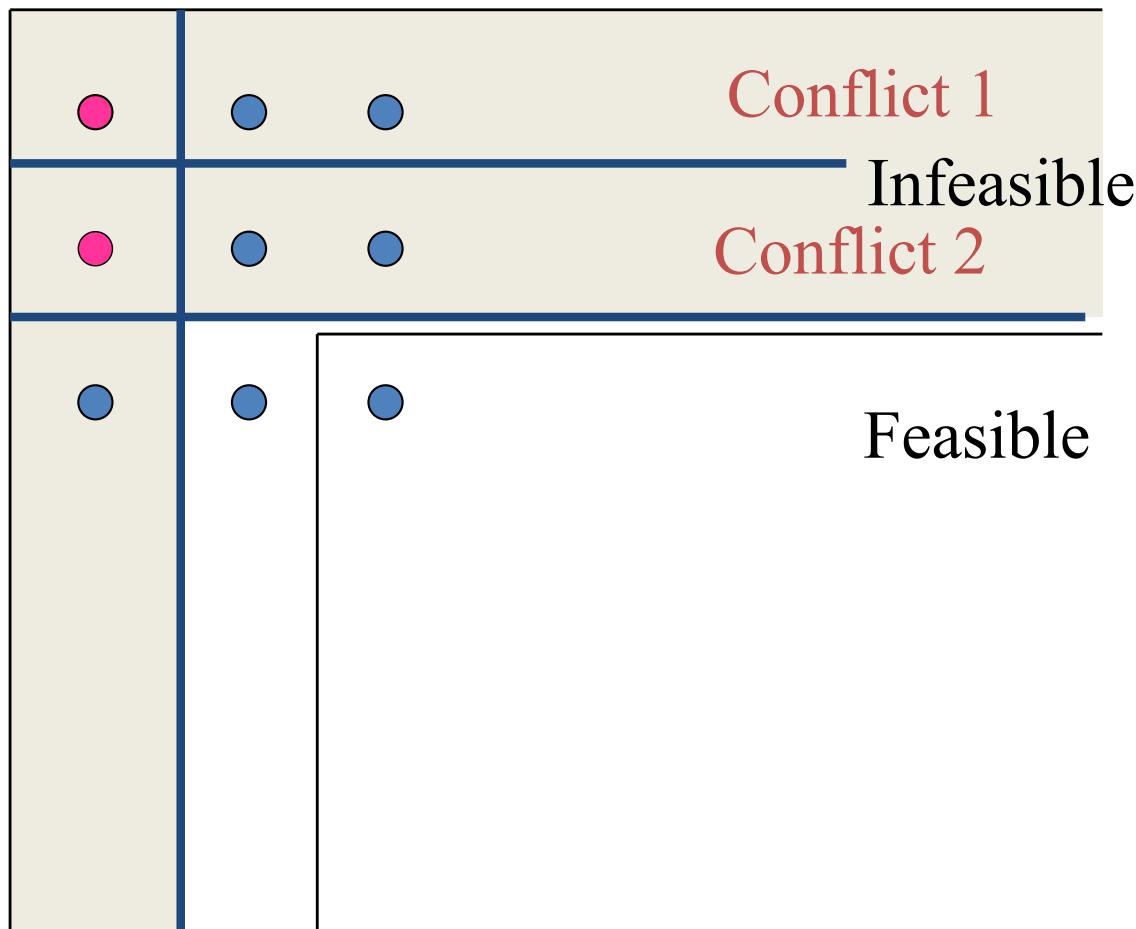
Conflict-directed A*

Increasing
Cost



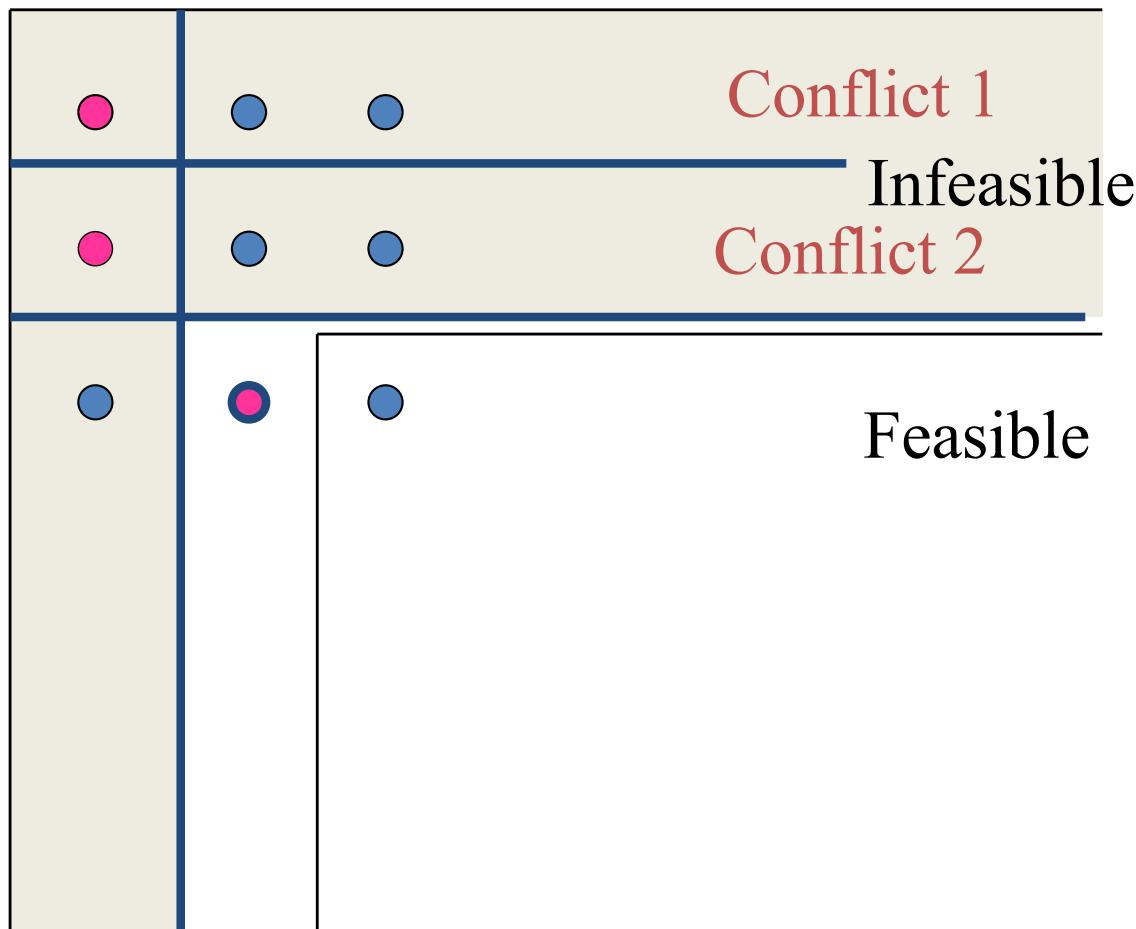
Conflict-directed A*

Increasing
Cost



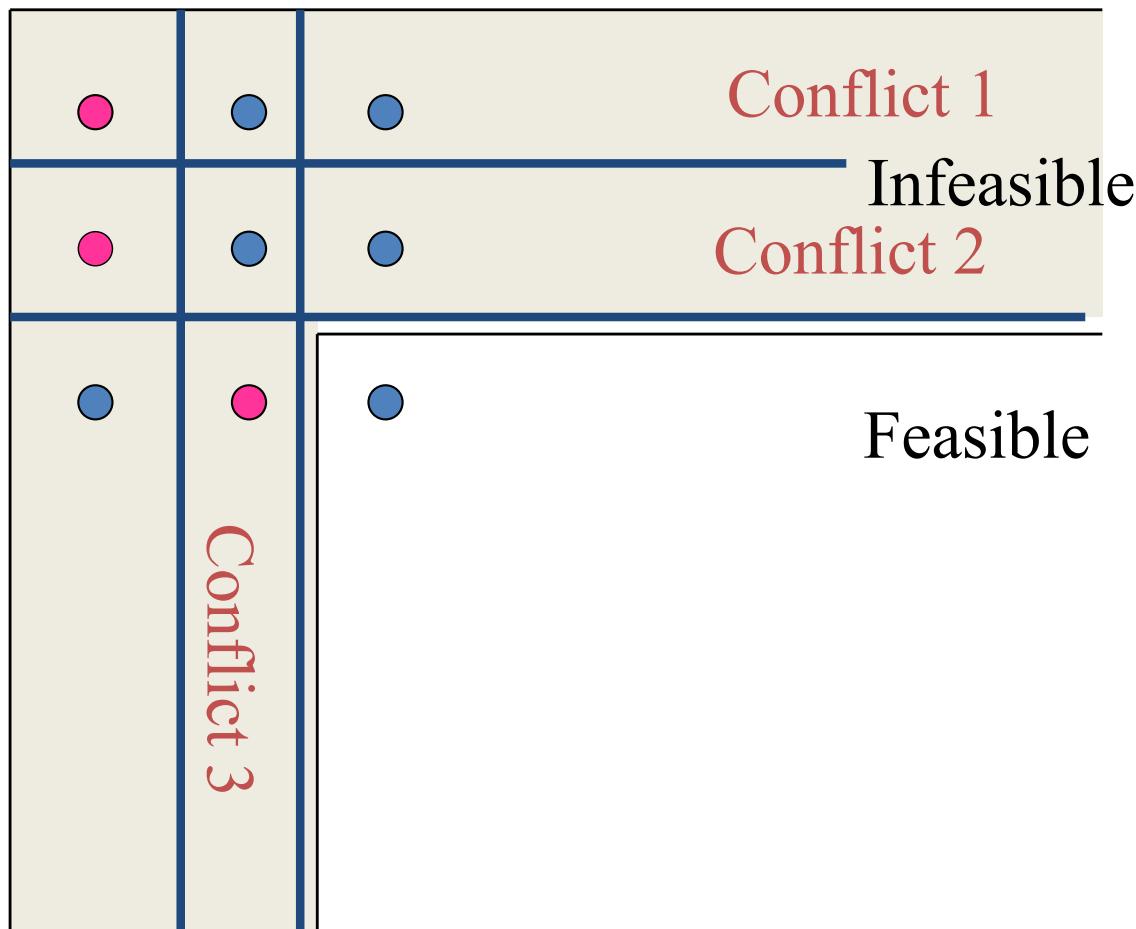
Conflict-directed A*

Increasing
Cost

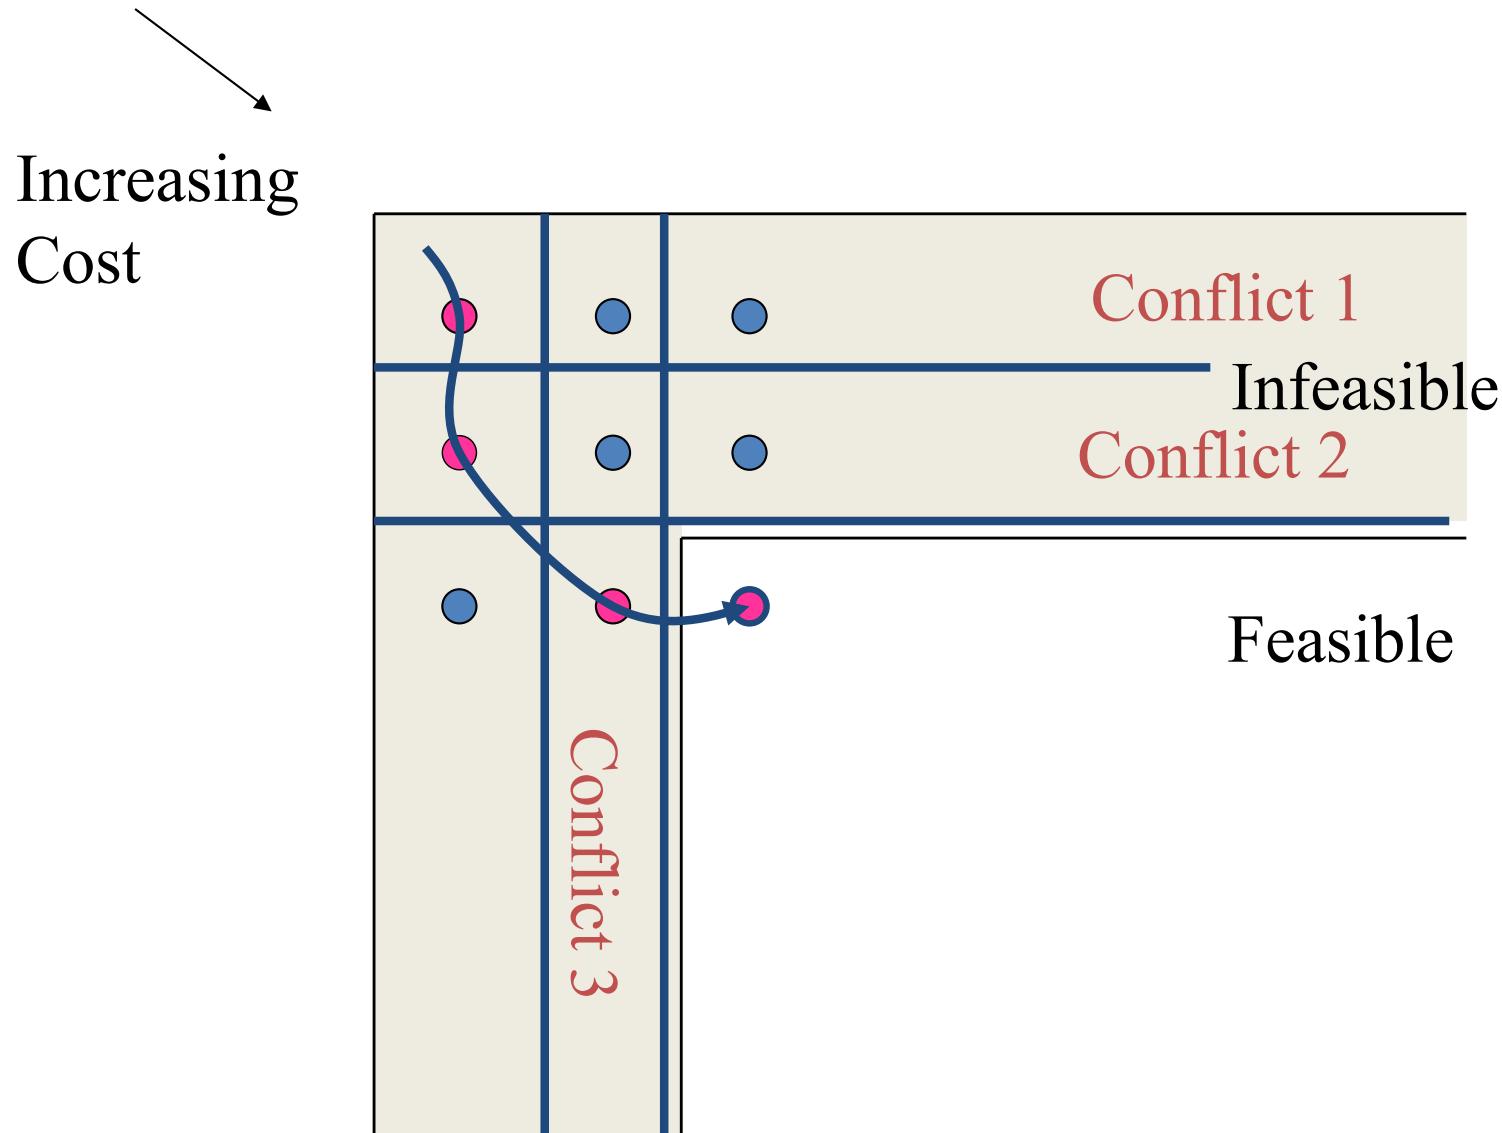


Conflict-directed A*

Increasing
Cost



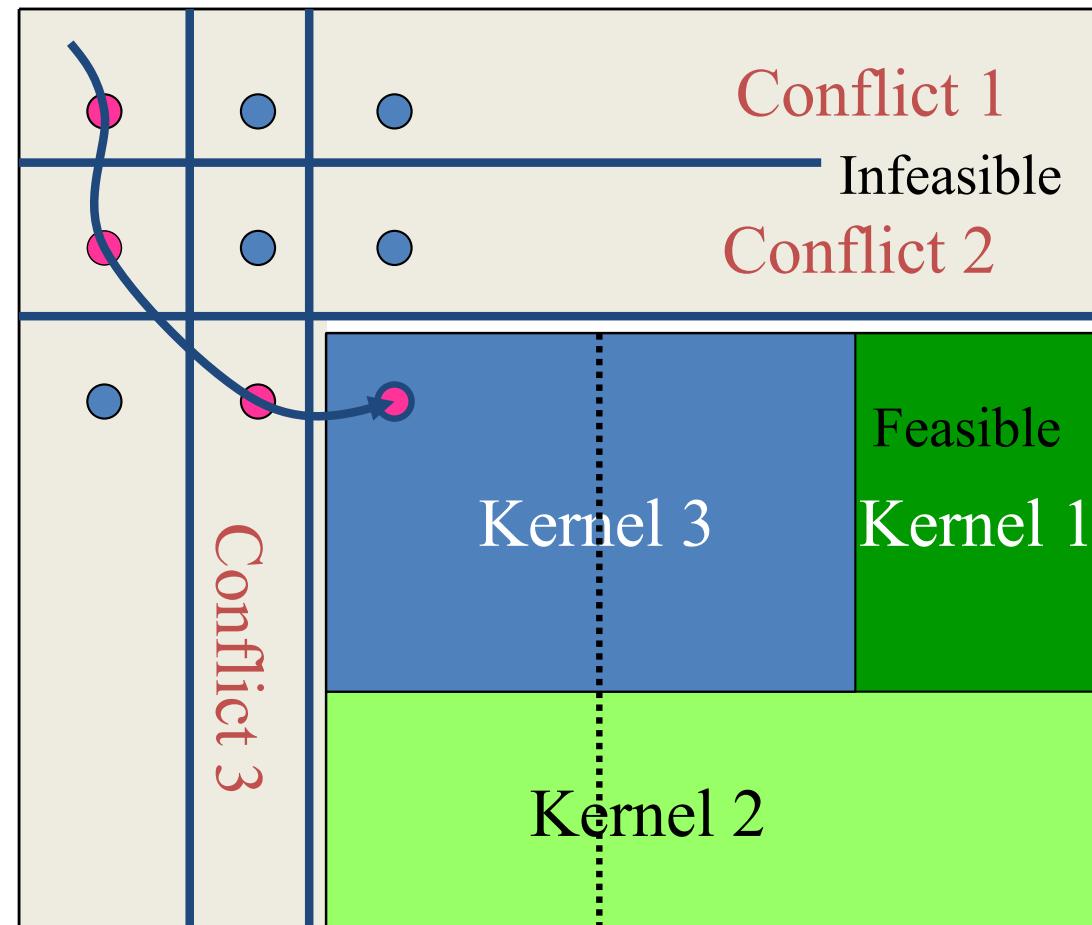
Conflict-directed A*



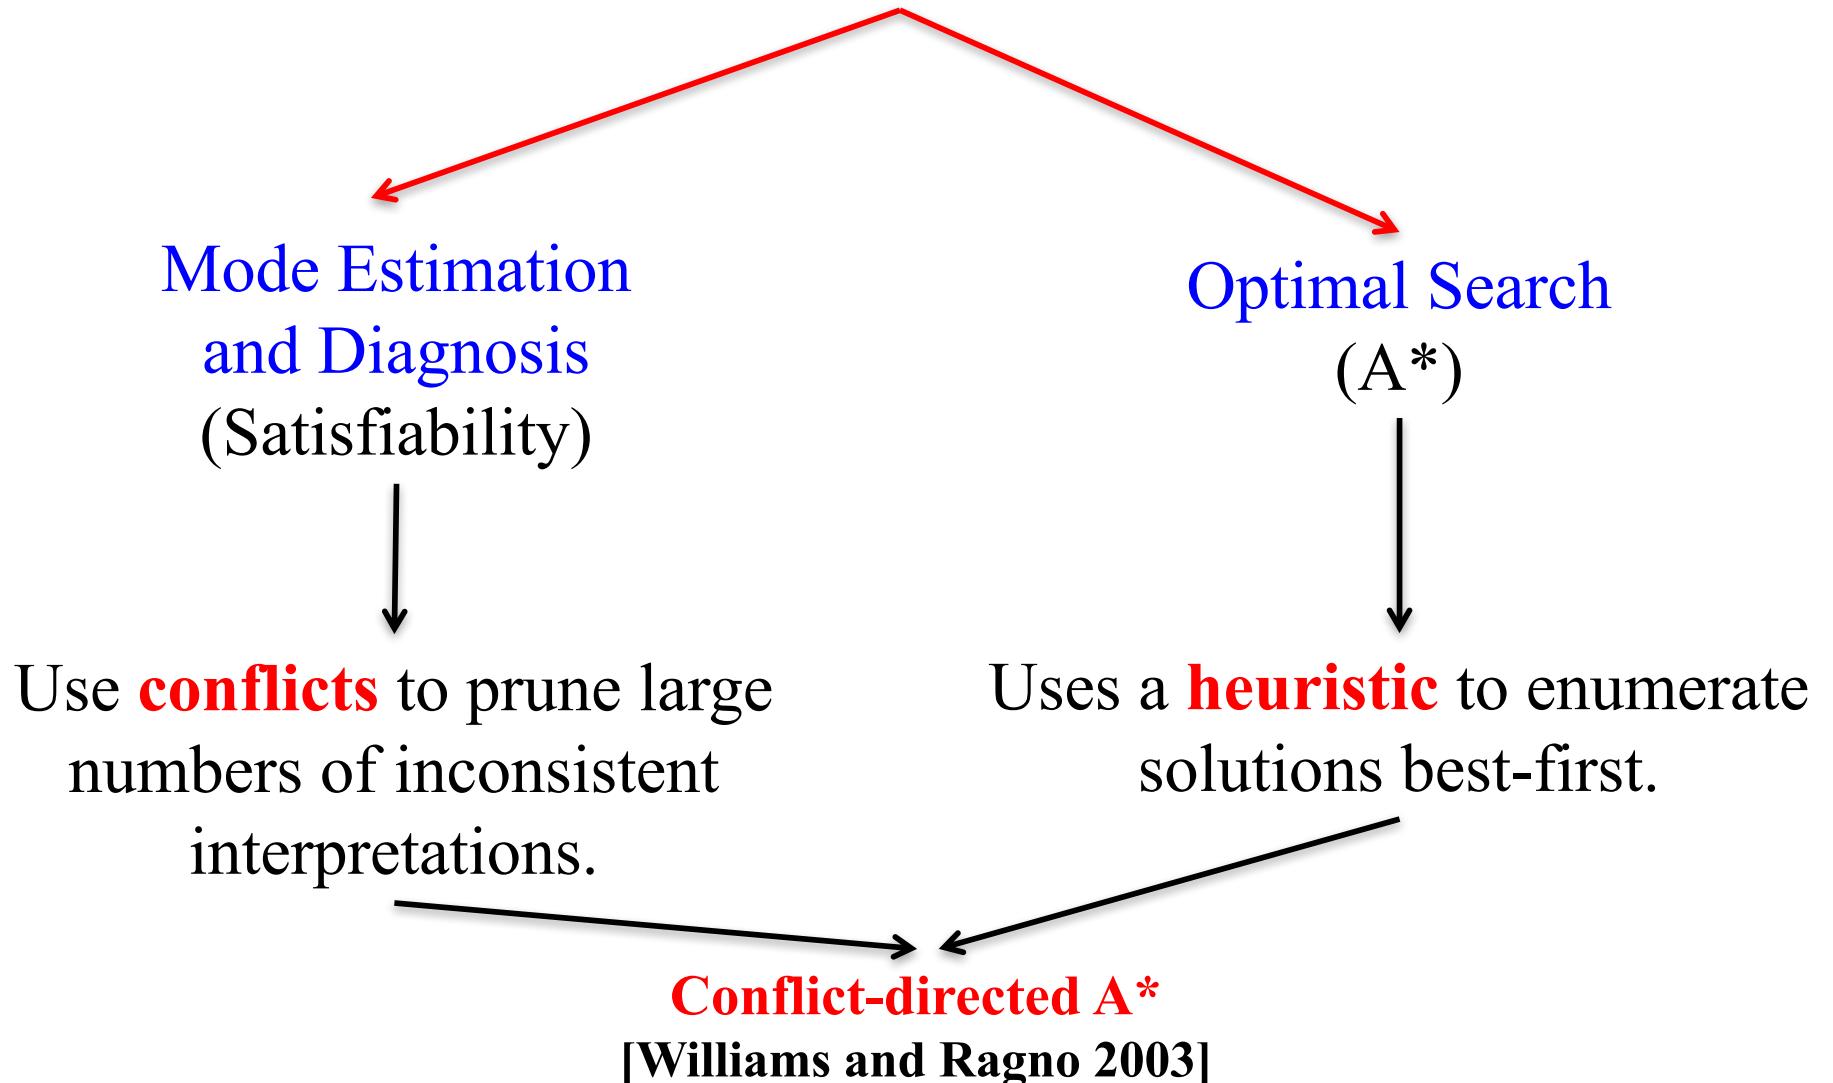
Conflict-directed A*

- Each feasible subregion described by a kernel assignment.
⇒ Approach: Use conflicts to search for kernel assignment containing the best cost candidate.

Increasing Cost

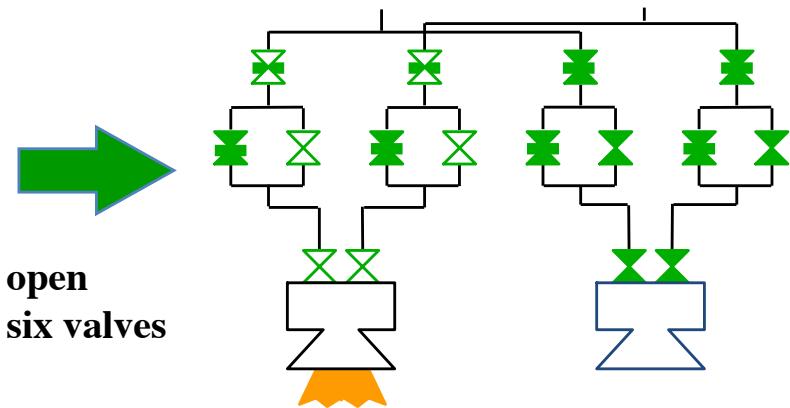


Enumerating Likely Mode Estimates

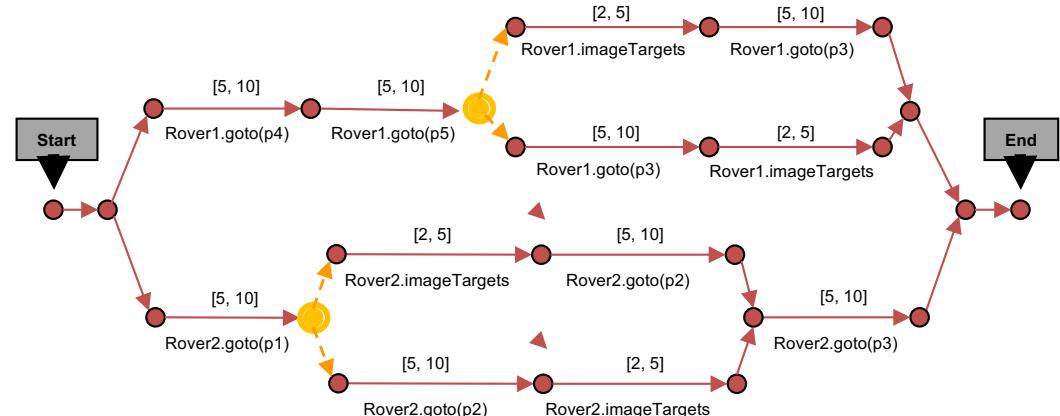


Use OpSat to . . .

Reconfigure Modes



Execute Programs with Choice



Livingstone
[Williams & Nayak, AAAI 96]

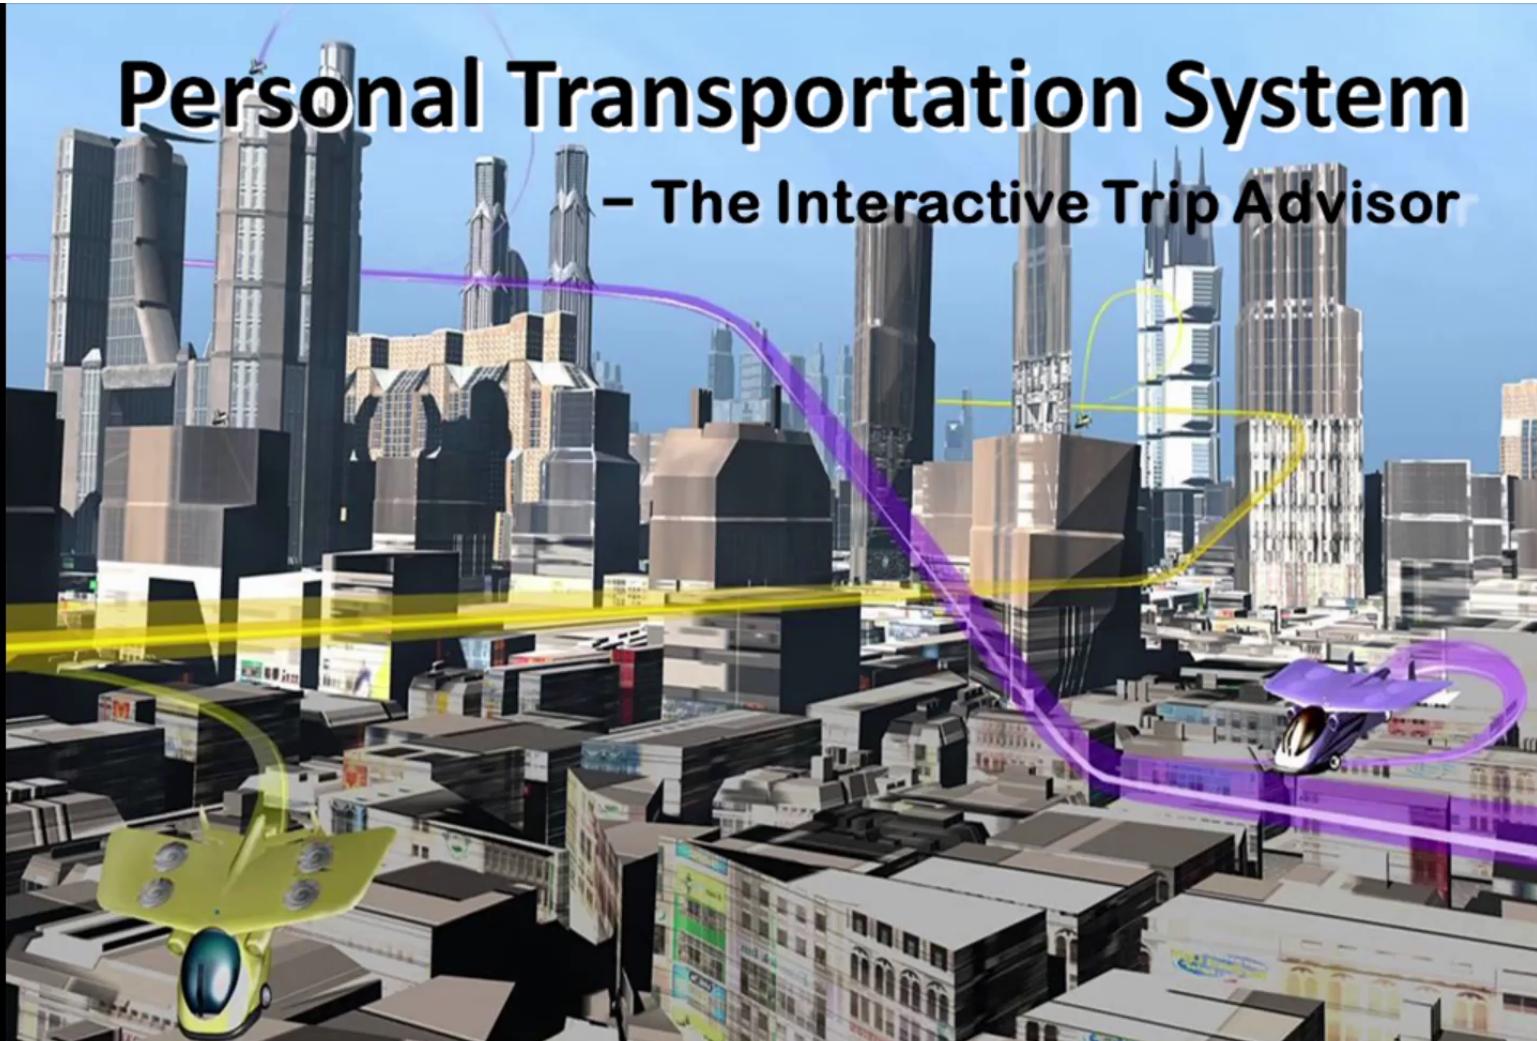
Kirk
[Kim, Williams & Abramson, IJCAI 01]

Outline

- Programs that Monitor State
 - Sub-goal monitoring
 - Model-based diagnosis and mode estimation
- Programs that Self-Diagnose (opt)

Personal Transportation System

– The Interactive Trip Advisor



Massachusetts
Institute of
Technology



Model-based Embedded & Robotic Systems



The Problem

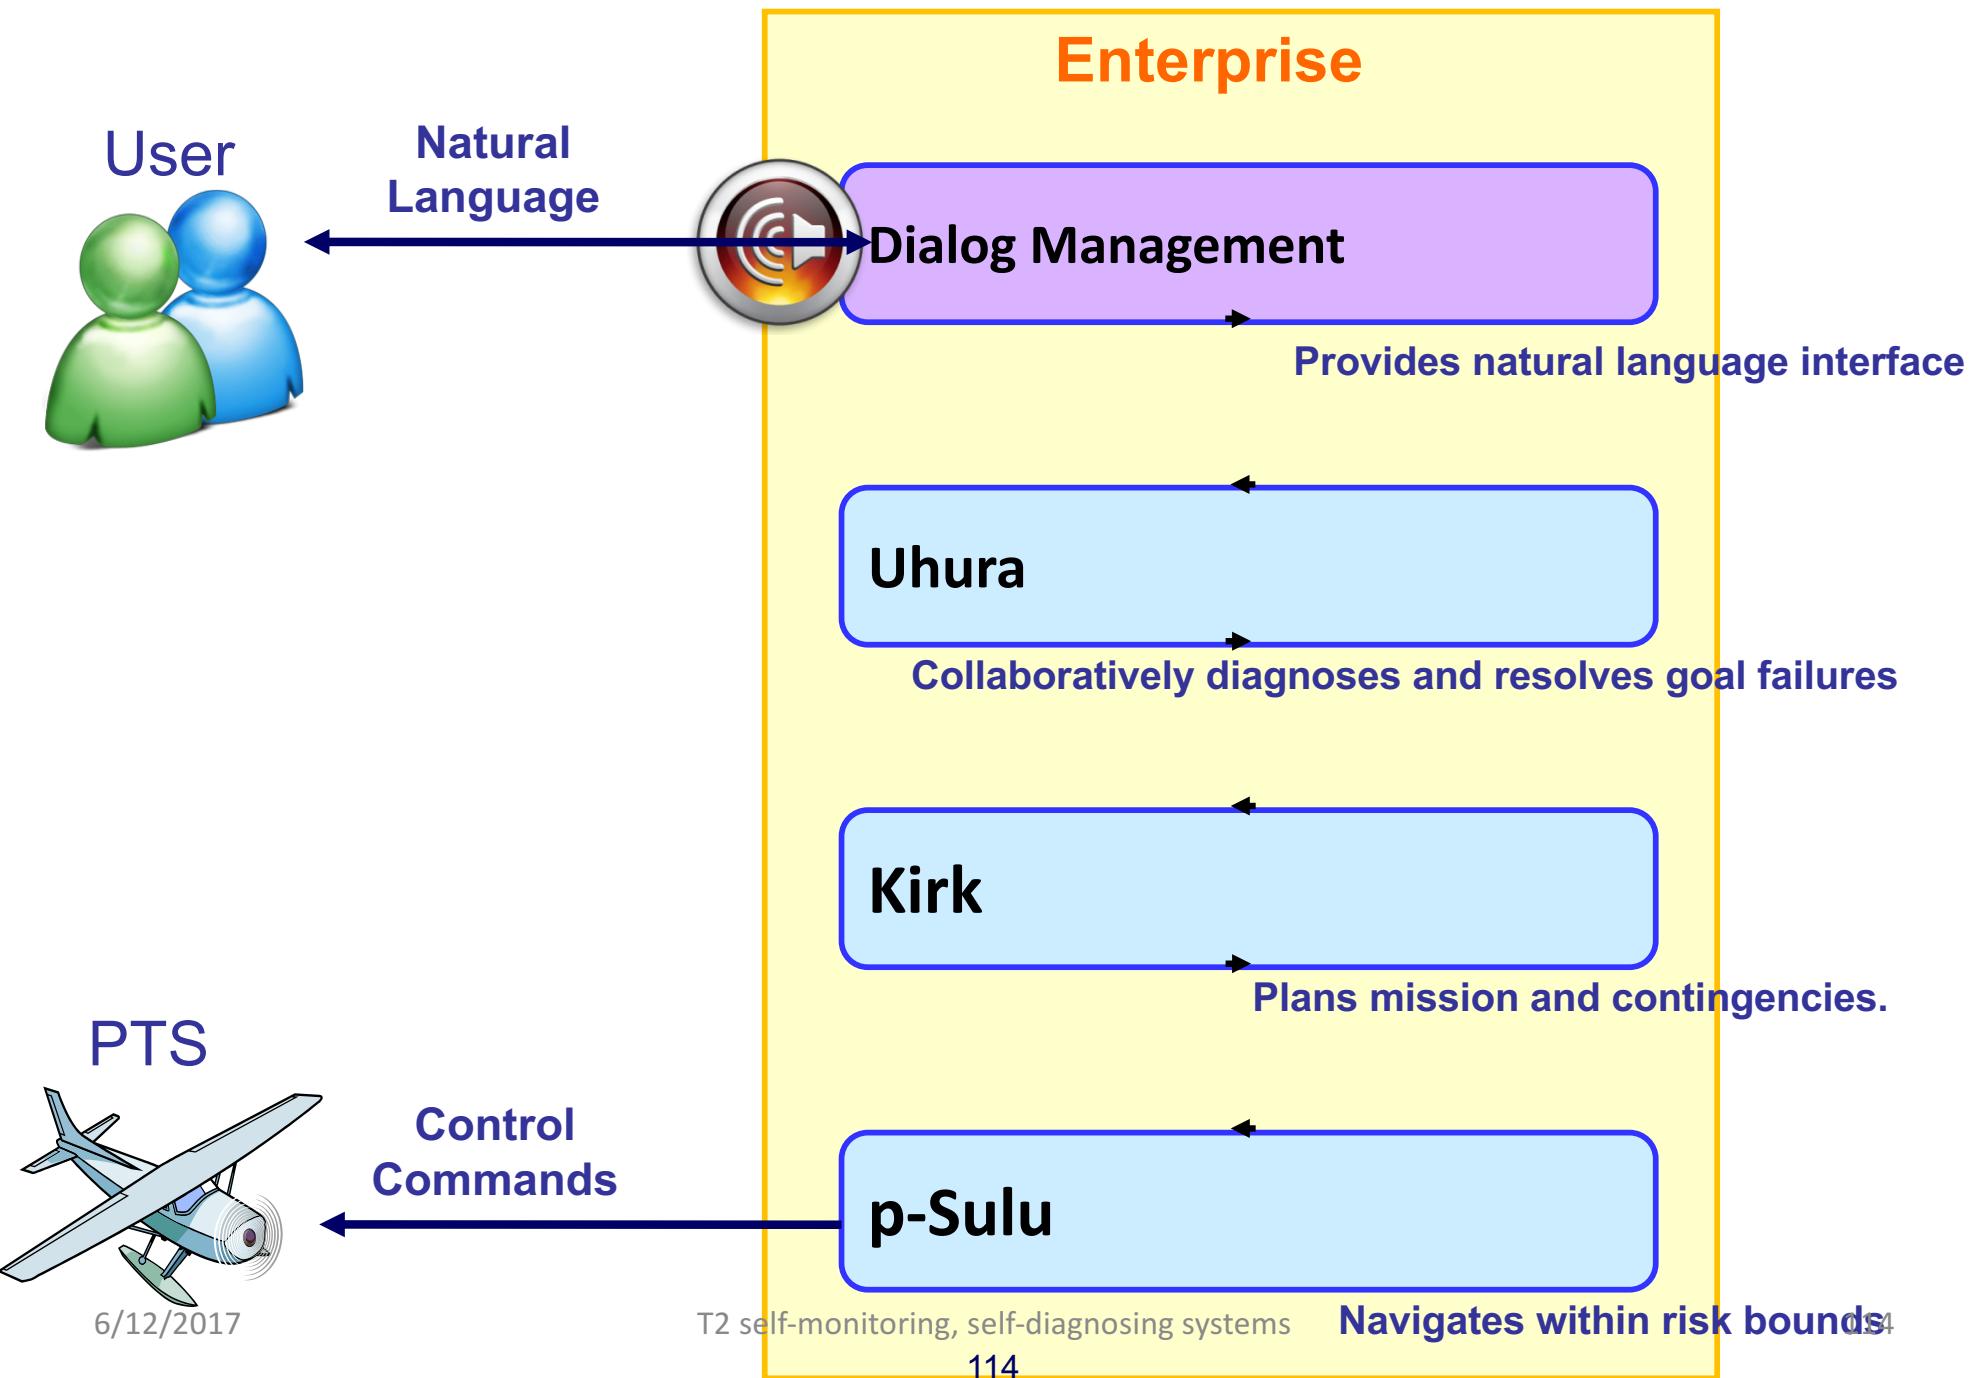
- As humans, we **often plan** to do **too much** . . .
- . . . our study plans, working schedules, and travel itineraries are often **over-subscribed**.
 - ‘I want to see the new movie premiere tonight!’
 - ‘Sorry, you cannot. You need submit your project by **mid night**, and it will keep you busy until **11pm**.’

Approach

View **human robot collaboration** as a form of collaborative diagnosis.

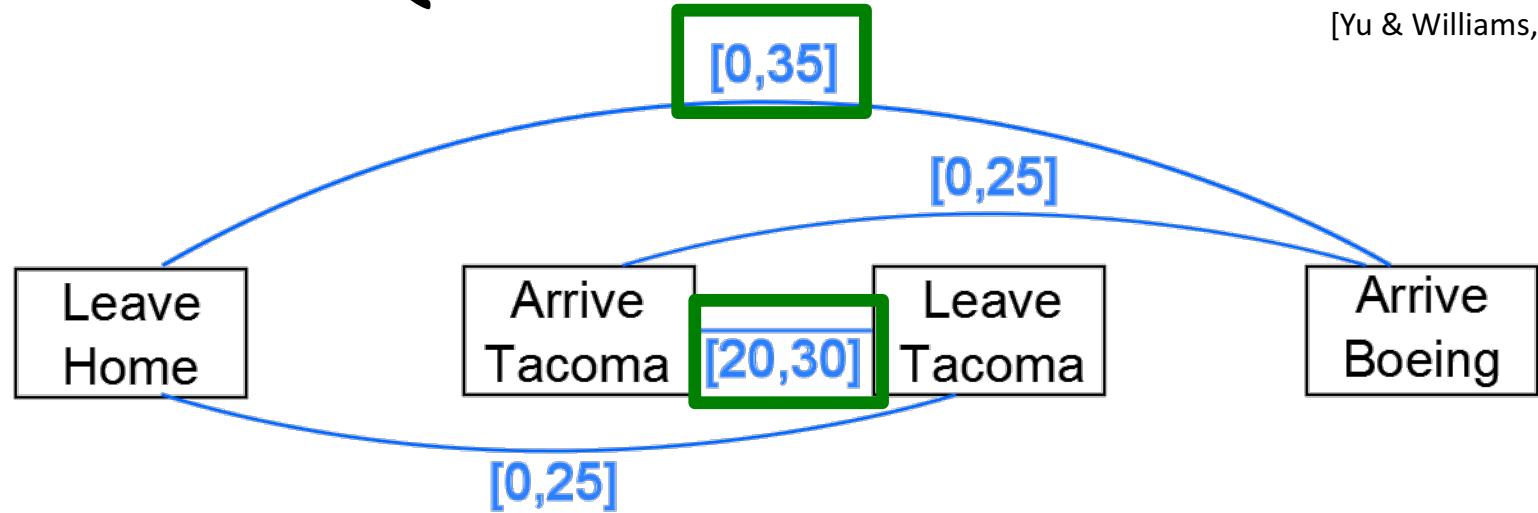
1. Identify **user goals** that are **over-subscribed**, within a temporal plan (QSP).
 2. Explain their source of **inconsistency**, using **conflicts**.
 3. Use constraint **suspension** and **kernels** to suggest goals to drop.
- Use **continuous relaxation** to **weaken** rather than drop user **goals**.
 - Generalize OpSat to **disjunctive linear programs**.





Uhura Translates User Goals to a Qualitative State Plan

[Yu & Williams, IJCAI 13]



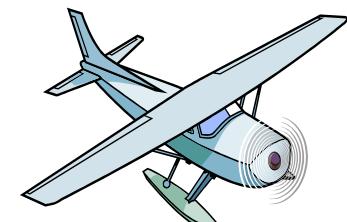
I want to go to Boeing in 35 minutes.
I want to stop at Tacoma airport for 20 to 30 minutes.
I've applied for a slot at Tacoma.
That's all.

User



6/12/2017

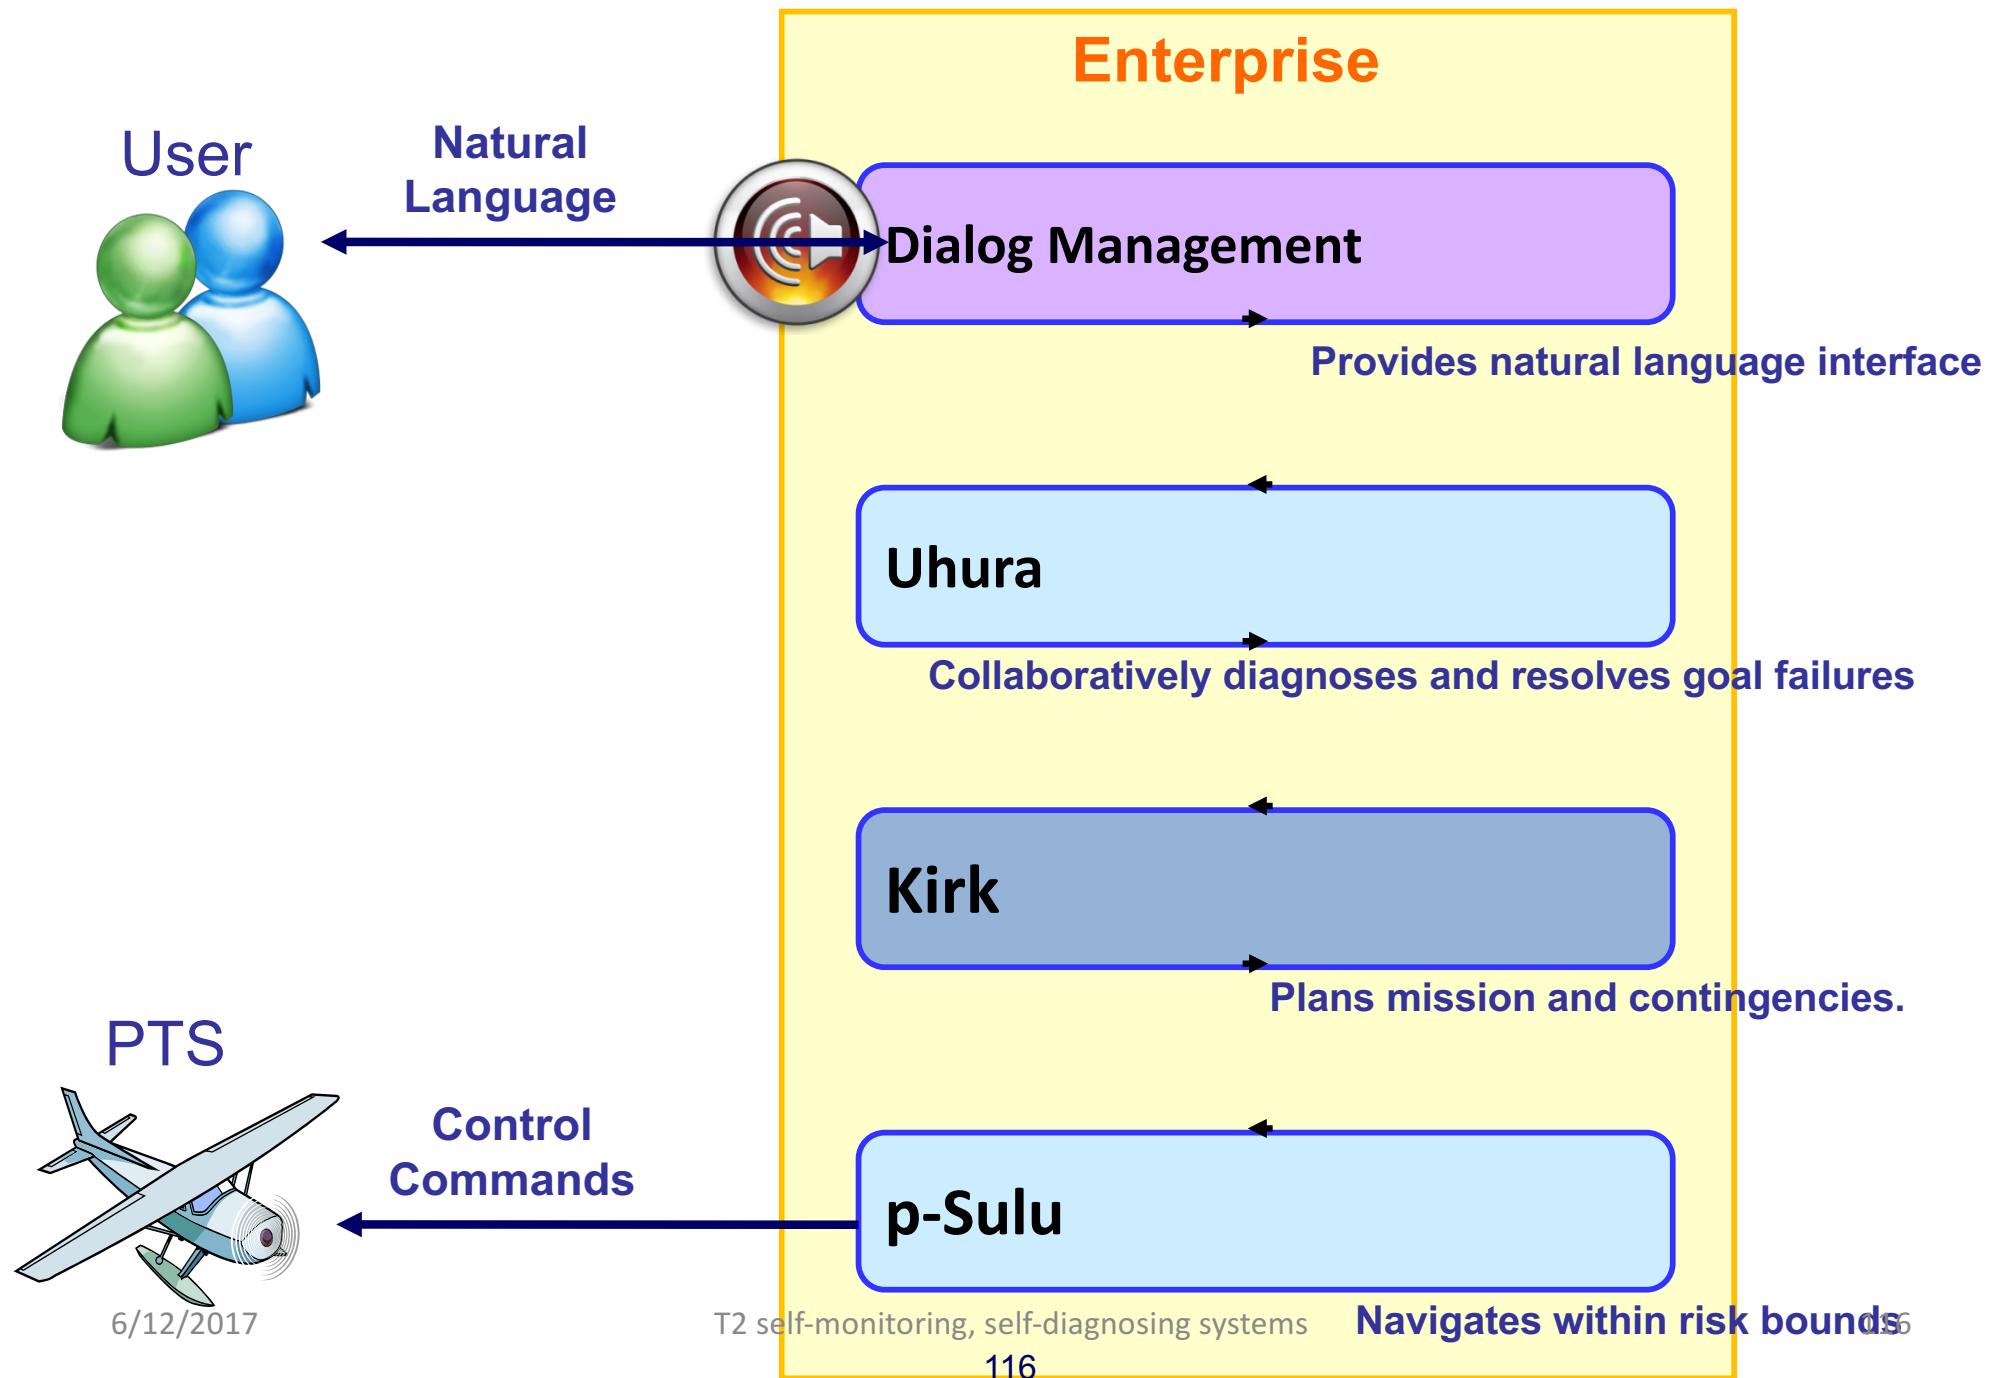
PAV



T2 self-monitoring, self-diagnosing systems

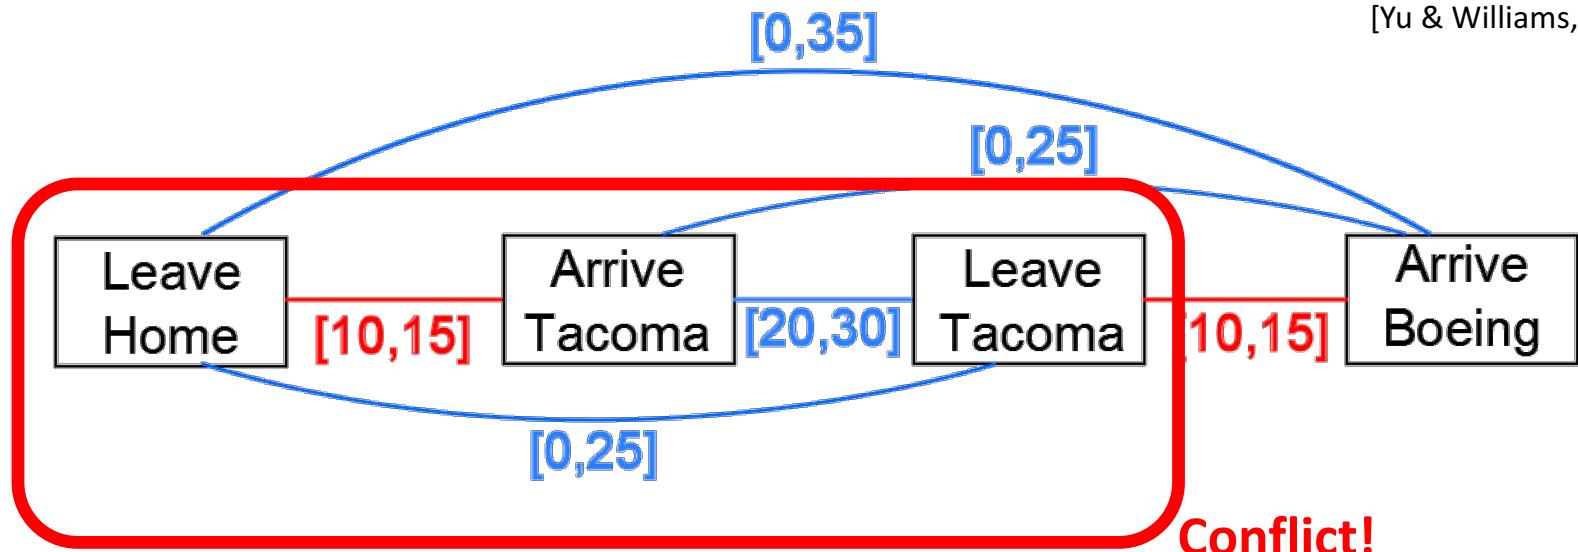
115

Kirk searches for an executable travel plan



Uhura Collaborates with User to Diagnose and Repair Plan Failure

[Yu & Williams, IJCAI 13]



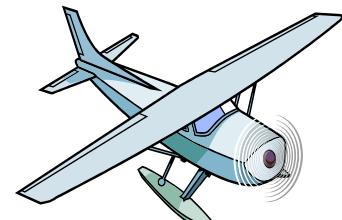
User



6/12/2017

I have a problem, navigating around the storm will delay us.

PAV

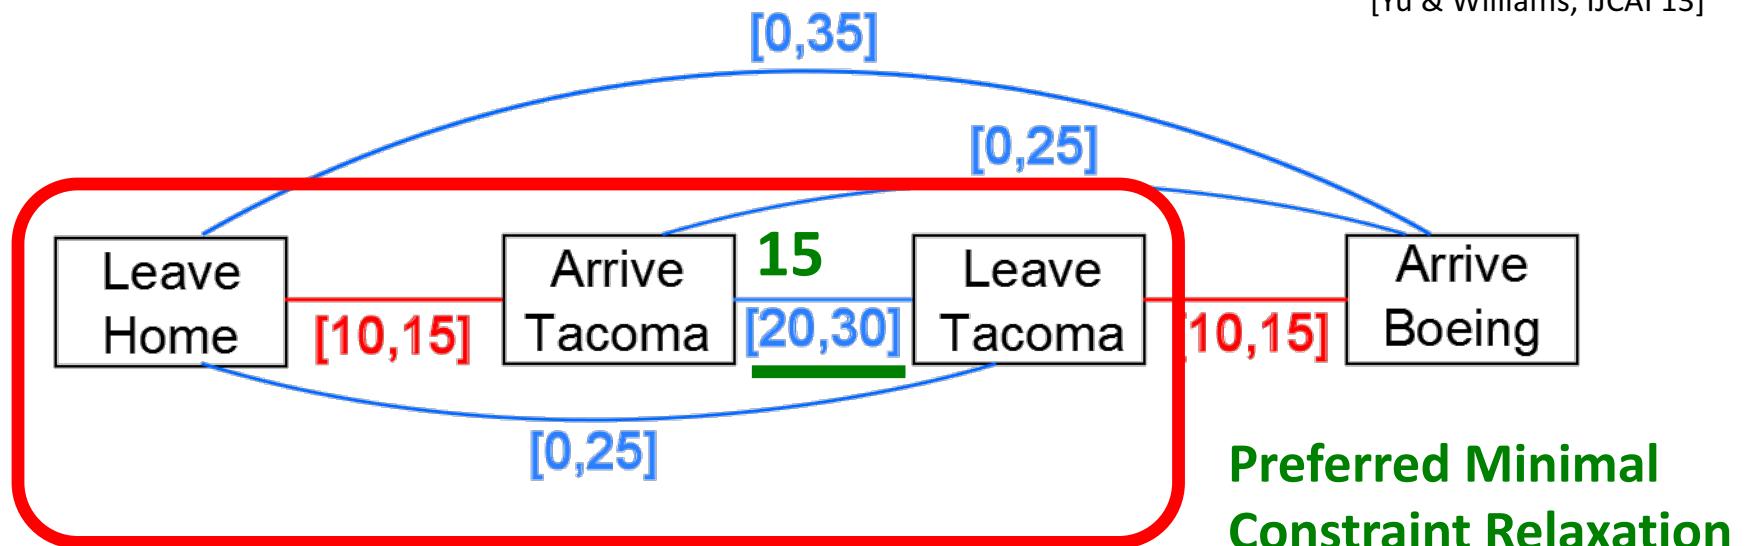


T2 self-monitoring, self-diagnosing systems

117

Uhura Collaborates with User to Diagnose and Repair Plan Failure

[Yu & Williams, IJCAI 13]



Can you shorten your stay at Tacoma to 15 minutes.

User



6/12/2017

PAV

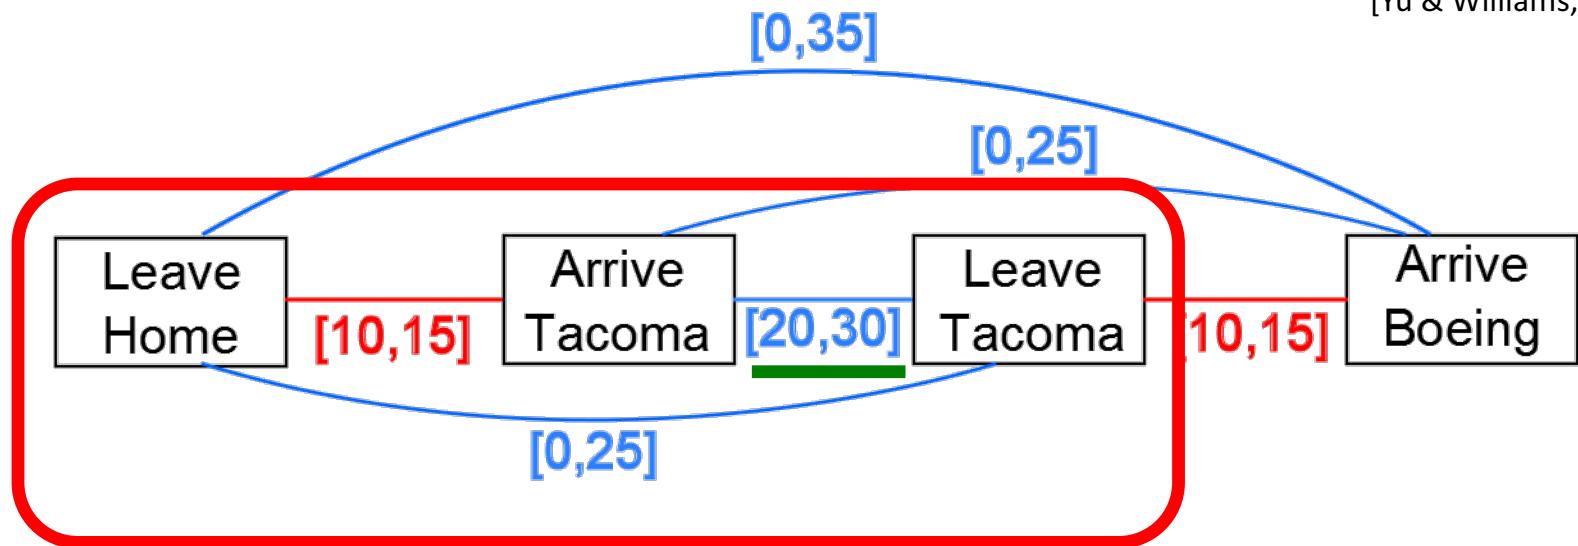


T2 self-monitoring, self-diagnosing systems

118

Uhura Collaborates with User to Diagnose and Repair Plan Failure

[Yu & Williams, IJCAI 13]



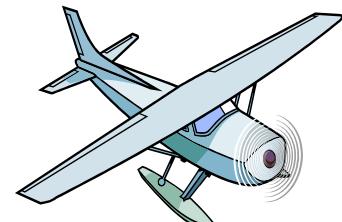
No, I cannot do that.

User



6/12/2017

PAV

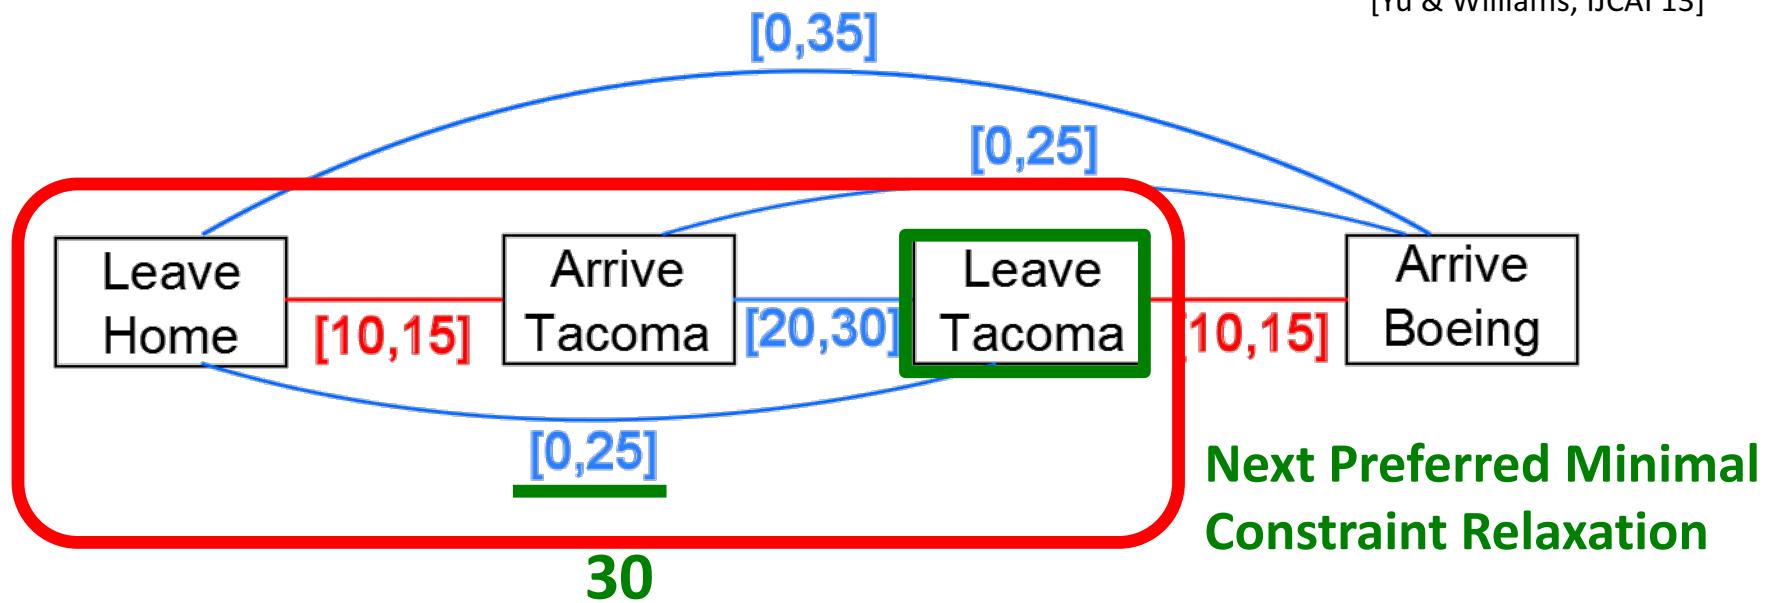


T2 self-monitoring, self-diagnosing systems

119

Uhura Collaborates with User to Diagnose and Repair Plan Failure

[Yu & Williams, IJCAI 13]

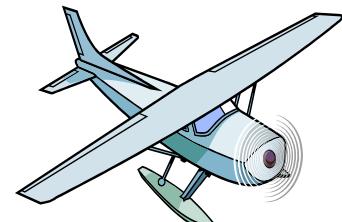


User



6/12/2017

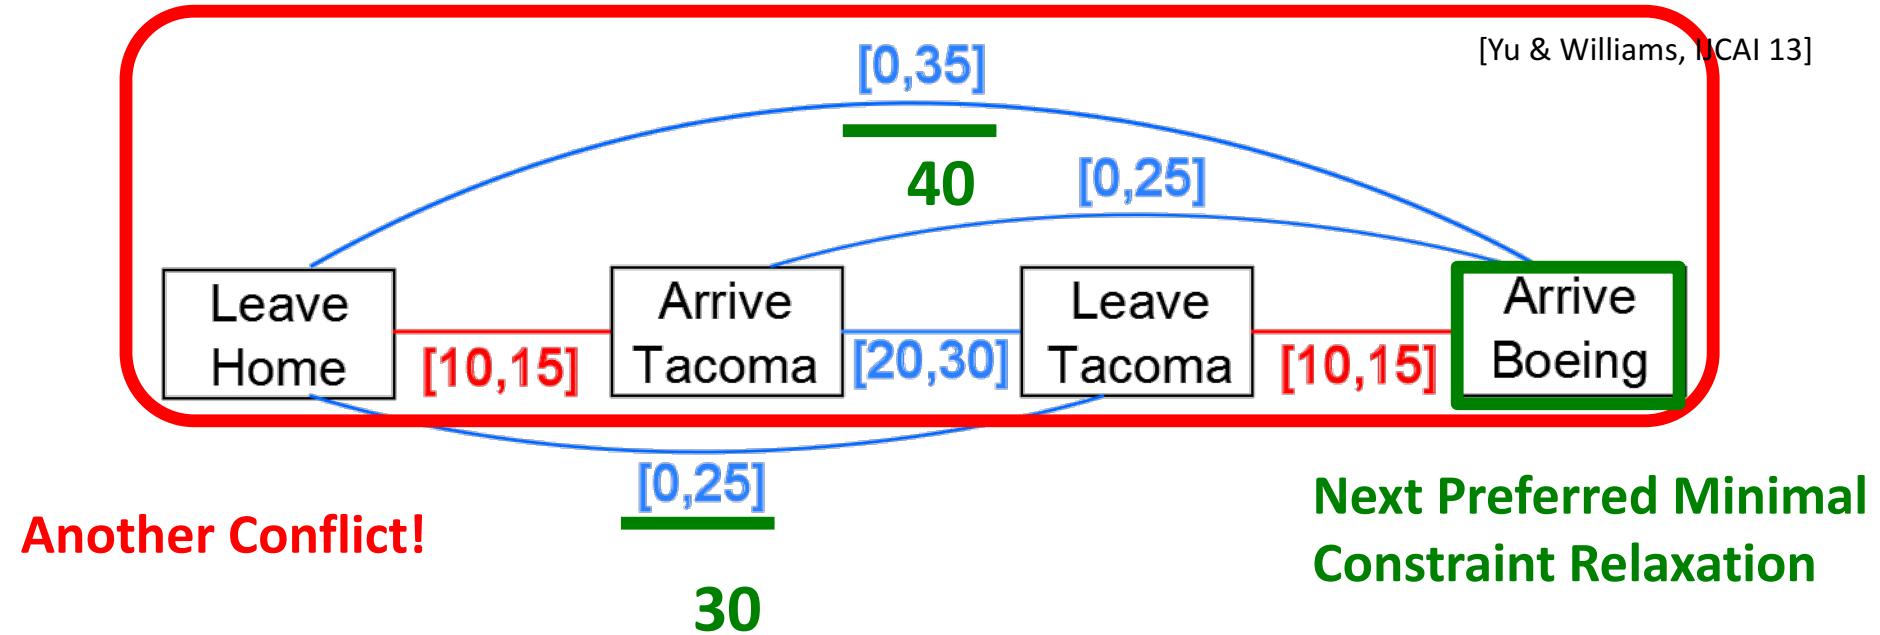
PAV



120

T2 self-monitoring, self-diagnosing systems

Uhura Collaborates with User to Diagnose and Repair Plan Failure

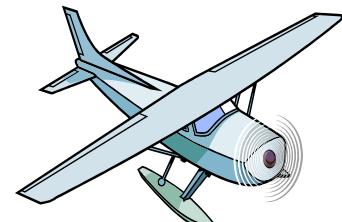


User



6/12/2017

PAV

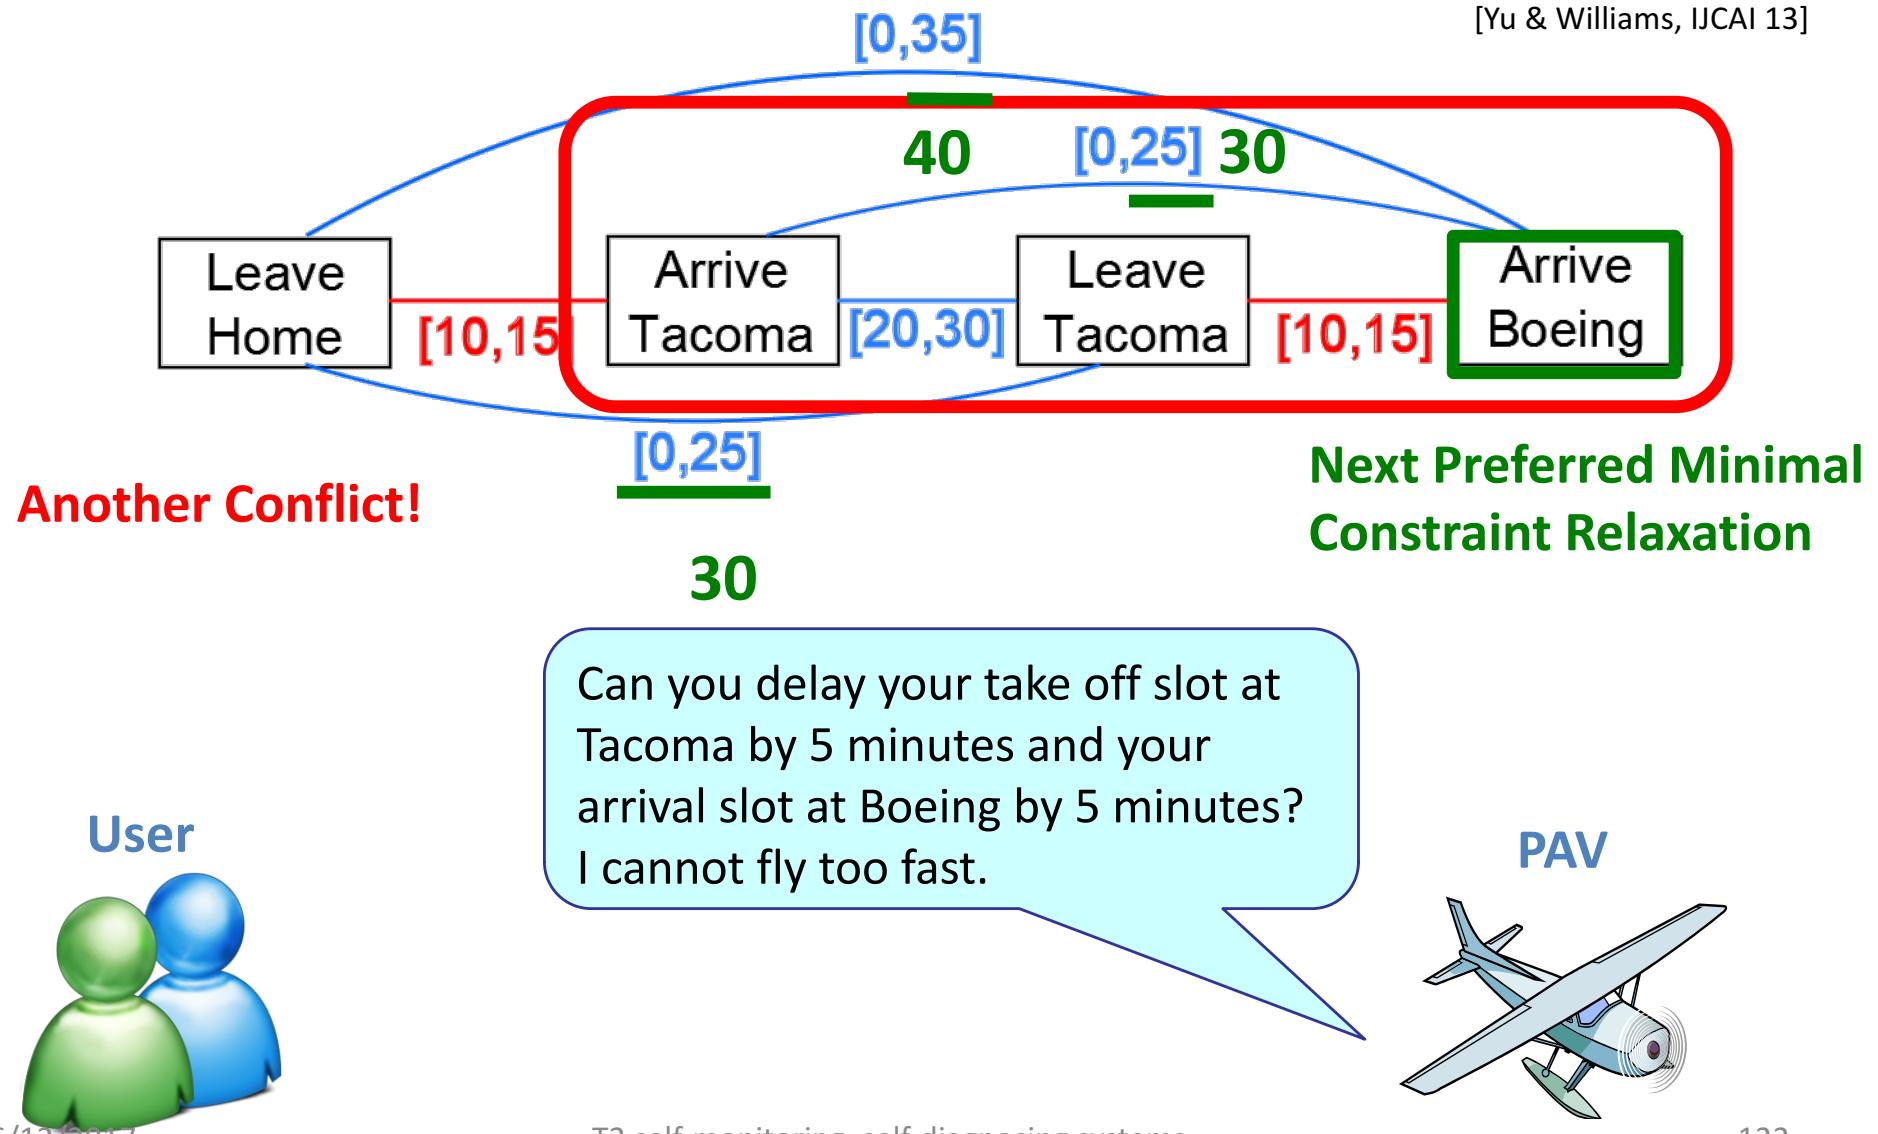


T2 self-monitoring, self-diagnosing systems

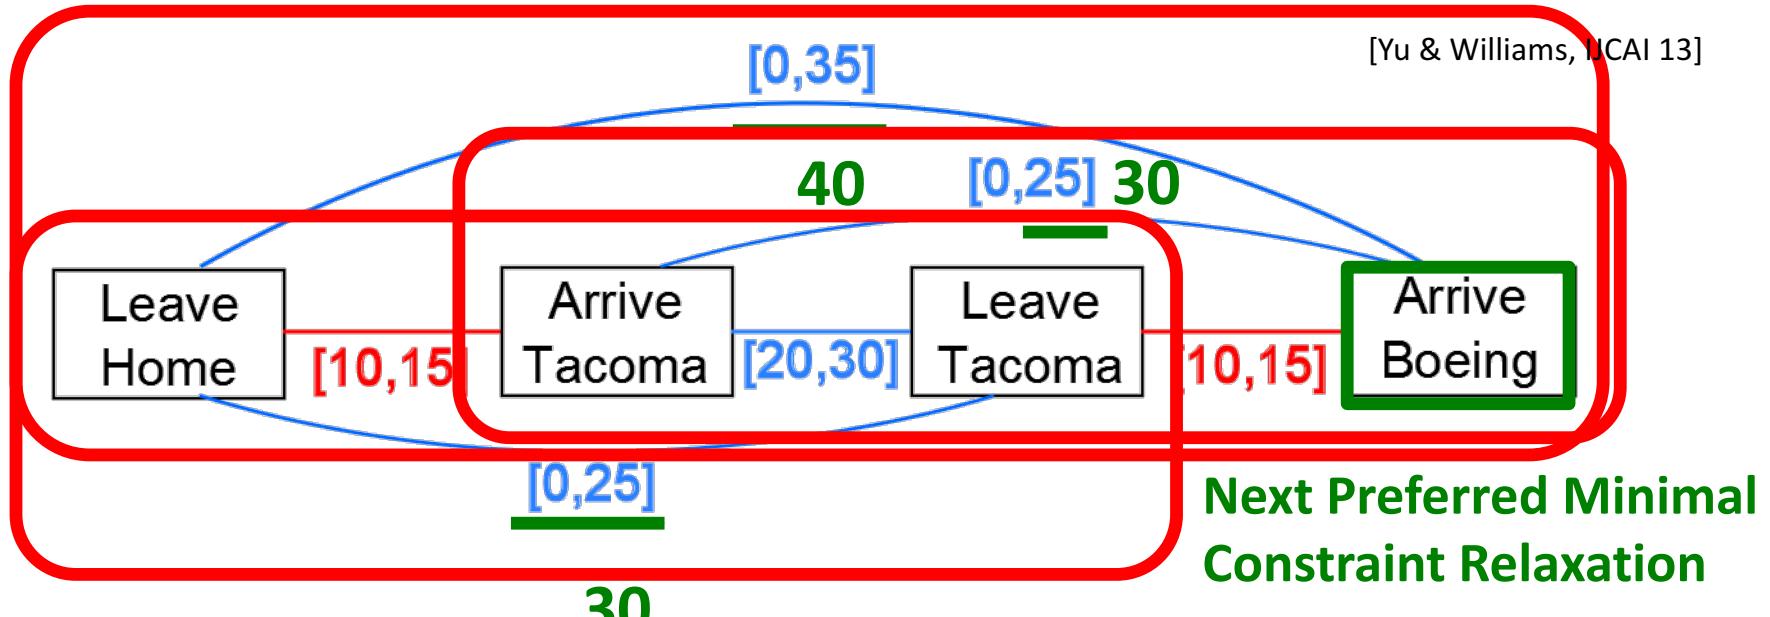
121

Uhura Collaborates with User to Diagnose and Repair Plan Failure

[Yu & Williams, IJCAI 13]



Uhura Collaborates with User to Diagnose and Repair Plan Failure

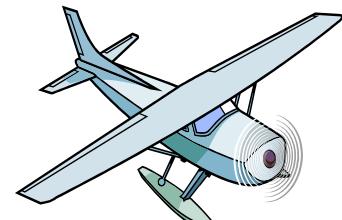


User



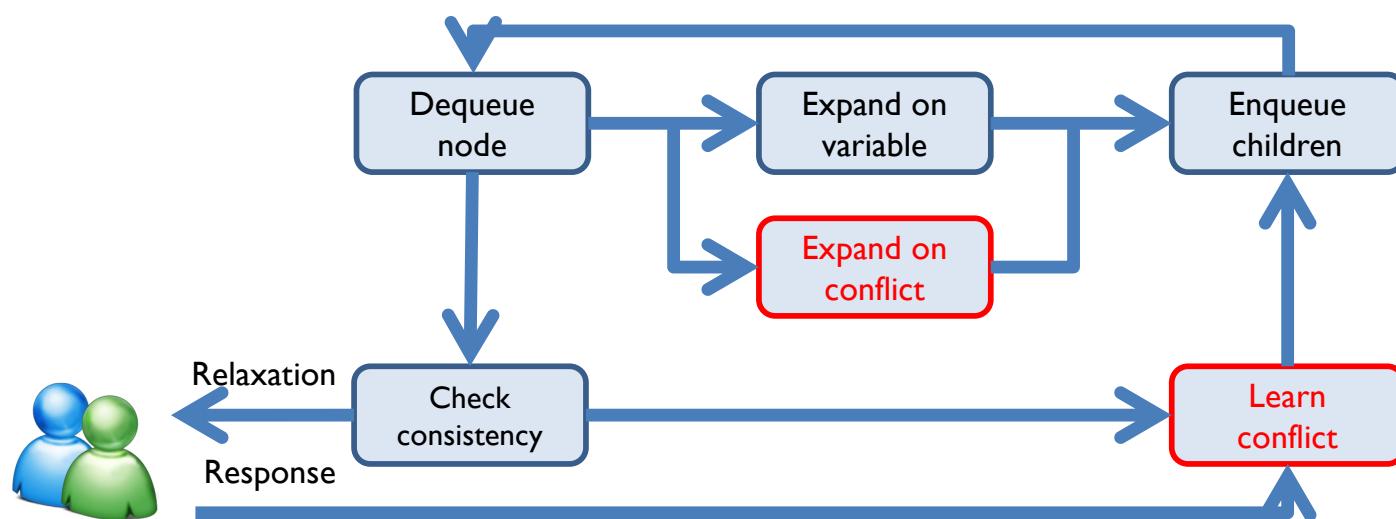
Yes, that's fine.

PAV



Best-first, Conflict-directed Relaxation

- **Generate:**
 - Enumerates **minimal relaxations** in **best-first order**.
 - Generalizes **conflict-directed A*** to **continuous conflicts**.
 - Framed as a **disjunctive linear program** with only **positive literals**.
- **Test:**
 - Scheduler with conflict extraction.



Goals can be relaxed in many ways



Self-Monitoring, Self-Diagnosing Systems

- Self-monitoring is essential to languages and architectures used to build robust robotic systems.
- Causal link monitoring enables robots to anticipate and adapt to sub-goal failures.
- Probabilistic mode estimation enable robots to diagnose component faults that lead to action failure.
- Conflict-direct search solves large decision problems by learning why hypotheses fail.
- Conflict-directed relaxation can be used to collaboratively resolve inconsistent plans and programs.

Questions?