



The Vehicle Routing Problem with Time Windows

Dr Philip Kilby | Team Leader, Optimisation Applications and Platforms
June 2017

www.data61.csiro.au



Outline

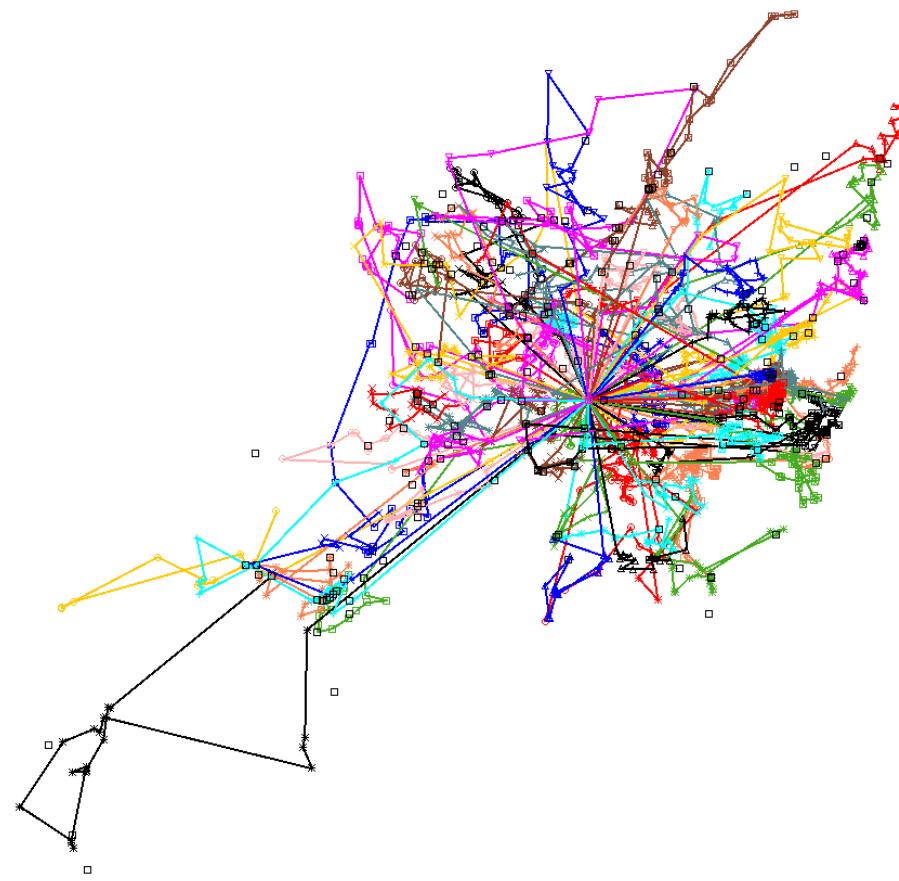


- Problem Description
- Solving the VRP
 - Construction
 - Local Search
 - Meta-heuristics
 - Variable Neighbourhood Search
 - Including Large Neighbourhood Search
- CP101
- A CP model for the VRP

Vehicle routing problem

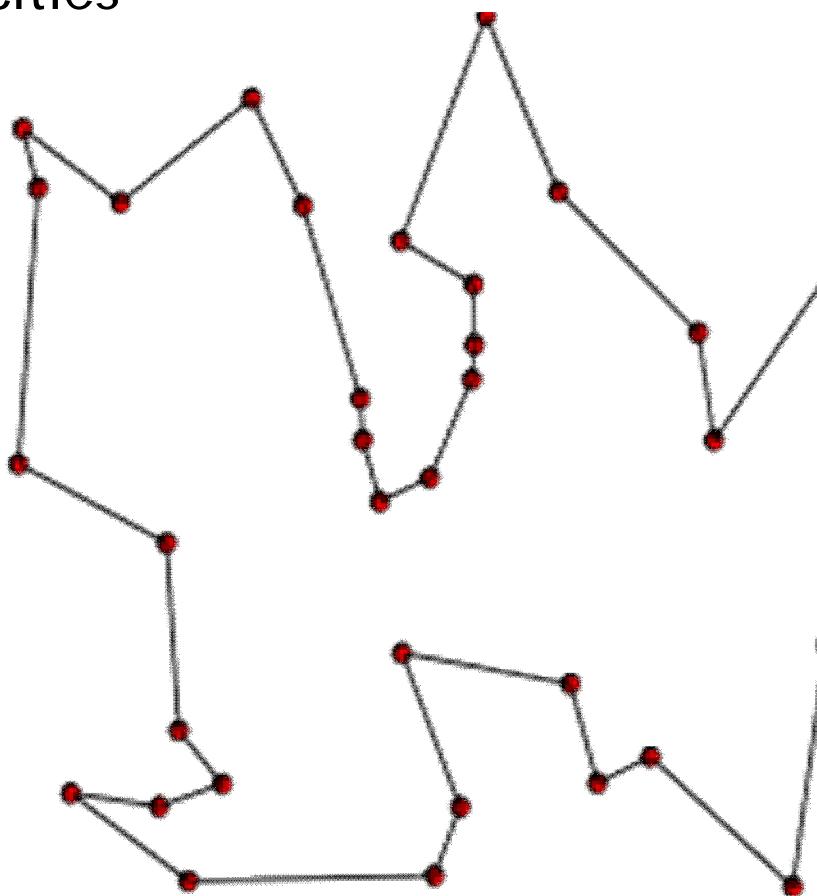


Given a set of customers, and a fleet of vehicles to make deliveries,
find a set of routes that services all customers at minimum cost



Travelling Salesman Problem

- Find the tour of minimum cost that visits all cities



Why study the VRP?



- It's hard: it exhibits all the difficulties of comb. opt.
- It's useful:
 - The logistics task is 9% of economic activity in Australia
 - Logistics accounts for 10% of the selling price of goods



Why study the VRP in Robotics?



Appears as a sub-problem



Task allocation to agents

- Multiple agents with multiple tasks
- Best allocation minimizes cost

Scheduling with setup costs

- Can be modelled as a VRPTW



Vehicle Routing Problem



For each customer, we know

- Quantity required
- The cost to travel to every other customer

For the vehicle fleet, we know

- The number of vehicles
- The capacity

We must determine which customers each vehicle serves, and in what order, to minimise **cost**

Vehicle Routing Problem



Objective function

- In academic studies, usually a combination:
 - First, minimise number of routes
 - Then minimise total distance or total time
- In real world
 - A combination of time and distance
 - Must include vehicle- and staff-dependent costs
 - Usually vehicle numbers are fixed
 - Includes “preferences” – like pretty routes

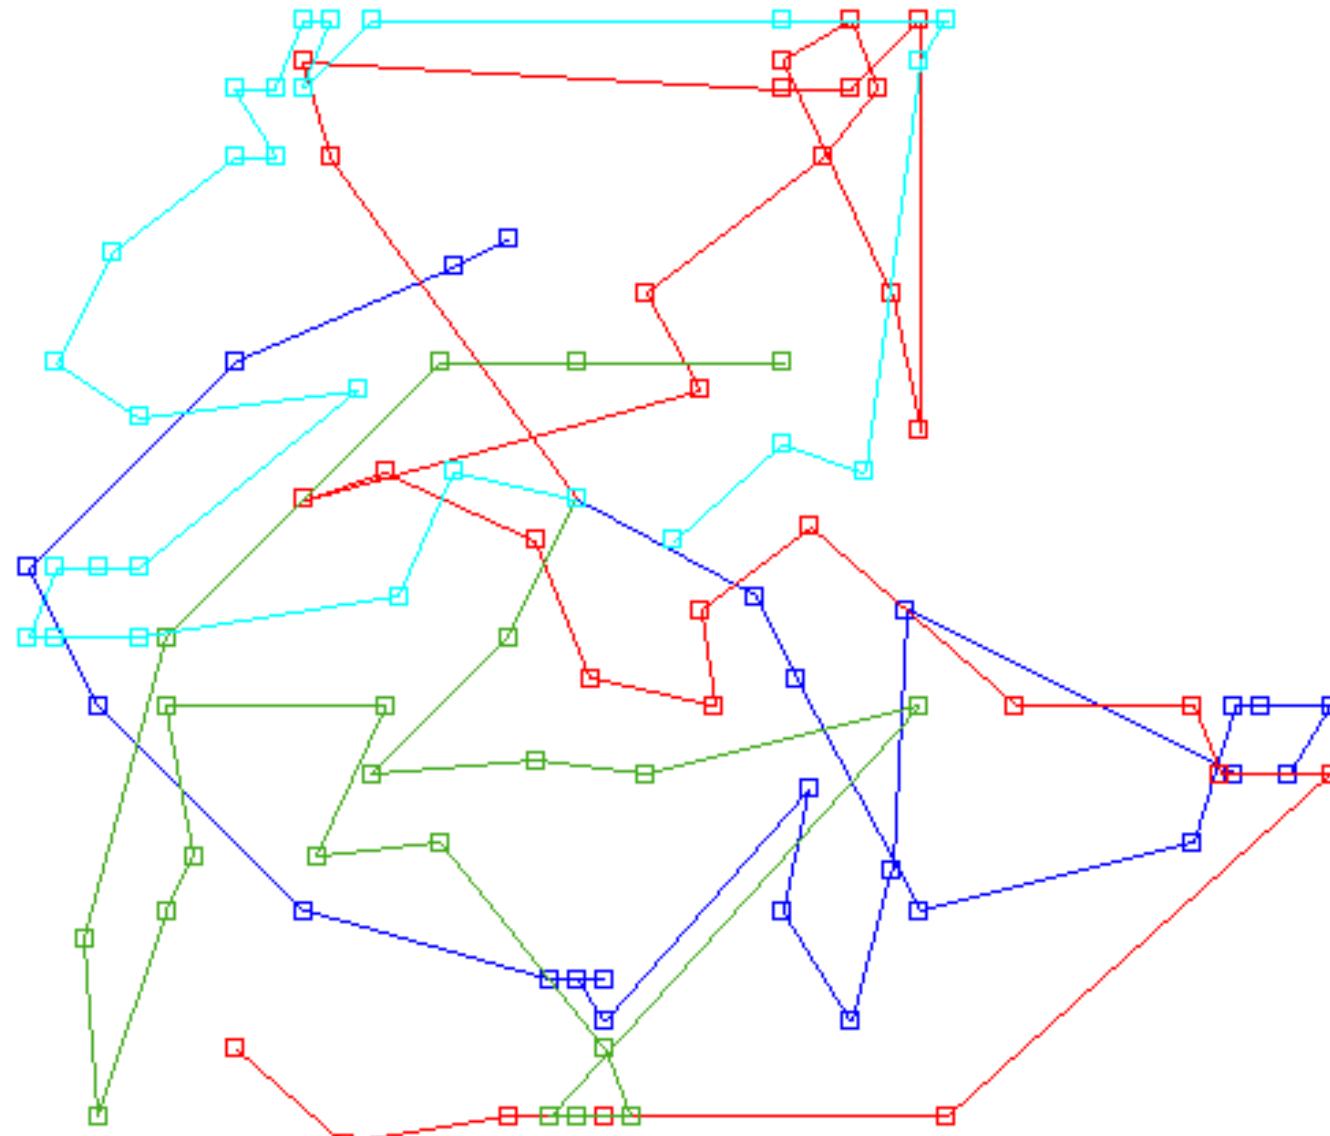
Time window constraints



Vehicle routing with Constraints

- Time Window constraints
 - A window during which service can start
 - E.g. only accept delivery 7:30am to 11:00am
 - Additional input data required
 - Duration of each customer visit
 - Time between *each pair* of customers
 - (Travel time can be vehicle-dependent or time-dependent)
 - Makes the route harder to visualise

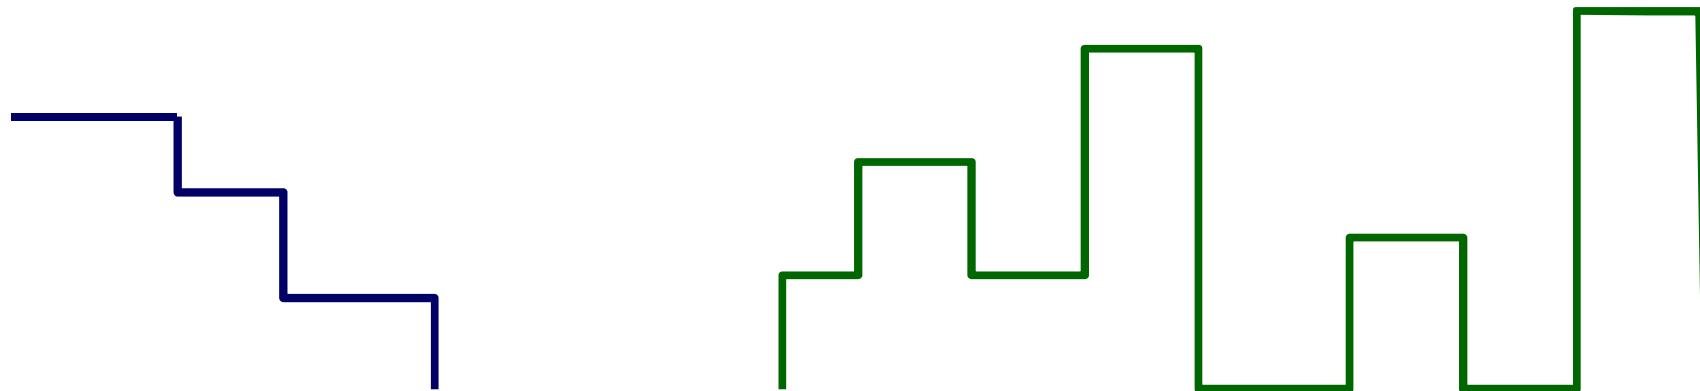
Time Window constraints



Pickup and Delivery problems



- Most routing considers delivery to/from a depot (depots)
- Pickup and Delivery problems consider FedEx style problem:
 - *pickup at location A, deliver to location B*
 - Load profile:



Other variants

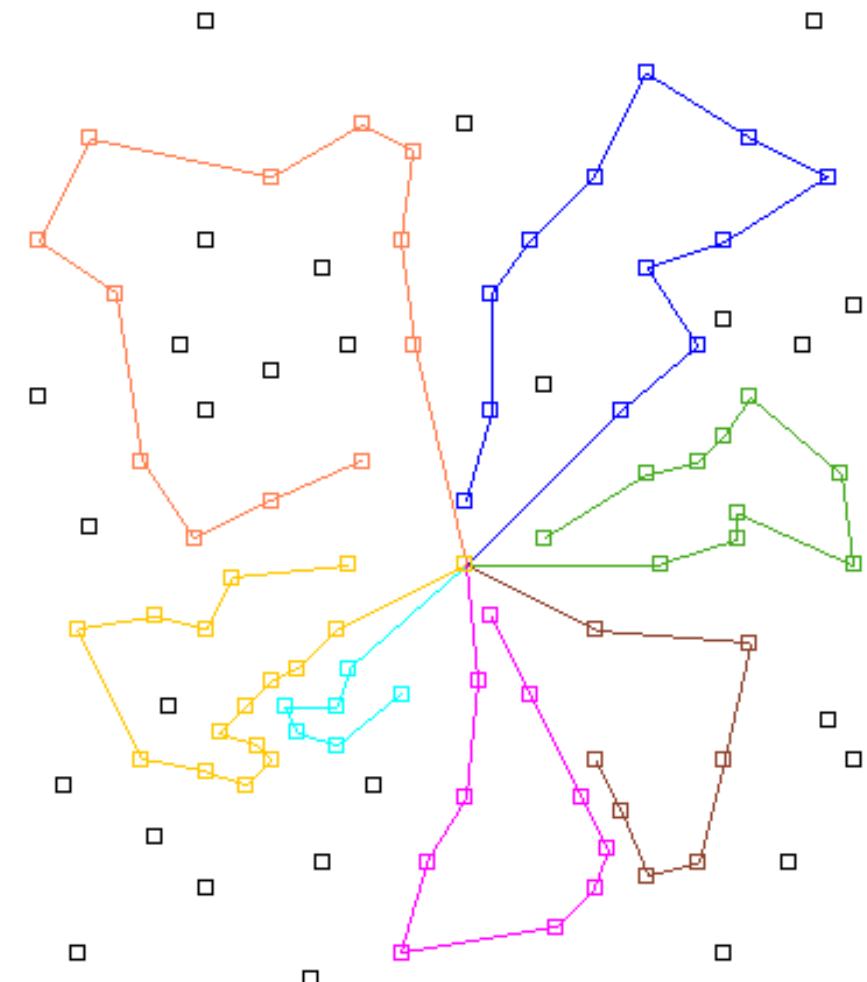


Profitable tour problem

- Not all visits need to be completed
- Known profit for each visit
- Choose a subset that gives maximum *profit* = (*revenue from visits*) – (*routing cost*)

Orienteering Problem

- Maximum revenue in limited time



VRP meets the real world



Many groups now looking at real-world constraints

Rich Vehicle Routing Problem

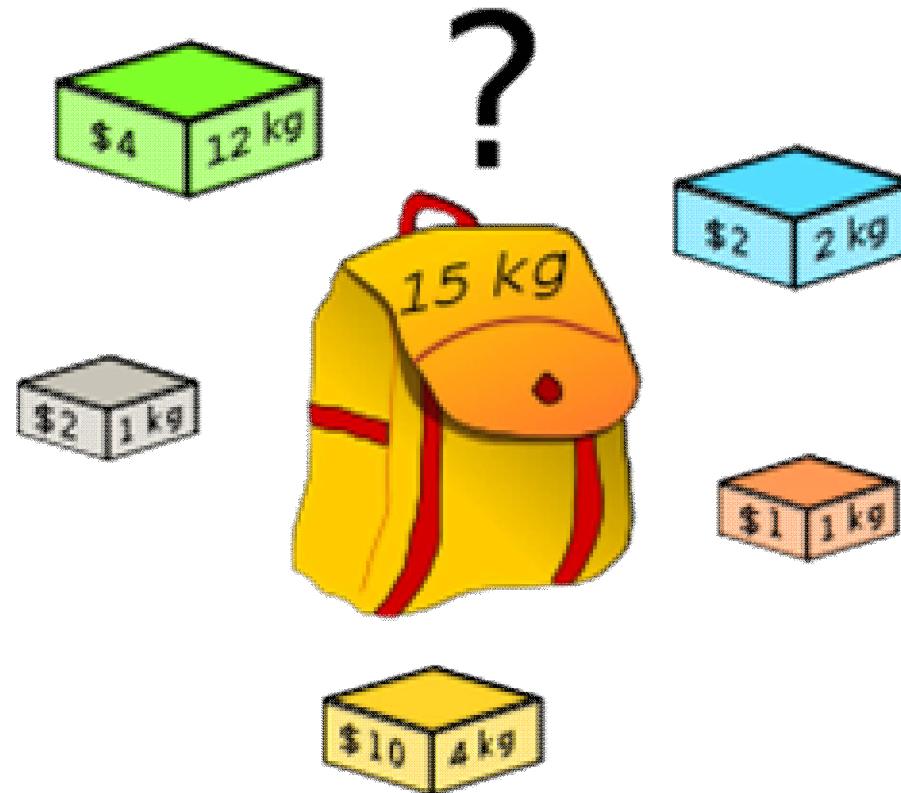
- Attempt to model constraints common to many real-life enterprises
 - Multiple Time windows
 - Multiple Commodities
 - Multiple Depots
 - Heterogeneous vehicles
 - Compatibility constraints
 - Goods for customer *A* {must | can't} travel with goods from customer *B*
 - Goods for customer *A* {must | can't} travel on vehicle *C*

VRP as an instance



VRP is a Combinatorial Optimization problem

- Others include
 - Scheduling
 - Assignment
 - Bin Packing





Solving VRPs

Solution Methods



Exact:

- Integer Programming or Mixed Integer Programming
- Constraint Programming

Heuristic:

- Construct
- Improve
 - Local Search
 - Meta-heuristics

Exact Methods



VRP:

- MIP: Can only solve problems with 100-150 customers
- CP: Similar size

ILP



$$\text{minimise} : \sum_{i,j} c_{ij} \sum_k x_{ijk}$$

$x_{ijk} = 1$ if vehicle k travels directly from i to j

subject to

$$\sum_i \sum_k x_{ijk} = 1 \quad \forall j$$

Exactly one vehicle in

$$\sum_j \sum_k x_{ijk} = 1 \quad \forall i$$

Exactly one vehicle out

$$\sum_j \sum_k x_{ihk} - \sum_j \sum_k x_{hjk} = 0 \quad \forall k, h$$

It's the same vehicle

$$\sum_i q_i \sum_j x_{ijk} \leq Q_k \quad \forall k$$

Capacity constraint

$$\sum x_{ijk} = |S| - 1 \quad S \subseteq P(N), 0 \notin S \quad \text{Subtour elimination}$$

$$x_{ijk} \in \{0,1\}$$

ILP



$$\text{minimise} : \sum_{i,j} c_{ij} \sum_k x_{ijk}$$

subject to

$$\sum_i \sum_k x_{ijk} = 1 \quad \forall j$$

$$\sum_j \sum_k x_{ijk} = 1 \quad \forall i$$

$$\sum_j \sum_k x_{ihk} - \sum_j \sum_k x_{hjk} = 0 \quad \forall k, h$$

$$\sum_i q_i \sum_j x_{ijk} \leq Q_k \quad \forall k$$

$$\sum x_{ijk} = |S| - 1 \quad S \subseteq P(N), 0 \notin S$$

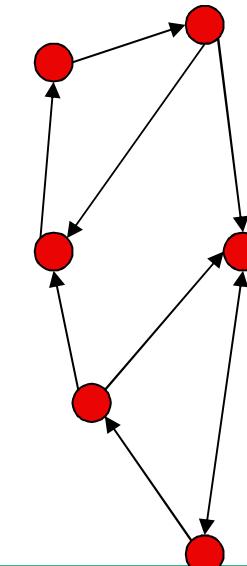
$$x_{ijk} \in \{0,1\}$$

Advantages

- Can find optimal solution

Disadvantages

- Only works for small problems
- One extra constraint → back to the drawing board
- S is huge



Depot

ILP – Column Generation



Columns
represent routes

Column/route cost c_k

Rows represent
customers

Array entry $a_{ik} = 1$ iff
customer i is
covered by route k

	89	76	99	45	32	
1	1	1	1	0	0	...
2	0	1	1	0	1	...
3	0	0	0	0	0	...
4	1	0	1	1	0	...
5	1	0	0	0	0	...

Column Generation



- Decision var x_k : Use column k ?
- Column only appears if feasible ordering is possible
- Cost of best ordering is c_k
- Best order stored separately
- Master problem at right

$$\begin{aligned} & \min && \sum c_k x_k \\ & \text{subject to} \end{aligned}$$

$$\lambda_i \longrightarrow \sum_k a_{ik} x_k = 1 \quad \forall i$$
$$x_k \in \{0,1\}$$



Heuristics for the VRP



Heuristics:

Often variants of

- Construct
- Improve

Heuristics for the VRP



Construction by Insertion

- Start with an empty solution
- Repeat
 - Choose which customer to insert
 - Choose where to insert it

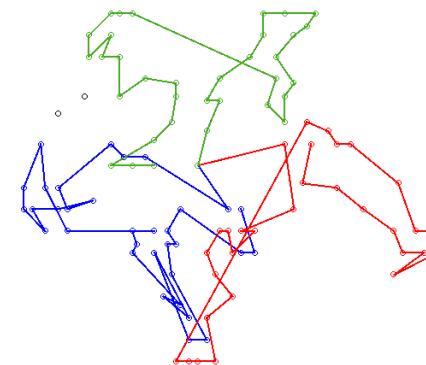
E.g. (Greedy)

- Choose the customer that increases the cost by the least
- Insert it in the position that increases the cost by the least

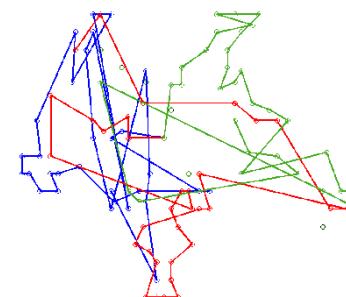
Solving the VRP the easy way



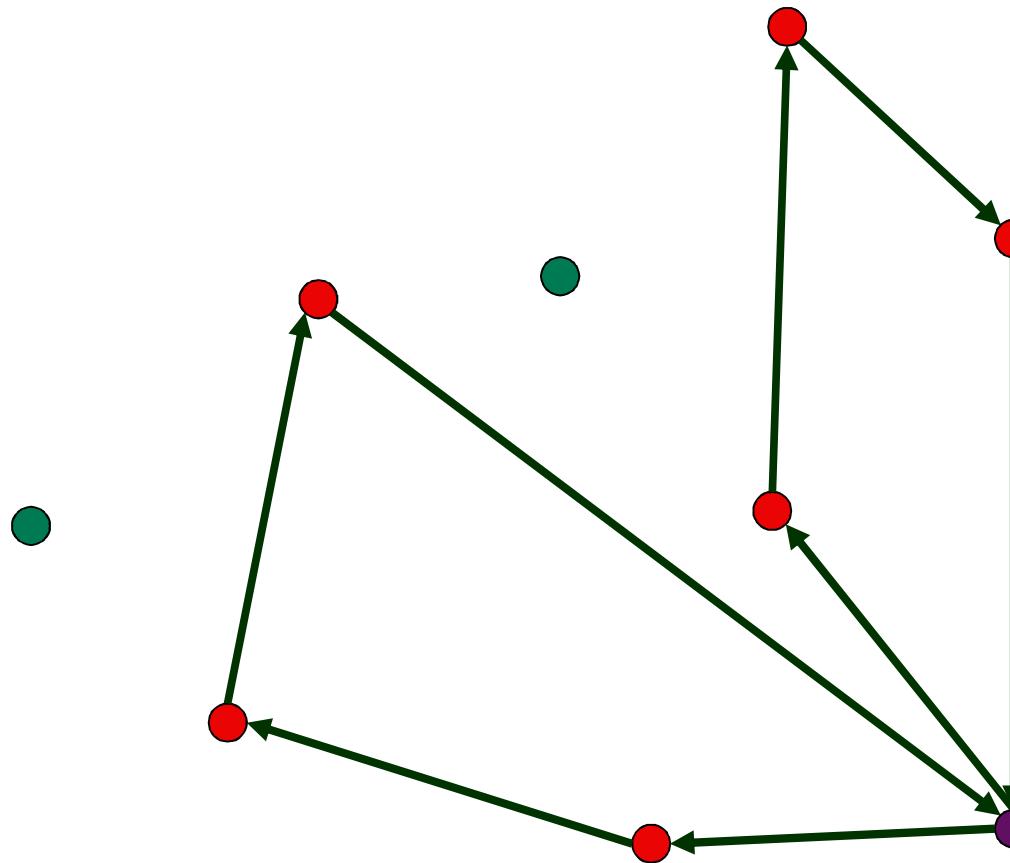
Insert methods



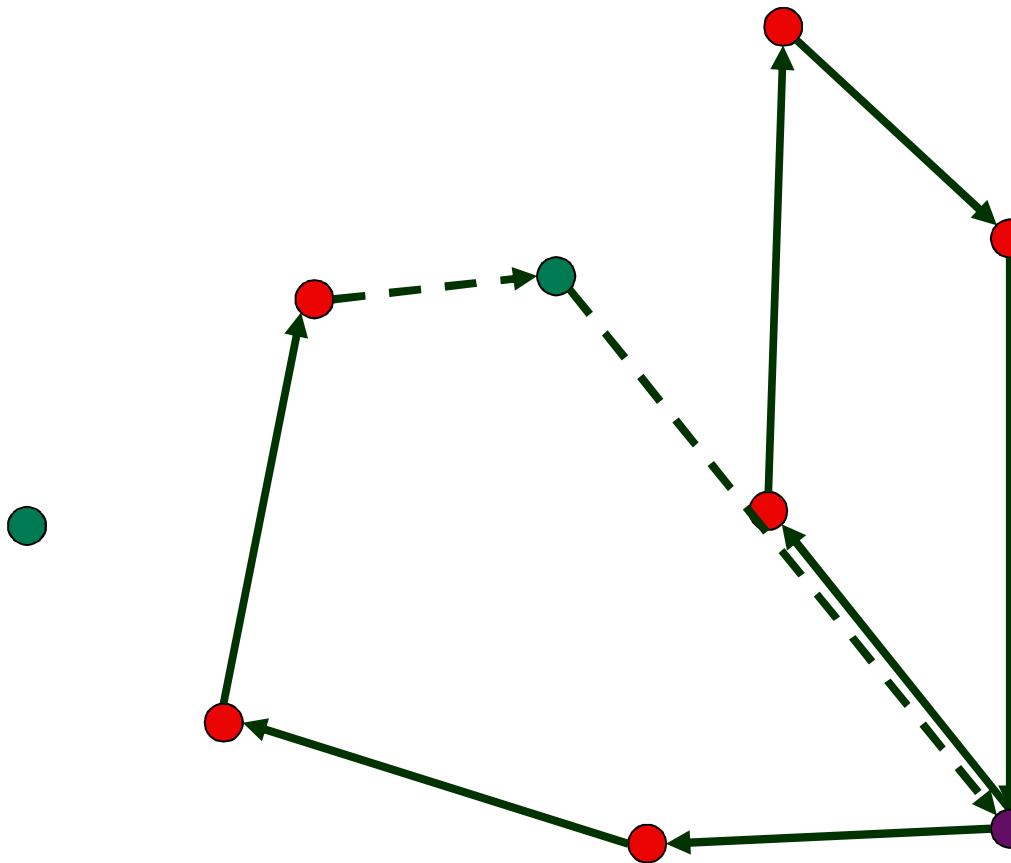
Order is important:



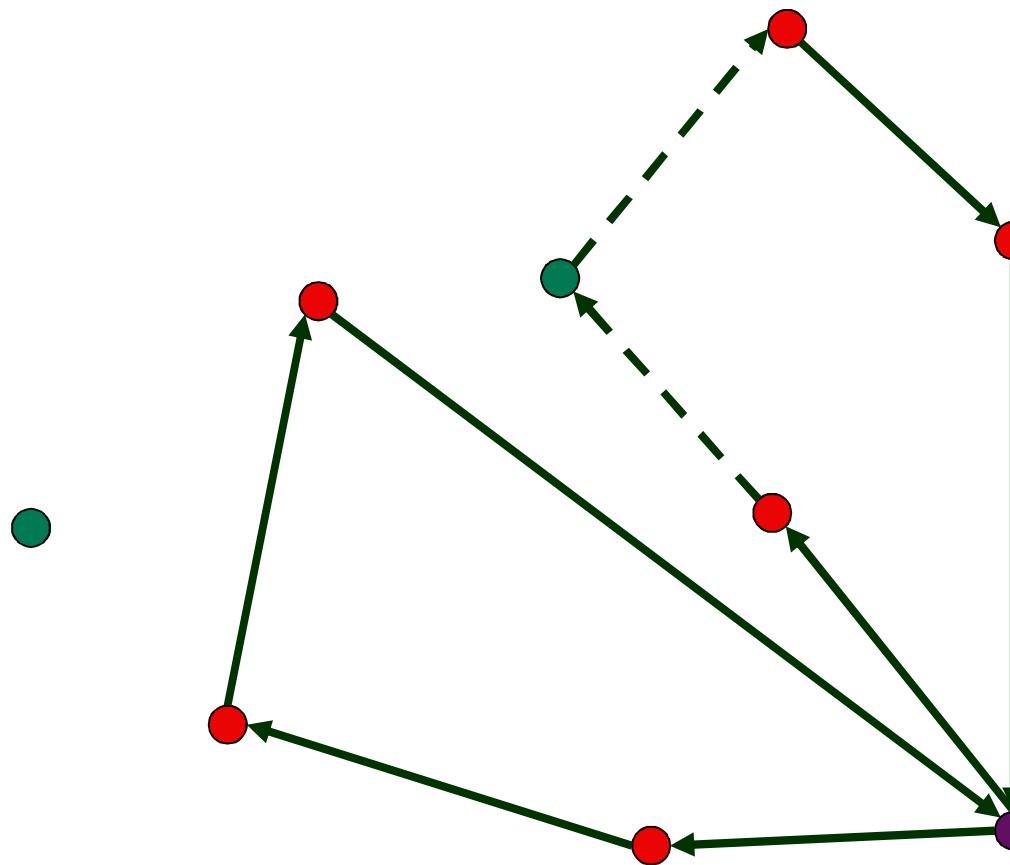
Regret



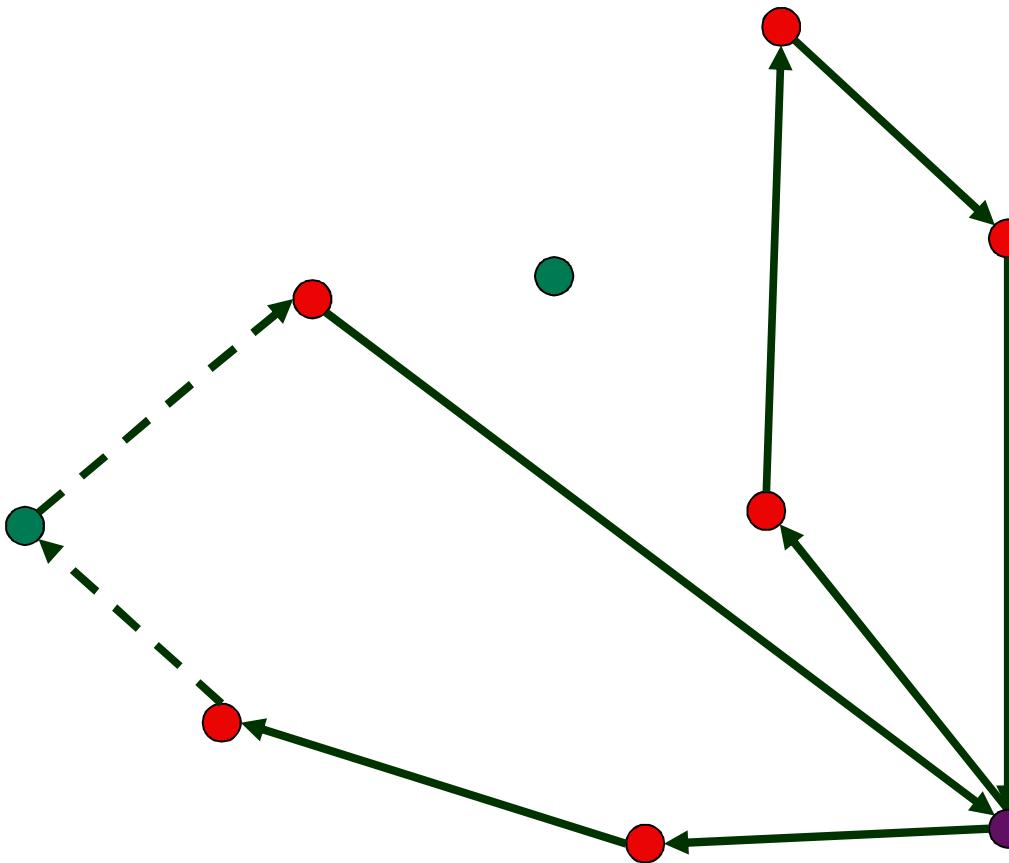
Regret



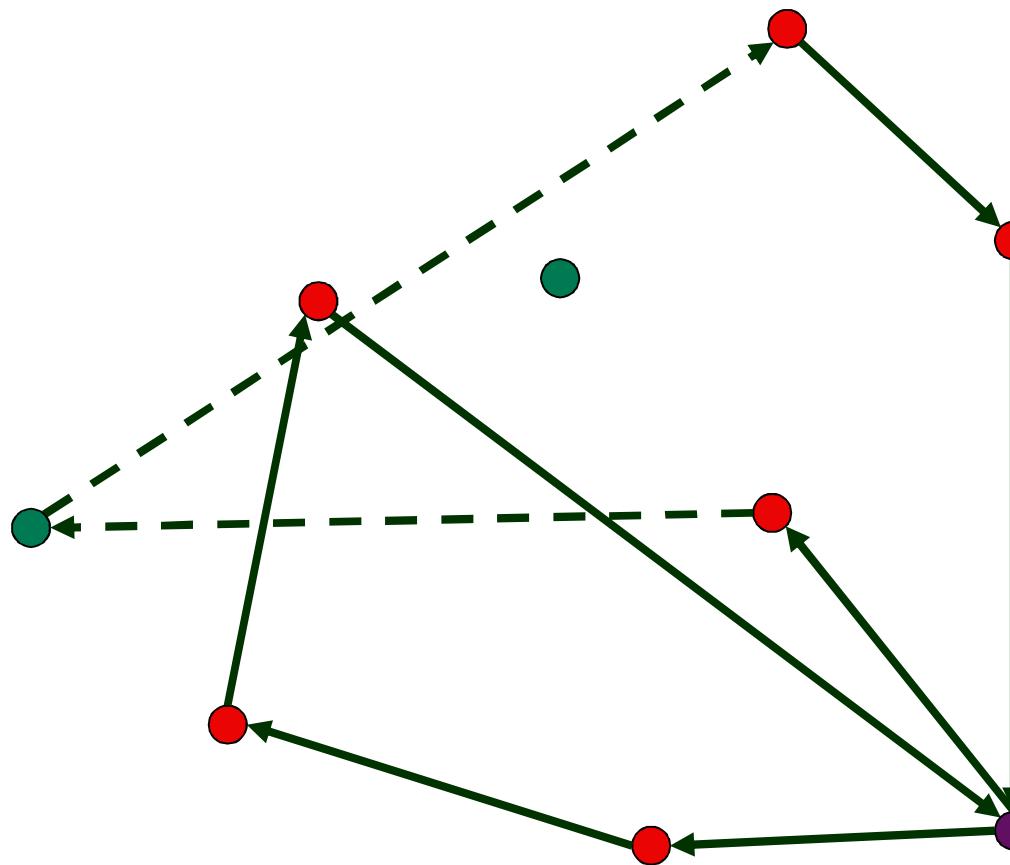
Regret



Regret



Regret



Regret

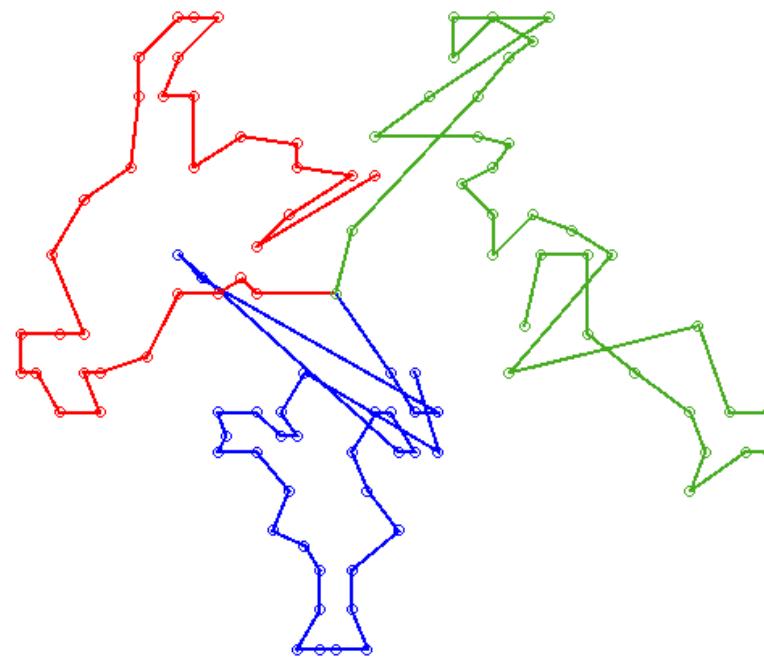


$$\begin{aligned}\text{Regret} &= C(\text{insert in 2^{nd}-best route}) - C(\text{insert in best route}) \\ &= f(2, i) - f(1, i)\end{aligned}$$

$$K\text{-Regret} = \sum_{k=1, K} (f(k, i) - f(1, i))$$

Insert customer with maximum regret

Insertion with Regret



Seeds

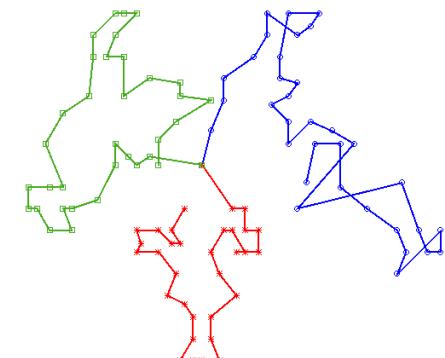


Initialise each route with one (or more) customer(s)

- Indicates the general area where a vehicle will be
- May indicate time it will be there
 - Depends on time window width

Distance-based seeding

- Find the customer (s_1) most distant from the depot
- Find the customer (s_2) most distant from s_1
- Find the customer (s_3) mist distance from s_1, s_2
- ...
- Continue until all vehicles have a seed



Implementation



- Heart of algorithm is deciding which customer to insert next, and where
- Data structure of “Insert Positions”
 - legal positions to insert a customer
 - Must calculate cost of insert
 - Must ensure feasibility of insert
- After each modification (customer insert)
 - Add new insert positions
 - Update cost of affected insert positions
 - Check legality of all insert positions
 - $O(1)$ check important for efficiency



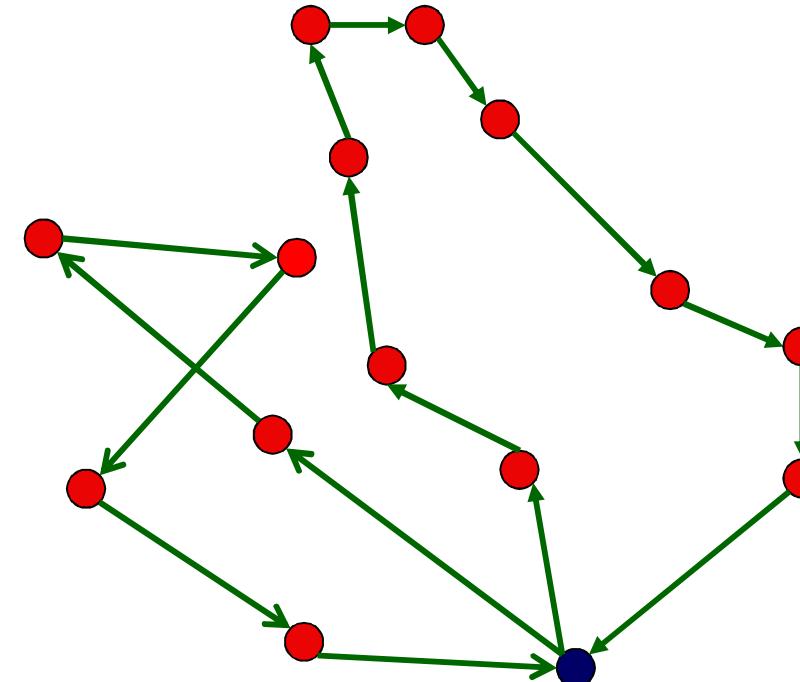
Local Search

Improvement Methods



Local Search

- Often defined using an “operator”

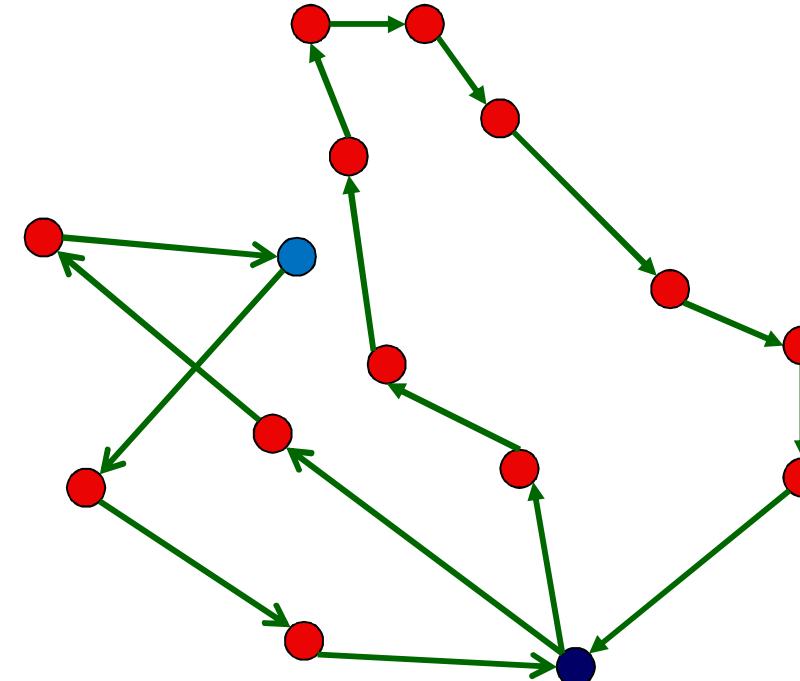


Improvement Methods



Local Search

- Often defined using an “operator”
 - e.g. 1-move

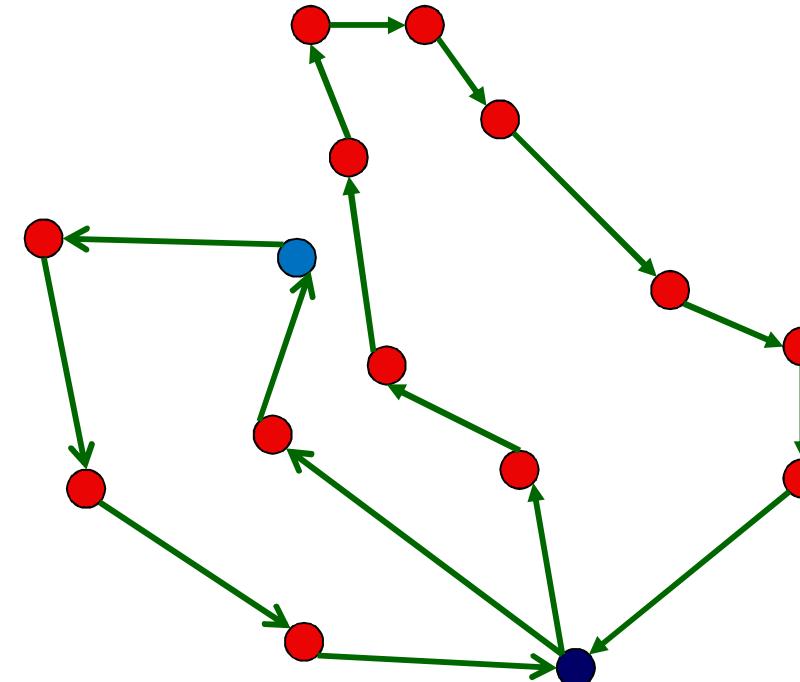


Improvement Methods



Local Search

- Often defined using an “operator”
 - e.g. 1-move

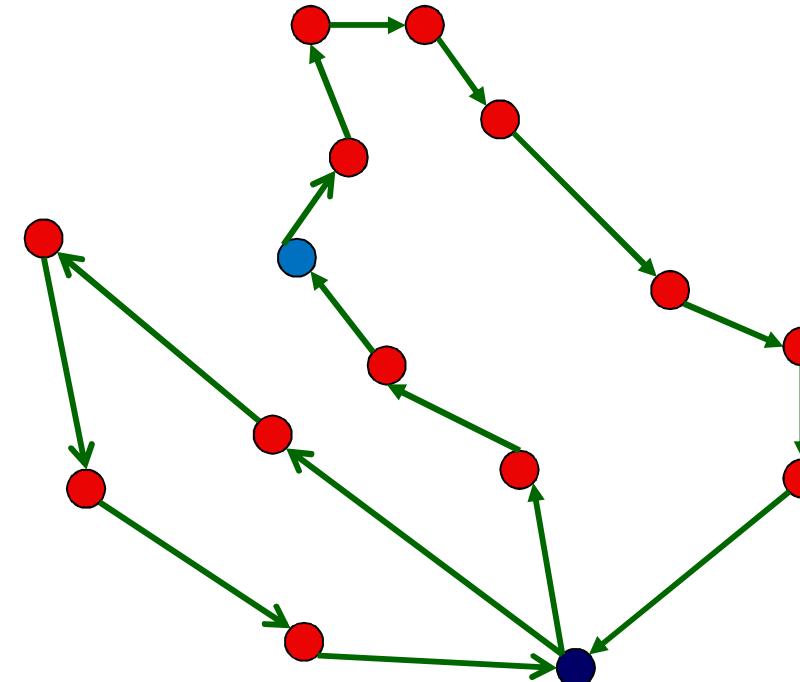


Improvement Methods



Local Search

- Often defined using an “operator”
 - e.g. 1-move

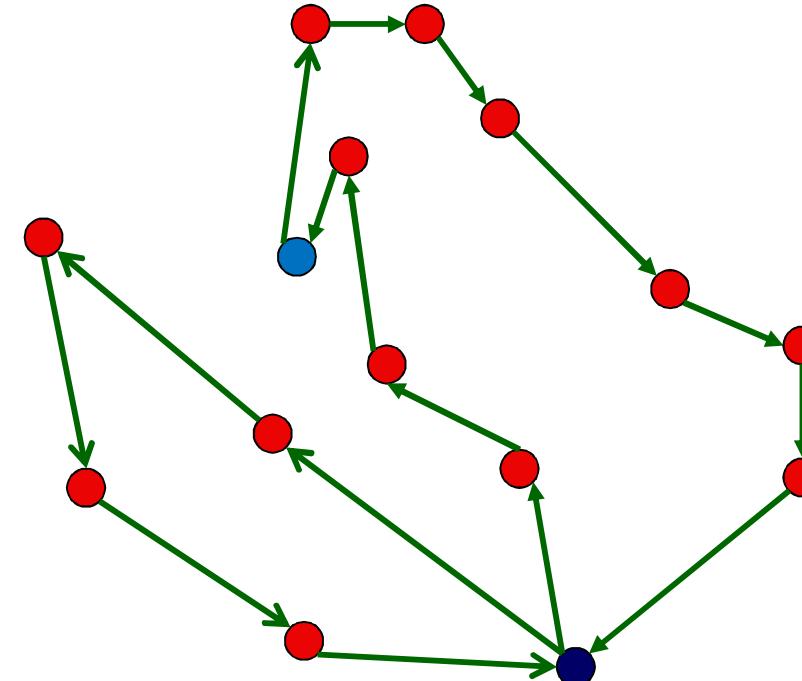


Improvement Methods



Local Search

- Often defined using an “operator”
 - e.g. 1-move

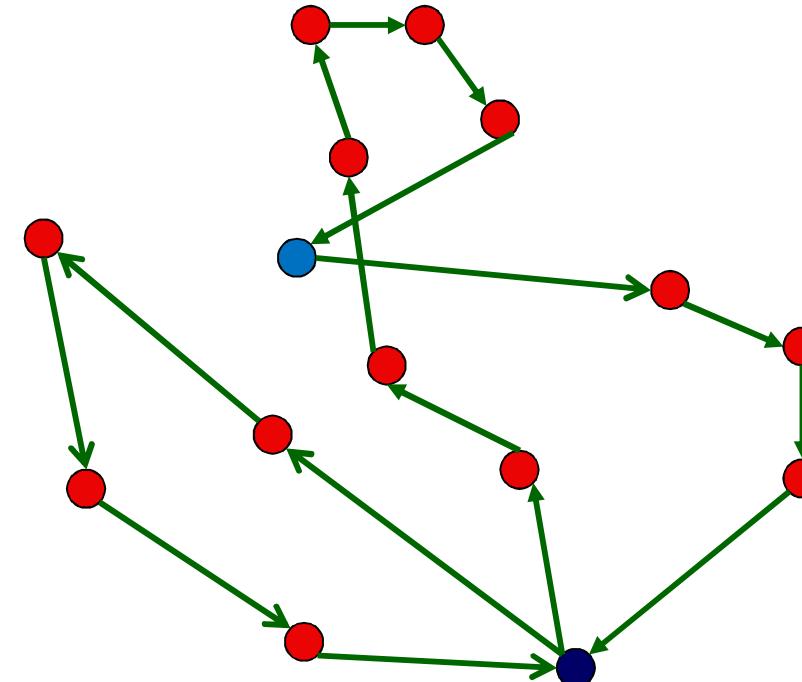


Improvement Methods



Local Search

- Often defined using an “operator”
 - e.g. 1-move

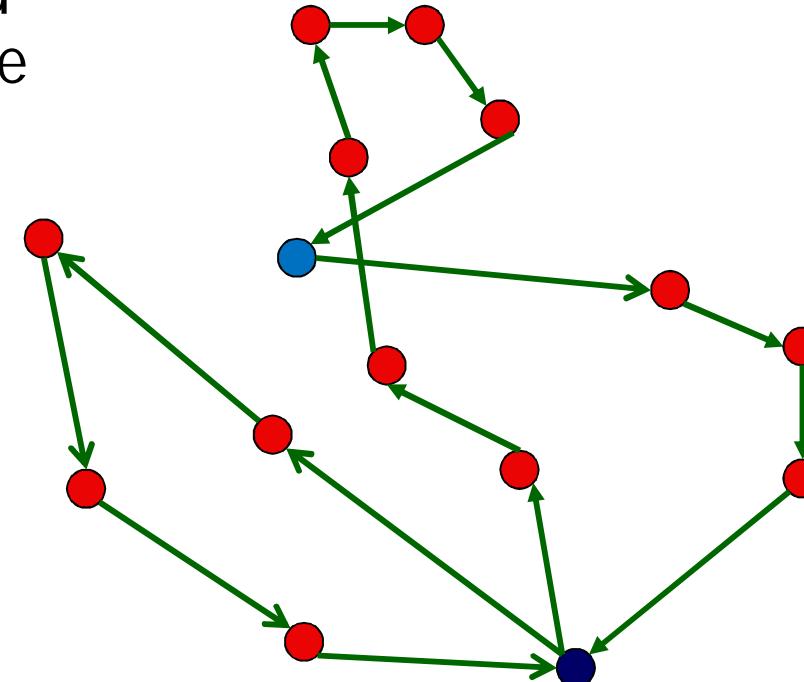


Improvement Methods



Local Search

- Often defined using an “operator”
 - e.g. 1-move
- Solutions that can be reached using the operator termed the *neighbourhood*
- Local Search explores the neighbourhood of the current solution

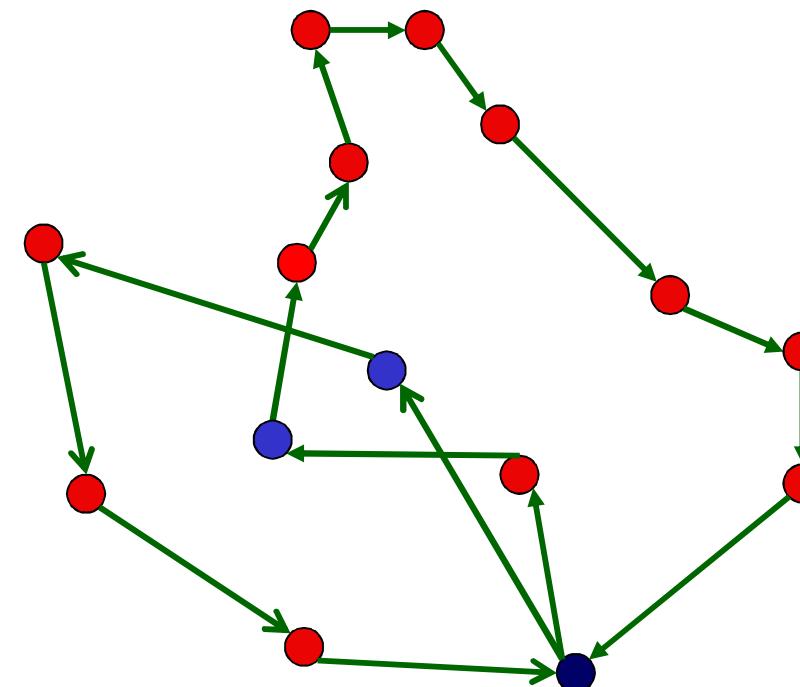


Local Search



Other Neighbourhoods for VRP:

- Swap 1-1

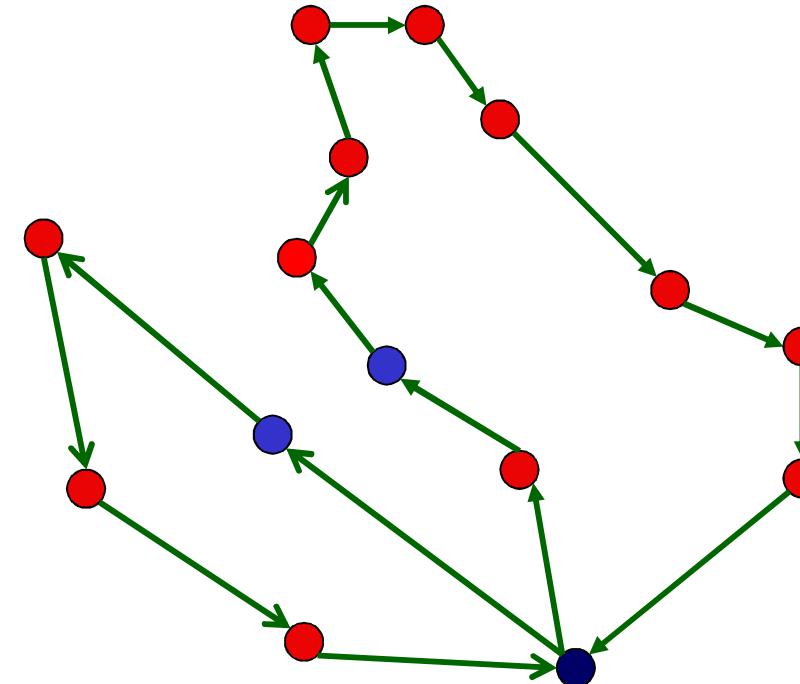


Local Search



Other Neighbourhoods for VRP:

- Swap 1-1

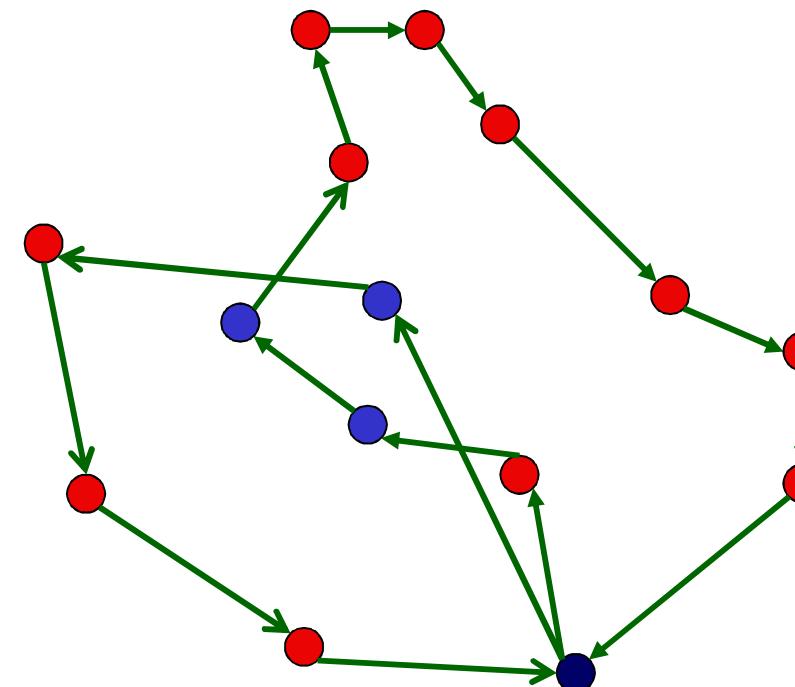


Local Search



Other Neighbourhoods for VRP:

- Swap 2-1

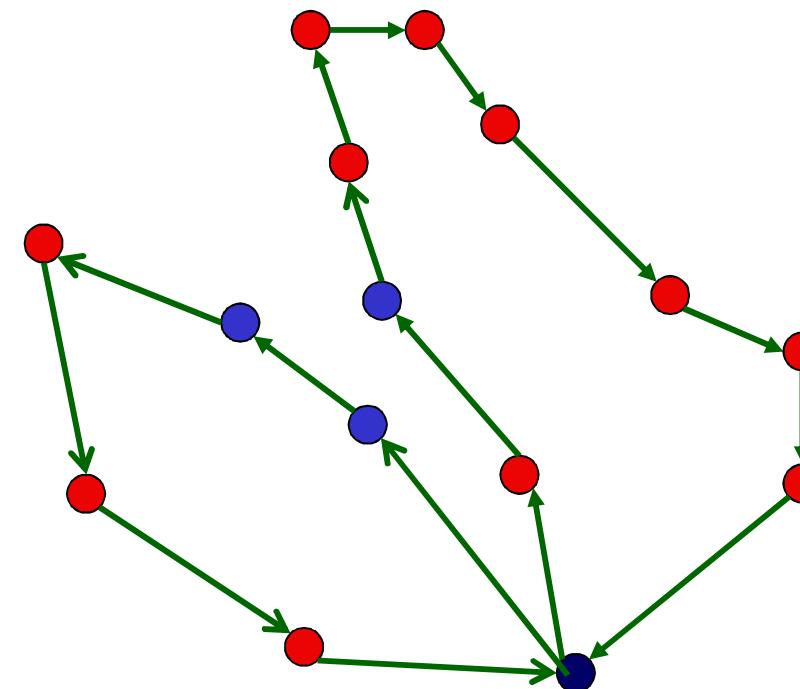


Local Search



Other Neighbourhoods for VRP:

- Swap 2-1

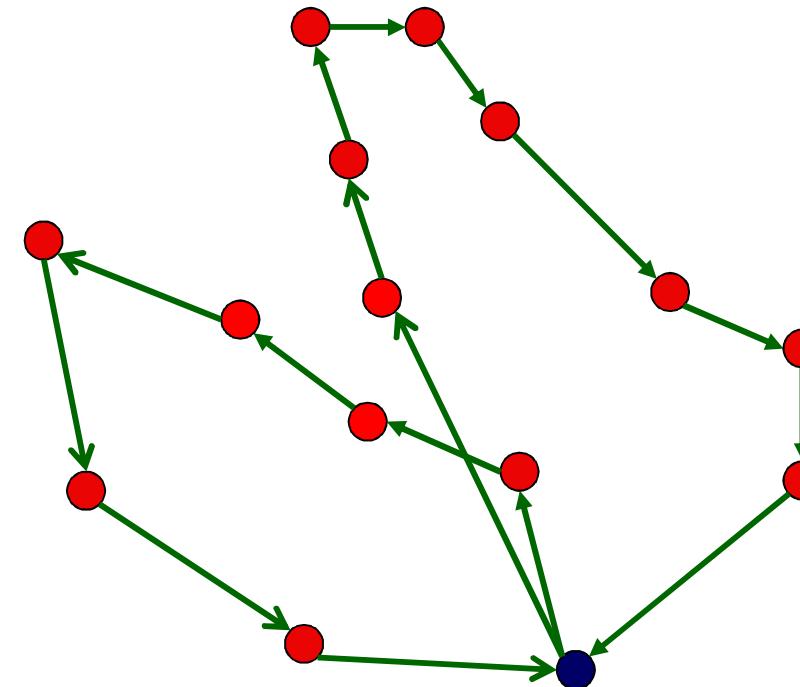


Local Search



Other Neighbourhoods for VRP:

- Swap tails

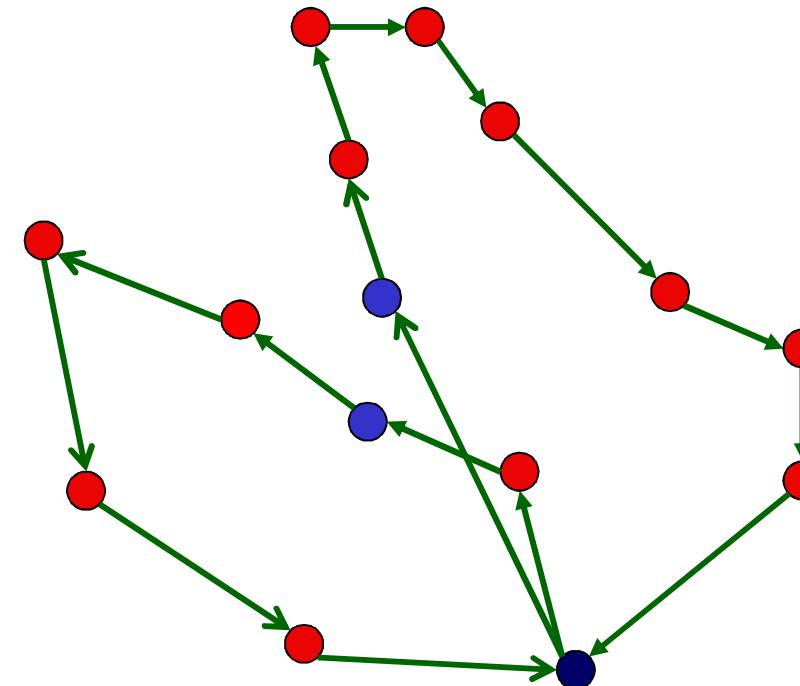


Local Search



Other Neighbourhoods for VRP:

- Swap tails

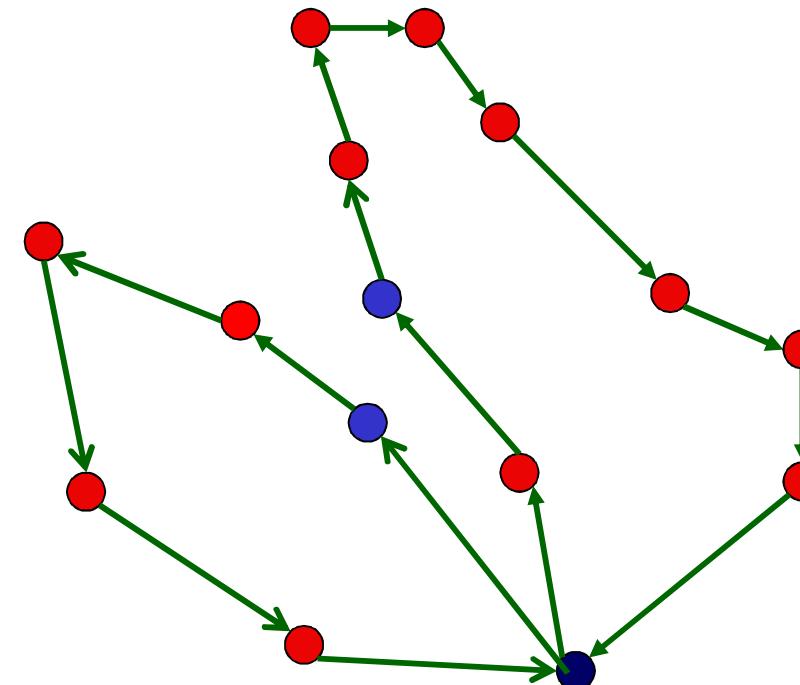


Local Search



Other Neighbourhoods for VRP:

- Swap tails

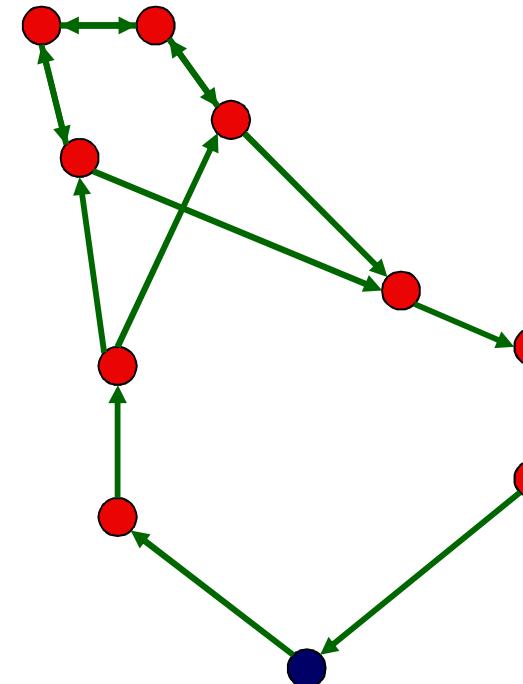


Improvement Methods



2-opt (3-opt, 4-opt...)

- Remove 2 arcs
- Replace with 2 others



Improvement Methods



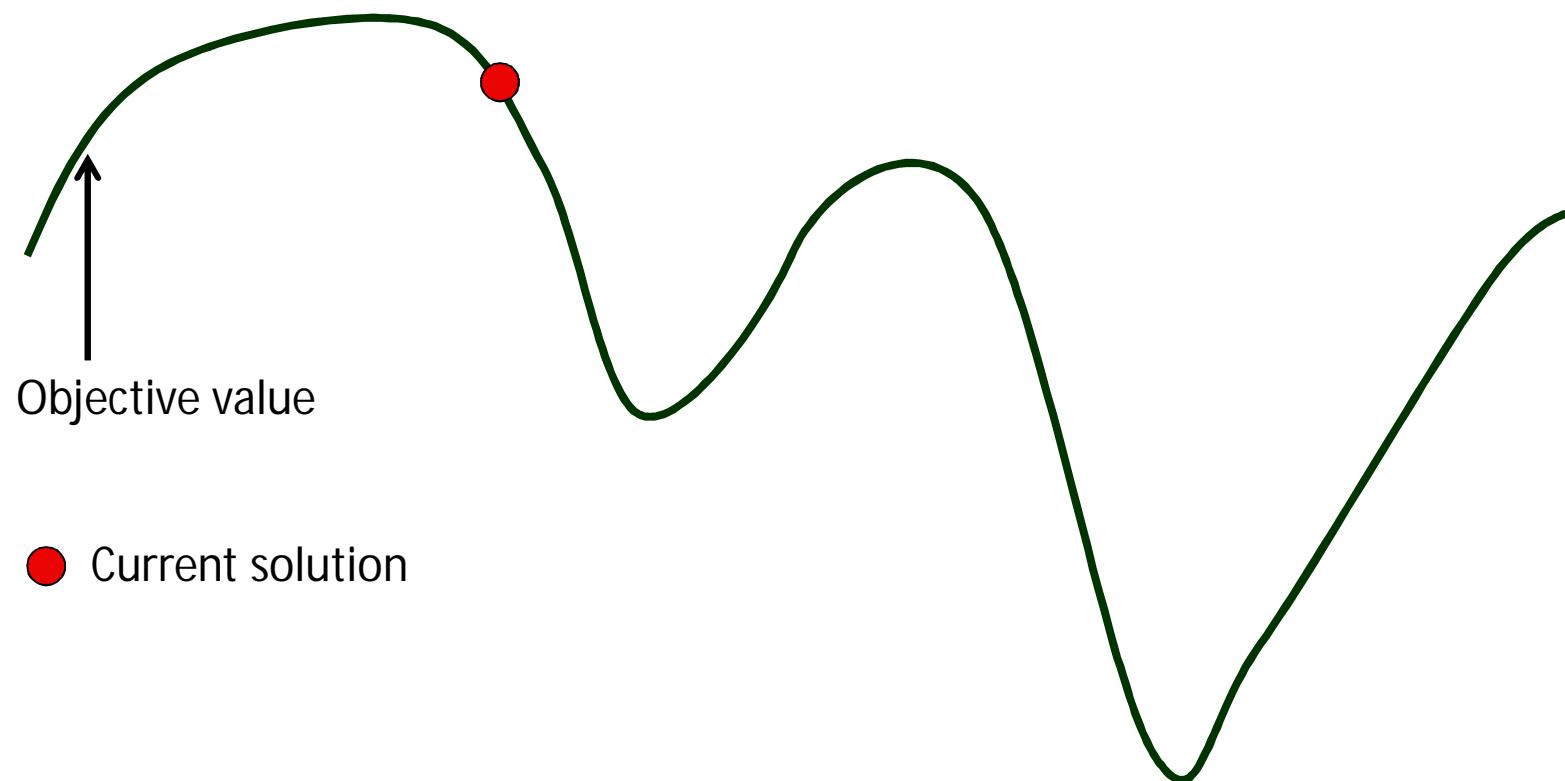
Or-opt

- Consider chains of length k
- k takes value $1 \dots n / 2$
- Remove the chain from its current position
- Consider placing in each other possible position
 - in forward orientation
 - and reverse orientation
- Very effective

Local Search



- Local minima



Local Search



Escaping local minima

Meta-heuristics

- Heuristic way of combining heuristics
- Designed to escape local minima

Local Search



Escaping local minima

- Define more (larger) neighbourhoods
 - 1-move (move 1 visit to another position)
 - 1-1 swap (swap visits in 2 routes)
 - 2-2 swap (swap 2 visits between 2 routes)
 - Tail exchange (swap final portion of routes)
 - 2-opt
 - Or-opt (all sizes 2 .. $n/2$)
 - 3-opt

Local Search



Variable Neighbourhood Search

- Consider multiple neighbourhoods
 - 1-move (move 1 visit to another position)
 - 1-1 swap (swap visits in 2 routes)
 - 2-2 swap (swap 2 visits between 2 routes)
 - 2-opt
 - Or-opt
 - Tail exchange (swap final portion of routes)
 - 3-opt
- Explore one neighbourhood completely
- If no improvement found, advance to next neighbourhood
- When an improvement is found, return to level 1

Local Search



Variable Neighbourhood Search

- For new constraints/new problems, add new neighbourhoods
- E.g. Orienteering problem
 - New neighbourhoods:
 - Unassign 1 customer (i.e. do not visit)
 - Unassign clusters of customer (e.g. sequences of customers)
 - Insert clusters of unassigned customers

Local Search



Many Meta-heuristics have been tried

- Simulated Annealing
- Tabu Search
- Genetic Algorithms
- Ants
- Bees
- Particle Swarms
- Large Neighbourhood Search

Large Neighbourhood Search



- Originally developed by Paul Shaw (1997)
- This version Ropke & Pisinger (2007)¹
- A.k.a “Record-to-record” search
- Destroy part of the solution
 - Remove visits from the solution
- Re-create solution
 - Use favourite construct method to re-insert customers
- If the solution is better, keep it
- Repeat

1: S Ropke and D Pisinger, *An Adaptive Large Neighborhood Search Heuristic for the Pickup and Delivery Problem with Time Windows*, Transportation Science **40**(4), pp 455-472, 2006

Large Neighbourhood Search



Destroy part of the solution (*Select* method)

- Remove some visits
- Move them to the “unassigned” lists

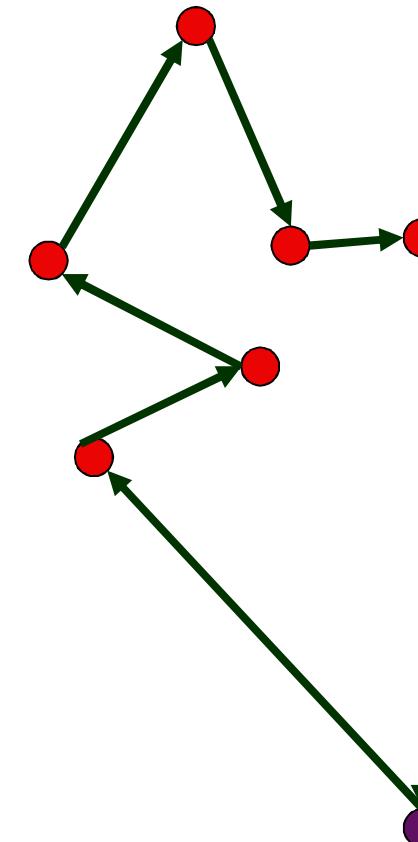
Large Neighbourhood Search



Destroy part of the solution (*Select* method)

Examples

- Remove a sequence of visits



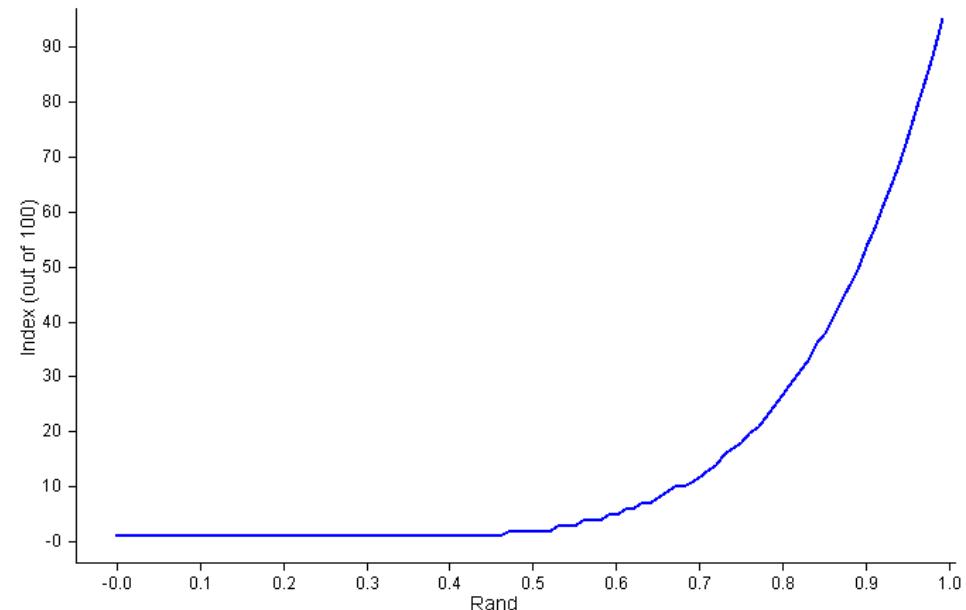
Large Neighbourhood Search



Destroy part of the solution (*Select* method)

Examples

- Choose longest (worst) arc in solution
 - Remove visits at each end
 - Remove nearby visits
- Actually, choose r^{th} worst
- $r = n * (\text{uniform}(0,1))^y$
- $y \sim 6$
 - $0.5^6 = 0.016$
 - $0.9^6 = 0.531$



Large Neighbourhood Search



Destroy part of the solution (*Select* method)

Examples

- Dump visits from k routes ($k = 1, 2, 3$)
 - Prefer routes that are close,
 - Better yet, overlapping

Large Neighbourhood Search



Destroy part of the solution (*Select* method)

Examples

- Choose first visit randomly
- Then, remove “related” visits
 - Based on distance, time compatibility, load

$$\begin{aligned} R_{ij} = & \varphi C_{ij} + \\ & \chi(|a_i - a_j|) + \\ & \psi(|q_i - q_j|) \end{aligned}$$

Large Neighbourhood Search



Destroy part of the solution (*Select* method)

Examples

- Dump visits from the smallest route
 - Good if saving vehicles
 - Sometimes fewer vehicles = reduced travel

Large Neighbourhood Search



Destroy part of the solution (*Select* method)

- Parameter: Max to dump
 - As a % of n ?
 - As a fixed number e.g. 100 for large problems
- Actual number is uniform rand (5, max)

Large Neighbourhood Search



Re-create solution

- Systematic search
 - Smaller problem, easier to solve
 - Can be very effective
- E.g.: CP Backtracking search
 - Constraint: objective must be less than current
 - (Implicitly) Look at all reconstructions
- Backtrack as soon as a better sol is found
- Backtrack anyway after *too many* failures

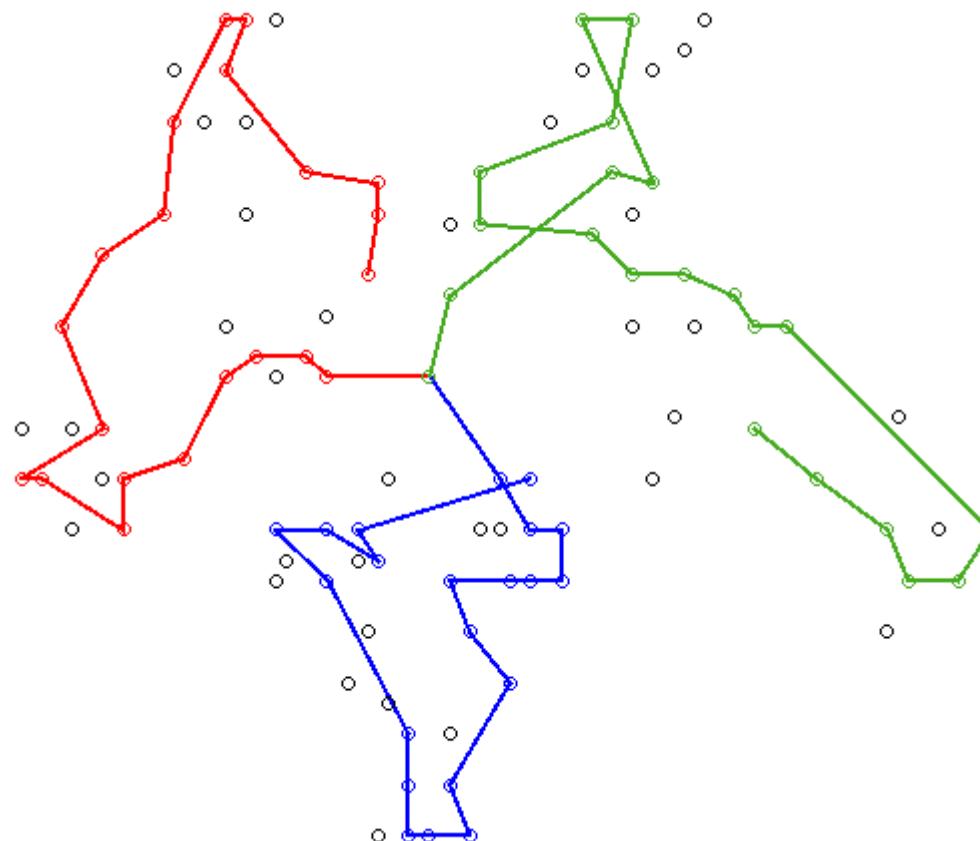
Large Neighbourhood Search



Re-create solution

- Use your favourite insert method
- Better still, use a portfolio
 - Ropke: Select amongst
 - Minimum Insert Cost
 - Regret
 - 3-regret
 - 4-regret

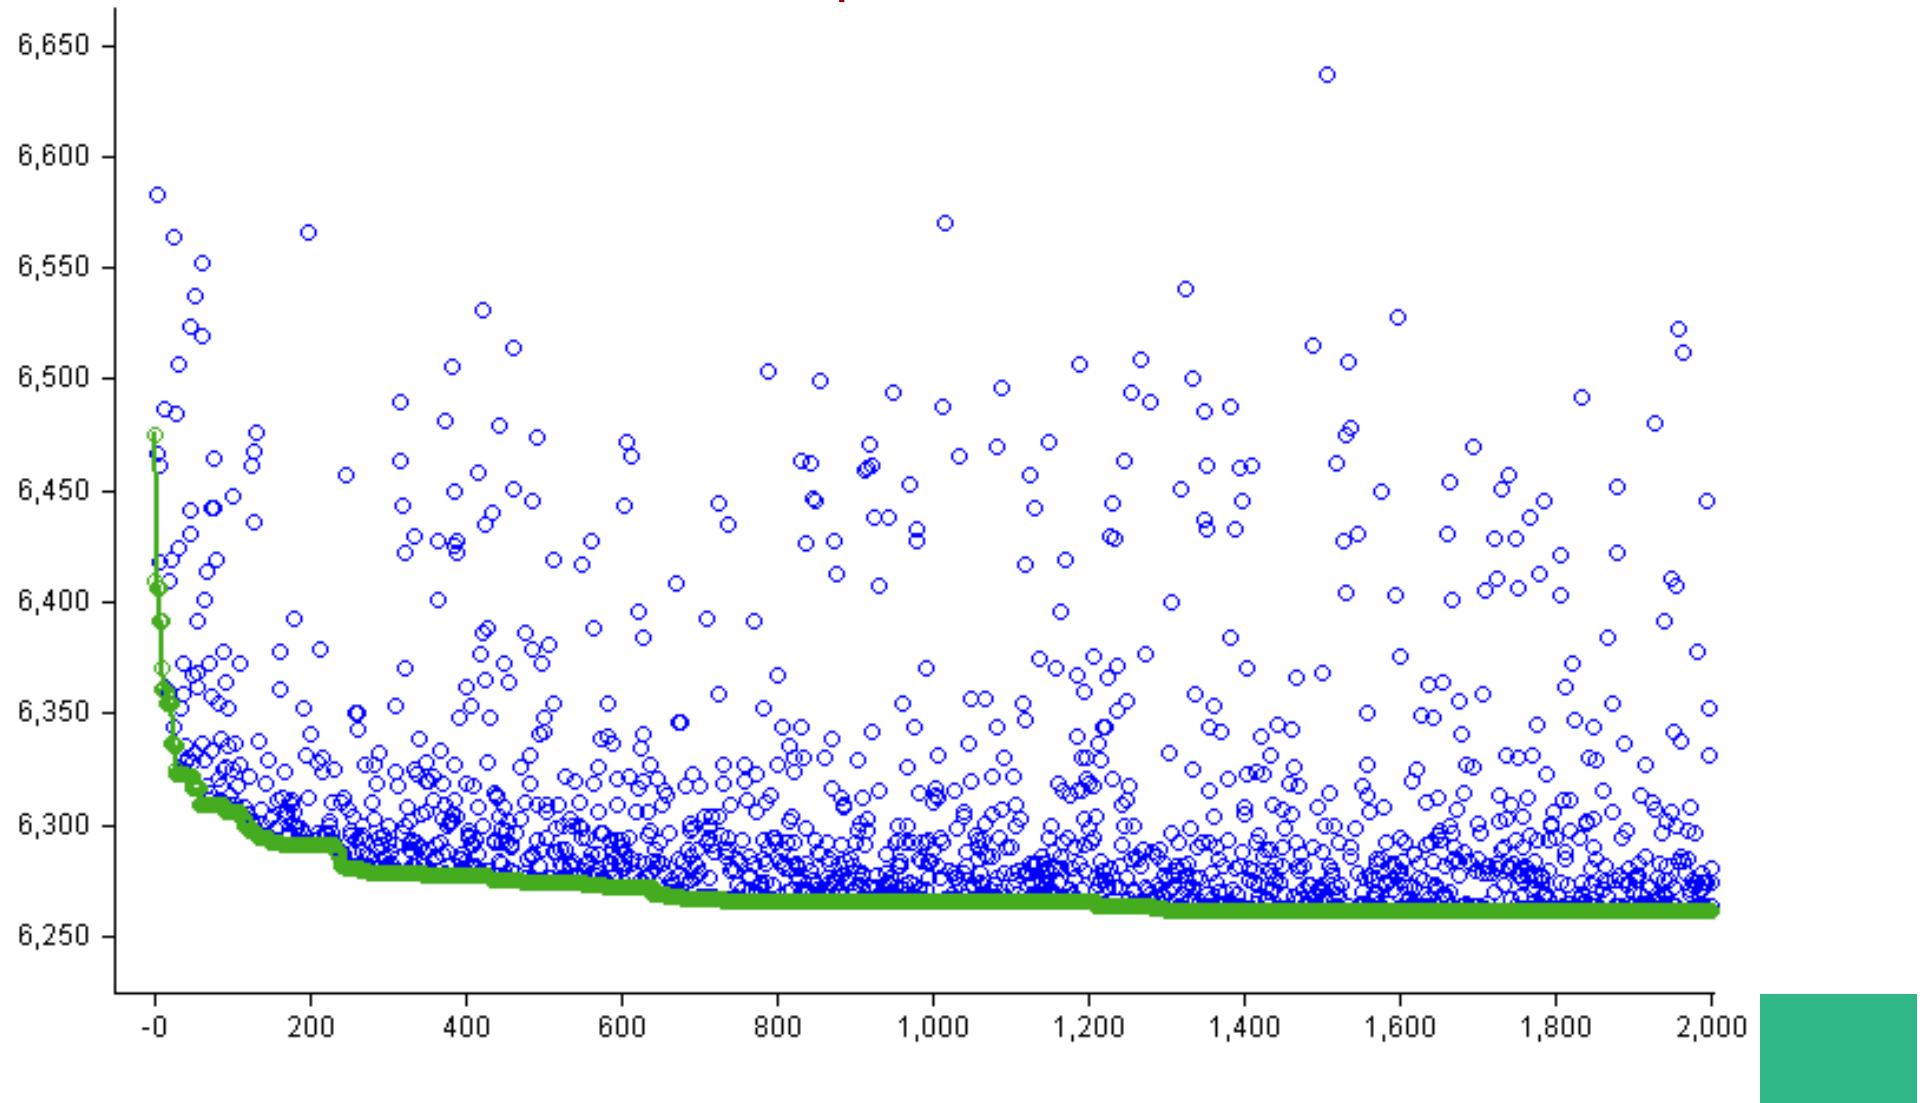
Large Neighbourhood Search



Large Neighbourhood Search



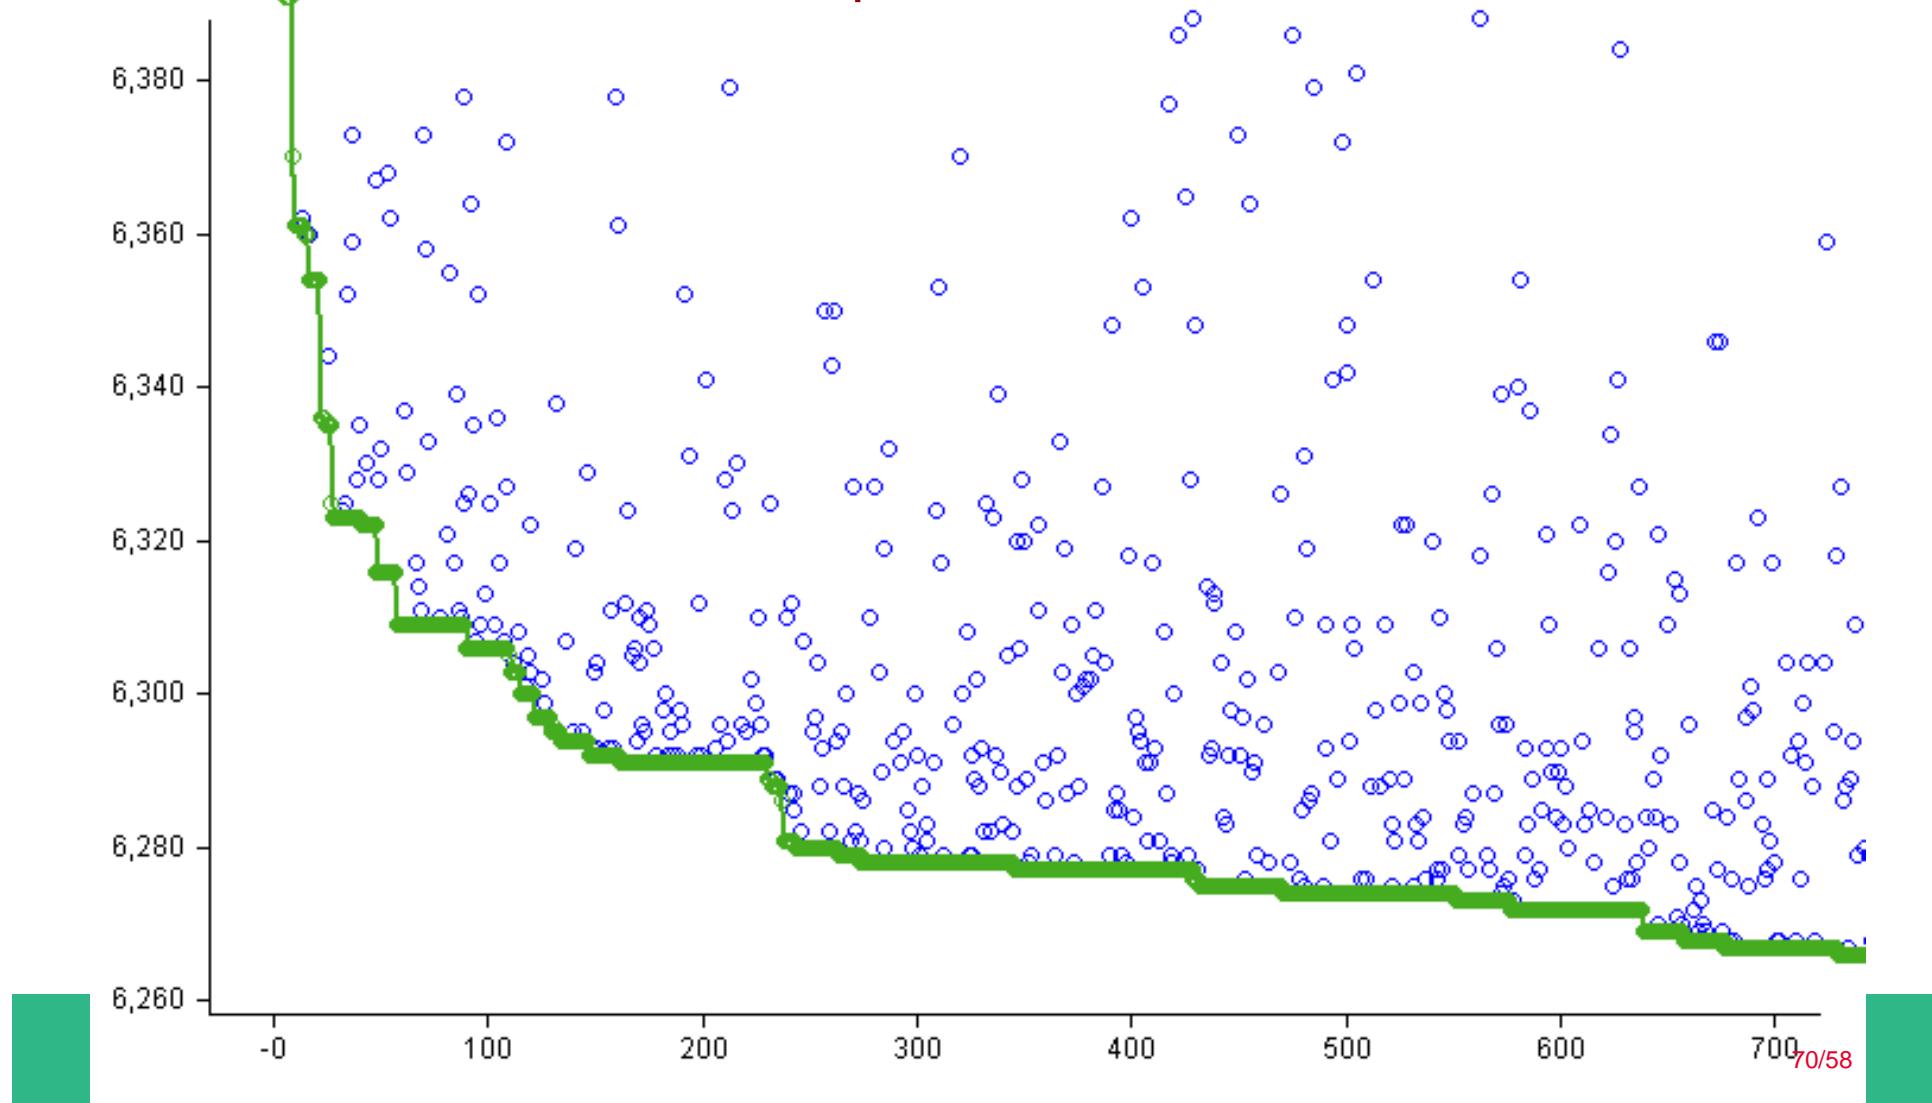
- If the solution is better, keep it



Large Neighbourhood Search



- If the solution is better, keep it



Large Neighbourhood Search

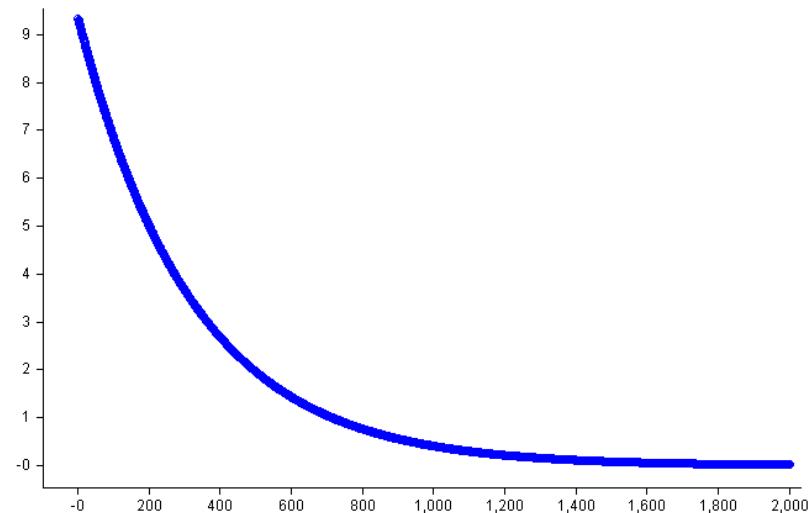


- If the solution is better, keep it
- Can use Hill-climbing
- Can use Simulated Annealing
- Can use Threshold Annealing
- ...

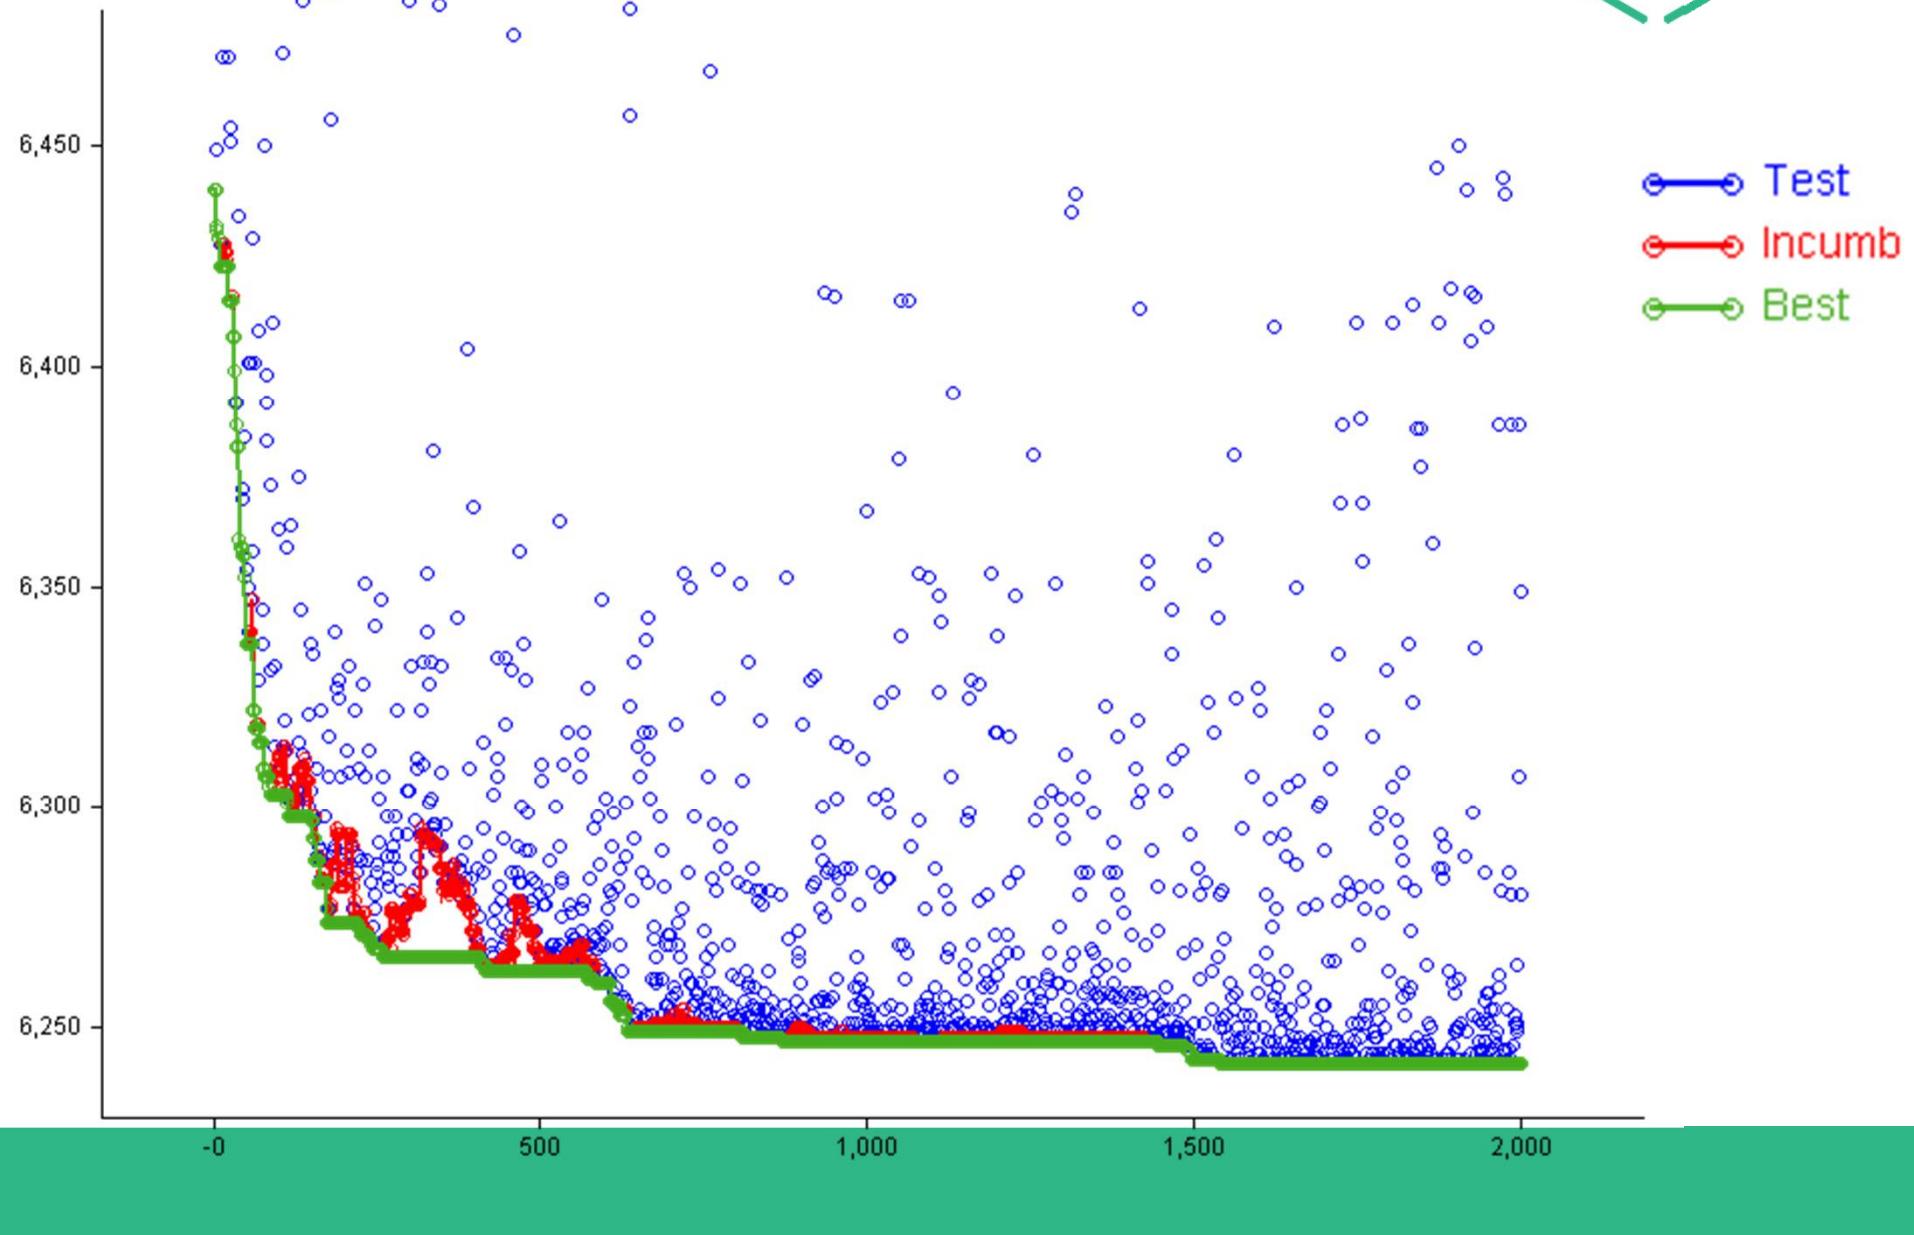
Large Neighbourhood Search



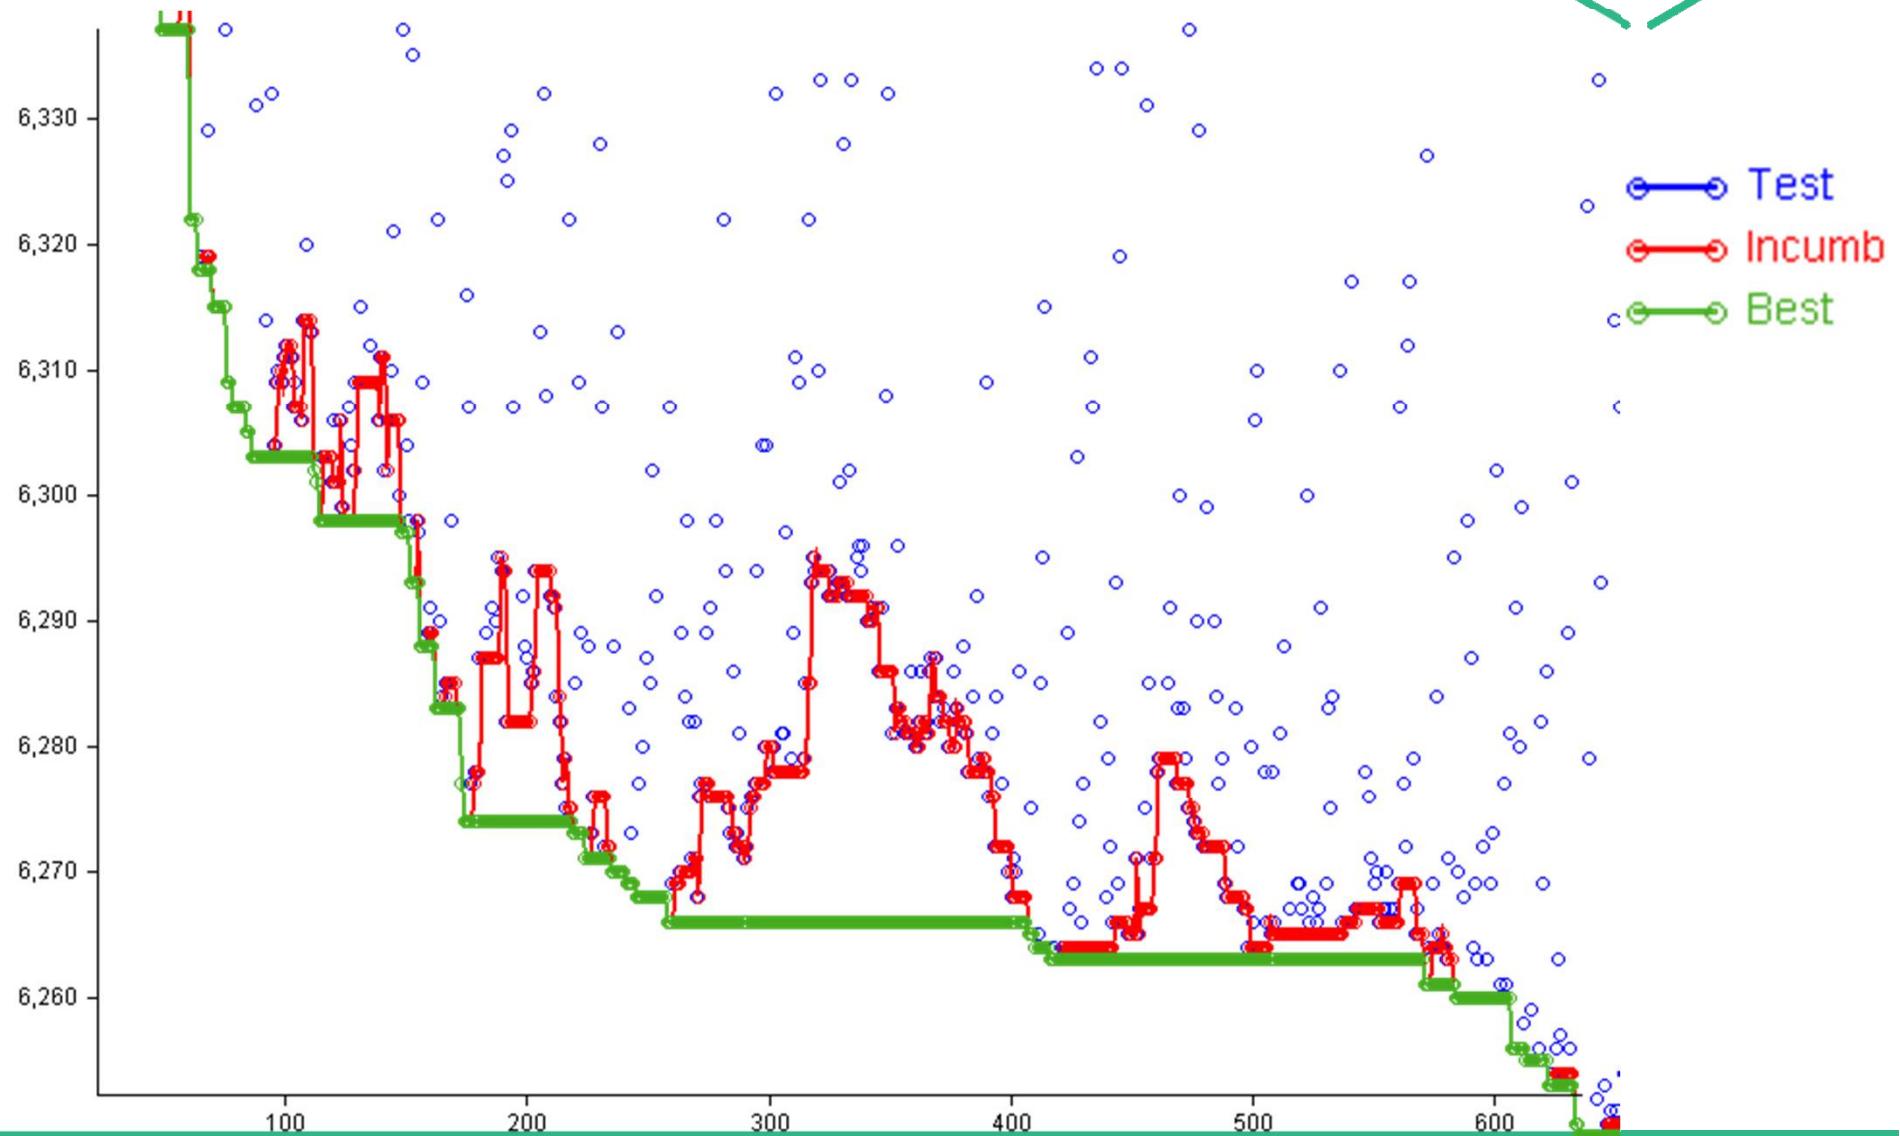
$$P(\text{accept increase } \Delta) = e^{-\Delta/T}$$



Large Neighbourhood Search



Large Neighbourhood Search

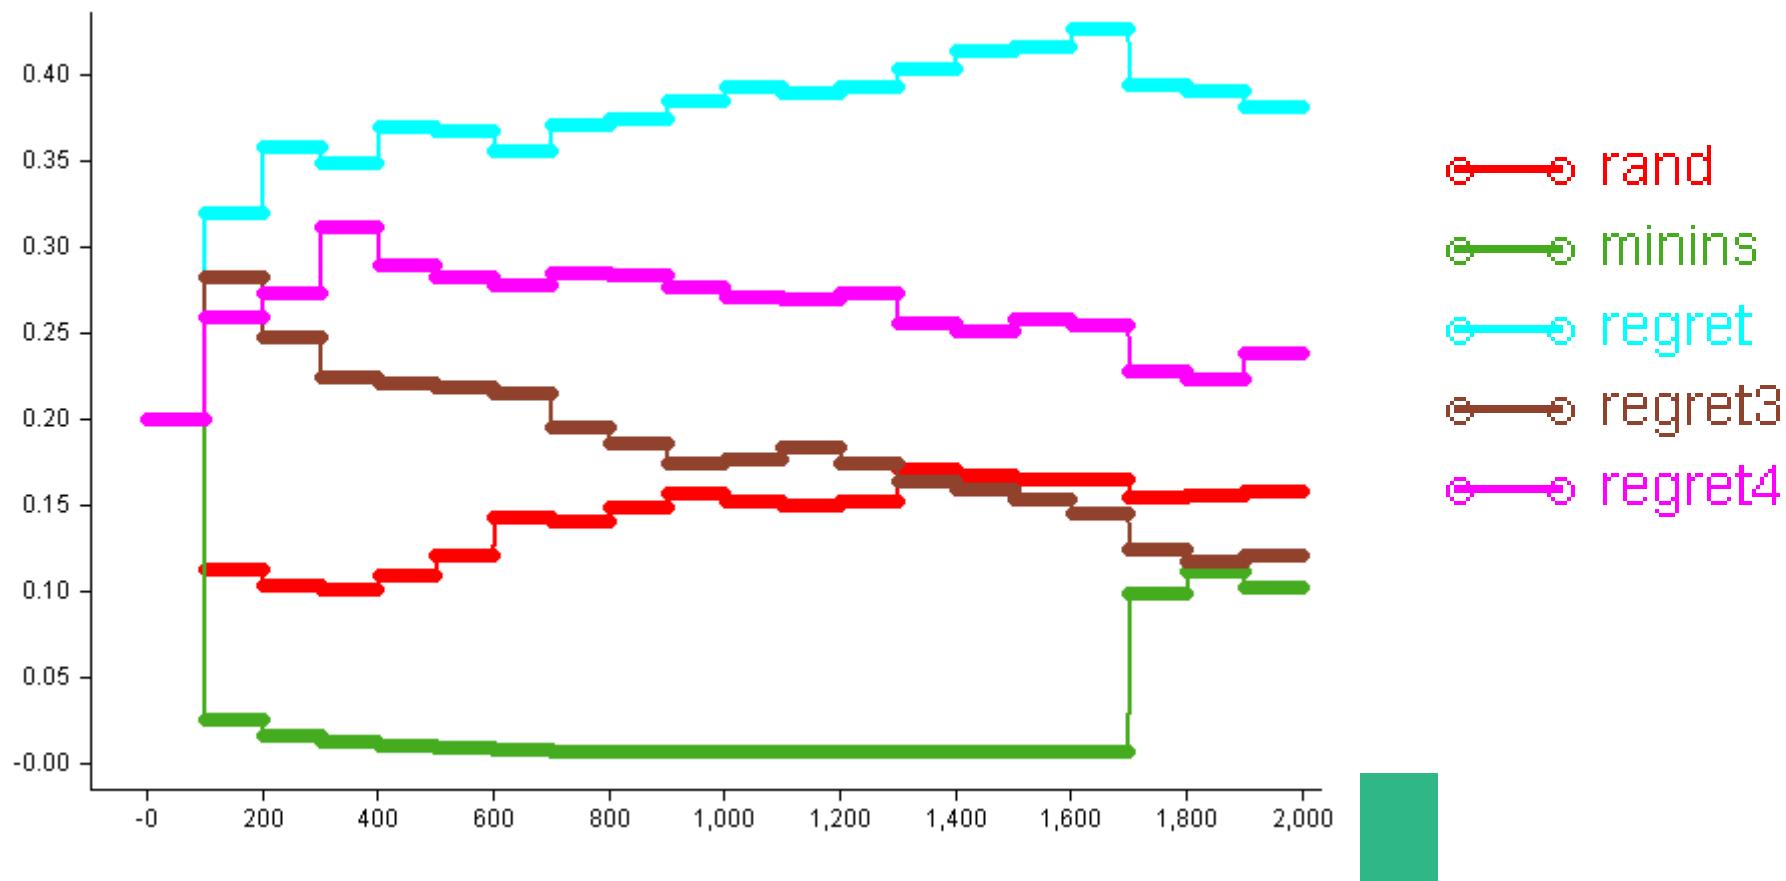


Large Neighbourhood Search



Adaptive

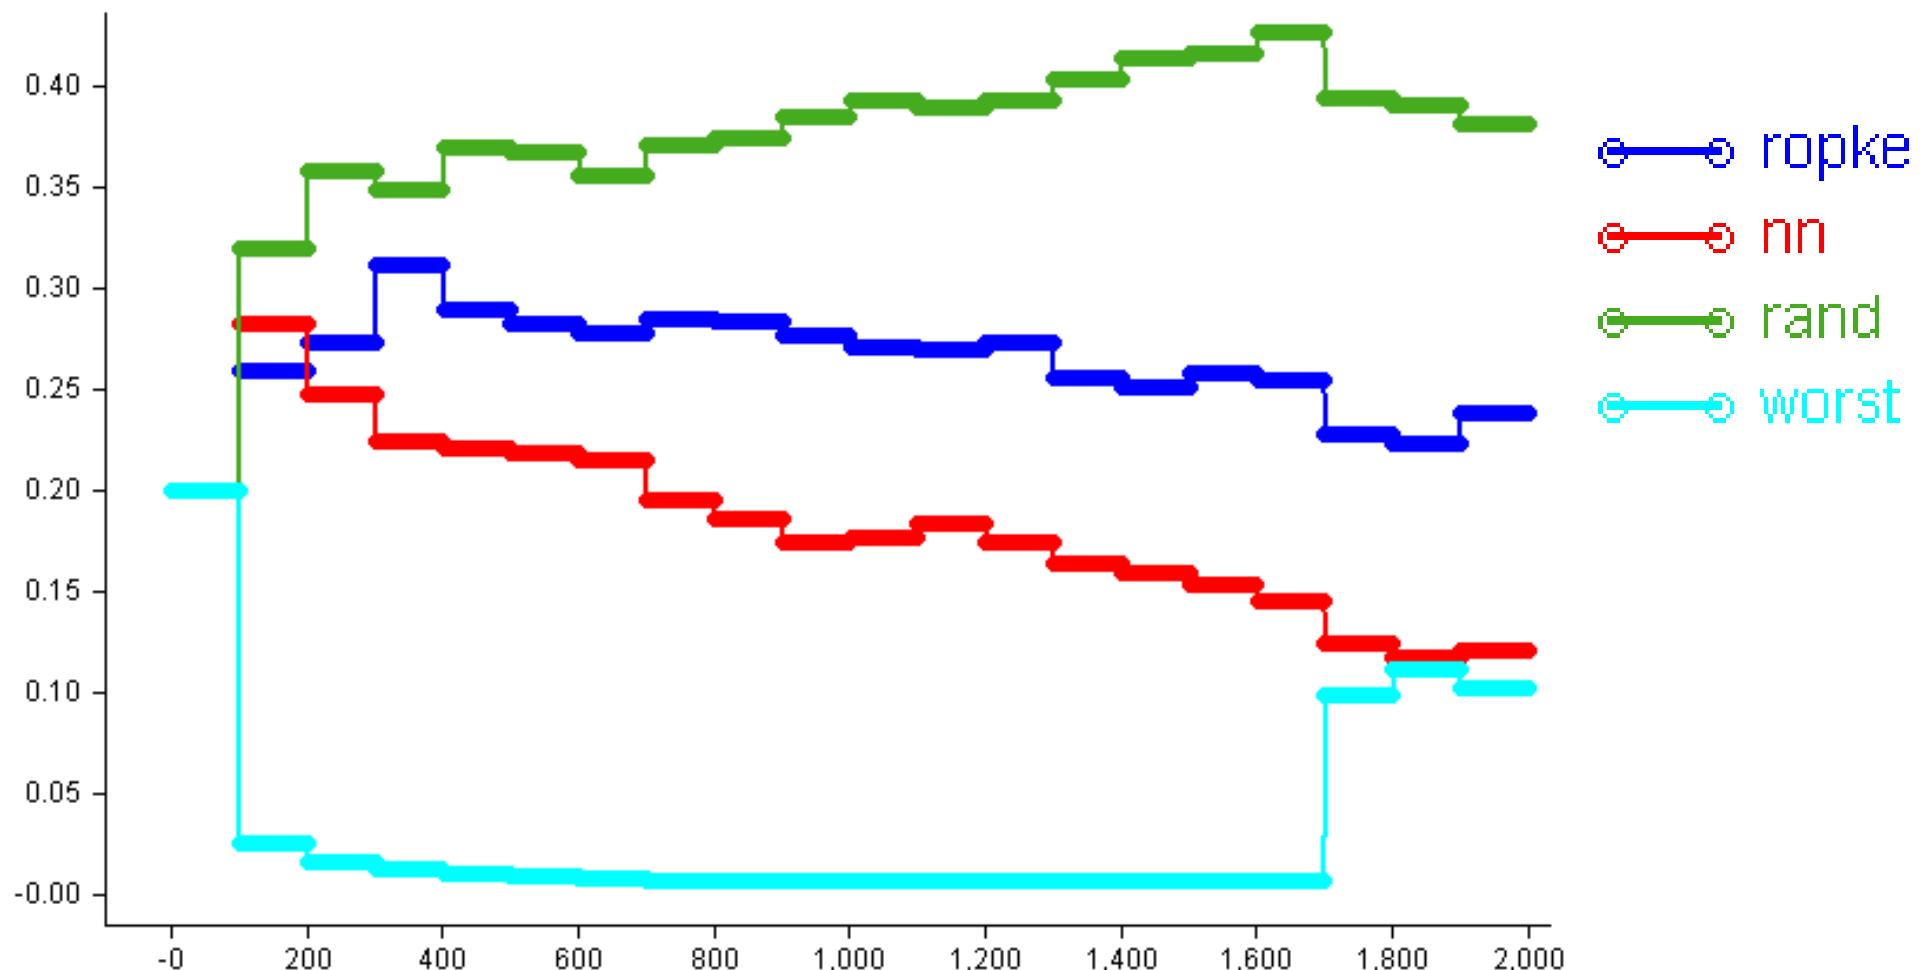
- Ropke adapts choice based on prior performance
 - “Good” methods are chosen more often



Large Neighbourhood Search



Adapting Select method



Large Neighbourhood Search



Ropke & Pisinger (with additions) can solve a variety of problems

- VRP
- VRP + Time Windows
- Pickup and Delivery
- Multiple Depots
- Multiple Commodities
- Heterogeneous Fleet
- Compatibility Constraints

Solution Methods



Summary so far:

- Introduced several successive insertion construction methods
 - Various ways to choose the next visit to insert
 - Various ways to choose where to insert
- Described two successful metaheuristics
 - Variable Neighbourhood Search
 - Large Neighbourhood Search

Solution Methods



What's wrong with that?

- New constraint → new code
 - Often right in the core
- New constraints interact
 - e.g. Multiple time windows mess up duration calculation
- Code is hard to understand, hard to maintain

Solution Methods



An alternative:

Constraint Programming

Constraint Programming



CP offers a language for representing problems

- Decision variables
- Constraints on variables

Also offers techniques for solving the problems

- Systematic search
- Heuristic Search

Variables are represented by their *domain*

- (Usually finite) set of feasible values
- E.g $x \in [0,100]$ or $x \in [0,1,3..15,16,18,55..99]$

Constraints link variables

- $x \leq 4y + 6z$
- $x^2 + y^2 = z^2$
- Cardinality ($X, 1, 4, 5$)
(In the set X the value '1' occurs at least 4 times, and no more than 5 times)
- AllDifferent (X) (All values in X are different)
- DriverBreak (30,120,240)
(A break of 30 minutes must be inserted after 120 minutes but no later than 240 minutes after start of route)

Propagators (efficiently) enforce constraints

- Wake when the domain of a linked variable is changed
- For each value in each variable
 - Ensure there is a set of feasible values of other variables that *supports* that value – e.g.

$$x < y$$

$$x = [3, 5, 7, 9]$$

$$y = [2, 4, 6, 8]$$

- The value '9' in x has no support in y
- The value '2' in y has no support in x
- After propagation: $x = [3, 5, 7]$

$$y = [4, 6, 8]$$

Eg Mutual Exclusion constraint in VRP

- (If any visit from the set D is assigned, then no others can be)
- Uses 'isAssigned' var
 - Domain [0,1]
- Attach propagator to the 'isAssigned' variable for each of the visits
- Propagator wakes when 'isAssigned' is bound to 1 for any visit
- Propagates by binding isAssigned to '0' for remaining visits.

CP101



- Typical execution:
 - Establish choice point (store all current domains)
 - Choose variable to instantiate
 - Choose value to assign, and assign it
 - Propagations fire until a *fixed point* is achieved, or an inconsistency is proved (empty domain)
 - If inconsistent,
 - Backtrack (restore to choicepoint)
 - Remove offending value from the variable's domain
 - Repeat until all variables are bound (assigned)
 - For complete search, store sol, then act like inconsistent

'Choose a variable to assign, choose value to assign'

- Very good fit for constructive route creation
- After each insert, propagators fire
- New variable domains give look-ahead to feasible future insertions
- Constraints guide insertion process

Step-to-new-solution does not work as well

- Local move operators can only use CP as a rule checker
 - Do not leverage full power of CP

Expressive Language (e.g. Minizinc)



```
string: Name;  
  
% Customers  
int: NumCusts;  
set of int: Customers =  
    1..NumCusts;  
  
% Locations  
int: NumLocs = NumCusts + 1;  
set of int: Locations =  
    1..NumLocs;  
  
% Vehicles  
int: NumVehicles;  
int: NumRoutes = NumVehicles;  
set of int: Vehicles =  
    1..NumVehicles;  
  
% Location data  
array[Locations] of float: locX;  
array[Locations] of float: locY;  
array [Locations, Locations] of  
int: dist;  
  
% Decision variables  
var int: obj;  
array[Visits] of var Visits:  
    routeOf;  
array[Visits] of var Visits:  
    succ;  
array[Visits] of var [0,1]:  
    isAssigned;  
  
constraint alldifferent (succ);  
constraint circuit (succ);  
  
constraint  
    obj = sum (i in Visits)  
        (dist[Loc[i],Loc[succ[i]]]);  
  
constraint  
    sum (i in Visits,j = routeOf[i])  
        (demand[i]) < j  
        for j in Vehicles;
```

Constraint Programming for the VRP



Constraint Programming

Advantages:

- Expressive language for formulating constraints
- Each constraint encapsulated
- Constraints interact naturally
- Constraints guide construction

Disadvantages:

- Can be slow
- No fine control of solving
 - (unless you use a low-level library like *gecode*)



Constraint Programming



Two ways to use constraint programming

- Rule Checker
- Properly

Rule Checker:

- Use favourite method to create/improve a solution
- Check it with CP
 - Very inefficient.



A CP Model for the VRP

Vocabulary



- A *solution* is made up of *routes* (one for each vehicle)
- A *route* is made up of a sequence of *visits*
- Some visits serve a customer (*customer visit*)

(Some tricks)

- Each route has a “start visit” and an “end visit”
- Start visit is first visit on a route – location is depot
- End visit is last visit on a route – location is depot
- Also have an additional route – the unassigned route
 - Where visits live that cannot be assigned

Model



A (rich) vehicle routing problem

- n customers (fixed in this model)
- v vehicles (fixed in this model)
- $m = v+1$ one route per vehicle plus “unassigned” route
- fixed locations
 - where things happen
 - one for each customer + one for (each?) depot
- c commodities (e.g. weight, volume, pallets)
 - Know demand from each customer for each commodity
- Know time between each location pair
- Know cost between each location pair
 - Both obey triangle inequality

Referencing



Sets

- $N = \{1 \dots n\}$ – customers
- $V = \{1 \dots v\}$ – vehicles/real routes
- $R = \{1 \dots m\}$ - routes include 'unassigned' route
- $S = \{n+1 \dots n+m\}$ – start visits
- $E = \{n+m+1 \dots n+2m\}$ – end visits
- $V = N \cup S \cup E$ – all visits
- $V^S = N \cup S$ – visits that have a sensible successor
- $V^E = N \cup E$ – visits that have a sensible predecessor

Referencing



Customers

- Each customer has an index in $N = \{1..n\}$
- Customers are 'named' in CP by their index

Routes

- Each route has an index in $R = \{1..m\}$
- Unassigned route has index m
- Routes are 'named' in CP by their index

Visits

- Customer visit index same as customer index
- Start visit for route k has index $n + k$; aka $start_k$
- End visit for route k has index $n + m + k$; aka end_k

Data



We know (note uppercase)

- V_i The 'value' of customer i
- D_{ik} Demand by customer i for commodity k
- E_i Earliest time to start service at i
- L_i Latest time to start service at i
- Q_{jk} Capacity of vehicle j for commodity k
- T_{ij} Travel time from visit i to visit j
- C_{ij} Cost (w.r.t. objective) of travel from i to j

Basic Variables



Successor variables: s_i

- s_i gives direct successor of i , i.e. the index of the next visit on the route that visits i
- $s_i \in V^E$ for i in V^S $s_i = 0$ for i in E

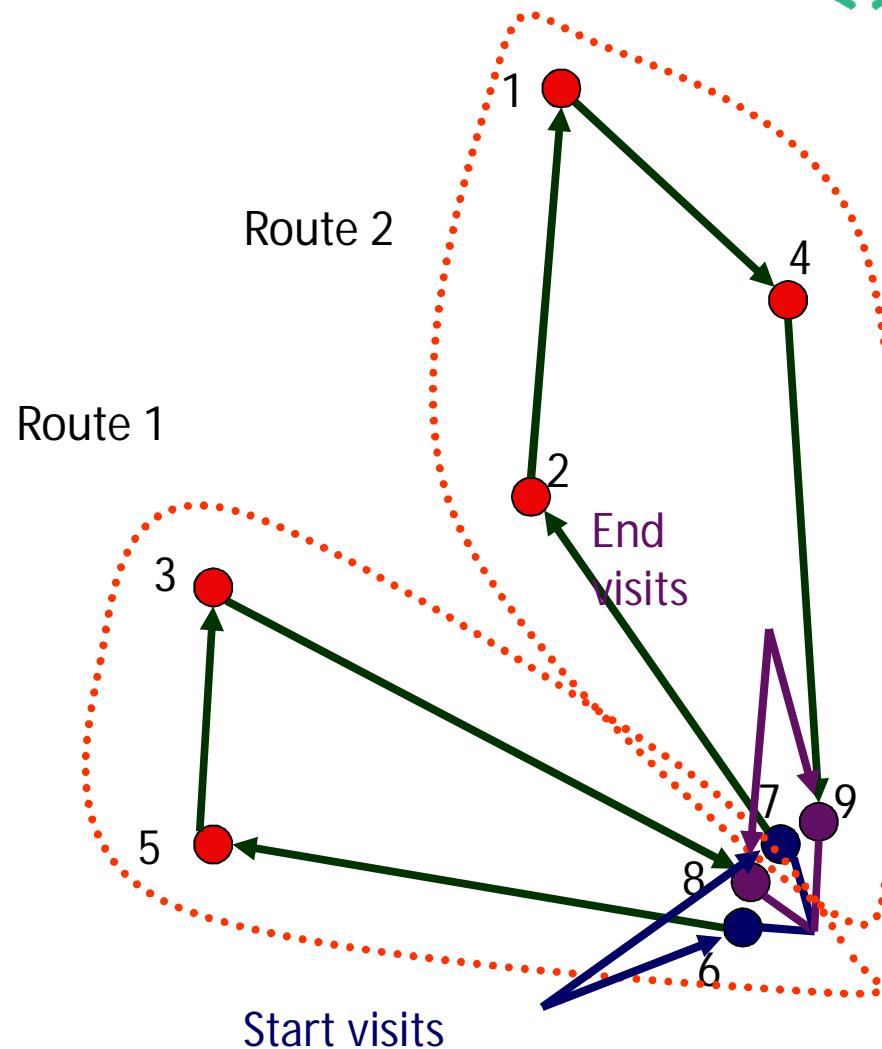
Predecessor variables p_i

- p_i gives the index of the previous visit in the route
- $p_i \in V^S$ for i in V^E $p_i = 0$ for i in S
- Redundant – but empirical evidence for its use
- Route variables r_i
- r_i gives the index of the route (vehicle) that visits i
- $r_i \in R$

Example



i	s_i	p_i	r_i
1	4	2	2
2	1	7	2
3	8	5	1
4	9	1	2
5	3	6	1
6	5	0	1
7	2	0	2
8	0	3	1
9	0	4	2



Other variables



Accumulation Variables

- q_{ik} Quantity of commodity k after visit i
- c_i Objective cost getting to i

- For problems with time constraints
 - a_i Arrival time at i
 - t_i Start time at i (time service starts)
 - d_i Departure time at i
-
- Actually, only t_i is required, but others allow for expressive constraints

What can we model?



- Basic VRP
- VRP with time windows
- Multi-depot
- Heterogeneous fleet
- Open VRP (vehicle not required to return to base)
 - Requires *anywhere* location
 - Route end visits located at *anywhere*
 - distance $i \rightarrow \text{anywhere} = 0$
- Compatibility
 - Customers on different / same vehicle
 - Customers on/not on given vehicle
- Pickup and Delivery problems

What can we model?



- Variable load/unload times
 - by changing departure time relative to start time
- Dispatch time constraints
 - e.g. limited docks
 - s_i for $i \in S$ is load-start time
- Depot close time
 - Time window on end visits
- Fleet size and mix
 - Add lots of vehicles
 - Need to introduce a 'fixed cost' for a vehicle
 - C_{ij} is increased by fixed cost for all $i \in S$, all $j \in N$

What can't we model



- Can't handle dynamic problems
 - Fixed domain for s , p , r vars
- Can't introduce new visits post-hoc
 - E.g. optional driver break must be allowed at start
- Can't handle multiple visits to same customer
 - 'Larger than truck-load' problems
 - If qty is fixed, can have multiple visits / cust
 - Heterogeneous fleet is a pain
- Can handle time- or vehicle-dependent travel times/costs with mods
- Can handle Soft Constraints with mods

Objective



Want to minimize

- sum of objective (c_{ij}) over used arcs, plus
- value of unassigned visits

minimize

$$\sum_{i \in E} c_i + \sum_{i | r_i = 0} v_i$$

Basic constraints



$$\begin{aligned} & \text{Path } (\mathcal{S}, \mathcal{E}, \{s_i \mid i \in V\}) \\ & \text{AllDifferent } (\{p_i \mid i \in V^{\mathcal{E}}\}) \end{aligned}$$

Accumulate obj.

$$c_{s_i} = c_i + C_{i,s_i} \quad \forall i \in V^{\mathcal{S}}$$

Accumulate time

$$a_{s_i} = d_i + T_{i,s_i} \quad \forall i \in V^{\mathcal{S}}$$

Time windows

$$t_i \geq a_i \quad \forall i \in V$$

$$t_i \leq L_i \quad \forall i \in V$$

$$t_i \geq E_i \quad \forall i \in V$$

$$t_i = 0 \quad \forall i \in S$$

Constraints

- Load

$$q_{s_i k} = q_{ik} + Q_{s_i k} \quad \forall i \in V^S, k \in C$$

$$q_{ik} \leq Q_{r_i k} \quad \forall i \in V, k \in C$$

$$q_{ik} \geq 0 \quad \forall i \in V, k \in C$$

$$q_{ik} = 0 \quad \forall i \in S, k \in C$$

- Consistency

$$s_{p_i} = i \quad \forall i \in V^S$$

$$p_{s_i} = i \quad \forall i \in V^E$$

$$r_i = r_{s_i} \quad \forall i \in V^S$$

$$r_{n+k} = k \quad \forall k \in M$$

$$r_{n+m+k} = k \quad \forall k \in M$$



Subtour elimination



- Most CP libraries have built-ins
 - MiniZinc: 'circuit'
 - Comet: 'circuit'
 - ILOG: Path constraint

Propagation – Cycles



'Path' constraint

- Propagates subtour elimination
- Also propagates cost
- $\text{path} (S, E, \text{succ}, P, z)$
 - succ array implies path
 - ensures path from nodes in S to nodes in T through nodes in P
 - variable z bounds cost of path
 - cost propagated incrementally based on shortest / longest paths



Large Neighbourhood Search revisited

Large Neighbourhood Search



Destroy & Re-create

- Destroy part of the solution
 - Remove visits from the solution
- Re-create solution
 - Use insert method to re-insert customers
 - Different insert methods lead to new (better?) solutions
- If the solution is better, keep it
- Repeat

Large Neighbourhood Search



Destroy part of the solution (*Select* method)

In CP terms, this means:

- Relax some variable assignments

In CP-VRP terms, this means

- Relax some *routeOf* and *successor* assignments

Large Neighbourhood Search



Re-create solution

- Use insert methods
- Uses full power of CP propagations

A MiniZinc VRP model



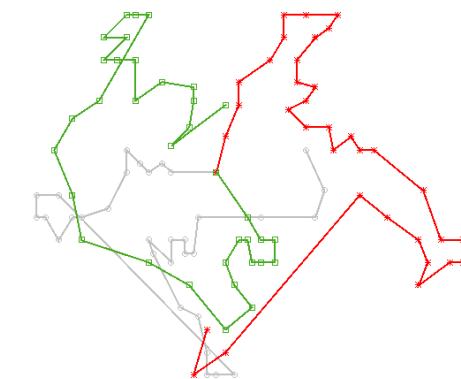
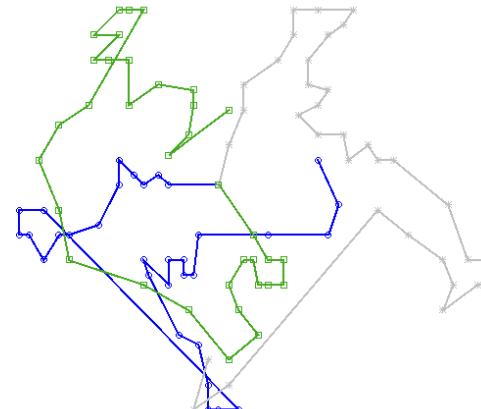
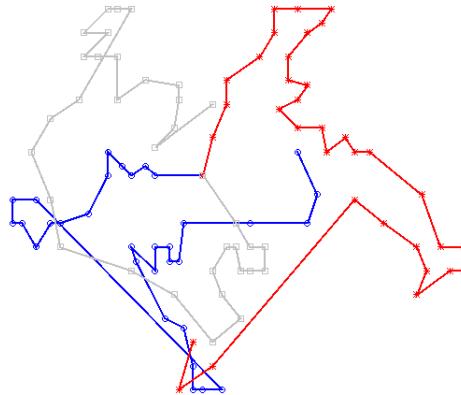
Advanced techniques – Recreate



Adaptive Decomposition

Decompose problem

- Only consider 2-3 routes
- Smaller problem is much easier to solve



Advanced techniques – Recreate



Adaptive Decomposition

Decompose problem

- Only consider 2-3 routes
- Smaller problem is much easier to solve

Adaptive

- Decompose in different ways
- Use problem features to determine decomposition

Conclusions



- Now you know
 - How to construct a solution to a VRP by successive insertion
 - How to improve the solution using
 - Variable Neighbourhood Search
 - Large Neighbourhood Search
- Argued that CP is “natural” for solving vehicle routing problems
 - Real problems often have unique constraints
 - Easy to change CP model to include new constraints
 - New constraints don’t change core solve method
 - Infrastructure for complete (completish) search in subproblems
- LNS is “natural” for CP
 - Insertion leverages propagation